

Controls Engineering in the FIRST Robotics Competition

Graduate-level control theory for high schoolers

Tyler Veness

Controls Engineering in the FIRST Robotics Competition

Graduate-level control theory for high schoolers

Tyler Veness

Copyright © 2017-2020 Tyler Veness

<https://github.com/calcogul/controls-engineering-in-frc>

Licensed under the Creative Commons Attribution-ShareAlike 4.0 International License (the “License”). You may not use this file except in compliance with the License. You may obtain a copy of the License at <https://creativecommons.org/licenses/by-sa/4.0/>. Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an “AS IS” BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

Generated from commit 5132b8b made on April 23, 2020. The latest version can be downloaded from <https://file.tavsys.net/control/controls-engineering-in-frc.pdf>.

Contents

Preface	xi
0 Notes to the reader	1
0.1 Prerequisites	1
0.2 Structure of this book	1
0.3 Ethos of this book	2
0.3.1 Role of classical control theory	2
0.3.2 An integrated approach to nonlinear control theory	3
0.4 Mindset of an egoless engineer	3
0.5 Request for feedback	3

I

Fundamentals of control theory

1 Control system basics	7
1.1 Nomenclature	7
1.2 What is gain?	8
1.3 Block diagrams	8
1.4 Why feedback control?	9
2 PID controllers	11
2.1 Proportional term	11
2.2 Derivative term	12
2.3 Integral term	14

2.4	PID controller definition	16
2.5	Response types	16
2.6	Manual tuning	17
2.7	Actuator saturation	17
2.8	Limitations	18
3	Transfer functions	19
3.1	Laplace transform	19
3.2	Parts of a transfer function	19
3.2.1	Poles and zeroes	19
3.2.2	Nonminimum phase zeroes	21
3.2.3	Pole-zero cancellation	21
3.3	Transfer functions in feedback	22
3.3.1	Root locus	22

II

Modern control theory

4	Linear algebra	27
4.1	Vectors	27
4.2	Linear combinations, span, and basis vectors	27
4.3	Linear transformations and matrices	27
4.4	Matrix multiplication as composition	27
4.5	The determinant	28
4.6	Inverse matrices, column space, and null space	28
4.7	Nonsquare matrices as transformations between dimensions	28
4.8	Eigenvectors and eigenvalues	28
4.9	Miscellaneous notation	28
5	State-space controllers	29
5.1	From PID control to model-based control	29
5.2	What is a dynamical system?	30
5.3	State-space notation	30
5.3.1	What is state-space?	30
5.3.2	Benefits over classical control	30
5.3.3	Definition	31
5.4	Controllability	32
5.5	Observability	32
5.6	Closed-loop controller	33
5.7	Pole placement	34
5.8	Linear-quadratic regulator	34
5.8.1	Bryson's rule	35
5.8.2	Pole placement vs LQR	36

5.9 Model augmentation	37
5.9.1 Plant augmentation	37
5.9.2 Controller augmentation	37
5.9.3 Observer augmentation	37
5.9.4 Output augmentation	38
5.9.5 Examples	38
5.10 Feedforward	39
5.10.1 Plant inversion	40
5.10.2 Unmodeled dynamics	41
5.11 Integral control	42
5.11.1 Plant augmentation	43
5.11.2 Input error estimation	43
6 Digital control	45
6.1 Phase loss	45
6.2 s-plane to z-plane	46
6.2.1 z-plane stability	46
6.2.2 z-plane behavior	47
6.2.3 Nyquist frequency	47
6.3 Discretization methods	48
6.4 Effects of discretization on controller performance	50
6.5 Matrix exponential	52
6.6 Taylor series	53
6.7 Zero-order hold for state-space	54
7 Nonlinear control	57
7.1 Introduction	57
7.2 Linearization	57
7.3 Lyapunov stability	58
7.4 Affine systems	59
7.4.1 Feedback linearization for reference tracking	59
7.5 Further reading	60
8 State-space applications	61
8.1 Elevator	61
8.1.1 Continuous state-space model	61
8.1.2 Model augmentation	62
8.1.3 Gravity feedforward	62
8.1.4 Simulation	63
8.1.5 Implementation	63
8.2 Flywheel	63
8.2.1 Continuous state-space model	63
8.2.2 Model augmentation	64
8.2.3 Simulation	64
8.2.4 Implementation	64
8.2.5 Flywheel model without encoder	64

8.2.6	Voltage compensation	66
8.3	Single-jointed arm	66
8.3.1	Continuous state-space model	66
8.3.2	Model augmentation	67
8.3.3	Gravity feedforward	67
8.3.4	Simulation	68
8.3.5	Implementation	68
8.4	Pendulum	69
8.4.1	State-space model	69
8.5	Differential drive	70
8.5.1	Velocity subspace state-space model	70
8.5.2	Linear time-varying model	72
8.5.3	Improving model accuracy	74
8.5.4	Cross track error controller	75
8.5.5	Explicit time-varying control law	76
8.5.6	Nonlinear observer design	78
8.6	Ramsete unicycle controller	82
8.6.1	Velocity and turning rate command derivation	82
8.6.2	Nonlinear controller equations	83
8.6.3	Linear reference tracker	84
8.7	Linear time-varying unicycle controller (cascaded)	87
8.7.1	Explicit time-varying control law	88

III

Estimation and localization

9	Stochastic control theory	93
9.1	Terminology	93
9.2	State observers	93
9.2.1	Luenberger observer	94
9.3	Introduction to probability	96
9.3.1	Random variables	96
9.3.2	Expected value	97
9.3.3	Variance	97
9.3.4	Joint probability density functions	98
9.3.5	Covariance	99
9.3.6	Correlation	99
9.3.7	Independence	99
9.3.8	Marginal probability density functions	100
9.3.9	Conditional probability density functions	100
9.3.10	Bayes's rule	100
9.3.11	Conditional expectation	100
9.3.12	Conditional variances	100
9.3.13	Random vectors	100
9.3.14	Covariance matrix	101
9.3.15	Relations for independent random vectors	101
9.3.16	Gaussian random variables	102

9.4	Linear stochastic systems	103
9.4.1	State vector expectation evolution	103
9.4.2	Error covariance matrix evolution	103
9.4.3	Measurement vector expectation	104
9.4.4	Measurement covariance matrix	104
9.5	Two-sensor problem	104
9.6	Kalman filter	105
9.6.1	Derivations	105
9.6.2	Predict and update equations	108
9.6.3	Setup	110
9.6.4	Simulation	111
9.6.5	Kalman filter as Luenberger observer	112
9.7	Kalman smoother	114
9.7.1	Derivations	114
9.7.2	State update equation	114
9.7.3	Error covariance equation	115
9.7.4	Optimal estimate	116
9.7.5	Predict and update equations	117
9.7.6	Example	118
9.8	Extended Kalman filter	120
9.9	Unscented Kalman filter	120
9.10	Multiple model adaptive estimation	121
10	Pose estimation	123
10.1	Euler integration	123
10.2	Pose exponential	123
10.2.1	What is a group?	123
10.2.2	What is a pose?	124
10.2.3	What is a twist?	124
10.2.4	Derivation	124
10.2.5	Lie groups	127
10.3	Pose correction	128

11	Calculus	131
11.1	Derivatives	131
11.1.1	Power rule	132
11.1.2	Product rule	132
11.1.3	Chain rule	132
11.2	Integrals	132
11.2.1	Change of variables	133
11.3	Tables	133
11.3.1	Common derivatives and integrals	133

12	Dynamics	135
12.1	Linear motion	135
12.2	Angular motion	135
12.3	Vectors	136
12.3.1	Basic vector operations	136
12.3.2	Cross product	136
12.4	Curvilinear motion	137
12.4.1	Differential drive	137
12.4.2	Mecanum drive	139
12.4.3	Swerve drive	142
13	Model examples	145
13.1	DC brushed motor	145
13.1.1	Equations of motion	146
13.1.2	Calculating constants	146
13.1.3	Current limiting	148
13.2	Elevator	148
13.2.1	Equations of motion	148
13.3	Flywheel	150
13.3.1	Equations of motion	150
13.3.2	Calculating constants	151
13.4	Single-jointed arm	152
13.4.1	Equations of motion	152
13.4.2	Calculating constants	153
13.5	Pendulum	154
13.5.1	Force derivation	154
13.5.2	Torque derivation	155
13.5.3	Energy derivation	156
13.6	Differential drive	157
13.6.1	Equations of motion	157
13.6.2	Calculating constants	160
14	System identification	161
14.1	1 DOF mechanism model	162
14.2	Drivetrain velocity model	163

15	Motion profiles	169
15.1	1 DOF motion profiles	169
15.1.1	Jerk	170
15.1.2	Profile selection	170
15.1.3	Profile equations	171
15.1.4	Other profile types	171
15.1.5	Further reading	172

15.2	2 DOF motion profiles	172
16	Trajectory optimization	173

VI

Appendices

A	Simplifying block diagrams	177
A.1	Cascaded blocks	177
A.2	Combining blocks in parallel	177
A.3	Removing a block from a feedforward loop	178
A.4	Eliminating a feedback loop	178
A.5	Removing a block from a feedback loop	179
B	Laplace domain analysis	181
B.1	Projections	181
B.2	Fourier transform	184
B.3	Laplace transform	186
B.4	Laplace transform definition	188
B.5	Case study: steady-state error	188
B.6	Case study: flywheel PID control	190
C	Robust control	195
C.1	Gain margin and phase margin	195
D	Installing Python packages	197
D.1	Windows instructions	197
D.2	Linux instructions	197
E	Linear-quadratic regulator	199
E.1	Derivation	199
E.2	Output feedback	202
E.3	Implicit model following	203
F	QR-weighted linear plant inversion	207
F.1	Necessary theorems	207
F.2	Setup	208
F.3	Minimization	208
G	Steady-state feedforward	211
G.1	Continuous case	211
G.2	Discrete case	212

G.3	Deriving steady-state input	213
G.3.1	Case study: second-order CIM motor model	213
H	Derivations	215
H.1	Transfer function in feedback	215
H.2	Zero-order hold for state-space	216
H.3	Kalman filter as Luenberger observer	217
H.3.1	Luenberger observer with separate prediction and update	218
H.4	Trapezoidal motion profile	218
	Glossary	220
	Bibliography	221
	Online	221
	Miscellaneous	222
	Index	223



Preface

Motivation

I am the software mentor for a FIRST Robotics Competition (FRC) team. My responsibilities for that include teaching programming, software engineering practices, and applications of control theory. The curriculum I developed as of the spring of 2017 (located at <https://csweb.frc3512.com/ci/>) taught rookies enough to be minimally competitive, but it provided no formal avenues of growth for veteran students.

Also, out of a six week build season, the software team usually only got a few days with the completed robot due to poor build schedule management. This led to two problems. First, two days is only enough time to verify basic software functionality, not test and tune feedback controllers. Second, this was also the first time the robot's electromechanical systems have been tested after integration, so any issues that arose consumed valuable software integration time while the team traced the problem to a mechanical, electrical, or software cause.

This book expands my curriculum to cover control theory topics I learned in my graduate-level engineering classes at University of California, Santa Cruz. It introduces state-space controllers and serves as a practical guide for formulating and implementing them. Since state-space control utilizes a system model, both problems mentioned earlier can be addressed. Basic software functionality can be tested against it and feedback controllers can be tuned automatically based on system constraints. This allows software teams to test their work much earlier in the build season in a controlled environment as well as save time during feedback controller design, implementation, and testing.

I originally started writing this book because I felt that the topic wasn't difficult, but the information for it wasn't easily accessible. When I was a high school robotics team member, I had to learn everything from scratch through various internet sources and eventually from college courses as part of my bachelor's degree. Control theory has a certain beauty to it, and I want more people to appreciate it the way I do. Therefore, my goal is to make the learning process quicker and easier for future team members by collating all the relevant information.

Intended audience

This guide is intended to make an advanced engineering topic approachable so it can be applied by those who aren't experts in control theory. My intended audience is high school students who are members of a FIRST Robotics Competition team. As such, they will likely have passing knowledge of PID control and have basic proficiency in programming. This guide will expand their incomplete understanding of control theory to include the fundamentals of classical control theory, enough linear algebra to understand the notation and mechanics of modern control, and finally how to apply modern control to design challenges they regularly see in their FRC robots from year to year.

Acknowledgements

I would like to thank my controls engineering instructors Dejan Milutinović and Gabriel Elkaim of University of California, Santa Cruz. They taught their classes from a pragmatic perspective focused on application and intuition that I appreciated. I would also like to thank Dejan Milutinović for introducing me to the field of control theory and both of my instructors for teaching me what it means to be a controls engineer.

Thanks to Austin Schuh from FRC team 971 for providing the final continuous state-space models used in the examples section.



0. Notes to the reader

0.1 Prerequisites

Knowledge of basic algebra and complex numbers is assumed. Some introductory physics and calculus will be taught as necessary.

0.2 Structure of this book

This book consists of five parts and a collection of appendices that address the four tasks a controls engineer carries out: derive a model of the system (kinematics), design a controller for the model (control theory), design an observer to estimate the current state of the model (localization), and plan how the controller is going to drive the model to a desired state (motion planning).

Part I, “Fundamentals of control theory,” introduces the basics of control theory, teaches the fundamentals of PID controller design, describes what a transfer function is, and shows how they can be used to analyze dynamical systems. Emphasis is placed on the geometric intuition of this analysis rather than the frequency domain math.

Part II, “Modern control theory,” first provides a crash course in the geometric intuition behind linear algebra and covers enough of the mechanics of evaluating matrix algebra for the reader to follow along in later chapters. It covers state-space representation, controllability, and observability. The intuition gained in part I and the notation of linear algebra are used to model and control linear multiple-input, multiple-output (MIMO) systems and covers discretization, LQR controller design, LQE observer design, and feedforwards. Then, these concepts are applied to design and implement controllers for real systems. The examples from part IV are converted to state-space representation, implemented, and tested with a discrete controller.

Part II also introduces the basics of nonlinear control system analysis with Lyapunov functions. It presents an example of a nonlinear controller for a unicycle-like vehicle as well as how to apply it to a two-wheeled vehicle. Since nonlinear control isn’t the focus of this book, we mention other books and resources for further reading.

Part III, “Estimation and localization,” introduces the field of stochastic control theory. The Luenberger observer and the probability theory behind the Kalman filter is taught with several examples of creative applications of Kalman filter theory.

Part IV, “System modeling,” introduces the basic calculus and physics concepts required to derive the models used in the previous chapters. It walks through the derivations for several common FRC subsystems. Then, methods for system identification are discussed for empirically measuring model parameters.

Part V, “Motion planning,” covers planning how the robot will get from its current state to some desired state in a manner achievable by its dynamics. It introduces motion profiles with one degree of freedom for simple maneuvers. Trajectory optimization methods are presented for generating profiles with higher degrees of freedom.

The appendices provide further enrichment that isn’t required for a passing understanding of the material. This includes derivations for many of the results presented and used in the mainmatter of the book.

The Python scripts used to generate the plots in the case studies double as reference implementations of the techniques discussed in their respective chapters. They are available in this book’s Git repository. Its location is listed on the copyright page.

0.3 Ethos of this book

This book is intended as both a tutorial for new students and as a reference manual for more experienced readers who need to review a thing or two. While it isn’t comprehensive, the reader will hopefully learn enough to either implement the concepts presented themselves or know where to look for more information.

Some parts are mathematically rigorous, but I believe in giving students a solid theoretical foundation with emphasis on intuition so they can apply it to new problems. To achieve deep understanding of the topics in this book, math is unavoidable. With that said, I try to provide practical and intuitive explanations whenever possible.

0.3.1 Role of classical control theory

The sections on classical control theory are only included to develop geometric intuition for the mathematical machinery of modern control theory. Many tools exclusive to classical control theory (root locus, Bode plots, Nyquist plots, etc.) aren’t useful for or relevant to the examples presented, so they serve only to complicate the learning process.

Classical control theory is interesting in that one can perform stability and robustness analyses and design reasonable controllers for systems on the back of a napkin. It’s also useful for controlling systems which don’t have a model. One can generate a Bode plot of a system by feeding in sine waves of increasing frequency and recording the amplitude of the output oscillations. This data can be used to create a transfer function or lead and lag compensators can be applied directly based on the Bode plot. However, computing power is much more plentiful nowadays; we should take advantage of this in the design phase and use the more modern tools that enables when it makes sense.

This book uses LQR and modern control over, say, loop shaping with Bode and Nyquist plots because we have accurate dynamical models to leverage, and LQR allows directly expressing what the author is concerned with optimizing: state excursion relative to control effort. Applying lead and lag compensators, while effective for robust controller design, doesn’t provide the same expressive power.

0.3.2 An integrated approach to nonlinear control theory

Most teaching resources separate linear and nonlinear control with the latter being reserved for a different course. Here, they are introduced together because the concepts of nonlinear control apply often, and it isn't that much of a leap (if Lyapunov stability isn't included). The control and estimation chapters cover relevant tools for dealing with nonlinearities like linearization when appropriate.

0.4 Mindset of an egoless engineer

Engineering has a mindset, not just a skillset. Engineers have a unique way of approaching problems, and the following maxim summarizes what I hope to teach my robotics students (with examples drawn from controls engineering).

“Engineer based on requirements, not an ideology.”

Engineering is filled with trade-offs. The tools should fit the job, and not every problem is a nail waiting to be struck by a hammer. Instead, assess the minimum requirements (min specs) for a solution to the task at hand and do only enough work to satisfy them; exceeding your specifications is a waste of time and money. If you require performance or maintainability above the min specs, your min specs were chosen incorrectly by definition.

Controls engineering is pragmatic in a similar respect: *solve. the. problem.* For control of nonlinear systems, plant inversion is elegant on paper but doesn't work with an inaccurate model, yet using a theoretically incorrect solution like linear approximations of the nonlinear system works well enough to be used industry-wide. There are more sophisticated controllers than PID, but we use PID anyway for its versatility and simplicity. Sometimes the inferior solutions are more effective or have a more desirable cost-benefit ratio than what the control system designer considers ideal or clean. Choose the tool that is most effective.

Solutions need to be good enough, but do not need to be perfect. We want to avoid integrators as they introduce instability, but we use them anyway because they work well for meeting tracking specifications. One should not blindly defend a design or follow an ideology, because there is always a case where its antithesis is a better option. The engineer should be able to determine when this is the case, set aside their ego, and do what will meet the specifications of their client (e.g., system response characteristics, maintainability, usability). Preferring one solution over another for pragmatic or technical reasons is fine, but the engineer should not care on a personal level which sufficient solution is chosen.

0.5 Request for feedback

While we have tried to write a book that makes the topics of control theory approachable, it still may be dense or fast-paced for some readers (it covers three classes of feedback control, two of which are for graduate students, in one short book). Please send us feedback, corrections, or suggestions through the GitHub link listed on the copyright page. New examples that demonstrate key concepts and make them more accessible are also appreciated.

This page intentionally left blank

Fundamentals of control theory

1	Control system basics	7
1.1	Nomenclature	
1.2	What is gain?	
1.3	Block diagrams	
1.4	Why feedback control?	
2	PID controllers	11
2.1	Proportional term	
2.2	Derivative term	
2.3	Integral term	
2.4	PID controller definition	
2.5	Response types	
2.6	Manual tuning	
2.7	Actuator saturation	
2.8	Limitations	
3	Transfer functions	19
3.1	Laplace transform	
3.2	Parts of a transfer function	
3.3	Transfer functions in feedback	

This page intentionally left blank



1. Control system basics

Control systems are all around us and we interact with them daily. A small list of ones you may have seen includes heaters and air conditioners with thermostats, cruise control and the anti-lock braking system (ABS) on automobiles, and fan speed modulation on modern laptops. [Control systems](#) monitor or control the behavior of [systems](#) like these and may consist of humans controlling them directly (manual control), or of only machines (automatic control).

How can we prove closed-loop [controllers](#) on an autonomous car, for example, will behave safely and meet the desired performance specifications in the presence of uncertainty? Control theory is an application of algebra and geometry used to analyze and predict the behavior of [systems](#), make them respond how we want them to, and make them [robust](#) to [disturbances](#) and uncertainty.

Controls engineering is, put simply, the engineering process applied to control theory. As such, it's more than just applied math. While control theory has some beautiful math behind it, controls engineering is an engineering discipline like any other that is filled with trade-offs. The solutions control theory gives should always be sanity checked and informed by our performance specifications. We don't need to be perfect; we just need to be good enough to meet our specifications (see section 0.4 for more on engineering).

1.1 Nomenclature

Most resources for advanced engineering topics assume a level of knowledge well above that which is necessary. Part of the problem is the use of jargon. While it efficiently communicates ideas to those within the field, new people who aren't familiar with it are lost. See the glossary for a list of words and phrases commonly used in control theory, their origins, and their meaning. Links to the glossary are provided for certain words throughout the book and will use [this color](#).

The [system](#) or collection of actuators being controlled by a [control system](#) is called the [plant](#). A [controller](#) is used to drive the [plant](#) from its current state to some desired state (the [reference](#)). Controllers which don't include information measured from the [plant](#)'s [output](#) are called [open-loop controllers](#).

Controllers which incorporate information fed back from the plant's output are called closed-loop controllers or feedback controllers. Figure 1.1 shows a plant in feedback with a controller.

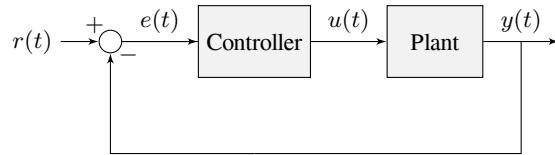


Figure 1.1: Control system nomenclature for a closed-loop system

$r(t)$	reference	$u(t)$	control input
$e(t)$	error	$y(t)$	output

Note that the input and output of a system are defined from the plant's point of view. The negative feedback controller shown is driving the difference between the reference and output, also known as the error, to zero.

1.2 What is gain?

Gain is a proportional value that shows the relationship between the magnitude of an input signal to the magnitude of an output signal at steady-state. Many systems contain a method by which the gain can be altered, providing more or less "power" to the system.

Figure 1.2 shows a system with a hypothetical input and output. Since the output is twice the amplitude of the input, the system has a gain of 2.

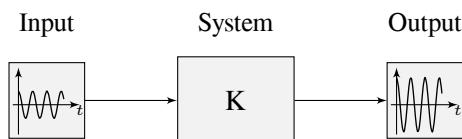


Figure 1.2: Demonstration of system with a gain of $K = 2$

1.3 Block diagrams

When designing or analyzing a control system, it is useful to model it graphically. Block diagrams are used for this purpose. They can be manipulated and simplified systematically (see appendix A). Figure 1.3 is an example of one.

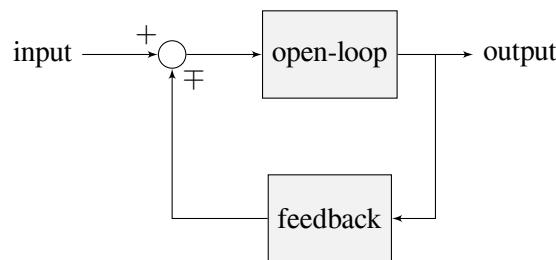


Figure 1.3: Block diagram with nomenclature

The open-loop gain is the total gain from the sum node at the input (the circle) to the output branch.

This would be the [system](#)'s gain if the feedback loop was disconnected. The [feedback gain](#) is the total gain from the output back to the input sum node. A sum node's output is the sum of its inputs.

Figure 1.4 is a block diagram with more formal notation in a feedback configuration.

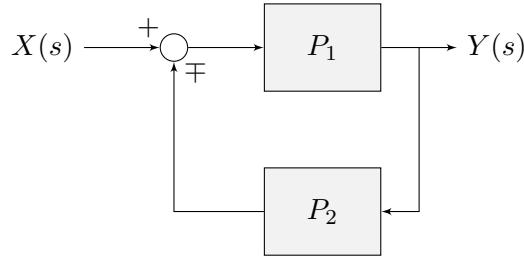


Figure 1.4: Feedback block diagram

\mp means “minus or plus” where a minus represents negative feedback.

1.4 Why feedback control?

Let's say we are controlling a DC brushed motor. With just a [mathematical model](#) and knowledge of all current [states](#) of the [system](#) (i.e., angular velocity), we can predict all future [states](#) given the future voltage [inputs](#). Why then do we need feedback control? If the [system](#) is [disturbed](#) in any way that isn't modeled by our equations, like a load was applied to the armature, or voltage sag in the rest of the circuit caused the commanded voltage to not match the actual applied voltage, the angular velocity of the motor will deviate from the [model](#) over time.

To combat this, we can take measurements of the [system](#) and the environment to detect this deviation and account for it. For example, we could measure the current position and estimate an angular velocity from it. We can then give the motor corrective commands as well as steer our [model](#) back to reality. This feedback allows us to account for uncertainty and be [robust](#) to it.

This page intentionally left blank

2. PID controllers

The PID controller is a commonly used feedback controller consisting of proportional, integral, and derivative terms, hence the name. This chapter will build up the definition of a PID controller term by term while trying to provide intuition for how each of them behaves.

First, we'll get some nomenclature for PID controllers out of the way. The [reference](#) is called the [setpoint](#) (the desired position) and the [output](#) is called the [process variable](#) (the measured position). Below are some common variable naming conventions for relevant quantities.

$r(t)$	setpoint	$u(t)$	control input
$e(t)$	error	$y(t)$	output

The [error](#) $e(t)$ is $r(t) - y(t)$.

For those already familiar with PID control, this book's interpretation won't be consistent with the classical intuition of "past", "present", and "future" error. We will be approaching it from the viewpoint of modern control theory with proportional controllers applied to different physical quantities we care about. This will provide a more complete explanation of the derivative term's behavior for constant and moving [setpoints](#), and this intuition will carry over to the modern control methods covered later in this book.

The proportional term drives the position error to zero, the derivative term drives the velocity error to zero, and the integral term accumulates the area between the [setpoint](#) and [output](#) plots over time (the integral of position [error](#)) and adds the current total to the [control input](#). We'll go into more detail on each of these.

2.1 Proportional term

The *Proportional* term drives the position error to zero.

Definition 2.1.1 — Proportional controller.

$$u(t) = K_p e(t) \quad (2.1)$$

where K_p is the proportional gain and $e(t)$ is the error at the current time t .

Figure 2.1 shows a block diagram for a system controlled by a P controller.

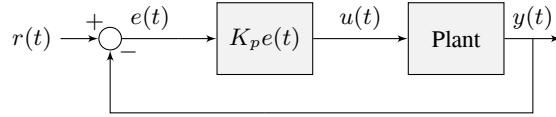


Figure 2.1: P controller block diagram

Proportional gains act like “software-defined springs” that pull the system toward the desired position. Recall from physics that we model springs as $F = -kx$ where F is the force applied, k is a proportional constant, and x is the displacement from the equilibrium point. This can be written another way as $F = k(0 - x)$ where 0 is the equilibrium point. If we let the equilibrium point be our feedback controller’s setpoint, the equations have a one-to-one correspondence.

$$F = k(r - x)$$

$$u(t) = K_p e(t) = K_p(r(t) - y(t))$$

so the “force” with which the proportional controller pulls the system’s output toward the setpoint is proportional to the error, just like a spring.

2.2 Derivative term

The *Derivative* term drives the velocity error to zero.

Definition 2.2.1 — PD controller.

$$u(t) = K_p e(t) + K_d \frac{de}{dt} \quad (2.2)$$

where K_p is the proportional gain, K_d is the derivative gain, and $e(t)$ is the error at the current time t .

Figure 2.2 shows a block diagram for a system controlled by a PD controller.

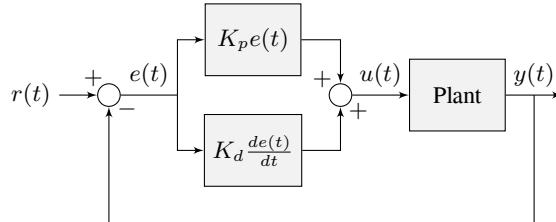


Figure 2.2: PD controller block diagram

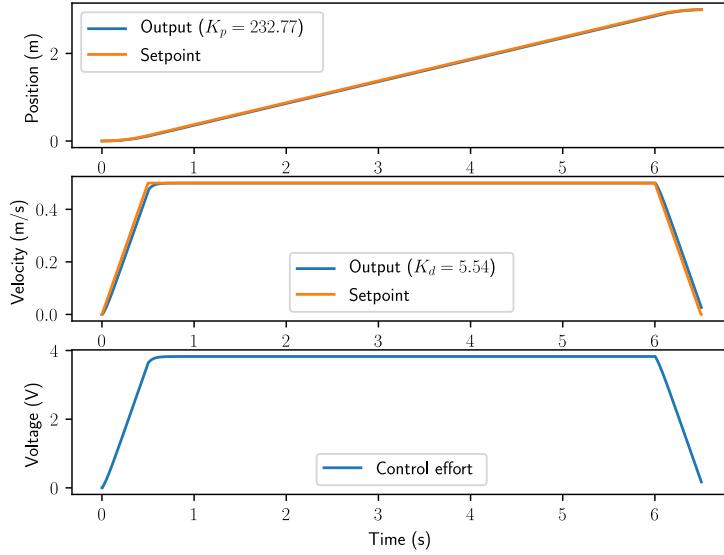


Figure 2.3: PD controller on an elevator

A PD controller has a proportional controller for position (K_p) and a proportional controller for velocity (K_d). The velocity **setpoint** is implicitly provided by how the position **setpoint** changes over time. Figure 2.3 shows an example for an elevator.

To prove a PD controller is just two proportional controllers, we will rearrange the equation for a PD controller.

$$u_k = K_p e_k + K_d \frac{e_k - e_{k-1}}{dt}$$

where u_k is the **control input** at timestep k and e_k is the **error** at timestep k . e_k is defined as $e_k = r_k - x_k$ where r_k is the **setpoint** and x_k is the current **state** at timestep k .

$$\begin{aligned} u_k &= K_p(r_k - x_k) + K_d \frac{(r_k - x_k) - (r_{k-1} - x_{k-1})}{dt} \\ u_k &= K_p(r_k - x_k) + K_d \frac{r_k - x_k - r_{k-1} + x_{k-1}}{dt} \\ u_k &= K_p(r_k - x_k) + K_d \frac{r_k - r_{k-1} - x_k + x_{k-1}}{dt} \\ u_k &= K_p(r_k - x_k) + K_d \frac{(r_k - r_{k-1}) - (x_k - x_{k-1})}{dt} \\ u_k &= K_p(r_k - x_k) + K_d \left(\frac{r_k - r_{k-1}}{dt} - \frac{x_k - x_{k-1}}{dt} \right) \end{aligned}$$

Notice how $\frac{r_k - r_{k-1}}{dt}$ is the **velocity of the setpoint**. By the same reasoning, $\frac{x_k - x_{k-1}}{dt}$ is the **system's velocity** at a given timestep. That means the K_d term of the PD controller is driving the estimated velocity to the **setpoint** velocity.

If the **setpoint** is constant, the **implicit velocity setpoint** is zero, so the K_d term slows the **system** down if it's moving. This acts like a “software-defined damper”. These are commonly seen on door closers, and their damping force increases linearly with velocity.

2.3 Integral term

The *Integral* term accumulates the area between the *setpoint* and *output* plots over time (i.e., the integral of position *error*) and adds the current total to the *control input*. Accumulating the area between two curves is called integration.

Definition 2.3.1 — PI controller.

$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau \quad (2.3)$$

where K_p is the proportional gain, K_i is the integral gain, $e(t)$ is the error at the current time t , and τ is the integration variable.

The integral integrates from time 0 to the current time t . We use τ for the integration because we need a variable to take on multiple values throughout the integral, but we can't use t because we already defined that as the current time.

Figure 2.4 shows a block diagram for a *system* controlled by a PI controller.

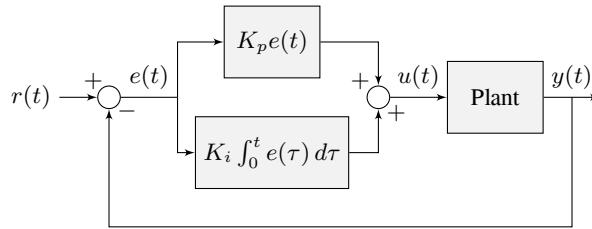


Figure 2.4: PI controller block diagram

When the *system* is close to the *setpoint* in steady-state, the proportional term may be too small to pull the *output* all the way to the *setpoint*, and the derivative term is zero. This can result in *steady-state error*, as shown in figure 2.5.

A common way of eliminating *steady-state error* is to integrate the *error* and add it to the *control input*. This increases the *control effort* until the *system* converges. Figure 2.5 shows an example of *steady-state error* for a flywheel, and figure 2.6 shows how an integrator added to the flywheel controller eliminates it. However, too high of an integral gain can lead to overshoot, as shown in figure 2.7.

R There are better approaches to fix *steady-state error* like using feedforwards or constraining when the integral control acts using other knowledge of the *system*. We will discuss these in more detail when we get to modern control theory.

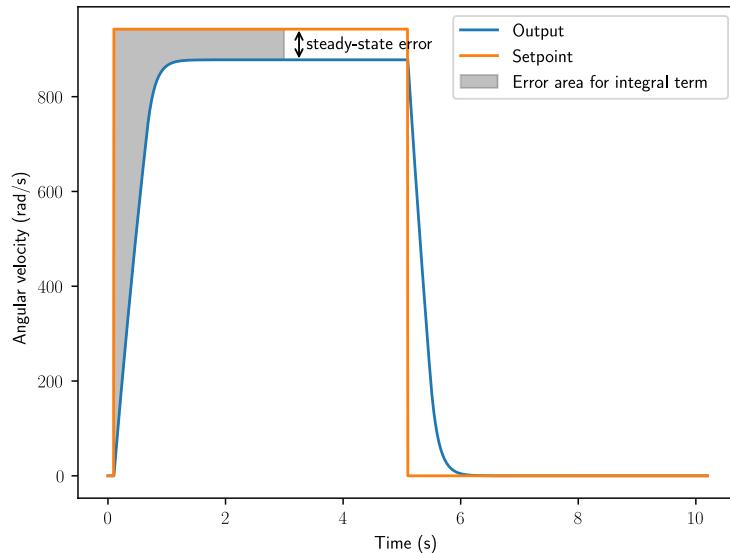


Figure 2.5: P controller on a flywheel with steady-state error

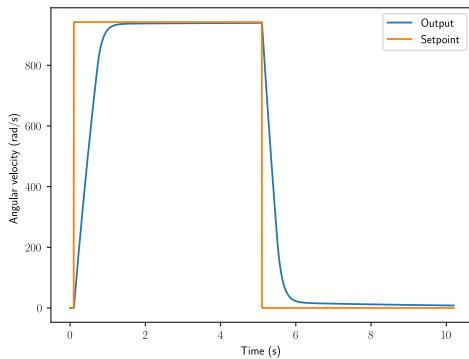
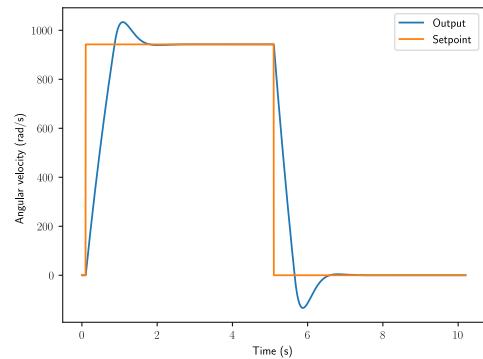


Figure 2.6: PI controller on a flywheel without steady-state error

Figure 2.7: PI controller on a flywheel with overshoot from large K_i gain

2.4 PID controller definition

When these three terms are combined, one gets the typical definition for a PID controller.

Definition 2.4.1 — PID controller.

$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau + K_d \frac{de}{dt} \quad (2.4)$$

where K_p is the proportional gain, K_i is the integral gain, K_d is the derivative gain, $e(t)$ is the error at the current time t , and τ is the integration variable.

Figure 2.8 shows a block diagram for a system controlled by a PID controller.

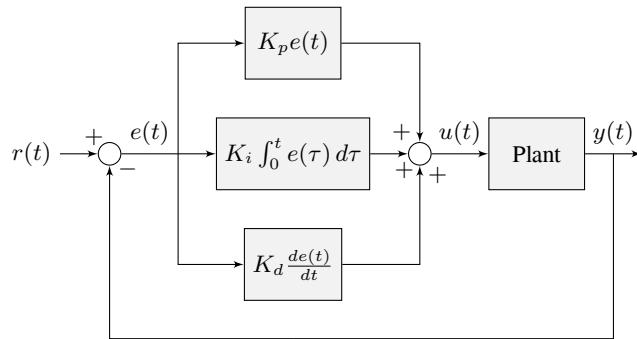


Figure 2.8: PID controller block diagram

2.5 Response types

A system driven by a PID controller generally has three types of responses: underdamped, overdamped, and critically damped. These are shown in figure 2.9.

For the step responses in figure 2.9, rise time is the time the system takes to initially reach the reference after applying the step input. Settling time is the time the system takes to settle at the reference after the step input is applied.

An *underdamped* response oscillates around the reference before settling. An *overdamped* response is slow to rise and does not overshoot the reference. A *critically damped* response has the fastest rise time without overshooting the reference.

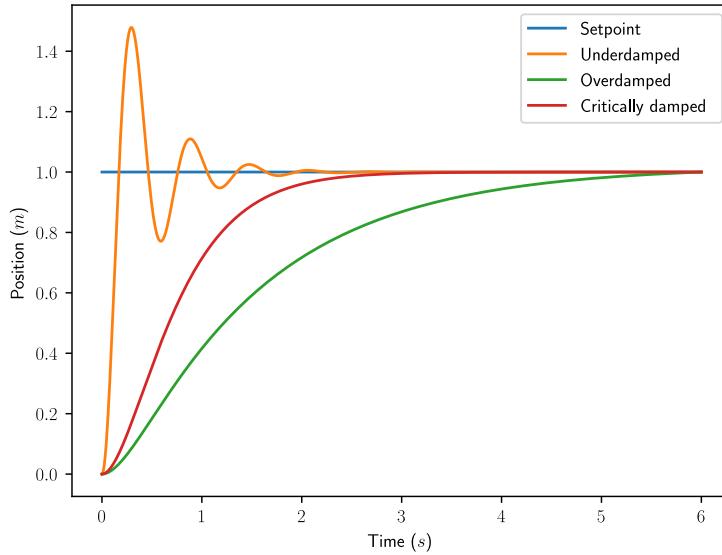


Figure 2.9: PID controller response types

2.6 Manual tuning

These steps apply to position PID controllers. Velocity PID controllers typically don't need K_d .

1. Set K_p , K_i , and K_d to zero.
2. Increase K_p until the [output](#) starts to oscillate around the [setpoint](#).
3. Increase K_d as much as possible without introducing jittering in the [system response](#).

If the [setpoint](#) follows a trapezoidal motion profile (see section 15.1), tuning becomes a lot easier. Plot the position [setpoint](#), velocity [setpoint](#), measured position, and measured velocity. The velocity [setpoint](#) can be obtained via numerical differentiation of the position [setpoint](#) (i.e., $v_{desired,k} = \frac{r_k - r_{k-1}}{\Delta t}$). Increase K_p until the position tracks well, then increase K_d until the velocity tracks well.

If the [controller](#) settles at an [output](#) above or below the [setpoint](#), one can increase K_i such that the [controller](#) reaches the [setpoint](#) in a reasonable amount of time. However, a steady-state feedforward is strongly preferred over integral control (especially for velocity PID control).



Note: Adding an integral gain to the [controller](#) is an incorrect way to eliminate [steady-state error](#). A better approach would be to tune it with an integrator added to the [plant](#), but this requires a [model](#). Since we are doing output-based rather than model-based control, our only option is to add an integrator to the [controller](#).

Beware that if K_i is too large, integral windup can occur. Following a large change in [setpoint](#), the integral term can accumulate an error larger than the maximal [control input](#). As a result, the system overshoots and continues to increase until this accumulated error is unwound.

2.7 Actuator saturation

A controller calculates its output based on the error between the [reference](#) and the current state. Plant in the real world don't have unlimited control authority available for the controller to apply. When the actuator limits are reached, the controller acts as if the gain has been temporarily reduced.

We'll try to explain this through a bit of math. Let's say we have a controller $u = k(r - x)$ where u is the [control effort](#), k is the [gain](#), r is the [reference](#), and x is the current [state](#). Let u_{max} be the limit of the actuator's output which is less than the uncapped value of u and k_{max} be the associated maximum gain. We will now compare the capped and uncapped controllers for the same [reference](#) and current [state](#).

$$\begin{aligned} u_{max} &< u \\ k_{max}(r - x) &< k(r - x) \\ k_{max} &< k \end{aligned}$$

For the inequality to hold, k_{max} must be less than the original value for k . This reduced gain is evident in a [system response](#) when there is a linear change in state instead of an exponential one as it approaches the [reference](#). This is due to the [control effort](#) no longer following a decaying exponential plot. Once the [system](#) is closer to the [reference](#), the controller will stop saturating and produce realistic controller values again.

2.8 Limitations

PID's heuristic method of tuning is a reasonable choice when there is no *a priori* knowledge of the [system](#) dynamics. However, controllers with much better response can be developed if a [dynamical model](#) of the [system](#) is known. Furthermore, PID only applies to single-input, single-output (SISO) [systems](#); we'll cover methods for multiple-input, multiple-output (MIMO) control in part II of this book.

3. Transfer functions

This chapter briefly discusses what transfer functions are, how the locations of poles and zeroes affect [system response](#) and stability, and how controllers affect pole locations. As long as you gain understanding of those concepts, don't worry too much about not being able to follow the math presented here; we won't use transfer functions in modern control theory (it's mainly provided for completeness). This chapter is intended to provide a framework within which to understand results from the mathematical machinery of modern control as well as vocabulary to communicate that understanding.

3.1 Laplace transform

For an introduction to Laplace transforms and the geometric intuition behind transfer functions, we recommend watching Zach Star's video "What does the Laplace Transform really tell us? A visual explanation (plus applications)" (21 minutes) [27]. An optional, more mathematical introduction is presented in appendix B for completeness.

3.2 Parts of a transfer function

A transfer function maps an input coordinate to an output coordinate in the Laplace domain. These can be obtained by applying the Laplace transform to a differential equation and rearranging the terms to obtain a ratio of the output variable to the input variable. Equation (3.1) is an example of a transfer function.

$$H(s) = \frac{\overbrace{(s - 9 + 9i)(s - 9 - 9i)}^{\text{zeroes}}}{\underbrace{s(s + 10)}_{\text{poles}}} \quad (3.1)$$

3.2.1 Poles and zeroes

The roots of factors in the numerator of a transfer function are called *zeroes* because they make the transfer function approach zero. Likewise, the roots of factors in the denominator of a transfer

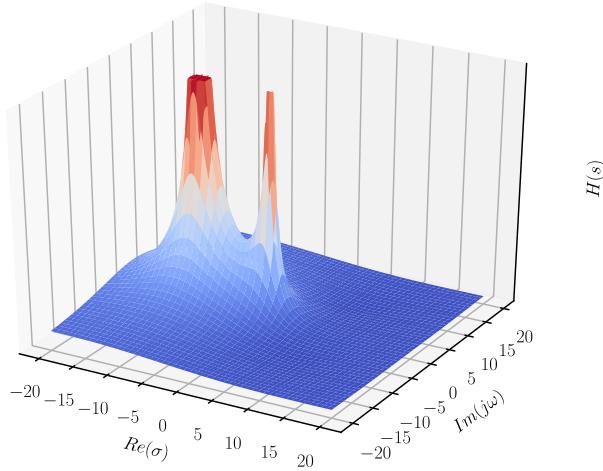


Figure 3.1: Equation (3.1) plotted in 3D

function are called *poles* because they make the transfer function approach infinity; on a 3D graph, these look like the poles of a circus tent (see figure 3.1).

When the factors of the denominator are broken apart using partial fraction expansion into something like $\frac{A}{s+a} + \frac{B}{s+b}$, the constants A and B are called residues, which determine how much each pole contributes to the [system response](#).

The factors representing poles are each the Laplace transform of a decaying exponential¹. That means the time domain responses of [systems](#) comprise decaying exponentials (e.g., $y = e^{-t}$).

R Imaginary poles and zeroes always come in complex conjugate pairs (e.g., $-2 + 3i$, $-2 - 3i$).

The locations of the closed-loop poles in the complex plane determine the stability of the [system](#). Each pole represents a frequency mode of the [system](#), and their location determines how much of each response is induced for a given input frequency. Figure 3.2 shows the [impulse responses](#) in the time domain for transfer functions with various pole locations. They all have an initial condition of 1.

Location	Stability
Left Half-plane (LHP)	Stable
Imaginary axis	Marginally stable
Right Half-plane (RHP)	Unstable

Table 3.1: Pole location and stability

When a [system](#) is stable, its output may oscillate but it converges to steady-state. When a [system](#) is

¹We are handwaving Laplace transform derivations because they are complicated and neither relevant nor useful.

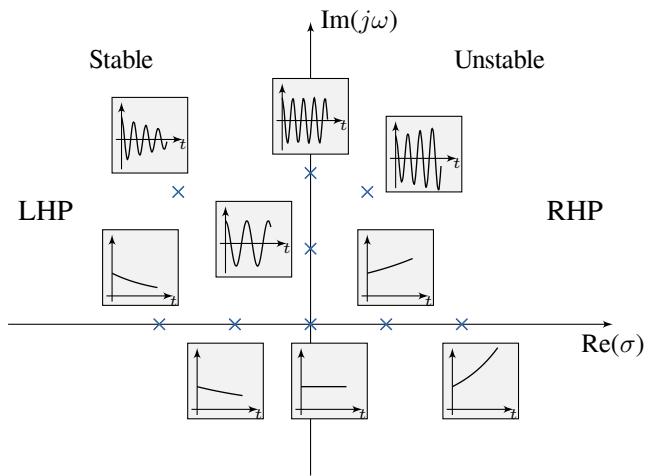


Figure 3.2: Impulse response vs pole location

marginally stable, its output oscillates at a constant amplitude forever. When a [system](#) is unstable, its output grows without bound.

3.2.2 Nonminimum phase zeroes

While poles in the RHP are unstable, the same is not true for zeroes. They can be characterized by the [system](#) initially moving in the wrong direction before heading toward the [reference](#). Since the poles always move toward the zeroes, zeroes impose a “speed limit” on the [system response](#) because it takes a finite amount of time to move the wrong direction, then change directions.

One example of this type of [system](#) is bicycle steering. Try riding a bicycle without holding the handle bars, then poke the right handle; the bicycle turns right. Furthermore, if one is holding the handlebars and wants to turn left, rotating the handlebars counterclockwise will make the bicycle fall toward the right. The rider has to lean into the turn and overpower the nonminimum phase dynamics to go the desired direction.

Another example is a Segway. To move forward by some distance, the Segway must first roll backward to rotate the Segway forward. Once the Segway starts falling in that direction, it begins rolling forward to avoid falling over until it reaches the target distance. At that point, the Segway increases its forward speed to pitch backward and slow itself down. To come to a stop, the Segway rolls backward again to level itself out.

3.2.3 Pole-zero cancellation

Pole-zero cancellation occurs when a pole and zero are located at the same place in the s -plane. This effectively eliminates the contribution of each to the [system](#) dynamics. By placing poles and zeroes at various locations (this is done by placing transfer functions in series), we can eliminate undesired [system](#) dynamics. While this may appear to be a useful design tool at first, there are major caveats. Most of these are due to [model](#) uncertainty resulting in poles which aren’t in the locations the controls designer expected.

Notch filters are typically used to dampen a specific range of frequencies in the [system response](#). If its band is made too narrow, it can still leave the undesirable dynamics, but now you can no longer measure them in the response. They are still happening, but they are what’s called *unobservable*.

Never pole-zero cancel unstable or nonminimum phase dynamics. If the [model](#) doesn’t quite reflect reality, an attempted pole cancellation by placing a nonminimum phase zero results in the pole still

moving to the zero placed next to it. You have the same dynamics as before, but the pole is also stuck where it is no matter how much **feedback gain** is applied. For an attempted nonminimum phase zero cancellation, you have effectively placed an unstable pole that's unobservable. This means the **system** will be going unstable and blowing up, but you won't be able to detect this and react to it.

Keep in mind when making design decisions that the **model** likely isn't perfect. The whole point of feedback control is to be robust to this kind of uncertainty.

3.3 Transfer functions in feedback

For **controllers** to regulate a **system** or **track** a reference, they must be placed in positive or negative feedback with the **plant** (whether to use positive or negative depends on the **plant** in question). Stable feedback loops attempt to make the **output** equal the **reference**.

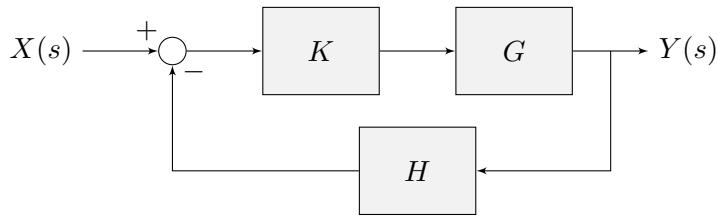


Figure 3.3: Feedback controller block diagram

$X(s)$	input	H	measurement transfer function
K	controller gain	$Y(s)$	output
G	plant transfer function		

The transfer function of figure 3.3, a **control system** diagram with feedback, from input to output is

$$G_{cl}(s) = \frac{Y(s)}{X(s)} = \frac{KG}{1 + KGH} \quad (3.2)$$

The numerator is the **open-loop gain** and the denominator is one plus the gain around the feedback loop, which may include parts of the **open-loop gain** (see appendix H.1 for a derivation). As another example, the transfer function from the input to the **error** is

$$G_{cl}(s) = \frac{E(s)}{X(s)} = \frac{1}{1 + KGH} \quad (3.3)$$

The roots of the denominator of $G_{cl}(s)$ are different from those of the open-loop transfer function $KG(s)$. These are called the **closed-loop poles**.

3.3.1 Root locus

In closed-loop, the poles can be moved around by adjusting the controller gain, but the zeroes stay put. The root locus shows where the poles will go as the gain for a P controller is increased and tells us for what range of gains the controller will be stable. As the controller gain is increased, poles can move toward negative infinity (figure 3.4), move toward each other then split toward asymptotes (figure 3.5), or move toward zeroes (figure 3.6). The **system** in figure 3.6 becomes unstable as the gain is increased.

We won't be using root locus plots for any of our control systems analysis later, but it does help provide an intuition for what **controllers** actually do to a **system**.

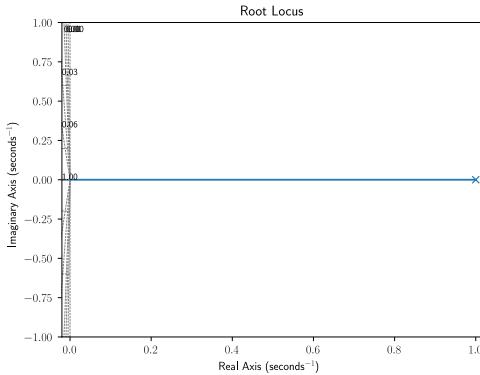


Figure 3.4: Root locus showing pole moving toward negative infinity

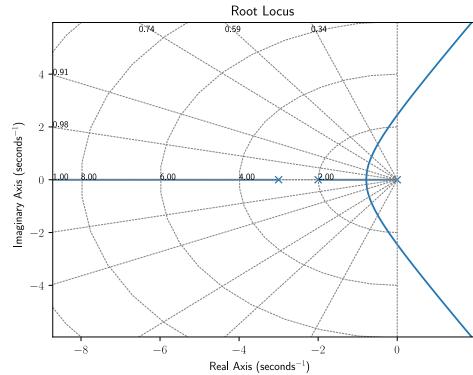


Figure 3.5: Root locus showing poles moving toward asymptotes

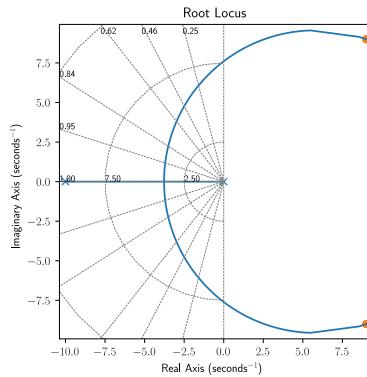


Figure 3.6: Root locus of equation (3.1) showing poles moving toward zeroes.

If poles are much farther left in the LHP than the typical [system](#) dynamics exhibit, they can be considered negligible. Every [system](#) has some form of unmodeled high frequency, nonlinear dynamics, but they can be safely ignored depending on the operating regime.

To demonstrate this, consider the transfer function for a second-order DC brushed motor (a CIM motor) from voltage to position

$$G(s) = \frac{K}{s((Js + b)(Ls + R) + K^2)}$$

where $J = 3.2284 \times 10^{-6} \text{ kg}\cdot\text{m}^2$, $b = 3.5077 \times 10^{-6} \text{ N}\cdot\text{m}\cdot\text{s}$, $K_e = K_t = 0.0181 \text{ V}/\text{rad/s}$, $R = 0.0902 \Omega$, and $L = 230 \times 10^{-6} \text{ H}$.

This [plant](#) has the root locus shown in figure 3.7. In proportional feedback, the [plant](#) is unstable for large values of K . However, if we remove the unstable pole by setting L in the transfer function to zero, we get the root locus in figure 3.8. For small values of K , both [systems](#) are stable and have nearly indistinguishable [step responses](#) due to the exceedingly small contribution from the fast pole (see figures 3.9 and 3.10). The high frequency dynamics only cause instability for large values of K that induce fast [system responses](#). In other words, the [system responses](#) of the second-order model and its first-order approximation are similar for low frequency operating regimes.

Why can't unstable poles close to the origin be ignored in the same way? The response of high

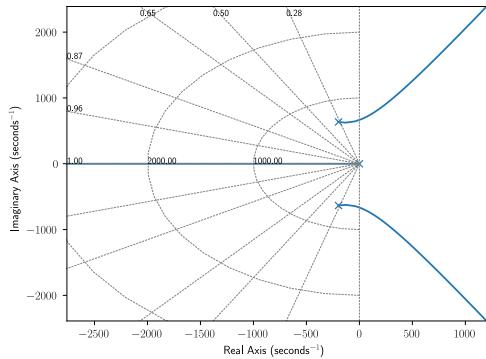


Figure 3.7: Root locus of second-order DC brushed motor plant

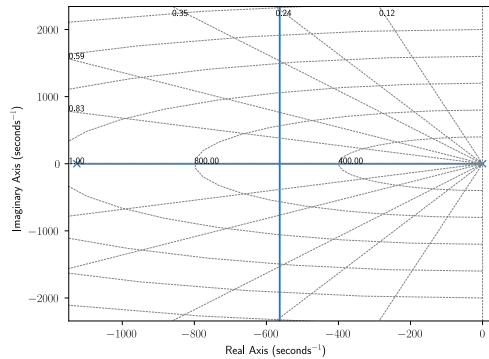


Figure 3.8: Root locus of first-order DC brushed motor plant

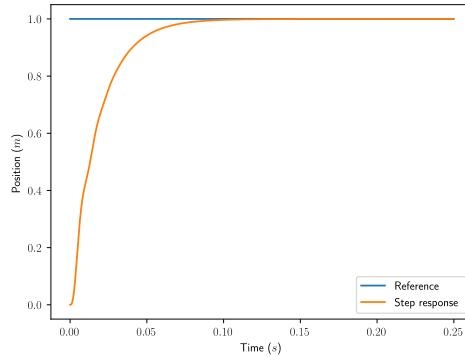


Figure 3.9: Step response of second-order DC brushed motor plant

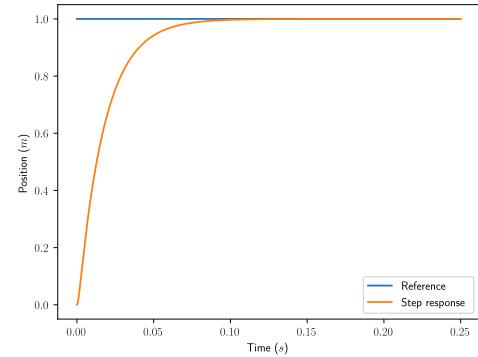


Figure 3.10: Step response of first-order DC brushed motor plant

frequency stable poles decays rapidly. Unstable poles, on the other hand, represent unstable dynamics which cause the [system output](#) to grow to infinity. Regardless of how slow these unstable dynamics are, they will eventually dominate the response.

Modern control theory

4	Linear algebra	27
4.1	Vectors	
4.2	Linear combinations, span, and basis vectors	
4.3	Linear transformations and matrices	
4.4	Matrix multiplication as composition	
4.5	The determinant	
4.6	Inverse matrices, column space, and null space	
4.7	Nonsquare matrices as transformations between dimensions	
4.8	Eigenvectors and eigenvalues	
4.9	Miscellaneous notation	
5	State-space controllers	29
5.1	From PID control to model-based control	
5.2	What is a dynamical system?	
5.3	State-space notation	
5.4	Controllability	
5.5	Observability	
5.6	Closed-loop controller	
5.7	Pole placement	
5.8	Linear-quadratic regulator	
5.9	Model augmentation	
5.10	Feedforward	
5.11	Integral control	
6	Digital control	45
6.1	Phase loss	
6.2	s-plane to z-plane	
6.3	Discretization methods	
6.4	Effects of discretization on controller performance	
6.5	Matrix exponential	
6.6	Taylor series	
6.7	Zero-order hold for state-space	
7	Nonlinear control	57
7.1	Introduction	
7.2	Linearization	
7.3	Lyapunov stability	
7.4	Affine systems	
7.5	Further reading	
8	State-space applications	61
8.1	Elevator	
8.2	Flywheel	
8.3	Single-jointed arm	
8.4	Pendulum	
8.5	Differential drive	
8.6	Ramsete unicycle controller	
8.7	Linear time-varying unicycle controller (cascaded)	

This page intentionally left blank

4. Linear algebra

Modern control theory borrows concepts from linear algebra. At first, linear algebra may appear very abstract, but there are simple geometric intuitions underlying it. First, watch 3Blue1Brown's preview video for the *Essence of linear algebra* video series (5 minutes) [4]. The goal here is to provide an intuitive, geometric understanding of linear algebra as a method of linear transformations.

We would normally include written material here for learning linear algebra, but 3Blue1Brown's animated videos are better at conveying the geometric intuition involved than anything we could write here with static text. Instead of making an inferior copy of his work, we'll provide bibliography entries for a selection of his videos.

4.1 Vectors

Watch the "Vectors, what even are they?" video from 3Blue1Brown's *Essence of linear algebra* series (5 minutes) [11].

4.2 Linear combinations, span, and basis vectors

Watch the "Linear combination, span, and basis vectors" video from 3Blue1Brown's *Essence of linear algebra* series (10 minutes) [6].

4.3 Linear transformations and matrices

Watch the "Linear transformations and matrices" video from 3Blue1Brown's *Essence of linear algebra* series (11 minutes) [7].

4.4 Matrix multiplication as composition

Watch the "Matrix multiplication as composition" video from 3Blue1Brown's *Essence of linear algebra* series (10 minutes) [8].

4.5 The determinant

Watch the “The determinant” video from 3Blue1Brown’s *Essence of linear algebra* series (10 minutes) [10].

4.6 Inverse matrices, column space, and null space

Watch the “Inverse matrices, column space, and null space” video from 3Blue1Brown’s *Essence of linear algebra* series (12 minutes) [5].

4.7 Nonsquare matrices as transformations between dimensions

Watch the “Nonsquare matrices as transformations between dimensions” video from 3Blue1Brown’s *Essence of linear algebra* series (4 minutes) [9].

4.8 Eigenvectors and eigenvalues

Watch the “Eigenvectors and eigenvalues” video from 3Blue1Brown’s *Essence of linear algebra* series (17 minutes) [3].

4.9 Miscellaneous notation

This book works with two-dimensional matrices in the sense that they only have rows and columns. The dimensionality of these matrices is specified by row first, then column. For example, a matrix with two rows and three columns would be a two-by-three matrix. A square matrix has the same number of rows as columns. Matrices commonly use capital letters while vectors use lowercase letters.

The matrix \mathbf{I} is known as the identity matrix, which is a square matrix with ones along its diagonal and zeroes elsewhere. For example

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

The matrix denoted by $\mathbf{0}_{m \times n}$ is a matrix filled with zeroes with m rows and n columns.

The T in \mathbf{A}^T denotes transpose, which flips the matrix across its diagonal such that the rows become columns and vice versa.

The † in \mathbf{B}^\dagger denotes the Moore-Penrose pseudoinverse given by $\mathbf{B}^\dagger = (\mathbf{B}^T \mathbf{B})^{-1} \mathbf{B}^T$. The pseudoinverse is used when the matrix is nonsquare and thus not invertible to produce a close approximation of an inverse in the least squares sense.

5. State-space controllers

 Chapters from here on use the `frcccontrol` Python package to demonstrate the concepts discussed and perform the complex math required. See appendix D for how to install it.

When we want to command a `system` to a set of `states`, we design a controller with certain `control laws` to do it. PID controllers use the system `outputs` with proportional, integral, and derivative `control laws`. In state-space, we also have knowledge of the system `states` so we can do better.

Modern control theory uses state-space representation to model and control systems. State-space representation models `systems` as a set of `state`, `input`, and `output` variables related by first-order differential equations that describe how the `system's state` changes over time given the current `states` and `inputs`.

5.1 From PID control to model-based control

As mentioned before, controls engineers have a more general framework to describe control theory than just PID control. PID controller designers are focused on fiddling with controller parameters relating to the current, past, and future `error` rather than the underlying system `states`. Integral control is a commonly used tool, and some people use integral action as the majority of the control action. While this approach works in a lot of situations, it is an incomplete view of the world.

Model-based control has a completely different mindset. Controls designers using model-based control care about developing an accurate `model` of the `system`, then driving the `states` they care about to zero (or to a `reference`). Integral control is added with u_{error} estimation if needed to handle `model` uncertainty, but we prefer not to use it because its response is hard to tune and some of its destabilizing dynamics aren't visible during simulation.

5.2 What is a dynamical system?

A dynamical system is a [system](#) whose motion varies according to a set of differential equations. A dynamical system is considered *linear* if the differential equations describing its dynamics consist only of linear operators. Linear operators are things like constant gain multiplications, derivatives, and integrals. You can define reasonably accurate linear [models](#) for pretty much everything you'll see in FRC with just those relations.

But let's say you have a DC brushed motor hooked up to a power supply and you applied a constant voltage to it from rest. The motor approaches a steady-state angular velocity, but the shape of the angular velocity curve over time isn't a line. In fact, it's a decaying exponential curve akin to

$$\omega = \omega_{max} (1 - e^{-t})$$

where ω is the angular velocity and ω_{max} is the maximum angular velocity. If DC brushed motors are said to behave linearly, then why is this?

Linearity refers to a [system](#)'s equations of motion, not its time domain response. The equation defining the motor's change in angular velocity over time looks like

$$\dot{\omega} = -a\omega + bV$$

where $\dot{\omega}$ is the derivative of ω with respect to time, V is the input voltage, and a and b are constants specific to the motor. This equation, unlike the one shown before, is actually linear because it only consists of multiplications and additions relating the [input](#) V and current [state](#) ω .

Also of note is that the relation between the input voltage and the angular velocity of the output shaft is a linear regression. You'll see why if you model a DC brushed motor as a voltage source and generator producing back-EMF (in the equation above, bV corresponds to the voltage source and $-a\omega$ corresponds to the back-EMF). As you increase the input voltage, the back-EMF increases linearly with the motor's angular velocity. If there was a friction term that varied with the angular velocity squared (air resistance is one example), the relation from input to output would be a curve. Friction that scales with just the angular velocity would result in a lower maximum angular velocity, but because that term can be lumped into the back-EMF term, the response is still linear.

5.3 State-space notation

5.3.1 What is state-space?

Recall from last chapter that 2D space has two axes: x and y . We represent locations within this space as a pair of numbers packaged in a vector, and each coordinate is a measure of how far to move along the corresponding axis. State-space is a Cartesian coordinate system with an axis for each [state](#) variable, and we represent locations within it the same way we do for 2D space: with a list of numbers in a vector. Each element in the vector corresponds to a [state](#) of the [system](#).

In addition to the [state](#), [inputs](#) and [outputs](#) are represented as vectors. Since the mapping from the current [states](#) and [inputs](#) to the change in [state](#) is a system of equations, it's natural to write it in matrix form.

5.3.2 Benefits over classical control

State-space notation provides a more convenient and compact way to model and analyze [systems](#) with multiple [inputs](#) and [outputs](#). For a [system](#) with p [inputs](#) and q [outputs](#), we would have to write $q \times p$ transfer functions to represent it. Not only is the resulting algebra unwieldy, but it only works

for linear [systems](#). Including nonzero initial conditions complicates the algebra even more. State-space representation uses the time domain instead of the Laplace domain, so it can model nonlinear [systems](#)¹ and trivially supports nonzero initial conditions.

If modern control theory is so great and classical control theory isn't needed to use it, why learn classical control theory at all? We teach classical control theory because it provides a framework within which to understand results from the mathematical machinery of modern control as well as vocabulary with which to communicate that understanding. For example, faster poles (poles moved to the left in the s-plane) mean faster decay, and oscillation means there is at least one pair of complex conjugate poles. Not only can you describe what happened succinctly, but you know why it happened from a theoretical perspective.

5.3.3 Definition

Below are the continuous and discrete versions of state-space notation.

Definition 5.3.1 — State-space notation.

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu} \quad (5.1)$$

$$\mathbf{y} = \mathbf{Cx} + \mathbf{Du} \quad (5.2)$$

$$\mathbf{x}_{k+1} = \mathbf{Ax}_k + \mathbf{Bu}_k \quad (5.3)$$

$$\mathbf{y}_k = \mathbf{Cx}_k + \mathbf{Du}_k \quad (5.4)$$

- | | | | |
|---|--------------------|---|---------------|
| A | system matrix | x | state vector |
| B | input matrix | u | input vector |
| C | output matrix | y | output vector |
| D | feedthrough matrix | | |

Matrix	Rows × Columns	Matrix	Rows × Columns
A	states × states	x	states × 1
B	states × inputs	u	inputs × 1
C	outputs × states	y	outputs × 1
D	outputs × inputs		

Table 5.1: State-space matrix dimensions

In the continuous case, the change in [state](#) and the [output](#) are linear combinations of the [state](#) vector and the [input](#) vector. The **A** and **B** matrices are used to map the [state](#) vector **x** and the [input](#) vector **u** to a change in the [state](#) vector $\dot{\mathbf{x}}$. The **C** and **D** matrices are used to map the [state](#) vector **x** and the [input](#) vector **u** to an [output](#) vector **y**.

¹This book focuses on analysis and control of linear [systems](#). See chapter 7 for more on nonlinear control.

5.4 Controllability

State controllability implies that it is possible – by admissible inputs – to steer the states from any initial value to any final value within some finite time window.

Theorem 5.4.1 — Controllability. A continuous time-invariant linear state-space model is controllable if and only if

$$\text{rank} ([\mathbf{B} \ \mathbf{AB} \ \mathbf{A}^2\mathbf{B} \ \cdots \ \mathbf{A}^{n-1}\mathbf{B}]) = n \quad (5.5)$$

where rank is the number of linearly independent rows in a matrix and n is the number of state variables.

The matrix in equation (5.5) being rank-deficient means the inputs cannot apply transforms along all axes in the state-space; the transformation the matrix represents is collapsed into a lower dimension.

The condition number of the controllability matrix \mathbb{C} is defined as $\frac{\sigma_{\max}(\mathbb{C})}{\sigma_{\min}(\mathbb{C})}$ where σ_{\max} is the maximum singular value² and σ_{\min} is the minimum singular value. As this number approaches infinity, one or more of the states becomes uncontrollable. This number can also be used to tell us which actuators are better than others for the given system; a lower condition number means that the actuators have more control authority.

5.5 Observability

Observability is a measure for how well internal states of a system can be inferred by knowledge of its external outputs. The observability and controllability of a system are mathematical duals (i.e., as controllability proves that an input is available that brings any initial state to any desired final state, observability proves that knowing enough output values provides enough information to predict the initial state of the system).

Theorem 5.5.1 — Observability. A continuous time-invariant linear state-space model is observable if and only if

$$\text{rank} \left(\begin{bmatrix} \mathbf{C} \\ \mathbf{CA} \\ \vdots \\ \mathbf{CA}^{n-1} \end{bmatrix} \right) = n \quad (5.6)$$

where rank is the number of linearly independent rows in a matrix and n is the number of state variables.

The matrix in equation (5.6) being rank-deficient means the outputs do not contain contributions from every state. That is, not all states are mapped to a linear combination in the output. Therefore, the outputs alone are insufficient to estimate all the states.

The condition number of the observability matrix \mathbb{O} is defined as $\frac{\sigma_{\max}(\mathbb{O})}{\sigma_{\min}(\mathbb{O})}$ where σ_{\max} is the maximum singular value² and σ_{\min} is the minimum singular value. As this number approaches infinity, one or more of the states becomes unobservable. This number can also be used to tell us which

²Singular values are a generalization of eigenvalues for nonsquare matrices.

sensors are better than others for the given [system](#); a lower condition number means the [outputs](#) produced by the sensors are better indicators of the [system state](#).

5.6 Closed-loop controller

With the [control law](#) $u = K(r - x)$, we can derive the closed-loop state-space equations. We'll discuss where this [control law](#) comes from in subsection 5.8.

First is the [state](#) update equation. Substitute the [control law](#) into equation (5.1).

$$\begin{aligned}\dot{x} &= Ax + BK(r - x) \\ \dot{x} &= Ax + BKr - BKx \\ \dot{x} &= (A - BK)x + BKr\end{aligned}\tag{5.7}$$

Now for the [output](#) equation. Substitute the [control law](#) into equation (5.2).

$$\begin{aligned}y &= Cx + D(K(r - x)) \\ y &= Cx + DKr - DKx \\ y &= (C - DK)x + DKr\end{aligned}\tag{5.8}$$

Now, we'll do the same for the discrete [system](#). We'd like to know whether the [system](#) defined by equation (5.3) operating with the [control law](#) $u_k = K(r_k - x_k)$ converges to the [reference](#) r_k .

$$\begin{aligned}x_{k+1} &= Ax_k + Bu_k \\ x_{k+1} &= Ax_k + B(K(r_k - x_k)) \\ x_{k+1} &= Ax_k + BKr_k - BKx_k \\ x_{k+1} &= Ax_k - BKx_k + BKr_k \\ x_{k+1} &= (A - BK)x_k + BKr_k\end{aligned}$$

Theorem 5.6.1 — Closed-loop state-space controller.

$$\dot{x} = (A - BK)x + BKr\tag{5.9}$$

$$y = (C - DK)x + DKr\tag{5.10}$$

$$x_{k+1} = (A - BK)x_k + BKr_k\tag{5.11}$$

$$y_k = (C - DK)x_k + DKr_k\tag{5.12}$$

A	system matrix	K	controller gain matrix
B	input matrix	x	state vector
C	output matrix	r	reference vector
D	feedthrough matrix	y	output vector

Matrix	Rows × Columns	Matrix	Rows × Columns
A	states × states	x	states × 1
B	states × inputs	u	inputs × 1
C	outputs × states	y	outputs × 1
D	outputs × inputs	r	states × 1
K	inputs × states		

Table 5.2: Controller matrix dimensions

Instead of commanding the [system](#) to a [state](#) using the vector **u** directly, we can now specify a vector of desired [states](#) through **r** and the [controller](#) will choose values of **u** for us over time to make the [system](#) converge to the [reference](#). For equation (5.9) to reach steady-state, the eigenvalues of **A** – **BK** must be in the left-half plane. For equation (5.11) to have a bounded output, the eigenvalues of **A** – **BK** must be within the unit circle.

The eigenvalues of **A** – **BK** are the poles of the closed-loop [system](#). Therefore, the rate of convergence and stability of the closed-loop [system](#) can be changed by moving the poles via the eigenvalues of **A** – **BK**. **A** and **B** are inherent to the [system](#), but **K** can be chosen arbitrarily by the controller designer.

5.7 Pole placement

This is the practice of placing the poles of a closed-loop [system](#) directly to produce a desired response. Python Control offers several pole placement algorithms for generating controller or observer gains from a set of poles. In general, pole placement should only be used if you know what you're doing. It's much easier to let LQR place the poles for you, which we'll discuss next.

5.8 Linear-quadratic regulator

Instead of placing the poles of a closed-loop [system](#) manually, the linear-quadratic regulator (LQR) places the poles for us based on acceptable relative [error](#) and [control effort](#) costs. This method of controller design uses a quadratic function for the cost-to-go defined as the sum of the [error](#) and [control effort](#) over time for the linear [system](#) $\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$.

$$J = \int_0^{\infty} (\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

where J represents a trade-off between [state](#) excursion and [control effort](#) with the weighting factors **Q** and **R**. LQR finds a [control law](#) **u** that minimizes the cost function. **Q** and **R** slide the cost along a Pareto boundary between state tracking and [control effort](#) (see figure 5.1). Pareto optimality for this problem means that an improvement in state [tracking](#) cannot be obtained without using more [control effort](#) to do so. Also, a reduction in [control effort](#) cannot be obtained without sacrificing state [tracking](#) performance. Pole placement, on the other hand, will have a cost anywhere on, above, or to the right of the Pareto boundary (no cost can be inside the boundary).

The minimum of LQR's cost function is found by setting the derivative of the cost function to zero and solving for the [control law](#) **u**. However, matrix calculus is used instead of normal calculus to take the derivative.

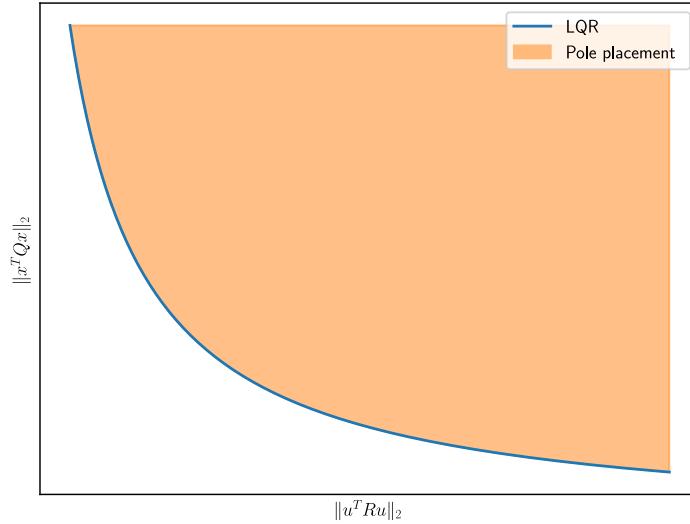


Figure 5.1: Pareto boundary for LQR

The feedback control law that minimizes J is shown in theorem 5.8.1.

Theorem 5.8.1 — Linear-quadratic regulator.

$$\begin{aligned} \min_{\mathbf{u}} \int_0^{\infty} (\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt \\ \text{subject to } \dot{\mathbf{x}} = \mathbf{A} \mathbf{x} + \mathbf{B} \mathbf{u} \end{aligned} \quad (5.13)$$

If the system is controllable, the optimal control policy \mathbf{u}^* that drives all the states to zero is $-\mathbf{K} \mathbf{x}$. To converge to nonzero states, a reference vector \mathbf{r} can be added to the state \mathbf{x} .

$$\mathbf{u} = \mathbf{K}(\mathbf{r} - \mathbf{x}) \quad (5.14)$$

This means that optimal control can be achieved with simply a set of proportional gains on all the states. To use the control law, we need knowledge of the full state of the system. That means we either have to measure all our states directly or estimate those we do not measure.

See appendix E for how \mathbf{K} is calculated in Python. If the result is finite, the controller is guaranteed to be stable and robust with a phase margin of 60 degrees [25].



LQR design's \mathbf{Q} and \mathbf{R} matrices don't need discretization, but the \mathbf{K} calculated for continuous time and discrete time systems will be different. The discrete time gains approach the continuous time gains as the sample period tends to zero.

5.8.1 Bryson's rule

The next obvious question is what values to choose for \mathbf{Q} and \mathbf{R} . While this can be more of an art than a science, Bryson's rule provides a good starting point. With Bryson's rule, the diagonals of the

Q and **R** matrices are chosen based on the maximum acceptable value for each state and actuator. The nondiagonal elements are zero. The balance between **Q** and **R** can be slid along the Pareto boundary using a weighting factor ρ .

$$J = \int_0^{\infty} \left(\rho \left[\left(\frac{x_1}{x_{1,max}} \right)^2 + \dots + \left(\frac{x_m}{x_{m,max}} \right)^2 \right] + \left[\left(\frac{u_1}{u_{1,max}} \right)^2 + \dots + \left(\frac{u_n}{u_{n,max}} \right)^2 \right] \right) dt$$

$$\mathbf{Q} = \begin{bmatrix} \frac{\rho}{x_{1,max}^2} & 0 & \dots & 0 \\ 0 & \frac{\rho}{x_{2,max}^2} & & \vdots \\ \vdots & \ddots & 0 & \\ 0 & \dots & 0 & \frac{\rho}{x_{m,max}^2} \end{bmatrix} \quad \mathbf{R} = \begin{bmatrix} \frac{1}{u_{1,max}^2} & 0 & \dots & 0 \\ 0 & \frac{1}{u_{2,max}^2} & & \vdots \\ \vdots & \ddots & 0 & \\ 0 & \dots & 0 & \frac{1}{u_{n,max}^2} \end{bmatrix}$$

Small values of ρ penalize control effort while large values of ρ penalize state excursions. Large values would be chosen in applications like fighter jets where performance is necessary. Spacecrafts would use small values to conserve their limited fuel supply.

5.8.2 Pole placement vs LQR

This example uses the following second-order model for a CIM motor (a DC brushed motor).

$$\mathbf{A} = \begin{bmatrix} -\frac{b}{J} & \frac{K_t}{J} \\ -\frac{K_e}{L} & -\frac{R}{L} \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 \\ \frac{1}{L} \end{bmatrix} \quad \mathbf{C} = [1 \ 0] \quad \mathbf{D} = [0]$$

Figure 5.2 shows the response using discrete poles³ placed at (0.1, 0) and (0.9, 0) and LQR with the following cost matrices.

$$\mathbf{Q} = \begin{bmatrix} \frac{1}{20^2} & 0 \\ 0 & \frac{1}{40^2} \end{bmatrix} \quad \mathbf{R} = \begin{bmatrix} \frac{1}{12^2} \end{bmatrix}$$

LQR selected poles at (0.593, 0) and (0.955, 0). Notice with pole placement that as the current pole moves left, the control effort becomes more aggressive.

³See section 6.2 for pole mappings of discrete systems (inside unit circle is stable). The pole mappings mentioned so far (LHP is stable) only apply to continuous systems.

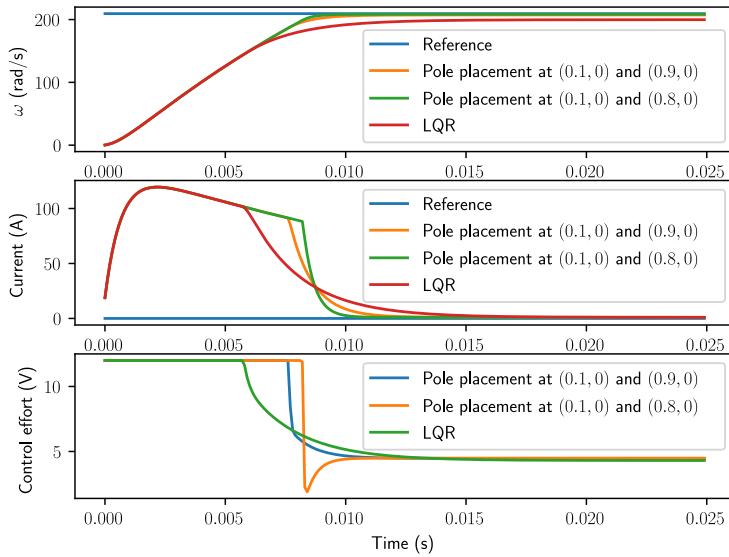


Figure 5.2: Second-order CIM motor response with pole placement and LQR

5.9 Model augmentation

This section will teach various tricks for manipulating state-space [models](#) with the goal of demystifying the matrix algebra at play. We will use the augmentation techniques discussed here in the section on integral control.

Matrix augmentation is the process of appending rows or columns to a matrix. In state-space, there are several common types of augmentation used: [plant](#) augmentation, controller augmentation, and [observer](#) augmentation.

5.9.1 Plant augmentation

Plant augmentation is the process of adding a state to a model's state vector and adding a corresponding row to the **A** and **B** matrices.

5.9.2 Controller augmentation

Controller augmentation is the process of adding a column to a controller's **K** matrix. This is often done in combination with [plant](#) augmentation to add controller dynamics relating to a newly added state.

5.9.3 Observer augmentation

Observer augmentation is closely related to [plant](#) augmentation. In addition to adding entries to the [observer](#) matrix **L**, the [observer](#) is using this augmented [plant](#) for estimation purposes. This is better explained with an example.

By augmenting the [plant](#) with a bias term with no dynamics (represented by zeroes in its rows in **A** and **B**), the [observer](#) will attempt to estimate a value for this bias term that makes the [model](#) best reflect the measurements taken of the real [system](#). Note that we're not collecting any data on this bias term directly; it's what's known as a hidden [state](#). Rather than our [inputs](#) and other [states](#) affecting it directly, the [observer](#) determines a value for it based on what is most likely given the [model](#) and

current measurements. We just tell the `plant` what kind of dynamics the term has and the `observer` will estimate it for us.

5.9.4 Output augmentation

Output augmentation is the process of adding rows to the \mathbf{C} matrix. This is done to help the controls designer visualize the behavior of internal states or other aspects of the `system` in MATLAB or Python Control. \mathbf{C} matrix augmentation doesn't affect `state` feedback, so the designer has a lot of freedom here. Noting that the `output` is defined as $\mathbf{y} = \mathbf{Cx} + \mathbf{Du}$, The following row augmentations of \mathbf{C} may prove useful. Of course, \mathbf{D} needs to be augmented with zeroes as well in these cases to maintain the correct matrix dimensionality.

Since $\mathbf{u} = -\mathbf{Kx}$, augmenting \mathbf{C} with $-\mathbf{K}$ makes the `observer` estimate the control input \mathbf{u} applied.

$$\begin{aligned} \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\ \begin{bmatrix} \mathbf{y} \\ \mathbf{u} \end{bmatrix} &= \begin{bmatrix} \mathbf{C} \\ -\mathbf{K} \end{bmatrix} \mathbf{x} + \begin{bmatrix} \mathbf{D} \\ \mathbf{0} \end{bmatrix} \mathbf{u} \end{aligned}$$

This works because \mathbf{K} has the same number of columns as `states`.

Various `states` can also be produced in the `output` with \mathbf{I} matrix augmentation.

5.9.5 Examples

Snippet 5.1 shows how one packs together the following augmented matrix in Python using concatenation.

$$\begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}$$

```
#!/usr/bin/env python3

import numpy as np


def main():
    A = np.array([[1, 2], [3, 4]])
    B = np.array([[5], [6]])
    C = np.array([[7, 8]])
    D = np.array([[9]])

    tmp = np.concatenate(
        (np.concatenate((A, B), axis=1), np.concatenate((C, D), axis=1)),
        axis=0
    )

    print("A =")
    print(A)
    print("B =")
    print(B)
    print("C =")
    print(C)
    print("D =")
    print(D)
    print("[A, B; C, D] =")
    print(tmp)

if __name__ == "__main__":
    main()
```

Snippet 5.1. Matrix augmentation example: concatenation

Snippet 5.2 shows how one packs together the same augmented matrix in Python using array slices.

```
#!/usr/bin/env python3

import numpy as np

def main():
    A = np.array([[1, 2], [3, 4]])
    B = np.array([[5], [6]])
    C = np.array([[7, 8]])
    D = np.array([[9]])

    tmp = np.zeros((3, 3))
    tmp[:2, :2] = A # tmp[0:2, 0:2] = A
    tmp[:2, 2:] = B # tmp[0:2, 2:3] = B
    tmp[2:, :2] = C # tmp[2:3, 0:2] = C
    tmp[2:, 2:] = D # tmp[2:3, 2:3] = D

    print("A =")
    print(A)
    print("B =")
    print(B)
    print("C =")
    print(C)
    print("D =")
    print(D)
    print("[A, B; C, D] =")
    print(tmp)

if __name__ == "__main__":
    main()
```

Snippet 5.2. Matrix augmentation example: array slices

Section 5.11 demonstrates `model` augmentation for different types of integral control.

5.10 Feedforward

So far, we've used feedback control for `reference tracking` (making a `system`'s output follow a desired `reference signal`). While this is effective, it's a reactionary measure; the `system` won't start applying `control effort` until the `system` is already behind. If we could tell the `controller` about the desired movement and required input beforehand, the `system` could react quicker and the feedback `controller` could do less work. A `controller` that feeds information forward into the `plant` like this is called a `feedforward controller`.

A `feedforward controller` injects information about the `system`'s dynamics (like a `model` does) or the desired movement. The feedforward handles parts of the control actions we already know must be applied to make a `system` track a `reference`, then feedback compensates for what we do not or cannot know about the `system`'s behavior at runtime.

There are two types of feedforwards: model-based feedforward and feedforward for unmodeled

dynamics. The first solves a mathematical model of the system for the inputs required to meet desired velocities and accelerations. The second compensates for unmodeled forces or behaviors directly so the feedback controller doesn't have to. Both types can facilitate simpler feedback controllers; we'll cover examples of each.

5.10.1 Plant inversion

Plant inversion is a method of model-based feedforward for state feedback. It solves the plant for the input that will make the plant track a desired output. This is called inversion because in a block diagram, the inverted plant feedforward and plant cancel out to produce a unity system from input to output.

While it can be an effective tool, the following should be kept in mind.

1. Don't invert an unstable plant. If the expected plant doesn't match the real plant exactly, the plant inversion will still result in an unstable system. Stabilize the plant first with feedback, then inject an inversion.
2. Don't invert a nonminimum phase system. The advice for pole-zero cancellation in subsection 3.2.3 applies here.

Necessary theorems

The following theorem will be needed to derive the linear plant inversion equation.

Theorem 5.10.1 $\frac{\partial \mathbf{x}^T \mathbf{A} \mathbf{x}}{\partial \mathbf{x}} = 2 \mathbf{A} \mathbf{x}$ where \mathbf{A} is symmetric.

Setup

Let's start with the equation for the reference dynamics

$$\mathbf{r}_{k+1} = \mathbf{A} \mathbf{r}_k + \mathbf{B} \mathbf{u}_k$$

where \mathbf{u}_k is the feedforward input. Note that this feedforward equation does not and should not take into account any feedback terms. We want to find the optimal \mathbf{u}_k such that we minimize the tracking error between \mathbf{r}_{k+1} and \mathbf{r}_k .

$$\mathbf{r}_{k+1} - \mathbf{A} \mathbf{r}_k = \mathbf{B} \mathbf{u}_k$$

To solve for \mathbf{u}_k , we need to take the inverse of the nonsquare matrix \mathbf{B} . This isn't possible, but we can find the pseudoinverse given some constraints on the state tracking error and control effort. To find the optimal solution for these sorts of trade-offs, one can define a cost function and attempt to minimize it. To do this, we'll first solve the expression for $\mathbf{0}$.

$$\mathbf{0} = \mathbf{B} \mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A} \mathbf{r}_k)$$

This expression will be the state tracking cost we use in the following cost function as an H_2 norm.

$$\mathbf{J} = (\mathbf{B} \mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A} \mathbf{r}_k))^T (\mathbf{B} \mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A} \mathbf{r}_k))$$

Minimization

Given theorem 5.10.1, find the minimum of \mathbf{J} by taking the partial derivative with respect to \mathbf{u}_k and setting the result to $\mathbf{0}$.

$$\frac{\partial \mathbf{J}}{\partial \mathbf{u}_k} = 2 \mathbf{B}^T (\mathbf{B} \mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A} \mathbf{r}_k))$$

$$\begin{aligned}
\mathbf{0} &= 2\mathbf{B}^T(\mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)) \\
\mathbf{0} &= 2\mathbf{B}^T\mathbf{B}\mathbf{u}_k - 2\mathbf{B}^T(\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \\
2\mathbf{B}^T\mathbf{B}\mathbf{u}_k &= 2\mathbf{B}^T(\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \\
\mathbf{B}^T\mathbf{B}\mathbf{u}_k &= \mathbf{B}^T(\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \\
\mathbf{u}_k &= (\mathbf{B}^T\mathbf{B})^{-1}\mathbf{B}^T(\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)
\end{aligned}$$

$(\mathbf{B}^T\mathbf{B})^{-1}\mathbf{B}^T$ is the Moore-Penrose pseudoinverse of \mathbf{B} denoted by \mathbf{B}^\dagger .

Theorem 5.10.2 — Linear plant inversion. Given the discrete model $\mathbf{x}_{k+1} = \mathbf{A}\mathbf{x}_k + \mathbf{B}\mathbf{u}_k$, the plant inversion feedforward is

$$\mathbf{u}_k = \mathbf{B}^\dagger(\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \quad (5.15)$$

where \mathbf{B}^\dagger is the Moore-Penrose pseudoinverse of \mathbf{B} , \mathbf{r}_{k+1} is the reference at the next timestep, and \mathbf{r}_k is the reference at the current timestep.

Discussion

Linear [plant](#) inversion in theorem 5.10.2 compensates for [reference](#) dynamics that don't follow how the [model](#) inherently behaves. If they do follow the [model](#), the feedforward has nothing to do as the [model](#) already behaves in the desired manner. When this occurs, $\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k$ will return a zero vector.

For example, a constant [reference](#) requires a feedforward that opposes [system](#) dynamics that would change the [state](#) over time. If the [system](#) has no dynamics, then $\mathbf{A} = \mathbf{I}$ and thus

$$\begin{aligned}
\mathbf{u}_k &= \mathbf{B}_\dagger(\mathbf{r}_{k+1} - \mathbf{I}\mathbf{r}_k) \\
\mathbf{u}_k &= \mathbf{B}_\dagger(\mathbf{r}_{k+1} - \mathbf{r}_k)
\end{aligned}$$

For a constant [reference](#), $\mathbf{r}_{k+1} = \mathbf{r}_k$.

$$\begin{aligned}
\mathbf{u}_k &= \mathbf{B}_\dagger(\mathbf{r}_k - \mathbf{r}_k) \\
\mathbf{u}_k &= \mathbf{B}_\dagger(\mathbf{0}) \\
\mathbf{u}_k &= \mathbf{0}
\end{aligned}$$

so no feedforward is required to hold a [system](#) with no dynamics at a constant [reference](#), as expected.

Figure 5.3 shows [plant](#) inversion applied to a second-order CIM motor model. [Plant](#) inversion accounts for the motor back-EMF and eliminates steady-state error.

5.10.2 Unmodeled dynamics

In addition to [plant](#) inversion, one can include feedforwards for unmodeled dynamics. Consider an elevator model which doesn't include gravity. A constant voltage offset can be used to compensate for this. The feedforward takes the form of a voltage constant because voltage is proportional to force applied, and the force is acting in only one direction at all times.

$$u_k = V_{app} \quad (5.16)$$

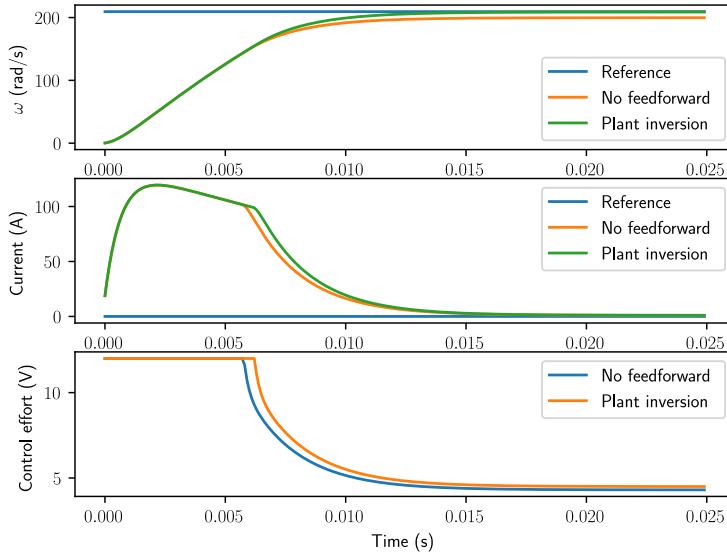


Figure 5.3: Second-order CIM motor response with plant inversion

where V_{app} is a constant. Another feedforward holds a single-jointed arm steady in the presence of gravity. It has the following form.

$$u_k = V_{app} \cos \theta \quad (5.17)$$

where V_{app} is the voltage required to keep the single-jointed arm level with the ground, and θ is the angle of the arm relative to the ground. Therefore, the force applied is greatest when the arm is parallel with the ground and zero when the arm is perpendicular to the ground (at that point, the joint supports all the weight).

Note that the elevator model could be augmented easily enough to include gravity and still be linear, but this wouldn't work for the single-jointed arm since a trigonometric function is required to model the gravitational force in the arm's rotating reference frame⁴.

5.11 Integral control

A common way of implementing integral control is to add an additional [state](#) that is the integral of the [error](#) of the variable intended to have zero [steady-state error](#).

There are two drawbacks to this method. First, there is integral windup on a unit [step input](#). That is, the integrator accumulates even if the [system](#) is [tracking](#) the [model](#) correctly. The second is demonstrated by an example from Jared Russell of FRC team 254. Say there is a position/velocity trajectory for some [plant](#) to follow. Without integral control, one can calculate a desired Kx to use as the [control input](#). As a result of using both desired position and velocity, [reference tracking](#) is good. With integral control added, the [reference](#) is always the desired position, but there is no way to tell the controller the desired velocity.

Consider carefully whether integral control is necessary. One can get relatively close without integral control, and integral adds all the issues listed above. Below, it is assumed that the controls designer

⁴While the applied torque of the motor is constant throughout the arm's range of motion, the torque caused by gravity in the opposite direction varies according to the arm's angle.

has determined that integral control will be worth the inconvenience.

We'll present two methods:

1. Augment the [plant](#) as described earlier. For an arm, one would add an “integral of position” state.
2. Estimate the “error” in the [control input](#) (the difference between what was applied versus what was observed to happen) via the [observer](#) and compensate for it. We'll call this “u error estimation”.

5.11.1 Plant augmentation

We want to augment the [system](#) with an integral term that integrates the [error](#) $e = r - y = r - Cx$.

$$\begin{aligned} \dot{x}_I &= \int e \, dt \\ \dot{x}_I &= e = r - Cx \end{aligned}$$

The [plant](#) is augmented as

$$\begin{aligned} \dot{\begin{bmatrix} x \\ x_I \end{bmatrix}} &= \begin{bmatrix} A & 0 \\ -C & 0 \end{bmatrix} \begin{bmatrix} x \\ x_I \end{bmatrix} + \begin{bmatrix} B \\ 0 \end{bmatrix} u + \begin{bmatrix} 0 \\ I \end{bmatrix} r \\ \dot{\begin{bmatrix} x \\ x_I \end{bmatrix}} &= \begin{bmatrix} A & 0 \\ -C & 0 \end{bmatrix} \begin{bmatrix} x \\ x_I \end{bmatrix} + \begin{bmatrix} B & 0 \\ 0 & I \end{bmatrix} \begin{bmatrix} u \\ r \end{bmatrix} \end{aligned}$$

The controller is augmented as

$$\begin{aligned} u &= K(r - x) - K_I x_I \\ u &= [K \quad K_I] \left(\begin{bmatrix} r \\ 0 \end{bmatrix} - \begin{bmatrix} x \\ x_I \end{bmatrix} \right) \end{aligned}$$

5.11.2 Input error estimation

Given the desired [input](#) produced by a [controller](#), unmodeled [disturbances](#) may cause the observed behavior of a [system](#) to deviate from its [model](#). Input error estimation estimates the difference between the desired [input](#) and a hypothetical [input](#) that makes the [model](#) match the observed behavior. This value can be added to the [control input](#) to make the [controller](#) compensate for unmodeled [disturbances](#) and make the [model](#) better predict the [system](#)'s future behavior.

First, we'll consider the one-dimensional case. Let u_{error} be the difference between the [input](#) actually applied to a [system](#) and the desired [input](#). The u_{error} term is then added to the [system](#) as follows.

$$\dot{x} = Ax + B(u + u_{error})$$

$u + u_{error}$ is the hypothetical [input](#) actually applied to the [system](#).

$$\dot{x} = Ax + Bu + Bu_{error}$$

The following equation generalizes this to a multiple-input [system](#).

$$\dot{x} = Ax + Bu + B_{error}u_{error}$$

where \mathbf{B}_{error} is a column vector that maps u_{error} to changes in the rest of the state the same way \mathbf{B} does for \mathbf{u} . \mathbf{B}_{error} is only a column of \mathbf{B} if u_{error} corresponds to an existing input within \mathbf{u} .

Given the above equation, we'll augment the plant as

$$\begin{aligned}\begin{bmatrix}\dot{\mathbf{x}} \\ u_{error}\end{bmatrix} &= \begin{bmatrix}\mathbf{A} & \mathbf{B}_{error} \\ \mathbf{0} & \mathbf{0}\end{bmatrix} \begin{bmatrix}\mathbf{x} \\ u_{error}\end{bmatrix} + \begin{bmatrix}\mathbf{B} \\ \mathbf{0}\end{bmatrix} \mathbf{u} \\ \mathbf{y} &= [\mathbf{C} \quad 0] \begin{bmatrix}\mathbf{x} \\ u_{error}\end{bmatrix} + \mathbf{D}\mathbf{u}\end{aligned}$$

Notice how the state is augmented with u_{error} . With this model, the observer will estimate both the state and the u_{error} term. The controller is augmented similarly. \mathbf{r} is augmented with a zero for the goal u_{error} term.

$$\begin{aligned}\mathbf{u} &= \mathbf{K}(\mathbf{r} - \mathbf{x}) - \mathbf{k}_{error}u_{error} \\ \mathbf{u} &= [\mathbf{K} \quad \mathbf{k}_{error}] \left(\begin{bmatrix}\mathbf{r} \\ 0\end{bmatrix} - \begin{bmatrix}\mathbf{x} \\ u_{error}\end{bmatrix} \right)\end{aligned}$$

where \mathbf{k}_{error} is a column vector with a 1 in a given row if u_{error} should be applied to that input or a 0 otherwise.

This process can be repeated for an arbitrary error which can be corrected via some linear combination of the inputs.

6. Digital control

The complex plane discussed so far deals with continuous [systems](#). In decades past, [plants](#) and controllers were implemented using analog electronics, which are continuous in nature. Nowadays, microprocessors can be used to achieve cheaper, less complex controller designs. [Discretization](#) converts the continuous [model](#) we've worked with so far from a differential equation like

$$\dot{x} = x - 3 \quad (6.1)$$

to a difference equation like

$$\begin{aligned} \frac{x_{k+1} - x_k}{\Delta T} &= x_k - 3 \\ x_{k+1} - x_k &= (x_k - 3)\Delta T \\ x_{k+1} &= x_k + (x_k - 3)\Delta T \end{aligned} \quad (6.2)$$

where x_k refers to the value of x at the k^{th} timestep. The difference equation is run with some update period denoted by T , by ΔT , or sometimes sloppily by dt^1 .

While higher order terms of a differential equation are derivatives of the [state](#) variable (e.g., \ddot{x} in relation to equation (6.1)), higher order terms of a difference equation are delayed copies of the [state](#) variable (e.g., x_{k-1} with respect to x_k in equation (6.2)).

6.1 Phase loss

However, [discretization](#) has drawbacks. Since a microcontroller performs discrete steps, there is a sample delay that introduces phase loss in the controller. Phase loss is the reduction of [phase margin](#)² that occurs in digital implementations of feedback controllers from sampling the continuous [system](#)

¹The discretization of equation (6.1) to equation (6.2) uses the forward Euler discretization method.

²See section C.1 for an explanation of phase margin.

at discrete time intervals. As the sample rate of the controller decreases, the [phase margin](#) decreases according to $-\frac{T}{2}\omega$ where T is the sample period and ω is the frequency of the [system](#) dynamics. Instability occurs if the [phase margin](#) of the [system](#) reaches zero. Large amounts of phase loss can make a stable controller in the continuous domain become unstable in the discrete domain. Here are a few ways to combat this.

- Run the controller with a high sample rate.
- Designing the controller in the analog domain with enough [phase margin](#) to compensate for any phase loss that occurs as part of [discretization](#).
- Convert the [plant](#) to the digital domain and design the controller completely in the digital domain.

6.2 s-plane to z-plane

Transfer functions are converted to impulse responses using the Z-transform. The s-plane's LHP maps to the inside of a unit circle in the z-plane. Table 6.1 contains a few common points and figure 6.1 shows the mapping visually.

s-plane	z-plane
(0, 0)	(1, 0)
imaginary axis	edge of unit circle
($-\infty$, 0)	(0, 0)

Table 6.1: Mapping from s-plane to z-plane

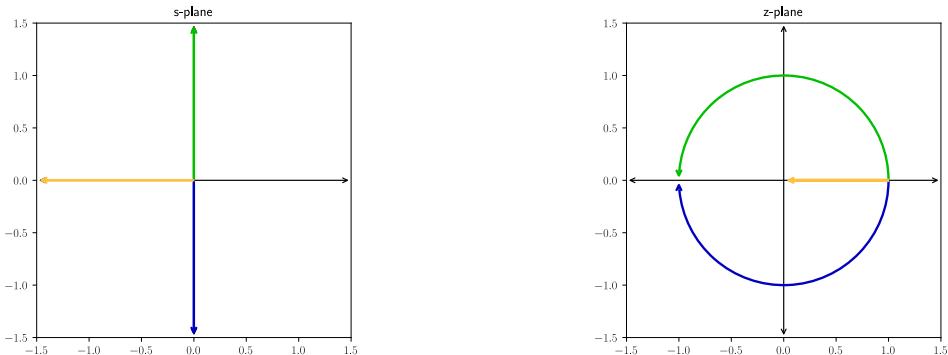


Figure 6.1: Mapping of axes from s-plane (left) to z-plane (right)

6.2.1 z-plane stability

Eigenvalues of a [system](#) that are within the unit circle are stable, but why is that? Let's consider a scalar equation $x_{k+1} = ax_k$. $a < 1$ makes x_{k+1} converge to zero. The same applies to a complex number like $z = x + yi$ for $x_{k+1} = zx_k$. If the magnitude of the complex number z is less than one, x_{k+1} will converge to zero. Values with a magnitude of 1 oscillate forever because x_{k+1} never decays.

6.2.2 z-plane behavior

As ω increases in $s = j\omega$, a pole in the z-plane moves around the perimeter of the unit circle. Once it hits $\frac{\omega_s}{2}$ (half the sampling frequency) at $(-1, 0)$, the pole wraps around. This is due to poles faster than the sample frequency folding down to below the sample frequency (that is, higher frequency signals *alias* to lower frequency ones).

You may notice that poles can be placed at $(0, 0)$ in the z-plane. This is known as a deadbeat controller. An N^{th} -order deadbeat controller decays to the [reference](#) in N timesteps. While this sounds great, there are other considerations like [control effort](#), [robustness](#), and [noise immunity](#). These will be discussed in more detail with LQR and LQE.

If poles from $(1, 0)$ to $(0, 0)$ on the x-axis approach infinity, then what do poles from $(-1, 0)$ to $(0, 0)$ represent? Them being faster than infinity doesn't make sense. Poles in this location exhibit oscillatory behavior similar to complex conjugate pairs. See figures 6.2 and 6.3. The jaggedness of these signals is due to the frequency of the [system](#) dynamics being above the Nyquist frequency (twice the sample frequency). The [discretized](#) signal doesn't have enough samples to reconstruct the continuous [system](#)'s dynamics.

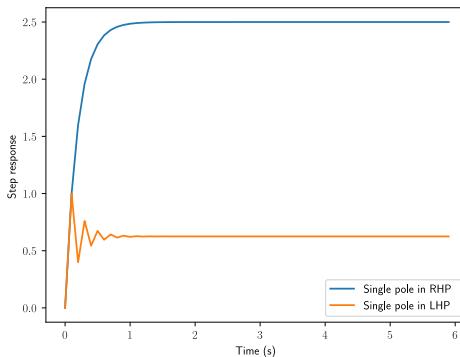


Figure 6.2: Single poles in various locations in z-plane

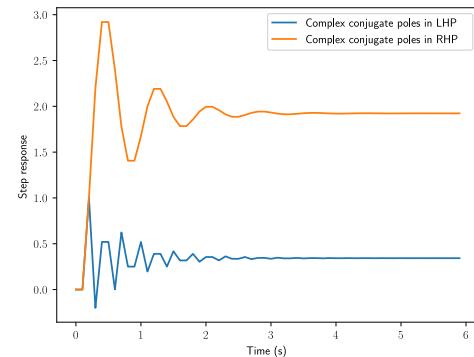


Figure 6.3: Complex conjugate poles in various locations in z-plane

6.2.3 Nyquist frequency

To completely reconstruct a signal, the Nyquist-Shannon sampling theorem states that it must be sampled at a frequency at least twice the maximum frequency it contains. The highest frequency a given sample rate can capture is called the Nyquist frequency, which is half the sample frequency. This is why recorded audio is sampled at 44.1 kHz . The maximum frequency a typical human can hear is about 20 kHz , so the Nyquist frequency is 40 kHz . (44.1 kHz in particular was chosen for unrelated historical reasons.)

Frequencies above the Nyquist frequency are folded down across it. The higher frequency and the folded down lower frequency are said to alias each other³. Figure 6.4 demonstrates aliasing.

The effect of these high-frequency aliases can be reduced with a low-pass filter (called an anti-aliasing filter in this application).

³The aliases of a frequency f can be expressed as $f_{\text{alias}}(N) \stackrel{\text{def}}{=} |f - Nf_s|$. For example, if a 200 Hz sine wave is sampled at 150 Hz , the [observer](#) will see a 50 Hz signal instead of a 200 Hz one.

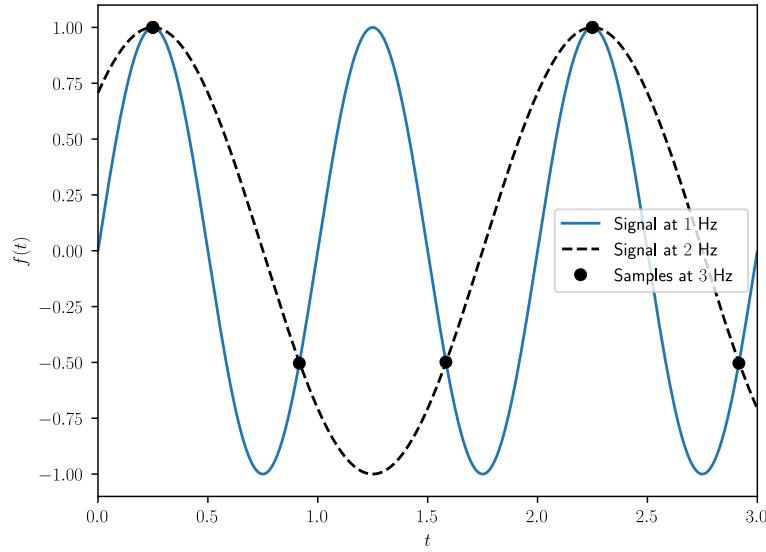


Figure 6.4: The samples of two sine waves can be identical when at least one of them is at a frequency above half the sample rate. In this case, the 2 Hz sine wave is above the Nyquist frequency 1.5 Hz .

6.3 Discretization methods

Discretization is done using a zero-order hold. That is, the [system state](#) is only updated at discrete intervals and it's held constant between samples (see figure 6.5). The exact method of applying this uses the matrix exponential, but this can be computationally expensive. Instead, approximations such as the following are used.

1. Forward Euler method. This is defined as $y_{n+1} = y_n + f(t_n, y_n)\Delta t$.
2. Backward Euler method. This is defined as $y_{n+1} = y_n + f(t_{n+1}, y_{n+1})\Delta t$.
3. Bilinear transform. The first-order bilinear approximation is $s = \frac{2}{T} \frac{1-z^{-1}}{1+z^{-1}}$.

where the function $f(t_n, y_n)$ is the slope of y at n and T is the sample period for the discrete [system](#). Each of these methods is essentially finding the area underneath a curve. The forward and backward Euler methods use rectangles to approximate that area while the bilinear transform uses trapezoids (see figures 6.6 and 6.7). Since these are approximations, there is distortion between the real discrete [system](#)'s poles and the approximate poles. This is in addition to the phase loss introduced by discretizing at a given sample rate in the first place. For fast-changing [systems](#), this distortion can quickly lead to instability.

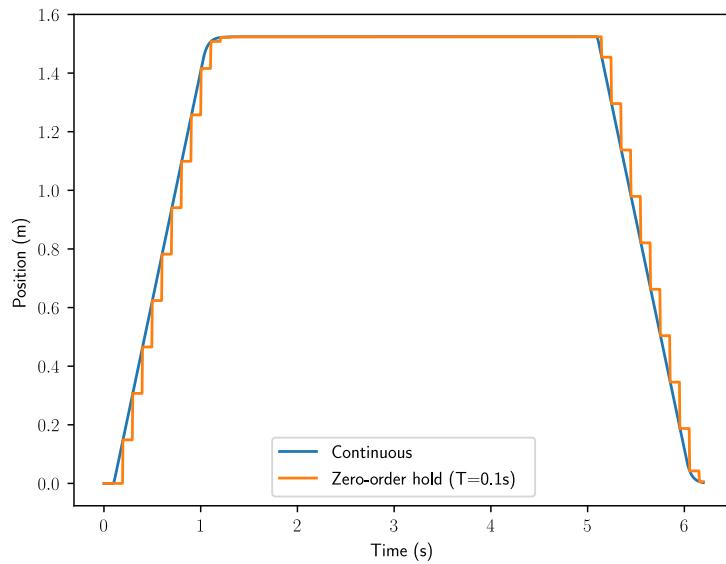


Figure 6.5: Zero-order hold of a system response

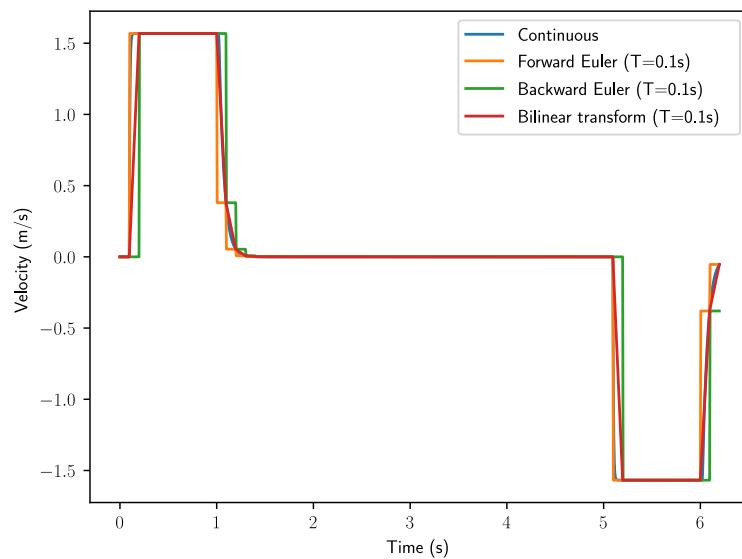


Figure 6.6: Discretization methods applied to velocity data

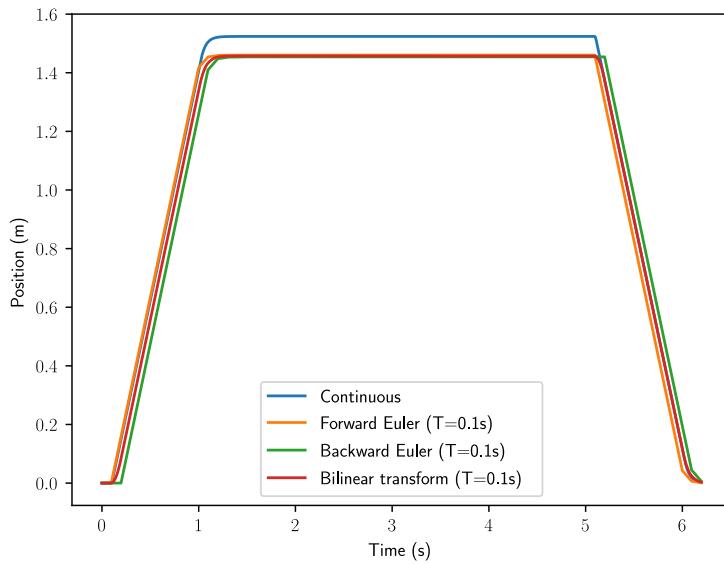


Figure 6.7: Position plot of discretization methods applied to velocity data

6.4 Effects of discretization on controller performance

Running a feedback controller at a faster update rate doesn't always mean better control. In fact, you may be using more computational resources than you need. However, here are some reasons for running at a faster update rate.

Firstly, if you have a discrete [model](#) of the [system](#), that [model](#) can more accurately approximate the underlying continuous [system](#). Not all controllers use a [model](#) though.

Secondly, the controller can better handle fast [system](#) dynamics. If the [system](#) can move from its initial state to the desired one in under 250ms, you obviously want to run the controller with a period less than 250ms. When you reduce the sample period, you're making the discrete controller more accurately reflect what the equivalent continuous controller would do (controllers built from analog circuit components like op-amps are continuous).

Running at a lower sample rate only causes problems if you don't take into account the response time of your [system](#). Some [systems](#) like heaters have [outputs](#) that change on the order of minutes. Running a control loop at 1kHz doesn't make sense for this because the [plant input](#) the controller computes won't change much, if at all, in 1ms.

Figures 6.8, 6.9, and 6.10 show simulations of the same controller for different sampling methods and sample rates, which have varying levels of fidelity to the real [system](#).

Forward Euler is numerically unstable for low sample rates. The bilinear transform is a significant improvement due to it being a second-order approximation, but zero-order hold performs best due to the matrix exponential including much higher orders (we'll cover the matrix exponential in the next section).

Table 6.2 compares the Taylor series expansions of the discretization methods presented so far (these are found using polynomial division). The bilinear transform does best with accuracy trailing off after the third-order term. Forward Euler has no second-order or higher terms, so it undershoots. Backward Euler has twice the second-order term and overshoots the remaining higher order terms as

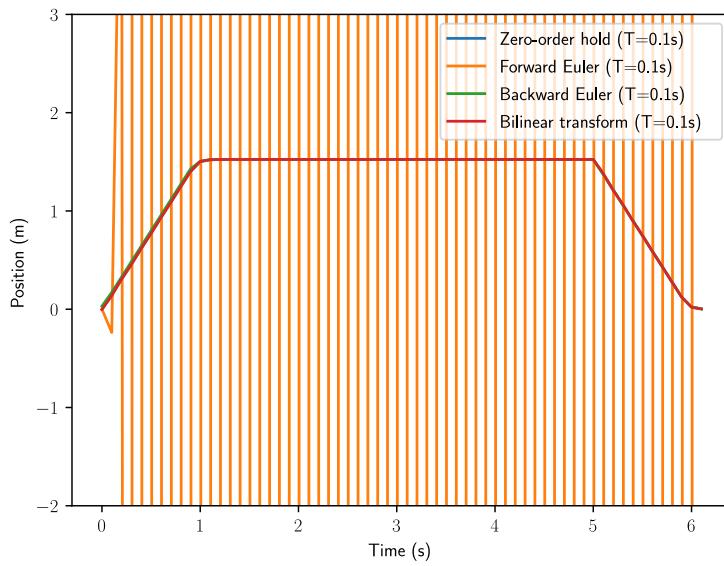


Figure 6.8: Sampling methods for system simulation with $T = 0.1s$

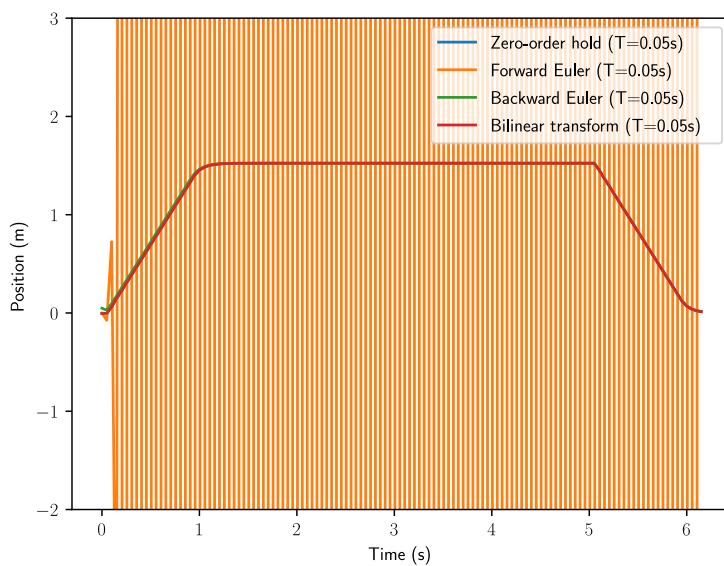


Figure 6.9: Sampling methods for system simulation with $T = 0.05s$

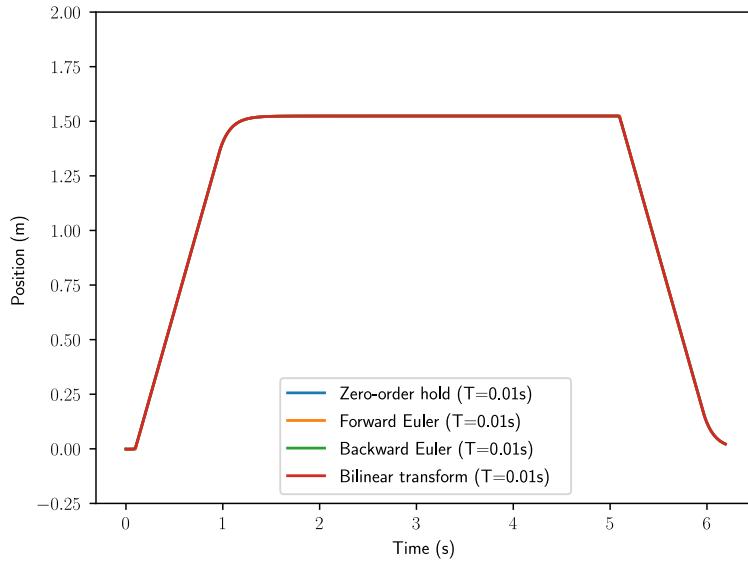


Figure 6.10: Sampling methods for system simulation with $T = 0.01s$

well.

Method	Conversion	Taylor series expansion
Zero-order hold	$z = e^{Ts}$	$z = 1 + Ts + \frac{1}{2}T^2s^2 + \frac{1}{6}T^3s^3 + \dots$
Bilinear	$z = \frac{1 + \frac{1}{2}Ts}{1 - \frac{1}{2}Ts}$	$z = 1 + Ts + \frac{1}{2}T^2s^2 + \frac{1}{4}T^3s^3 + \dots$
Forward Euler	$z = 1 + Ts$	$z = 1 + Ts$
Reverse Euler	$z = \frac{1}{1-Ts}$	$z = 1 + Ts + T^2s^2 + T^3s^3 + \dots$

Table 6.2: Taylor series expansions of discretization methods (scalar case). The zero-order hold discretization method is exact.

6.5 Matrix exponential

The matrix exponential (and [system discretization](#) in general) is typically solved with a computer. Python Control's `StateSpace.sample()` with the “zoh” method (the default) does this.

Definition 6.5.1 — Matrix exponential. Let \mathbf{X} be an $n \times n$ matrix. The exponential of \mathbf{X} denoted by $e^{\mathbf{X}}$ is the $n \times n$ matrix given by the following power series.

$$e^{\mathbf{X}} = \sum_{k=0}^{\infty} \frac{1}{k!} \mathbf{X}^k \quad (6.3)$$

where \mathbf{X}^0 is defined to be the identity matrix \mathbf{I} with the same dimensions as \mathbf{X} .

To understand why the matrix exponential is used in the [discretization](#) process, consider the set of

differential equations $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x}$ we use to describe [systems](#) ([systems](#) also have a $\mathbf{B}\mathbf{u}$ term, but we'll ignore it for clarity). The solution to this type of differential equation uses an exponential. Since we are using matrices and vectors here, we use the matrix exponential.

$$\mathbf{x}(t) = e^{\mathbf{A}t}\mathbf{x}_0$$

where \mathbf{x}_0 contains the initial conditions. If the initial [state](#) is the current system [state](#), then we can describe the [system's state](#) over time as

$$\mathbf{x}_{k+1} = e^{\mathbf{A}T}\mathbf{x}_k$$

where T is the time between samples \mathbf{x}_k and \mathbf{x}_{k+1} .

6.6 Taylor series

 Watch the “Taylor series” video from 3Blue1Brown’s *Essence of calculus* series (22 minutes) [13] for an explanation of how the Taylor series expansion works.

The definition for the matrix exponential and the approximations below all use the *Taylor series expansion*. The Taylor series is a method of approximating a function like e^t via the summation of weighted polynomial terms like t^k . e^t has the following Taylor series around $t = 0$.

$$e^t = \sum_{n=0}^{\infty} \frac{t^n}{n!}$$

where a finite upper bound on the number of terms produces an approximation of e^t . As n increases, the polynomial terms increase in power and the weights by which they are multiplied decrease. For e^t and some other functions, the Taylor series expansion equals the original function for all values of t as the number of terms approaches infinity⁴. Figure 6.11 shows the Taylor series expansion of e^t around $t = 0$ for a varying number of terms.

We'll expand the first few terms of the Taylor series expansion in equation (6.3) for $\mathbf{X} = \mathbf{A}T$ so we can compare it with other methods.

$$\sum_{k=0}^3 \frac{1}{k!} (\mathbf{A}T)^k = \mathbf{I} + \mathbf{A}T + \frac{1}{2}\mathbf{A}^2T^2 + \frac{1}{6}\mathbf{A}^3T^3$$

Table 6.3 compares the Taylor series expansions of the [discretization](#) methods for the matrix case. These use a more complex formula which we won't present here.

Each of them has different stability properties. The bilinear transform preserves the (in)stability of the continuous time [system](#).

⁴Functions for which their Taylor series expansion converges to and also equals it are called analytic functions.

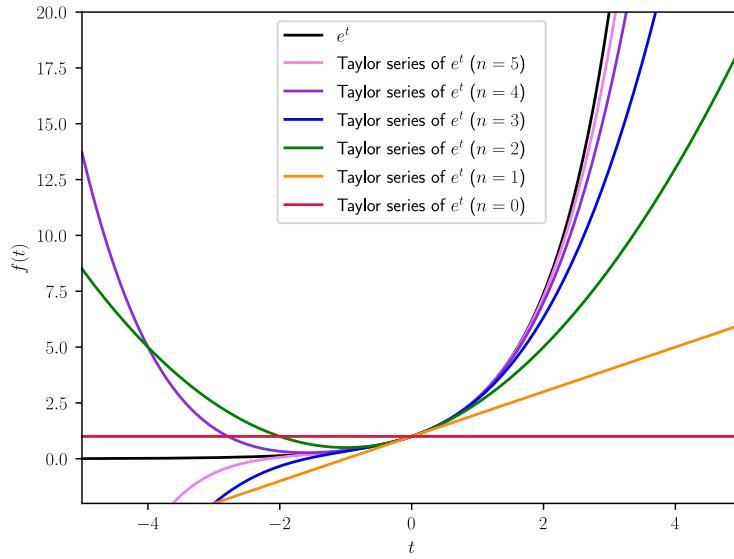


Figure 6.11: Taylor series expansions of e^t around $t = 0$ for n terms

Method	Conversion	Taylor series expansion
Zero-order hold	$\mathbf{A}_d = e^{\mathbf{A}_c T}$	$\mathbf{A}_d = \mathbf{I} + \mathbf{A}_c T + \frac{1}{2}\mathbf{A}_c^2 T^2 + \frac{1}{6}\mathbf{A}_c^3 T^3 + \dots$
Bilinear	$\mathbf{A}_d = (\mathbf{I} + \frac{1}{2}\mathbf{A}_c T) (\mathbf{I} - \frac{1}{2}\mathbf{A}_c T)^{-1}$	$\mathbf{A}_d = \mathbf{I} + \mathbf{A}_c T + \frac{1}{2}\mathbf{A}_c^2 T^2 + \frac{1}{4}\mathbf{A}_c^3 T^3 + \dots$
Forward Euler	$\mathbf{A}_d = \mathbf{I} + \mathbf{A}_c T$	$\mathbf{A}_d = \mathbf{I} + \mathbf{A}_c T$
Reverse Euler	$\mathbf{A}_d = (\mathbf{I} - \mathbf{A}_c T)^{-1}$	$\mathbf{A}_d = \mathbf{I} + \mathbf{A}_c T + \mathbf{A}_c^2 T^2 + \mathbf{A}_c^3 T^3 + \dots$

Table 6.3: Taylor series expansions of discretization methods (matrix case). The zero-order hold discretization method is exact.

6.7 Zero-order hold for state-space

Given the following continuous time state space model

$$\begin{aligned}\dot{\mathbf{x}} &= \mathbf{A}_c \mathbf{x} + \mathbf{B}_c \mathbf{u} + \mathbf{w} \\ \mathbf{y} &= \mathbf{C}_c \mathbf{x} + \mathbf{D}_c \mathbf{u} + \mathbf{v}\end{aligned}$$

where \mathbf{w} is the process noise, \mathbf{v} is the measurement noise, and both are zero-mean white noise sources with covariances of \mathbf{Q}_c and \mathbf{R}_c respectively. \mathbf{w} and \mathbf{v} are defined as normally distributed random variables.

$$\begin{aligned}\mathbf{w} &\sim N(0, \mathbf{Q}_c) \\ \mathbf{v} &\sim N(0, \mathbf{R}_c)\end{aligned}$$

The model can be discretized as follows

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{A}_d \mathbf{x}_k + \mathbf{B}_d \mathbf{u}_k + \mathbf{w}_k \\ \mathbf{y}_k &= \mathbf{C}_d \mathbf{x}_k + \mathbf{D}_d \mathbf{u}_k + \mathbf{v}_k\end{aligned}$$

with covariances

$$\mathbf{w}_k \sim N(0, \mathbf{Q}_d)$$

$$\mathbf{v}_k \sim N(0, \mathbf{R}_d)$$

Theorem 6.7.1 — Zero-order hold for state-space.

$$\mathbf{A}_d = e^{\mathbf{A}_c T} \quad (6.4)$$

$$\mathbf{B}_d = \int_0^T e^{\mathbf{A}_c \tau} d\tau \mathbf{B}_c = \mathbf{A}_c^{-1} (\mathbf{A}_d - \mathbf{I}) \mathbf{B}_c \quad (6.5)$$

$$\mathbf{C}_d = \mathbf{C}_c \quad (6.6)$$

$$\mathbf{D}_d = \mathbf{D}_c \quad (6.7)$$

$$\mathbf{Q}_d = \int_{\tau=0}^T e^{\mathbf{A}_c \tau} \mathbf{Q}_c e^{\mathbf{A}_c^T \tau} d\tau \quad (6.8)$$

$$\mathbf{R}_d = \frac{1}{T} \mathbf{R}_c \quad (6.9)$$

where a subscript of d denotes discrete, a subscript of c denotes the continuous version of the corresponding matrix, T is the sample period for the discrete system, and $e^{\mathbf{A}_c T}$ is the matrix exponential of \mathbf{A}_c .

See appendix H.2 for derivations.

To compute \mathbf{A}_d and \mathbf{B}_d in one step, one can utilize the following property.

$$e^{\begin{bmatrix} \mathbf{A}_c & \mathbf{B}_c \\ \mathbf{0} & \mathbf{0} \end{bmatrix} T} = \begin{bmatrix} \mathbf{A}_d & \mathbf{B}_d \\ \mathbf{0} & \mathbf{I} \end{bmatrix}$$

\mathbf{Q}_d can be computed as

$$\Phi = e^{\begin{bmatrix} -\mathbf{A}_c & \mathbf{Q}_c \\ \mathbf{0} & \mathbf{A}_c^T \end{bmatrix} T} = \begin{bmatrix} -\mathbf{A}_d & \mathbf{A}_d^{-1} \mathbf{Q}_d \\ \mathbf{0} & \mathbf{A}_d^T \end{bmatrix}$$

where $\mathbf{Q}_d = \Phi_{2,2}^T \Phi_{1,2}$ [22].

To see why \mathbf{R}_c is being divided by T , consider the discrete white noise sequence \mathbf{v}_k and the (non-physically realizable) continuous white noise process \mathbf{v} . Whereas $\mathbf{R}_{d,k} = E[\mathbf{v}_k \mathbf{v}_k^T]$ is a covariance matrix, $\mathbf{R}_c(t)$ defined by $E[\mathbf{v}(t) \mathbf{v}^T(\tau)] = \mathbf{R}_c(t) \delta(t - \tau)$ is a spectral density matrix (the Dirac function $\delta(t - \tau)$ has units of 1/sec). The covariance matrix $\mathbf{R}_c(t) \delta(t - \tau)$ has infinite-valued elements. The discrete white noise sequence can be made to approximate the continuous white noise process by shrinking the pulse lengths (T) and increasing their amplitude, such that $\mathbf{R}_d \rightarrow \frac{1}{T} \mathbf{R}_c$.

That is, in the limit as $T \rightarrow 0$, the discrete noise sequence tends to one of infinite-valued pulses of zero duration such that the area under the "impulse" autocorrelation function is $\mathbf{R}_d T$. This is equal to the area \mathbf{R}_c under the continuous white noise impulse autocorrelation function.

This page intentionally left blank

7. Nonlinear control

While many tools exist for designing controllers for linear [systems](#), all [systems](#) in reality are inherently nonlinear. We'll briefly mention some considerations for nonlinear [systems](#).

7.1 Introduction

Recall from linear [system](#) theory that we defined [systems](#) as having the following form:

$$\begin{aligned}\dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} + \mathbf{Gw} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} + \mathbf{v}\end{aligned}$$

In this equation, \mathbf{A} and \mathbf{B} are constant matrices, which means they are both time-invariant and linear (all transformations on the [system state](#) are linear ones, and those transformations remain the same for all time). In nonlinear and time-variant [systems](#), the [state](#) evolution and [output](#) are defined by arbitrary functions of the current [states](#) and [inputs](#).

$$\begin{aligned}\dot{\mathbf{x}} &= f(\mathbf{x}, \mathbf{u}, \mathbf{w}) \\ \mathbf{y} &= h(\mathbf{x}, \mathbf{u}, \mathbf{v})\end{aligned}$$

Nonlinear functions come up regularly when attempting to control the [pose](#) of a vehicle in the global coordinate frame instead of the vehicle's rotating local coordinate frame. Converting from one to the other requires applying a rotation matrix, which consists of sine and cosine operations. These functions are nonlinear.

7.2 Linearization

One way to control nonlinear [systems](#) is to [linearize](#) the [model](#) around a reference point. Then, all the powerful tools that exist for linear controls can be applied. This is done by taking the Jacobians of f and h .

$$\mathbf{A} = \frac{\partial f(\mathbf{x}, \mathbf{u}, \mathbf{w})}{\partial \mathbf{x}} \quad \mathbf{B} = \frac{\partial f(\mathbf{x}, \mathbf{u}, \mathbf{w})}{\partial \mathbf{u}} \quad \mathbf{C} = \frac{\partial h(\mathbf{x}, \mathbf{u}, \mathbf{v})}{\partial \mathbf{x}} \quad \mathbf{D} = \frac{\partial h(\mathbf{x}, \mathbf{u}, \mathbf{v})}{\partial \mathbf{u}}$$

Let $f(\mathbf{x}, \mathbf{u})$ be defined as

$$\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u}) = \begin{bmatrix} f_1(\mathbf{x}, \mathbf{u}) \\ f_2(\mathbf{x}, \mathbf{u}) \\ \vdots \\ f_m(\mathbf{x}, \mathbf{u}) \end{bmatrix}$$

The subscript denotes a row of f , where each row represents the dynamics of a state.

The Jacobian is the partial derivative of a vector-valued function with respect to one of the vector arguments. The Jacobian of f has as many rows as f , and the columns are filled with partial derivatives of f 's rows with respect to each of the argument's elements. For example, the Jacobian of f with respect to \mathbf{x} is

$$\mathbf{A} = \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{x}_1} & \frac{\partial f_1}{\partial \mathbf{x}_2} & \cdots & \frac{\partial f_1}{\partial \mathbf{x}_m} \\ \frac{\partial f_2}{\partial \mathbf{x}_1} & \frac{\partial f_2}{\partial \mathbf{x}_2} & \cdots & \frac{\partial f_2}{\partial \mathbf{x}_m} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial \mathbf{x}_1} & \frac{\partial f_m}{\partial \mathbf{x}_2} & \cdots & \frac{\partial f_m}{\partial \mathbf{x}_m} \end{bmatrix}$$

$\frac{\partial f_1}{\partial \mathbf{x}_1}$ is the partial derivative of the first row of f with respect to the first state, and so on for all rows of f and states. This has n^2 permutations and thus produces a square matrix.

The Jacobian of f with respect to \mathbf{u} is

$$\mathbf{B} = \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{u}_1} & \frac{\partial f_1}{\partial \mathbf{u}_2} & \cdots & \frac{\partial f_1}{\partial \mathbf{u}_n} \\ \frac{\partial f_2}{\partial \mathbf{u}_1} & \frac{\partial f_2}{\partial \mathbf{u}_2} & \cdots & \frac{\partial f_2}{\partial \mathbf{u}_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial \mathbf{u}_1} & \frac{\partial f_m}{\partial \mathbf{u}_2} & \cdots & \frac{\partial f_m}{\partial \mathbf{u}_n} \end{bmatrix}$$

$\frac{\partial f_1}{\partial \mathbf{u}_1}$ is the partial derivative of the first row of f with respect to the first input, and so on for all rows of f and inputs. This has $m \times n$ permutations and can produce a nonsquare matrix if $m \neq n$.

Linearization of a nonlinear equation is a Taylor series expansion to only the first-order terms (that is, terms whose variables have exponents on the order of x^1). This is where the small angle approximations for $\sin \theta$ and $\cos \theta$ (θ and 1 respectively) come from.

Higher order partial derivatives can be added to better approximate the nonlinear dynamics. We typically only linearize around equilibrium points¹ because we are interested in how the system behaves when perturbed from equilibrium. An FAQ on this goes into more detail [17]. To be clear though, linearizing the system around the current state as the system evolves does give a closer approximation over time.

Note that linearization with static matrices (that is, with a time-invariant linear system) only works if the original system in question is feedback linearizable.

7.3 Lyapunov stability

Lyapunov stability is a fundamental concept in nonlinear control, so we're going to give a brief overview of what it is so students can research it further.

¹Equilibrium points are points where $\dot{\mathbf{x}} = \mathbf{0}$. At these points, the system is in steady-state.

Since the [state](#) evolution in nonlinear [systems](#) is defined by a function rather than a constant matrix, the [system](#)'s poles as determined by [linearization](#) move around. Nonlinear control uses Lyapunov stability to determine if nonlinear [systems](#) are stable. From a linear control theory point of view, Lyapunov stability says the [system](#) is stable if, for a given initial condition, all possible eigenvalues of \mathbf{A} from that point on remain in the left-half plane. However, nonlinear control uses a different definition.

Lyapunov stability means that the [system](#) trajectory can be kept arbitrarily close to the origin by starting sufficiently close to it. Lyapunov's direct method uses a function consisting of the energy in a [system](#) or derivatives of the [system](#)'s [state](#) to prove stability around an equilibrium point. This is done by showing that the function, and thus its inputs, decay to some ground state.

More than one Lyapunov function can prove stability, and if one function doesn't prove it, another candidate should be tried. For this reason, we refer to these functions as *Lyapunov candidate functions*.

7.4 Affine systems

Let $\mathbf{x} = \mathbf{x}_0 + \delta\mathbf{x}$ and $\mathbf{u} = \mathbf{u}_0 + \delta\mathbf{u}$ where $\delta\mathbf{x}$ and $\delta\mathbf{u}$ are perturbations from $(\mathbf{x}_0, \mathbf{u}_0)$. A first-order linearization of $\dot{\mathbf{x}} = f(\mathbf{x}, \mathbf{u})$ around $(\mathbf{x}_0, \mathbf{u}_0)$ gives

$$\begin{aligned}\dot{\mathbf{x}} &\approx f(\mathbf{x}_0, \mathbf{u}_0) + \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} \bigg|_{\mathbf{x}_0, \mathbf{u}_0} \delta\mathbf{x} + \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}} \bigg|_{\mathbf{x}_0, \mathbf{u}_0} \delta\mathbf{u} \\ \dot{\mathbf{x}} &= f(\mathbf{x}_0, \mathbf{u}_0) + \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} \bigg|_{\mathbf{x}_0, \mathbf{u}_0} \delta\mathbf{x} + \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}} \bigg|_{\mathbf{x}_0, \mathbf{u}_0} \delta\mathbf{u}\end{aligned}$$

An affine system is a linear system with a constant offset in the dynamics. If $(\mathbf{x}_0, \mathbf{u}_0)$ is an equilibrium point, $f(\mathbf{x}_0, \mathbf{u}_0) = \mathbf{0}$, the resulting [model](#) is linear, and LQR works as usual. If $(\mathbf{x}_0, \mathbf{u}_0)$ is, say, the current operating point rather than an equilibrium point, the easiest way to correctly apply LQR is

1. Find a control input \mathbf{u}_0 that makes $(\mathbf{x}_0, \mathbf{u}_0)$ an equilibrium point.
2. Obtain an LQR for the linearized system.
3. Add \mathbf{u}_0 to the LQR's control input.

For a control-affine [system](#) (a nonlinear [system](#) with linear control inputs) $\dot{\mathbf{x}} = f(\mathbf{x}) + \mathbf{B}\mathbf{u}$, \mathbf{u}_0 can be derived via plant inversion as follows.

$$\begin{aligned}\dot{\mathbf{x}} &= f(\mathbf{x}_0) + \mathbf{B}\mathbf{u}_0 \\ \mathbf{0} &= f(\mathbf{x}_0) + \mathbf{B}\mathbf{u}_0 \\ \mathbf{B}\mathbf{u}_0 &= -f(\mathbf{x}_0) \\ \mathbf{u}_0 &= -\mathbf{B}^\dagger f(\mathbf{x}_0)\end{aligned}\tag{7.1}$$

7.4.1 Feedback linearization for reference tracking

Feedback linearization lets us erase the nonlinear dynamics of a system so we can apply our own (usually linear) dynamics for [reference](#) tracking. To do this, we will perform a similar procedure as in subsection 5.10.1 and solve for \mathbf{u} given the [reference](#) dynamics in $\dot{\mathbf{r}}$.

$$\begin{aligned}\dot{\mathbf{r}} &= f(\mathbf{x}) + \mathbf{B}\mathbf{u} \\ \mathbf{B}\mathbf{u} &= \dot{\mathbf{r}} - f(\mathbf{x})\end{aligned}$$

$$\mathbf{u} = \mathbf{B}^\dagger(\dot{\mathbf{r}} - f(\mathbf{x})) \quad (7.2)$$

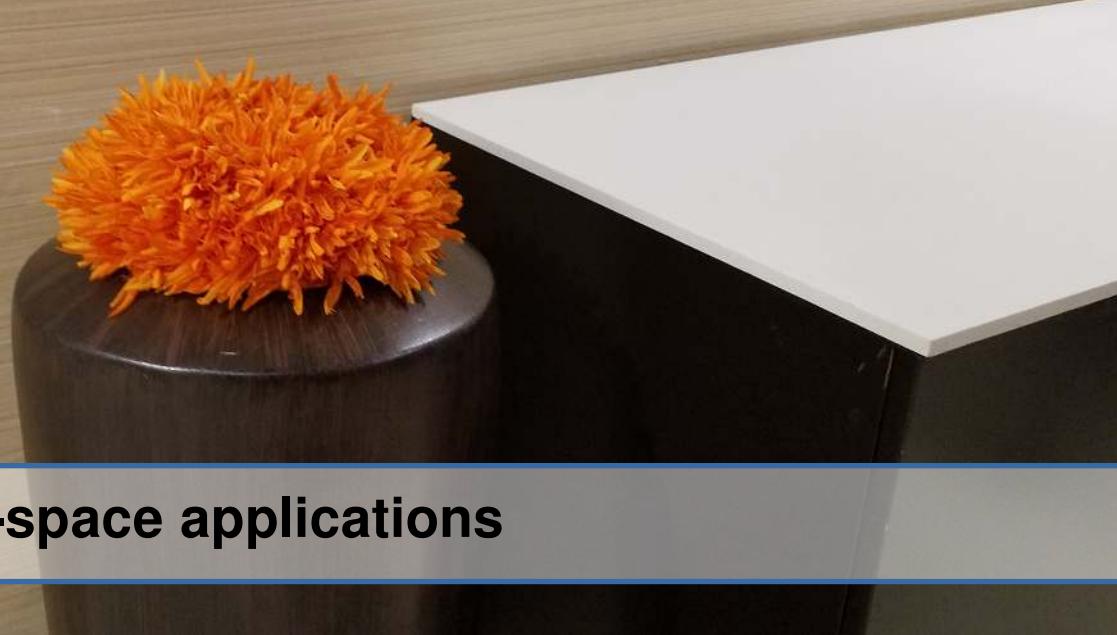


To use equation (7.2) in a discrete controller, one can approximate $\dot{\mathbf{r}}$ with $\frac{\mathbf{r}_{k+1} - \mathbf{r}_k}{T}$ where T is the time period between the two [references](#).

7.5 Further reading

For a more complex type of trajectory tracking, read *A guiding vector field algorithm for path following control of nonholonomic mobile robots* [19].

To learn more about nonlinear control, we recommend watching the underactuated robotics lectures by MIT [28]. Links to it and the associated lecture note [29] are in the bibliography. The books *Nonlinear Dynamics and Chaos* by Steven Strogatz and *Applied Nonlinear Control* by Jean-Jacques Slotine are also good references.



8. State-space applications

Up to now, we've just been teaching what tools are available. Now, we'll go into specifics on how to apply them and provide advice on certain applications.

The code shown in each example can be obtained from frcccontrol's Git repository at <https://github.com/calcmogul/frcccontrol/tree/master/examples>. See appendix D for setup instructions.

8.1 Elevator

8.1.1 Continuous state-space model

The position and velocity derivatives of the elevator can be written as

$$\dot{x} = v \quad (8.1)$$

$$\dot{v} = a \quad (8.2)$$

where by equation (13.15),

$$a = \frac{GK_t}{Rrm}V - \frac{G^2K_t}{Rr^2mK_v}v$$

Substitute this into equation (8.2).

$$\begin{aligned} \dot{v} &= \frac{GK_t}{Rrm}V - \frac{G^2K_t}{Rr^2mK_v}v \\ \dot{v} &= -\frac{G^2K_t}{Rr^2mK_v}v + \frac{GK_t}{Rrm}V \end{aligned} \quad (8.3)$$

Factor out v and V into column vectors.

$$\dot{[v]} = \left[-\frac{G^2K_t}{Rr^2mK_v} \right] [v] + \left[\frac{GK_t}{Rrm} \right] [V]$$

Augment the matrix equation with the position state x , which has the model equation $\dot{x} = v$. The matrix elements corresponding to v will be 1, and the others will be 0 since they don't appear, so $\dot{x} = 0x + 1v + 0V$. The existing rows will have zeroes inserted where x is multiplied in.

$$\begin{bmatrix} \dot{x} \\ v \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & -\frac{G^2 K_t}{Rr^2 m K_v} \end{bmatrix} \begin{bmatrix} x \\ v \end{bmatrix} + \begin{bmatrix} 0 \\ \frac{G K_t}{R r m} \end{bmatrix} [V]$$

Theorem 8.1.1 — Elevator state-space model.

$$\begin{aligned} \dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\ \mathbf{x} &= \begin{bmatrix} x \\ v \end{bmatrix} \quad \mathbf{y} = x \quad \mathbf{u} = V \\ \mathbf{A} &= \begin{bmatrix} 0 & 1 \\ 0 & -\frac{G^2 K_t}{Rr^2 m K_v} \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 \\ \frac{G K_t}{R r m} \end{bmatrix} \quad \mathbf{C} = [1 \quad 0] \quad \mathbf{D} = 0 \end{aligned} \quad (8.4)$$

8.1.2 Model augmentation

As per subsection 5.11.2, we will now augment the [model](#) so a u_{error} state is added to the [control input](#).

The [plant](#) and [observer](#) augmentations should be performed before the [model](#) is [discretized](#). After the controller gain is computed with the unaugmented discrete [model](#), the controller may be augmented. Therefore, the [plant](#) and [observer](#) augmentations assume a continuous [model](#) and the [controller](#) augmentation assumes a discrete [controller](#).

$$\begin{aligned} \mathbf{x}_{aug} &= \begin{bmatrix} x \\ v \\ u_{error} \end{bmatrix} \quad \mathbf{y} = x \quad \mathbf{u} = V \\ \mathbf{A}_{aug} &= \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{0}_{1 \times 2} & 0 \end{bmatrix} \quad \mathbf{B}_{aug} = \begin{bmatrix} \mathbf{B} \\ 0 \end{bmatrix} \quad \mathbf{C}_{aug} = [\mathbf{C} \quad 0] \quad \mathbf{D}_{aug} = \mathbf{D} \end{aligned} \quad (8.5)$$

$$\mathbf{K}_{aug} = [\mathbf{K} \quad 1] \quad \mathbf{r}_{aug} = \begin{bmatrix} \mathbf{r} \\ 0 \end{bmatrix} \quad (8.6)$$

This will compensate for unmodeled dynamics such as gravity. However, using a constant voltage feedforward to counteract gravity is preferred over u_{error} estimation in this case because it results in a simpler controller with similar performance.

8.1.3 Gravity feedforward

Input voltage is proportional to force and gravity is a constant force, so a constant voltage feedforward can compensate for gravity. We'll model gravity as a disturbance described by $-mg$. To compensate for it, we want to find a voltage that is equal and opposite to it. The bottom row of the continuous elevator model contains the acceleration terms.

$$B_{uff} = -(unmodeled\ dynamics)$$

where B is the motor acceleration term from \mathbf{B} and u_{ff} is the voltage feedforward.

$$B_{uff} = -(-mg)$$

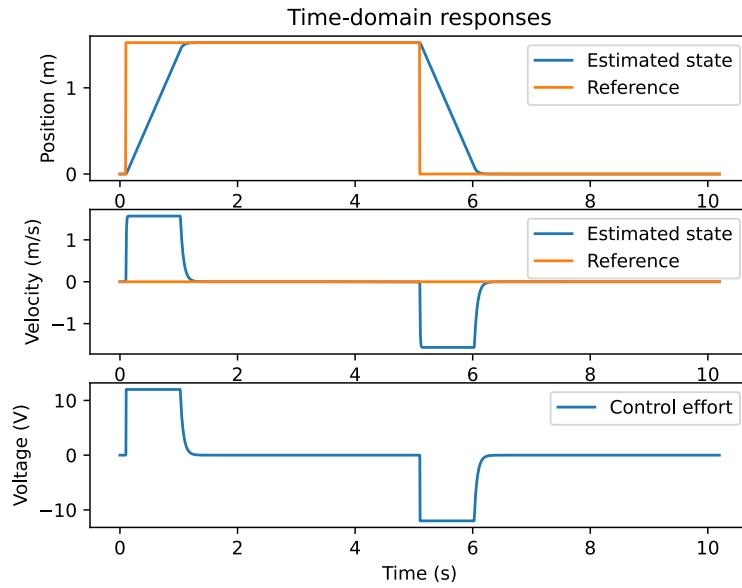


Figure 8.1: Elevator response

$$\begin{aligned}
 Bu_{ff} &= mg \\
 \frac{GK_t}{Rrm}u_{ff} &= mg \\
 u_{ff} &= \frac{Rrm^2g}{GK_t}
 \end{aligned}$$

8.1.4 Simulation

Python Control will be used to [discretize](#) the [model](#) and simulate it. One of the frcontrol examples¹ creates and tests a controller for it. Figure 8.1 shows the closed-loop [system](#) response.

8.1.5 Implementation

The script linked above also generates two files: ElevatorCoeffs.h and ElevatorCoeffs.cpp. These can be used with the WPILib StateSpacePlant, StateSpaceController, and StateSpaceObserver classes in C++ and Java. A C++ implementation of this elevator controller is available online².

8.2 Flywheel

8.2.1 Continuous state-space model

By equation (13.18)

$$\dot{\omega} = -\frac{G^2 K_t}{K_v R J} \omega + \frac{G K_t}{R J} V$$

Factor out ω and V into column vectors.

$$\dot{[\omega]} = \left[-\frac{G^2 K_t}{K_v R J} \right] [\omega] + \left[\frac{G K_t}{R J} \right] [V]$$

¹<https://github.com/calcmogul/frcontrol/blob/master/examples/elevator.py>

²<https://github.com/calcmogul/allwpilib/tree/state-space/wpilibcExamples/src/main/cpp/examples/StateSpaceElevator>

Theorem 8.2.1 — Flywheel state-space model.

$$\begin{aligned}
 \dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\
 \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\
 \mathbf{x} &= \omega \quad \mathbf{y} = \omega \quad \mathbf{u} = V \\
 \mathbf{A} &= -\frac{G^2 K_t}{K_v R J} \quad \mathbf{B} = \frac{G K_t}{R J} \quad \mathbf{C} = 1 \quad \mathbf{D} = 0
 \end{aligned} \tag{8.7}$$

8.2.2 Model augmentation

As per subsection 5.11.2, we will now augment the `model` so a u_{error} state is added to the `control input`.

The `plant` and `observer` augmentations should be performed before the `model` is `discretized`. After the controller gain is computed with the unaugmented discrete `model`, the controller may be augmented. Therefore, the `plant` and `observer` augmentations assume a continuous `model` and the `controller` augmentation assumes a discrete controller.

$$\begin{aligned}
 \mathbf{x} &= \begin{bmatrix} \omega \\ u_{error} \end{bmatrix} \quad \mathbf{y} = \omega \quad \mathbf{u} = V \\
 \mathbf{A}_{aug} &= \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ 0 & 0 \end{bmatrix} \quad \mathbf{B}_{aug} = \begin{bmatrix} \mathbf{B} \\ 0 \end{bmatrix} \quad \mathbf{C}_{aug} = [\mathbf{C} \ 0] \quad \mathbf{D}_{aug} = \mathbf{D}
 \end{aligned} \tag{8.8}$$

$$\mathbf{K}_{aug} = [\mathbf{K} \ 1] \quad \mathbf{r}_{aug} = \begin{bmatrix} \mathbf{r} \\ 0 \end{bmatrix} \tag{8.9}$$

This will compensate for unmodeled dynamics such as projectiles slowing down the flywheel.

8.2.3 Simulation

Python Control will be used to `discretize` the `model` and simulate it. One of the `frcontrol` examples³ creates and tests a controller for it. Figure 8.2 shows the closed-loop `system` response.

Notice how the `control effort` when the `reference` is reached is nonzero. This is a plant inversion feedforward compensating for the `system` dynamics attempting to slow the flywheel down when no voltage is applied.

8.2.4 Implementation

The script linked above also generates two files: `FlywheelCoeffs.h` and `FlywheelCoeffs.cpp`. These can be used with the WPILib `StateSpacePlant`, `StateSpaceController`, and `StateSpaceObserver` classes in C++ and Java. A C++ implementation of this flywheel controller is available online⁴.

8.2.5 Flywheel model without encoder

In the FIRST Robotics Competition, we can get the current drawn for specific channels on the power distribution panel. We can theoretically use this to estimate the angular velocity of a DC motor without an encoder. We'll start with the flywheel model derived earlier as equation (13.18).

$$\dot{\omega} = \frac{G K_t}{R J} V - \frac{G^2 K_t}{K_v R J} \omega$$

³<https://github.com/calcmogul/frcontrol/blob/master/examples/flywheel.py>

⁴<https://github.com/calcmogul/allwpilib/tree/state-space/wpilibcExamples/src/main/cpp/examples/StateSpaceFlywheel>

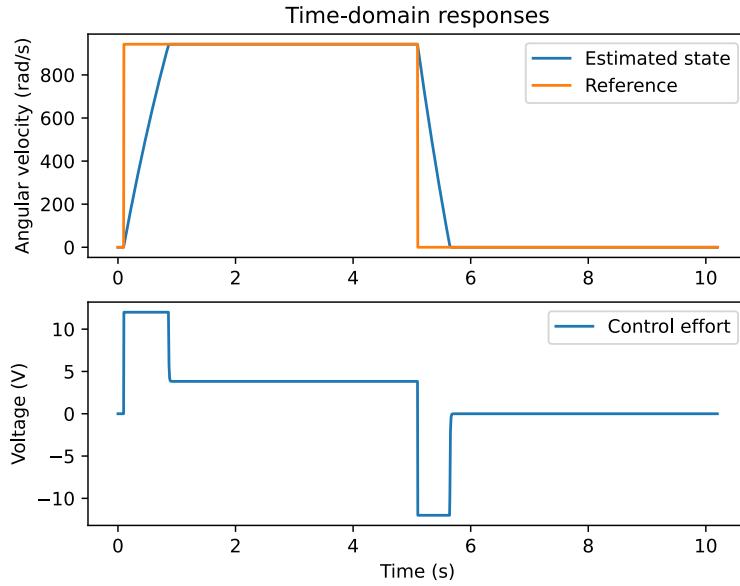


Figure 8.2: Flywheel response

$$\dot{\omega} = -\frac{G^2 K_t}{K_v R J} \omega + \frac{G K_t}{R J} V$$

Next, we'll derive the current I as an output.

$$\begin{aligned} V &= IR + \frac{\omega}{K_v} \\ IR &= V - \frac{\omega}{K_v} \\ I &= -\frac{1}{K_v R} \omega + \frac{1}{R} V \end{aligned}$$

Therefore,

Theorem 8.2.2 — Flywheel state-space model without encoder.

$$\begin{aligned} \dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\ \mathbf{x} &= \omega \quad \mathbf{y} = I \quad \mathbf{u} = V \\ \mathbf{A} &= -\frac{G^2 K_t}{K_v R J} \quad \mathbf{B} = \frac{G K_t}{R J} \quad \mathbf{C} = -\frac{1}{K_v R} \quad \mathbf{D} = \frac{1}{R} \end{aligned} \tag{8.10}$$

Notice that in this model, the output doesn't provide any direct measurements of the state. To estimate the full state (also known as full observability), we only need the outputs to collectively include linear combinations of every state⁵. We'll revisit this in chapter 9 with an example that uses range measurements to estimate an object's orientation.

⁵While the flywheel model's outputs are a linear combination of both the states and inputs, inputs don't provide new information about the states. Therefore, they don't affect whether the system is observable.

The effectiveness of this [model's observer](#) is heavily dependent on the quality of the current sensor used. If the sensor's noise isn't zero-mean, the [observer](#) won't converge to the true [state](#).

8.2.6 Voltage compensation

To improve controller [tracking](#), one may want to use the voltage renormalized to the power rail voltage to compensate for voltage drop when current spikes occur. This can be done as follows.

$$V = V_{cmd} \frac{V_{nominal}}{V_{rail}} \quad (8.11)$$

where V is the [controller](#)'s new input voltage, V_{cmd} is the old input voltage, $V_{nominal}$ is the rail voltage when effects like voltage drop due to current draw are ignored, and V_{rail} is the real rail voltage.

To drive the [model](#) with a more accurate voltage that includes voltage drop, the reciprocal can be used.

$$V = V_{cmd} \frac{V_{rail}}{V_{nominal}} \quad (8.12)$$

where V is the [model](#)'s new input voltage. Note that if both the [controller](#) compensation and [model](#) compensation equations are applied, the original voltage is obtained. The [model](#) input only drops from ideal if the compensated [controller](#) voltage saturates.

8.3 Single-jointed arm

8.3.1 Continuous state-space model

The angle and angular rate derivatives of the arm can be written as

$$\dot{\theta}_{arm} = \omega_{arm} \quad (8.13)$$

$$\dot{\omega}_{arm} = \dot{\omega}_{arm} \quad (8.14)$$

By equation (13.22)

$$\dot{\omega}_{arm} = -\frac{G^2 K_t}{K_v R J} \omega_{arm} + \frac{G K_t}{R J} V$$

Factor out ω_{arm} and V into column vectors.

$$\dot{[\omega_{arm}]} = \left[-\frac{G^2 K_t}{K_v R J} \right] [\omega_{arm}] + \left[\frac{G K_t}{R J} \right] [V]$$

Augment the matrix equation with the angle state $\dot{\theta}_{arm}$, which has the model equation $\dot{\theta}_{arm} = \omega_{arm}$. The matrix elements corresponding to ω_{arm} will be 1, and the others will be 0 since they don't appear, so $\dot{\theta}_{arm} = 0\theta_{arm} + 1\omega_{arm} + 0V$. The existing rows will have zeroes inserted where θ_{arm} is multiplied in.

$$\dot{[\theta_{arm} \omega_{arm}]} = \begin{bmatrix} 0 & 1 \\ 0 & -\frac{G^2 K_t}{K_v R J} \end{bmatrix} [\theta_{arm} \omega_{arm}] + \begin{bmatrix} 0 \\ \frac{G K_t}{R J} \end{bmatrix} [V]$$

Theorem 8.3.1 — Single-jointed arm state-space model.

$$\begin{aligned}\dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\ \mathbf{x} &= \begin{bmatrix} \theta_{arm} \\ \omega_{arm} \end{bmatrix} \quad \mathbf{y} = \theta_{arm} \quad \mathbf{u} = V \\ \mathbf{A} &= \begin{bmatrix} 0 & 1 \\ 0 & -\frac{G^2 K_t}{K_v R J} \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 \\ \frac{G K_t}{R J} \end{bmatrix} \quad \mathbf{C} = [1 \ 0] \quad \mathbf{D} = 0\end{aligned}\quad (8.15)$$

8.3.2 Model augmentation

As per subsection 5.11.2, we will now augment the [model](#) so a u_{error} state is added to the [control input](#).

The [plant](#) and [observer](#) augmentations should be performed before the [model](#) is [discretized](#). After the [controller](#) gain is computed with the unaugmented discrete [model](#), the controller may be augmented. Therefore, the [plant](#) and [observer](#) augmentations assume a continuous [model](#) and the [controller](#) augmentation assumes a discrete [controller](#).

$$\begin{aligned}\mathbf{x}_{aug} &= \begin{bmatrix} \mathbf{x} \\ u_{error} \end{bmatrix} \quad \mathbf{y} = \theta_{arm} \quad \mathbf{u} = V \\ \mathbf{A}_{aug} &= \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{0}_{1 \times 2} & 0 \end{bmatrix} \quad \mathbf{B}_{aug} = \begin{bmatrix} \mathbf{B} \\ 0 \end{bmatrix} \quad \mathbf{C}_{aug} = [\mathbf{C} \ 0] \quad \mathbf{D}_{aug} = \mathbf{D}\end{aligned}\quad (8.16)$$

$$\mathbf{K}_{aug} = [\mathbf{K} \ 1] \quad \mathbf{r}_{aug} = \begin{bmatrix} \mathbf{r} \\ 0 \end{bmatrix}\quad (8.17)$$

This will compensate for unmodeled dynamics such as gravity or other external loading from lifted objects. However, if only gravity compensation is desired, a feedforward of the form $u_{ff} = V_{gravity} \cos \theta$ is preferred where $V_{gravity}$ is the voltage required to hold the arm level with the ground and θ is the angle of the arm with the ground.

8.3.3 Gravity feedforward

Input voltage is proportional to torque and gravity is a constant force, but the torque applied against the motor varies according to the arm's angle. We'll use sum of torques to find a compensating torque.

We'll model gravity as a disturbance described by $-mg$ where m is the arm's mass. To compensate for it, we want to find a torque that is equal and opposite to the torque applied to the arm by gravity. The bottom row of the continuous elevator model contains the angular acceleration terms, so B_{uff} is angular acceleration caused by the motor; JB_{uff} is the torque.

$$\begin{aligned}JB_{uff} &= -(\mathbf{r} \times \mathbf{F}) \\ JB_{uff} &= -(rF \cos \theta)\end{aligned}$$

Torque is usually written as $rF \sin \theta$ where θ is the angle between the \mathbf{r} and \mathbf{F} vectors, but θ in this case is being measured from the horizontal axis rather than the vertical one, so the force vector is $\frac{\pi}{4}$ radians out of phase. Thus, an angle of 0 results in the maximum torque from gravity being applied rather than the minimum.

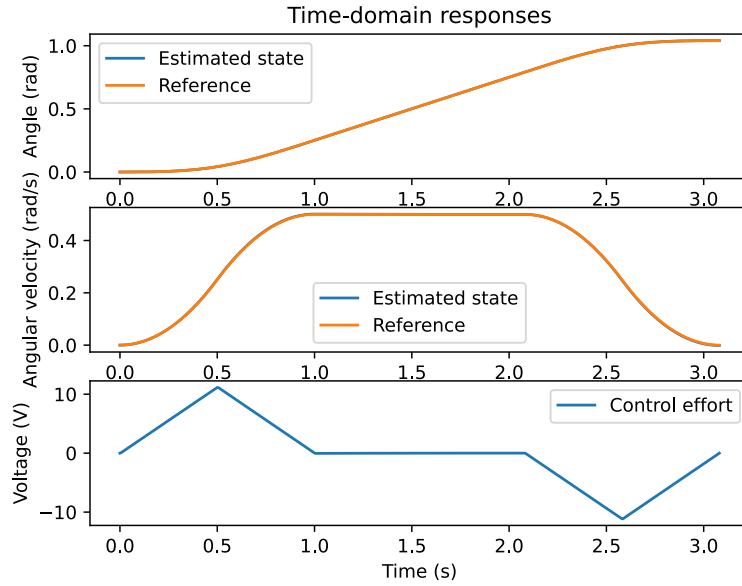


Figure 8.3: Single-jointed arm response

The force of gravity mg is applied at the center of the arm's mass. For a uniform beam, this is halfway down its length, or $\frac{L}{2}$ where L is the length of the arm.

$$J B_{eff} = - \left(\left(\frac{L}{2} \right) (-mg) \cos \theta \right)$$

$$J B_{eff} = mg \frac{L}{2} \cos \theta$$

$B = \frac{GK_t}{RJ}$, so

$$J \frac{GK_t}{RJ} u_{eff} = mg \frac{L}{2} \cos \theta$$

$$u_{eff} = \frac{RJ}{JGK_t} mg \frac{L}{2} \cos \theta$$

$$u_{eff} = \frac{L}{2} \frac{Rmg}{GK_t} \cos \theta$$

$\frac{L}{2}$ can be adjusted according to the location of the arm's center of mass.

8.3.4 Simulation

Python Control will be used to [discretize](#) the [model](#) and simulate it. One of the frcontrol examples⁶ creates and tests a controller for it. Figure 8.3 shows the closed-loop [system](#) response.

8.3.5 Implementation

The script linked above also generates two files: SingleJointedArmCoeffs.h and SingleJointedArmCoeffs.cpp. These can be used with the WPILib StateSpacePlant, StateSpaceController, and StateS-

⁶https://github.com/calcmogul/frcontrol/blob/master/examples/single_jointed_arm.py

paceObserver classes in C++ and Java. A C++ implementation of this single-jointed arm controller is available online⁷.

8.4 Pendulum

8.4.1 State-space model

Below is the [model](#) for a pendulum

$$\ddot{\theta} = -\frac{g}{l} \sin \theta$$

where θ is the angle of the pendulum and l is the length of the pendulum.

Since state-space representation requires that only single derivatives be used, they should be broken up as separate [states](#). We'll reassign $\dot{\theta}$ to be ω so the derivatives are easier to keep straight for state-space representation.

$$\dot{\omega} = -\frac{g}{l} \sin \theta$$

Now separate the [states](#).

$$\begin{aligned}\dot{\theta} &= \omega \\ \dot{\omega} &= -\frac{g}{l} \sin \theta\end{aligned}$$

This makes our state vector $[\theta \quad \omega]^T$ and our nonlinear model

$$f(\mathbf{x}, \mathbf{u}) = \begin{bmatrix} \omega \\ -\frac{g}{l} \sin \theta \end{bmatrix}$$

Linearization around $\theta = 0$

To apply our tools for linear control theory, the [model](#) must be a linear combination of the [states](#) and [inputs](#) (addition and multiplication by constants). Since this [model](#) is nonlinear on account of the sine function, we should [linearize](#) it.

Linearization finds a tangent line to the nonlinear dynamics at a desired point in the state-space. The Taylor series is a way to approximate arbitrary functions with polynomials, so we can use it for linearization.

The taylor series expansion for $\sin \theta$ around $\theta = 0$ is $\theta - \frac{1}{6}\theta^3 + \frac{1}{120}\theta^5 - \dots$. We'll take just the first-order term θ to obtain a linear function.

$$\begin{aligned}\dot{\theta} &= \omega \\ \dot{\omega} &= -\frac{g}{l}\theta\end{aligned}$$

Now write the [model](#) in state-space representation. We'll write out the system of equations with the zeroed variables included to assist with this.

$$\begin{aligned}\dot{\theta} &= 0\theta + 1\omega \\ \dot{\omega} &= -\frac{g}{l}\theta + 0\omega\end{aligned}$$

⁷<https://github.com/calcmogul/allwpilib/tree/state-space/wpilibcExamples/src/main/cpp/examples/StateSpaceSingleJointedArm>

Factor out θ and ω into a column vector.

$$\begin{bmatrix} \dot{\theta} \\ \omega \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ -\frac{g}{l} & 0 \end{bmatrix} \begin{bmatrix} \theta \\ \omega \end{bmatrix} \quad (8.18)$$

Linearization with the Jacobian

Here's the original nonlinear model in state-space representation.

$$f(\mathbf{x}, \mathbf{u}) = \begin{bmatrix} \omega \\ -\frac{g}{l} \sin \theta \end{bmatrix}$$

If we want to linearize around an arbitrary point, we can take the Jacobian with respect to \mathbf{x} .

$$\frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} = \begin{bmatrix} 0 & 1 \\ -\frac{g}{l} \cos \theta & 0 \end{bmatrix}$$

For full [state](#) feedback, knowledge of all [states](#) is required. If not all [states](#) are measured directly, an estimator can be used to supplement them.

We may only be measuring θ in the pendulum example, not $\dot{\theta}$, so we'll need to estimate the latter. The \mathbf{C} matrix the [observer](#) would use in this case is

$$\mathbf{C} = [1 \ 0]$$

Therefore, the measurement vector is

$$\begin{aligned} \mathbf{y} &= \mathbf{Cx} \\ \mathbf{y} &= [1 \ 0] \begin{bmatrix} \theta \\ \omega \end{bmatrix} \\ \mathbf{y} &= 1\theta + 0\omega \\ \mathbf{y} &= \theta \end{aligned}$$

Theorem 8.4.1 — Linear time-varying pendulum state-space model with no input.

$$\begin{aligned} \dot{\mathbf{x}} &= \mathbf{Ax} \\ \mathbf{y} &= \mathbf{Cx} \\ \mathbf{x} &= \begin{bmatrix} \theta \\ \omega \end{bmatrix} \quad \mathbf{y} = \theta \\ \mathbf{A} &= \begin{bmatrix} 0 & 1 \\ -\frac{g}{l} \cos \theta & 0 \end{bmatrix} \quad \mathbf{C} = [1 \ 0] \end{aligned} \quad (8.19)$$

8.5 Differential drive

8.5.1 Velocity subspace state-space model

By equations (13.34) and (13.35)

$$\dot{v}_l = \left(\frac{1}{m} + \frac{r_b^2}{J} \right) (C_1 v_l + C_2 V_l) + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) (C_3 v_r + C_4 V_r)$$

$$\dot{v}_r = \left(\frac{1}{m} - \frac{r_b^2}{J} \right) (C_1 v_l + C_2 V_l) + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) (C_3 v_r + C_4 V_r)$$

Regroup the terms into states v_l and v_r and inputs V_l and V_r .

$$\begin{aligned}\dot{v}_l &= \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_1 v_l + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_2 V_l + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_3 v_r + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_4 V_r \\ \dot{v}_r &= \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_1 v_l + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_2 V_l + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_3 v_r + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_4 V_r\end{aligned}$$

$$\begin{aligned}\dot{v}_l &= \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_1 v_l + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_3 v_r + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_2 V_l + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_4 V_r \\ \dot{v}_r &= \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_1 v_l + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_3 v_r + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_2 V_l + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_4 V_r\end{aligned}$$

Factor out v_l and v_r into a column vector and V_l and V_r into a column vector.

$$\begin{bmatrix} \dot{v}_l \\ \dot{v}_r \end{bmatrix} = \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_3 \\ \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_3 \end{bmatrix} \begin{bmatrix} v_l \\ v_r \end{bmatrix} + \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_4 \end{bmatrix} \begin{bmatrix} V_l \\ V_r \end{bmatrix}$$

Theorem 8.5.1 — Differential drive velocity state-space model.

$$\begin{aligned}\dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\ \mathbf{x} &= \begin{bmatrix} v_l \\ v_r \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} v_l \\ v_r \end{bmatrix} \quad \mathbf{u} = \begin{bmatrix} V_l \\ V_r \end{bmatrix} \\ \mathbf{A} &= \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_3 \\ \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_3 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J} \right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J} \right) C_4 \end{bmatrix} \\ \mathbf{C} &= \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad \mathbf{D} = \mathbf{0}_{2 \times 2}\end{aligned}\tag{8.20}$$

where $C_1 = -\frac{G_l^2 K_t}{K_v R r^2}$, $C_2 = \frac{G_l K_t}{R r}$, $C_3 = -\frac{G_r^2 K_t}{K_v R r^2}$, and $C_4 = \frac{G_r K_t}{R r}$.

Simulation

Python Control will be used to [discretize](#) the [model](#) and simulate it. One of the frcontrol examples⁸ creates and tests a controller for it. Figure 8.4 shows the closed-loop [system](#) response.

Given the high inertia in drivetrains, it's better to drive the reference with a motion profile instead of a [step input](#) for reproducibility.

⁸https://github.com/calcmogul/frcontrol/blob/master/examples/differential_drive.py

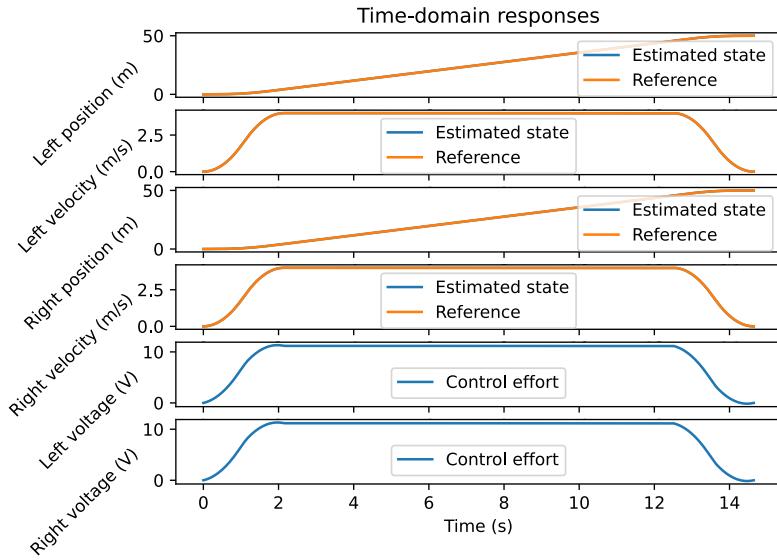


Figure 8.4: Drivetrain response

8.5.2 Linear time-varying model

The model in theorem 8.5.1 is linear, but only includes the velocity dynamics, not the dynamics of the drivetrain's global pose. The change in global pose is defined by these three equations.

$$\begin{aligned}\dot{x} &= \frac{v_l + v_r}{2} \cos \theta = \frac{v_r}{2} \cos \theta + \frac{v_l}{2} \cos \theta \\ \dot{y} &= \frac{v_l + v_r}{2} \sin \theta = \frac{v_r}{2} \sin \theta + \frac{v_l}{2} \sin \theta \\ \dot{\theta} &= \frac{v_r - v_l}{2r_b} = \frac{v_r}{2r_b} - \frac{v_l}{2r_b}\end{aligned}$$

Next, we'll augment the linear subspace's state with the global pose x , y , and θ . Here's the model as a vector function where $\mathbf{x} = [x \ y \ \theta \ v_l \ v_r]^T$ and $\mathbf{u} = [V_l \ V_r]^T$.

$$f(\mathbf{x}, \mathbf{u}) = \begin{bmatrix} \frac{v_r}{2} \cos \theta + \frac{v_l}{2} \cos \theta \\ \frac{v_r}{2} \sin \theta + \frac{v_l}{2} \sin \theta \\ \frac{v_r}{2r_b} - \frac{v_l}{2r_b} \\ \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 v_l + \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 v_r + \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 V_l + \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 V_r \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 v_l + \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 v_r + \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 V_l + \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 V_r \end{bmatrix} \quad (8.21)$$

As mentioned in chapter 7, one can approximate a nonlinear system via linearizations around points of interest in the state-space and design controllers for those linearized subspaces. If we sample linearization points progressively closer together, we converge on a control policy for the original nonlinear system. Since the linear [plant](#) being controlled varies with time, its controller is called a linear time-varying (LTV) controller.

If we use LQRs for the linearized subspaces, the nonlinear control policy will also be locally optimal. We'll be taking this approach with a differential drive. To create an LQR, we need to linearize

equation (8.21).

$$\frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} = \begin{bmatrix} 0 & 0 & -\frac{v_l+v_r}{2} \sin \theta & \frac{1}{2} \cos \theta & \frac{1}{2} \cos \theta \\ 0 & 0 & \frac{v_l+v_r}{2} \cos \theta & \frac{1}{2} \sin \theta & \frac{1}{2} \sin \theta \\ 0 & 0 & 0 & -\frac{1}{2r_b} & \frac{1}{2r_b} \\ 0 & 0 & 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ 0 & 0 & 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix}$$

$$\frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 \end{bmatrix}$$

Therefore,

Theorem 8.5.2 — Linear time-varying differential drive state-space model.

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$$

$$\mathbf{y} = \mathbf{Cx} + \mathbf{Du}$$

$$\mathbf{x} = [x \ y \ \theta \ v_l \ v_r]^T \quad \mathbf{y} = [\theta \ v_l \ v_r]^T \quad \mathbf{u} = [V_l \ V_r]^T$$

$$\mathbf{A} = \begin{bmatrix} 0 & 0 & -vs & \frac{1}{2}c & \frac{1}{2}c \\ 0 & 0 & vc & \frac{1}{2}s & \frac{1}{2}s \\ 0 & 0 & 0 & -\frac{1}{2r_b} & \frac{1}{2r_b} \\ 0 & 0 & 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ 0 & 0 & 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 \end{bmatrix}$$

$$\mathbf{C} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix} \quad \mathbf{D} = \mathbf{0}_{3 \times 2} \quad (8.22)$$

where $v = \frac{v_l+v_r}{2}$, $c = \cos \theta$, $s = \sin \theta$, $C_1 = -\frac{G_l^2 K_t}{K_v R r^2}$, $C_2 = \frac{G_l K_t}{R r}$, $C_3 = -\frac{G_r^2 K_t}{K_v R r^2}$, and $C_4 = \frac{G_r K_t}{R r}$. The constants C_1 through C_4 are from the derivation in section 13.6.

The LQR gain should be recomputed around the current operating point regularly due to the high nonlinearity of this system. A less computationally expensive controller will be developed in later sections.

We can also use this in an extended Kalman filter as is since the measurement model ($\mathbf{y} = \mathbf{Cx} + \mathbf{Du}$) is linear.

With this controller, θ becomes uncontrollable at $v = 0$ due to the x and y dynamics being equivalent to a unicycle; it can't change its heading unless it's rolling (just like a bicycle). However, a differential drive *can* rotate in place. To address this controller's limitation at $v = 0$, one can temporarily switch to an LQR of just θ , v_l , and v_r ; linearize the controller around a slightly nonzero state; or plan a new trajectory after the previous one completes that provides nonzero wheel velocities to rotate the robot.

8.5.3 Improving model accuracy

Figures 8.5 and 8.6 demonstrate the tracking behavior of the linearized differential drive controller.

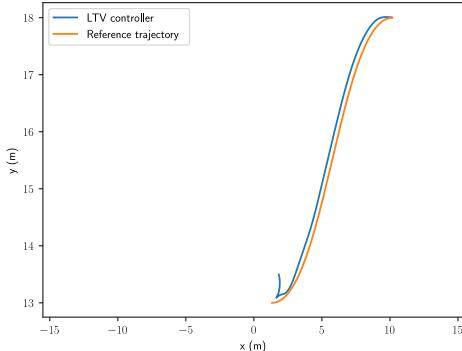


Figure 8.5: Linear time-varying differential drive controller x-y plot (first order)

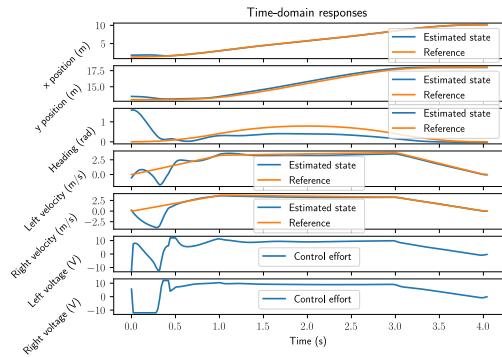


Figure 8.6: Linear time-varying differential drive controller response (first order)

The linearized differential drive model doesn't track well because the first-order linearization of \mathbf{A} doesn't capture the full heading dynamics, making the [model](#) update inaccurate. This linearization inaccuracy is evident in the Hessian matrix (second partial derivative with respect to the state vector) being nonzero.

$$\frac{\partial^2 f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}^2} = \begin{bmatrix} 0 & 0 & -\frac{v_l + v_r}{2} \cos \theta & 0 & 0 \\ 0 & 0 & -\frac{v_l + v_r}{2} \sin \theta & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

The second-order Taylor series expansion of the [model](#) around \mathbf{x}_0 would be

$$f(\mathbf{x}, \mathbf{u}_0) \approx f(\mathbf{x}_0, \mathbf{u}_0) + \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}}(\mathbf{x} - \mathbf{x}_0) + \frac{1}{2} \frac{\partial^2 f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}^2}(\mathbf{x} - \mathbf{x}_0)^2$$

To include higher-order dynamics in the linearized differential drive model integration, we recommend using the fourth-order Runge-Kutta (RK4) integration method on equation (8.21). See snippet 8.1 for an implementation of RK4.

```
def runge_kutta(f, x, u, dt):
    """Fourth order Runge-Kutta integration.

    Keyword arguments:
    f -- vector function to integrate
    x -- vector of states
    u -- vector of inputs (constant for dt)
    dt -- time for which to integrate
    """
    half_dt = dt * 0.5
    k1 = f(x, u)
    k2 = f(x + half_dt * k1, u)
    k3 = f(x + half_dt * k2, u)
    k4 = f(x + dt * k3, u)
    return x + dt / 6.0 * (k1 + 2.0 * k2 + 2.0 * k3 + k4)
```

Snippet 8.1. Fourth-order Runge-Kutta integration in Python

Figures 8.7 and 8.8 show a simulation using RK4 instead of the first-order model.

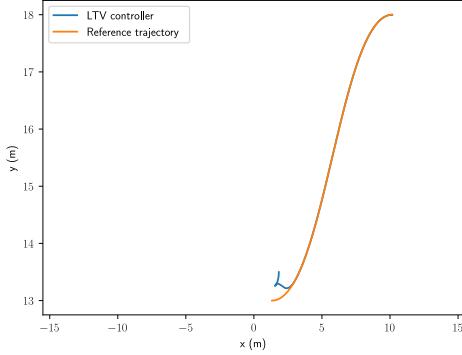


Figure 8.7: Linear time-varying differential drive controller (global reference frame formulation) x-y plot

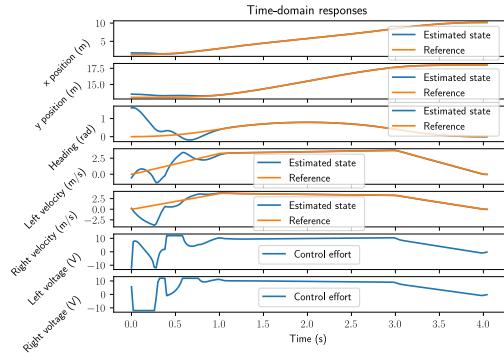


Figure 8.8: Linear time-varying differential drive controller (global reference frame formulation) response

8.5.4 Cross track error controller

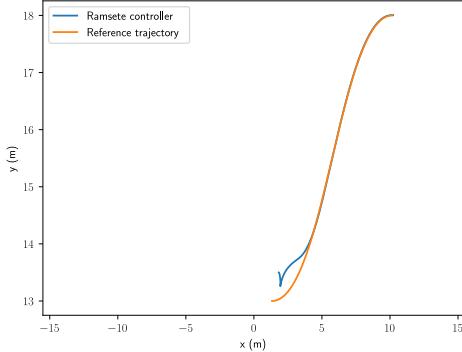


Figure 8.9: Ramsete nonlinear controller x-y plot

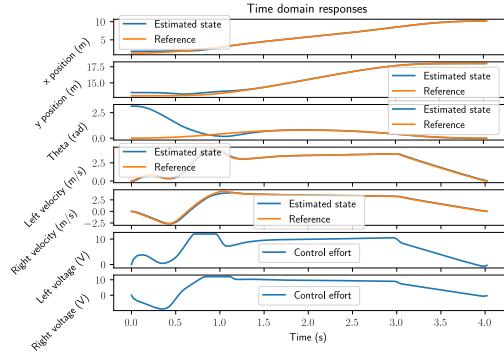


Figure 8.10: Ramsete nonlinear controller response

The tracking performance of the linearized differential drive controller (figures 8.7 and 8.8) and Ramsete (figures 8.9 and 8.10) for a given trajectory are similar, but the former's performance-effort trade-off can be tuned more intuitively via the Q and R gains. However, if the x and y components of the controller will fight each other, and it will take longer to converge to the path. This can be fixed by applying a counterclockwise rotation matrix to the global tracking error to transform it into the robot's coordinate frame.

$$R \begin{bmatrix} e_x \\ e_y \\ e_\theta \end{bmatrix} = \begin{bmatrix} \cos \theta & \sin \theta & 0 \\ -\sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}^G \begin{bmatrix} e_x \\ e_y \\ e_\theta \end{bmatrix}$$

where the superscript R represents the robot's coordinate frame and the superscript G represents the global coordinate frame.

With this transformation, the x and y error cost in LQR penalize the error ahead of the robot and cross-track error respectively instead of global pose error. Since the cross-track error is always measured from the robot's coordinate frame, the `model` used to compute the LQR should be linearized around $\theta = 0$ at all times.

$$\mathbf{A} = \begin{bmatrix} 0 & 0 & -\frac{v_l+v_r}{2} \sin 0 & \frac{1}{2} \cos 0 & \frac{1}{2} \cos 0 \\ 0 & 0 & \frac{v_l+v_r}{2} \cos 0 & \frac{1}{2} \sin 0 & \frac{1}{2} \sin 0 \\ 0 & 0 & 0 & -\frac{1}{2r_b} & \frac{1}{2r_b} \\ 0 & 0 & 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ 0 & 0 & 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} 0 & 0 & 0 & \frac{1}{2} & \frac{1}{2} \\ 0 & 0 & \frac{v_l+v_r}{2} & 0 & 0 \\ 0 & 0 & 0 & -\frac{1}{2r_b} & \frac{1}{2r_b} \\ 0 & 0 & 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ 0 & 0 & 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix}$$

This `model` results in figures 8.11 and 8.12, which show slightly better tracking performance than the previous formulation.

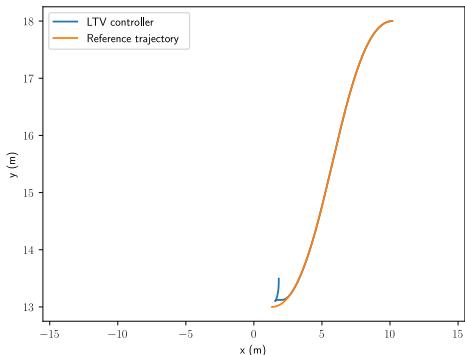


Figure 8.11: Linear time-varying differential drive controller x-y plot

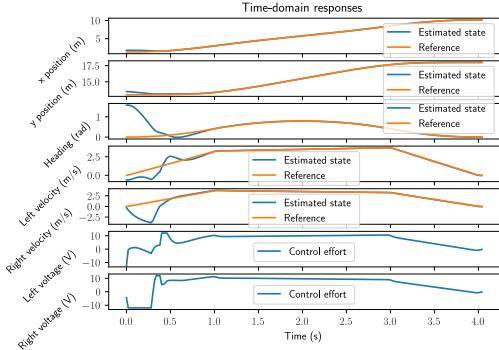


Figure 8.12: Linear time-varying differential drive controller response

8.5.5 Explicit time-varying control law

Another downside of this controller over Ramsete is that the user must generate controller gains for every state they visit in the state-space (an implicit control law) whereas Ramsete has two closed-form functions for its explicit control law.

A possible solution to this is fitting a vector function of the states to the linearized differential drive controller gains. Fortunately, the formulation of the controller dealing with cross-track error only requires linearization around different values of v rather than v and θ ; only a two-dimensional function is needed rather than a three-dimensional one. See figures 8.13 through 8.17 for plots of each controller gain for a range of velocities.

With the exception of the x gain plot, all functions are a variation of a square root. $v = 0$ and $v = 1$ were used to scale each function to produce a close approximation. The sign function⁹ is used for symmetry around the origin in the regression for y .

Theorem 8.5.3 — Linear time-varying differential drive controller.

$$\mathbf{x} = [x \ y \ \theta \ v_l \ v_r]^T \quad \mathbf{y} = [\theta \ v_l \ v_r]^T \quad \mathbf{u} = [V_l \ V_r]^T$$

The following \mathbf{A} and \mathbf{B} matrices of a continuous system are used to compute the LQR.

$$\mathbf{A} = \begin{bmatrix} 0 & 0 & 0 & \frac{1}{2} & \frac{1}{2} \\ 0 & 0 & v & 0 & 0 \\ 0 & 0 & 0 & -\frac{1}{2r_b} & \frac{1}{2r_b} \\ 0 & 0 & 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ 0 & 0 & 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 0 & 0 \\ \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 \end{bmatrix} \quad (8.23)$$

where $v = \frac{v_l+v_r}{2}$, $C_1 = -\frac{G_l^2 K_t}{K_v R r^2}$, $C_2 = \frac{G_l K_t}{R r}$, $C_3 = -\frac{G_r^2 K_t}{K_v R r^2}$, and $C_4 = \frac{G_r K_t}{R r}$. The constants C_1 through C_4 are from the derivation in section 13.6.

The locally optimal controller for this model is $\mathbf{u} = \mathbf{K}(v)(\mathbf{r} - \mathbf{x})$ where

$$\mathbf{K}(v) = \begin{bmatrix} k_x & k_{1,2}(v) \operatorname{sgn}(v) & k_{\theta,1} \sqrt{|v|} & k_{1,4}(v) & k_{2,4}(v) \\ k_x & -k_{1,2}(v) \operatorname{sgn}(v) & -k_{\theta,1} \sqrt{|v|} & k_{2,4}(v) & k_{1,4}(v) \end{bmatrix} \quad (8.24)$$

$$k_{1,2}(v) = k_{y,0} + (k_{y,1} - k_{y,0}) \sqrt{|v|} \quad (8.25)$$

$$k_{1,4}(v) = k_{v^+,0} + (k_{v^+,1} - k_{v^+,0}) \sqrt{|v|} \quad (8.26)$$

$$k_{2,4}(v) = k_{v^-,0} - (k_{v^+,1} - k_{v^+,0}) \sqrt{|v|} \quad (8.27)$$

Using \mathbf{K} computed via LQR at $v = 0$

$$k_x = \mathbf{K}_{1,1} \quad k_{y,0} = \mathbf{K}_{1,2} \quad k_{v^+,0} = \mathbf{K}_{1,4} \quad k_{v^-,0} = \mathbf{K}_{2,4}$$

Using \mathbf{K} computed via LQR at $v = 1$

$$k_{y,1} = \mathbf{K}_{1,2} \quad k_{\theta,1} = \mathbf{K}_{1,3} \quad k_{v^+,1} = \mathbf{K}_{1,4}$$

Figures 8.18 through 8.21 show the responses of the exact and approximate solutions are the same.

⁹The sign function is defined as follows:

$$\operatorname{sgn}(x) = \begin{cases} -1 & x < 0 \\ 0 & x = 0 \\ 1 & x > 0 \end{cases}$$

8.5.6 Nonlinear observer design

Encoder position augmentation

Estimation of the global pose can be significantly improved if encoder position measurements are used instead of velocity measurements. By augmenting the plant with the line integral of each wheel's velocity over time, we can provide a mapping from model states to position measurements. We can augment the linear subspace of the model as follows.

Augment the matrix equation with position states x_l and x_r , which have the model equations $\dot{x}_l = v_l$ and $\dot{x}_r = v_r$. The matrix elements corresponding to v_l in the first equation and v_r in the second equation will be 1, and the others will be 0 since they don't appear, so $\dot{x}_l = 1v_l + 0v_r + 0x_l + 0x_r + 0V_l + 0V_r$ and $\dot{x}_r = 0v_l + 1v_r + 0x_l + 0x_r + 0V_l + 0V_r$. The existing rows will have zeroes inserted where x_l and x_r are multiplied in.

$$\begin{bmatrix} \dot{x}_l \\ \dot{x}_r \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} v_l \\ v_r \end{bmatrix} + \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} V_l \\ V_r \end{bmatrix}$$

This produces the following linear subspace over $\mathbf{x} = [v_l \ v_r \ x_l \ x_r]^T$.

$$\mathbf{A} = \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 & 0 & 0 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \quad (8.28)$$

The measurement model for the complete nonlinear model is now $\mathbf{y} = [\theta \ x_l \ x_r]^T$ instead of $\mathbf{y} = [\theta \ v_l \ v_r]^T$.

U error estimation

As per subsection 5.11.2, we will now augment the [model](#) so u_{error} states are added to the [control inputs](#).

The [plant](#) and [observer](#) augmentations should be performed before the [model](#) is [discretized](#). After the [controller](#) gain is computed with the unaugmented discrete [model](#), the controller may be augmented. Therefore, the [plant](#) and [observer](#) augmentations assume a continuous [model](#) and the [controller](#) augmentation assumes a discrete [controller](#).

The three u_{error} states we'll be adding are $u_{error,l}$, $u_{error,r}$, and $u_{error,heading}$ for left voltage error, right voltage error, and heading error respectively. The left and right wheel positions are filtered encoder positions and are not adjusted for heading error. The turning angle computed from the left and right wheel positions is adjusted by the gyroscope heading. The heading u_{error} state is the heading error between what the wheel positions imply and the gyroscope measurement.

The full state is thus $\mathbf{x} = [x \ y \ \theta \ v_l \ v_r \ x_l \ x_r \ u_{error,l} \ u_{error,r} \ u_{error,heading}]^T$.

The complete nonlinear model is as follows. Let $v = \frac{v_l+v_r}{2}$. The three u_{error} states augment the linear subspace, so the nonlinear pose dynamics are the same.

$$\begin{bmatrix} \dot{x} \\ \dot{y} \\ \dot{\theta} \end{bmatrix} = \begin{bmatrix} v \cos \theta \\ v \sin \theta \\ \frac{v_r}{2r_b} - \frac{v_l}{2r_b} \end{bmatrix} \quad (8.29)$$

The left and right voltage error states should be mapped to the corresponding velocity states, so the system matrix should be augmented with \mathbf{B} .

The heading u_{error} is measuring counterclockwise encoder understeer relative to the gyroscope heading, so it should add to the left position and subtract from the right position for clockwise correction of encoder positions. That corresponds to the following input mapping vector.

$$\mathbf{B}_\theta = \begin{bmatrix} 0 \\ 0 \\ 1 \\ -1 \end{bmatrix}$$

Now we'll augment the linear system matrix horizontally to accomodate the u_{error} states.

$$\begin{bmatrix} \dot{v}_l \\ \dot{v}_r \\ \dot{x}_l \\ \dot{x}_r \end{bmatrix} = [\mathbf{A} \quad \mathbf{B} \quad \mathbf{B}_\theta] \begin{bmatrix} v_l \\ v_r \\ x_l \\ x_r \\ u_{error,l} \\ u_{error,r} \\ u_{error,heading} \end{bmatrix} + \mathbf{B}\mathbf{u} \quad (8.30)$$

\mathbf{A} and \mathbf{B} are the linear subspace from equation (8.28).

The u_{error} states have no dynamics. The observer selects them to minimize the difference between the expected and actual measurements.

$$\begin{bmatrix} \dot{u}_{error,l} \\ \dot{u}_{error,r} \\ \dot{u}_{error,heading} \end{bmatrix} = \mathbf{0}_{3 \times 1} \quad (8.31)$$

The controller is augmented as follows.

$$\mathbf{K}_{error} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix} \quad \mathbf{K}_{aug} = [\mathbf{K} \quad \mathbf{K}_{error}] \quad \mathbf{r}_{aug} = \begin{bmatrix} \mathbf{r} \\ 0 \\ 0 \\ 0 \end{bmatrix} \quad (8.32)$$

This controller augmentation compensates for unmodeled dynamics like:

1. Understeer caused by wheel friction inherent in skid-steer robots
2. Battery voltage drop under load, which reduces the available control authority

 The process noise for the voltage error states should be how much the voltage can be expected to drop. The heading error state should be the encoder [model](#) uncertainty.

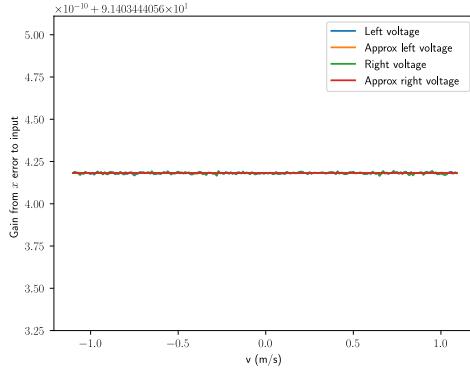


Figure 8.13: Linear time-varying differential drive controller LQR gain regression fit (x)

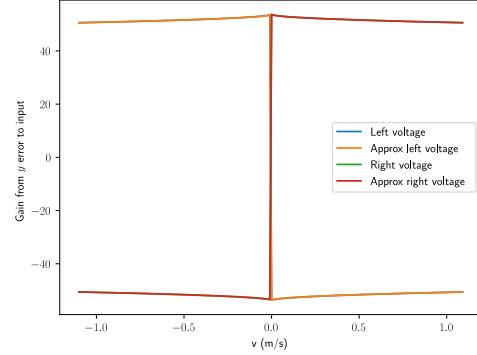


Figure 8.14: Linear time-varying differential drive controller LQR gain fit regression fit (y)

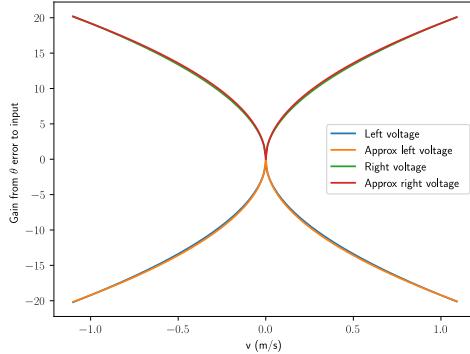


Figure 8.15: Linear time-varying differential drive controller LQR gain regression fit (θ)

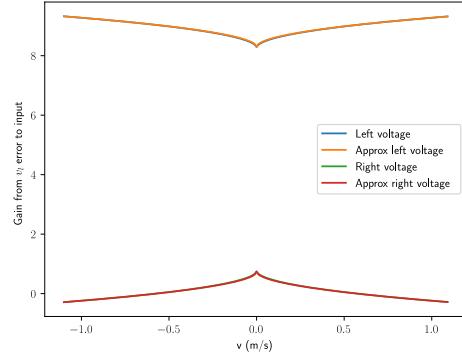


Figure 8.16: Linear time-varying differential drive controller LQR gain regression fit (v_l)

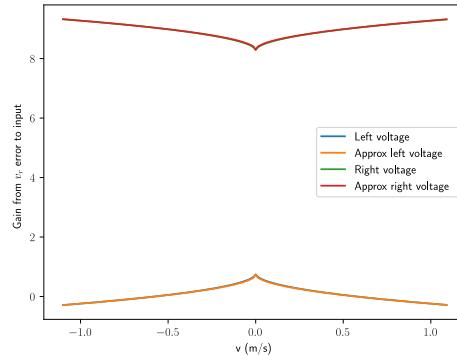


Figure 8.17: Linear time-varying differential drive controller LQR gain regression fit (v_r)

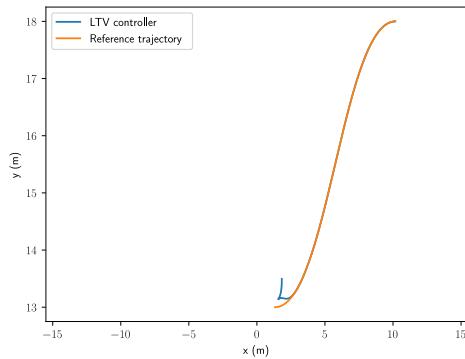


Figure 8.18: Linear time-varying differential drive controller x-y plot (approximate)

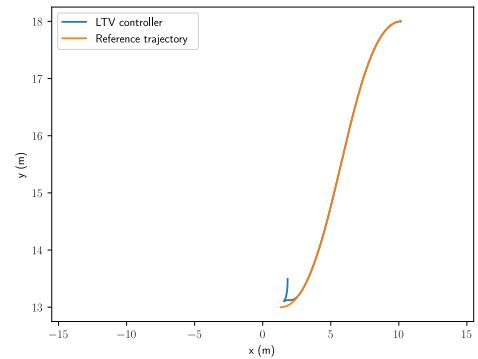


Figure 8.19: Linear time-varying differential drive controller x-y plot (exact)

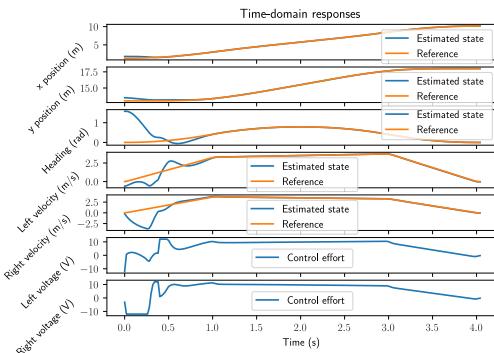


Figure 8.20: Linear time-varying differential drive controller response (approximate)

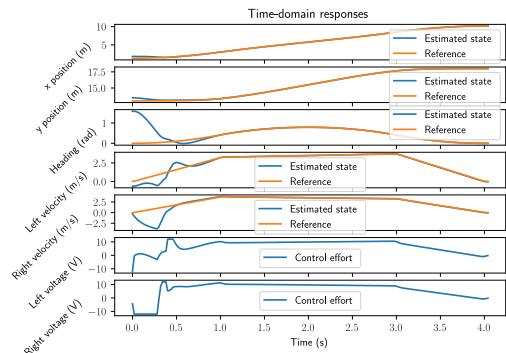


Figure 8.21: Linear time-varying differential drive controller response (exact)

8.6 Ramsete unicycle controller

Ramsete is a nonlinear time-varying feedback controller for unicycle [models](#) that drives the [model](#) to a desired [pose](#) along a two-dimensional trajectory. Why would we need a nonlinear control law in addition to the linear ones we have used so far? If we use the original approach with an LQR [controller](#) for left and right position and velocity [states](#), the [controller](#) only deals with the local [pose](#). If the robot deviates from the path, there is no way for the [controller](#) to correct and the robot may not reach the desired global [pose](#). This is due to multiple endpoints existing for the robot which have the same encoder path arc lengths.

Instead of using wheel path arc lengths (which are in the robot's local coordinate frame), nonlinear controllers like pure pursuit and Ramsete use global pose. The [controller](#) uses this extra information to guide a linear [reference tracker](#) like an LQR [controller](#) back in by adjusting the [references](#) of the LQR [controller](#).

The paper *Control of Wheeled Mobile Robots: An Experimental Overview* describes a nonlinear controller for a wheeled vehicle with unicycle-like kinematics; a global [pose](#) consisting of x , y , and θ ; and a desired [pose](#) consisting of x_d , y_d , and θ_d [24]. We'll call it Ramsete because that's the acronym for the title of the book it came from in Italian ("Robotica Articolata e Mobile per i SERVizi e le TEcnologie").

8.6.1 Velocity and turning rate command derivation

The [state](#) tracking [error](#) e in the vehicle's coordinate frame is defined as

$$\begin{bmatrix} e_x \\ e_y \\ e_\theta \end{bmatrix} = \begin{bmatrix} \cos \theta & \sin \theta & 0 \\ -\sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_d - x \\ y_d - y \\ \theta_d - \theta \end{bmatrix}$$

where e_x is the tracking [error](#) in x , e_y is the tracking [error](#) in y , and e_θ is the tracking [error](#) in θ . The 3×3 matrix is a rotation matrix that transforms the [error](#) in the [pose](#) (represented by $x_d - x$, $y_d - y$, and $\theta_d - \theta$) from the global coordinate frame into the vehicle's coordinate frame.

We will use the following control laws u_1 and u_2 for velocity and turning rate respectively.

$$\begin{aligned} u_1 &= -k_1 e_x \\ u_2 &= -k_3 e_\theta - k_2 v_d \operatorname{sinc}(e_\theta) e_y \end{aligned} \tag{8.33}$$

where k_1 , k_2 , and k_3 are time-varying gains and $\operatorname{sinc}(e_\theta)$ is defined as $\frac{\sin e_\theta}{e_\theta}$. This choice of control law may seem arbitrary, and that's because it is. The only requirement on our choice is that there exist an associated Lyapunov candidate function that proves the control law is globally asymptotically stable. We'll provide a sketch of a proof in theorem 8.6.1.

Our velocity and turning rate commands for the vehicle will use the following nonlinear transformation of the inputs.

$$v = v_d \cos e_\theta - u_1$$

$$\omega = \omega_d - u_2$$

Substituting the control laws u_1 and u_2 into these equations gives

$$\begin{aligned} v &= v_d \cos e_\theta + k_1 e_x \\ \omega &= k_3 e_\theta + \omega_d + k_2 v_d \operatorname{sinc}(e_\theta) e_y \end{aligned}$$

Theorem 8.6.1 Assuming that v_d and ω_d are bounded with bounded derivatives, and that $v_d(t) \rightarrow 0$ or $\omega_d(t) \rightarrow 0$ when $t \rightarrow \infty$, the control laws in equation (8.33) globally asymptotically stabilize the origin $\mathbf{e} = 0$.

Proof:

To prove convergence, the paper previously mentioned uses the following Lyapunov function.

$$V = \frac{k_2}{2}(e_x^2 + e_y^2) + \frac{e_\theta^2}{2}$$

where k_2 is a tuning constant, e_x is the tracking error in x , e_y is the tracking error in y , and e_θ is the tracking error in θ .

The time derivative along the solutions of the closed-loop system is nonincreasing since

$$\dot{V} = -k_1 k_2 e_x^2 - k_3 e_\theta^2 \leq 0$$

Thus, $\|e(t)\|$ is bounded, $\dot{V}(t)$ is uniformly continuous, and $V(t)$ tends to some limit value. Using the Barbalat lemma, $\dot{V}(t)$ tends to zero [24].

8.6.2 Nonlinear controller equations

Let $k_2 = b$ and $k = k_1(v_d, \omega_d) = k_3(v_d, \omega_d) = 2\zeta\sqrt{\omega_d^2 + bv_d^2}$.

Theorem 8.6.2 — Ramsete unicycle controller.

$$\begin{bmatrix} e_x \\ e_y \\ e_\theta \end{bmatrix} = \begin{bmatrix} \cos \theta & \sin \theta & 0 \\ -\sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_d - x \\ y_d - y \\ \theta_d - \theta \end{bmatrix} \quad (8.34)$$

$$v = v_d \cos e_\theta + k e_x \quad (8.35)$$

$$\omega = \omega_d + k e_\theta + b v_d \operatorname{sinc}(e_\theta) e_y \quad (8.36)$$

$$k = 2\zeta\sqrt{\omega_d^2 + bv_d^2} \quad (8.37)$$

$$\operatorname{sinc}(e_\theta) = \frac{\sin e_\theta}{e_\theta} \quad (8.38)$$

v	velocity command	v_d	desired velocity
ω	turning rate command	ω_d	desired turning rate
x	actual x position in global coordinate frame	x_d	desired x position
y	actual y position in global coordinate frame	y_d	desired y position
θ	actual angle in global coordinate frame	θ_d	desired angle

b and ζ are tuning parameters where $b > 0$ and $\zeta \in (0, 1)$. Larger values of b make convergence more aggressive (like a proportional term), and larger values of ζ provide more damping.

v and ω should be the references for a reference tracker for the drivetrain. This can be a linear controller (see subsection 8.6.3). x , y , and θ are obtained via a pose estimator (see chapter 10 for

how to implement one). The desired velocity, turning rate, and pose can be varied over time according to a desired trajectory.

With this controller, θ becomes uncontrollable at $v = 0$. This is fine for controlling unicycles because that's already an inherent structural limitation; unicycles can't change their heading unless they are rolling (just like bicycles). However, differential drives *can* rotate in place. To address this controller's limitation at $v = 0$, one can temporarily switch to an LQR of just θ , v_l , and v_r ; or plan a new trajectory that provides nonzero desired turning rates to rotate the robot.

8.6.3 Linear reference tracker

We need a velocity reference tracker for the nonlinear controller's commands. Starting from equation (8.20), we'll derive two models for this purpose and compare their responses with a straight profile and as part of a nonlinear trajectory follower.

Left/right velocity reference tracker

$$\mathbf{x} = \begin{bmatrix} x_l \\ v_l \\ x_r \\ v_r \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} x_l \\ x_r \end{bmatrix} \quad \mathbf{u} = \begin{bmatrix} V_l \\ V_r \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ 0 & 0 & 0 & 1 \\ 0 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & 0 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 0 & 0 \\ \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 \\ 0 & 0 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 \end{bmatrix}$$

$$\mathbf{C} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \quad \mathbf{D} = \mathbf{0}_{2 \times 2}$$

where $C_1 = -\frac{G_l^2 K_t}{K_v R r^2}$, $C_2 = \frac{G_l K_t}{R r}$, $C_3 = -\frac{G_r^2 K_t}{K_v R r^2}$, and $C_4 = \frac{G_r K_t}{R r}$.

To obtain a left/right velocity reference tracker, we can just remove the position states from the model.

Theorem 8.6.3 — Left/right velocity reference tracker.

$$\mathbf{x} = \begin{bmatrix} v_l \\ v_r \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} v_l \\ v_r \end{bmatrix} \quad \mathbf{u} = \begin{bmatrix} V_l \\ V_r \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_3 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_1 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_3 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_4 \\ \left(\frac{1}{m} - \frac{r_b^2}{J}\right) C_2 & \left(\frac{1}{m} + \frac{r_b^2}{J}\right) C_4 \end{bmatrix} \quad (8.39)$$

$$\mathbf{C} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad \mathbf{D} = \mathbf{0}_{2 \times 2}$$

where $C_1 = -\frac{G_l^2 K_t}{K_v R r^2}$, $C_2 = \frac{G_l K_t}{R r}$, $C_3 = -\frac{G_r^2 K_t}{K_v R r^2}$, and $C_4 = \frac{G_r K_t}{R r}$.

See https://github.com/calcmogul/controls-engineering-in-frc/blob/master/modern-control-theory/ss-applications/ramsete_decoupled.py for an implementation.

Linear/angular velocity reference tracker

Since the Ramsete controller produces velocity and turning rate commands, it would be more convenient if the [states](#) are velocity and turning rate instead of left and right wheel velocity. We'll create a second model that has velocity and angular velocity states and see how well it performs.

$$\begin{aligned}\dot{v}_l &= \left(\frac{1}{m} + \frac{r_b^2}{J} \right) (C_1 v_l + C_2 V_l) + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) (C_3 v_r + C_4 V_r) \\ \dot{v}_r &= \left(\frac{1}{m} - \frac{r_b^2}{J} \right) (C_1 v_l + C_2 V_l) + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) (C_3 v_r + C_4 V_r)\end{aligned}$$

Substitute in equations (12.1) and (12.2).

$$\dot{v}_c + \dot{\omega} r_b = \left(\frac{1}{m} + \frac{r_b^2}{J} \right) (C_1(v_c - \omega r_b) + C_2 V_l) + \left(\frac{1}{m} - \frac{r_b^2}{J} \right) (C_3(v_c + \omega r_b) + C_4 V_r) \quad (8.40)$$

$$\dot{v}_c - \dot{\omega} r_b = \left(\frac{1}{m} - \frac{r_b^2}{J} \right) (C_1(v_c - \omega r_b) + C_2 V_l) + \left(\frac{1}{m} + \frac{r_b^2}{J} \right) (C_3(v_c + \omega r_b) + C_4 V_r) \quad (8.41)$$

Now, we need to solve for \dot{v}_c and $\dot{\omega}$. First, we'll add equation (8.40) to equation (8.41).

$$\begin{aligned}2\dot{v}_c &= \frac{2}{m} (C_1(v_c - \omega r_b) + C_2 V_l) + \frac{2}{m} (C_3(v_c + \omega r_b) + C_4 V_r) \\ \dot{v}_c &= \frac{1}{m} (C_1(v_c - \omega r_b) + C_2 V_l) + \frac{1}{m} (C_3(v_c + \omega r_b) + C_4 V_r) \\ \dot{v}_c &= \frac{1}{m} (C_1 + C_3) v_c + \frac{1}{m} (-C_1 + C_3) \omega r_b + \frac{1}{m} C_2 V_l + \frac{1}{m} C_4 V_r \\ \dot{v}_c &= \frac{1}{m} (C_1 + C_3) v_c + \frac{r_b}{m} (-C_1 + C_3) \omega + \frac{1}{m} C_2 V_l + \frac{1}{m} C_4 V_r\end{aligned}$$

Next, we'll subtract equation (8.41) from equation (8.40).

$$\begin{aligned}2\dot{\omega} r_b &= \frac{2r_b^2}{J} (C_1(v_c - \omega r_b) + C_2 V_l) - \frac{2r_b^2}{J} (C_3(v_c + \omega r_b) + C_4 V_r) \\ \dot{\omega} &= \frac{r_b}{J} (C_1(v_c - \omega r_b) + C_2 V_l) - \frac{r_b}{J} (C_3(v_c + \omega r_b) + C_4 V_r) \\ \dot{\omega} &= \frac{r_b}{J} (C_1 - C_3) v_c + \frac{r_b}{J} (-C_1 - C_3) \omega r_b + \frac{r_b}{J} C_2 V_l - \frac{r_b}{J} C_4 V_r \\ \dot{\omega} &= \frac{r_b}{J} (C_1 - C_3) v_c + \frac{r_b^2}{J} (-C_1 - C_3) \omega + \frac{r_b}{J} C_2 V_l - \frac{r_b}{J} C_4 V_r\end{aligned}$$

Now, just convert the two equations to state-space notation.

Theorem 8.6.4 — Linear/angular velocity reference tracker.

$$\mathbf{x} = \begin{bmatrix} v \\ \omega \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} v_l \\ v_r \end{bmatrix} \quad \mathbf{u} = \begin{bmatrix} V_l \\ V_r \end{bmatrix}$$

$$\begin{aligned}\mathbf{A} &= \begin{bmatrix} \frac{1}{m}(C_1 + C_3) & \frac{r_b}{m}(-C_1 + C_3) \\ \frac{r_b}{J}(C_1 - C_3) & \frac{r_b^2}{J}(-C_1 - C_3) \end{bmatrix} & \mathbf{B} &= \begin{bmatrix} \frac{1}{m}C_2 & \frac{1}{m}C_4 \\ \frac{r_b}{J}C_2 & -\frac{r_b}{J}C_4 \end{bmatrix} \\ \mathbf{C} &= \begin{bmatrix} 1 & -r_b \\ 1 & r_b \end{bmatrix} & \mathbf{D} &= \mathbf{0}_{2 \times 2}\end{aligned} \quad (8.42)$$

where $C_1 = -\frac{G_l^2 K_t}{K_v R r^2}$, $C_2 = \frac{G_l K_t}{R r}$, $C_3 = -\frac{G_r^2 K_t}{K_v R r^2}$, and $C_4 = \frac{G_r K_t}{R r}$.

See https://github.com/calcogul/controls-engineering-in-frc/blob/master/modern-control-theory/ss-applications/ramsete_coupled.py for an implementation.

Reference tracker comparison

Figures 8.22 and 8.23 shows how well each [reference](#) tracker performs.

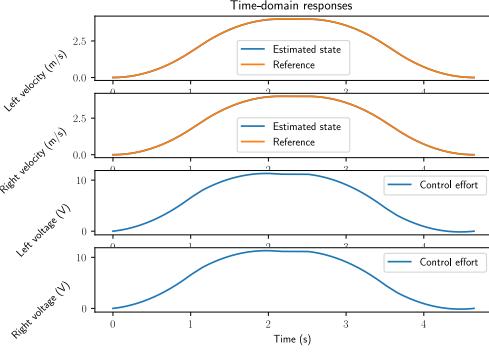


Figure 8.22: Left/right velocity reference tracker response to a motion profile

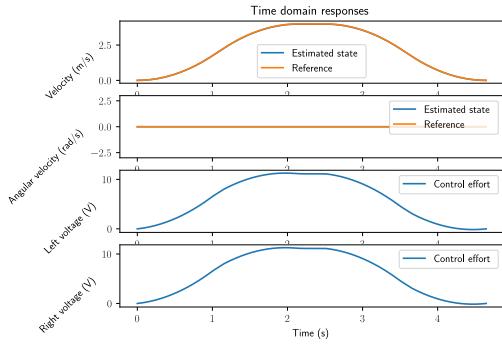


Figure 8.23: Linear/angular velocity reference tracker response to a motion profile

For a simple s-curve motion profile, they behave similarly.

Figure 8.24 demonstrates the Ramsete controller with the left/right velocity [reference](#) tracker for a typical FRC drivetrain with the [reference](#) tracking behavior shown in figure 8.25 and figure 8.26 demonstrates the Ramsete controller with the velocity / angular velocity [reference](#) tracker for a typical FRC drivetrain with the [reference](#) tracking behavior shown in figure 8.27.

The Ramsete [controller](#) behaves relatively the same in each case, but the left/right velocity [reference](#) tracker tracks the [references](#) produced by the Ramsete [controller](#) better and with smaller [control efforts](#) overall. This is likely due to the second [controller](#) having coupled dynamics. That is, the control inputs don't act independently on the system [states](#). In the coupled [controller](#), v and ω require opposing control actions to reach their respective [references](#), which results in them fighting each other. This hypothesis is supported by the condition number of the coupled [controller](#)'s controllability matrix being larger (2.692 for coupled and 1.314 for decoupled).

Therefore, theorem 8.6.3 is a superior [reference](#) tracker to theorem 8.6.4 due to its decoupled dynamics, even though conversions are required between it and the Ramsete [controller](#)'s commands.

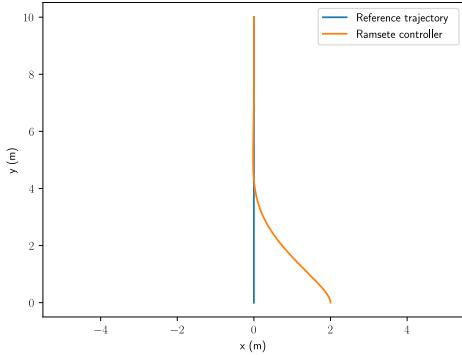


Figure 8.24: Ramsete controller response with left/right velocity reference tracker ($b = 2$, $\zeta = 0.7$)

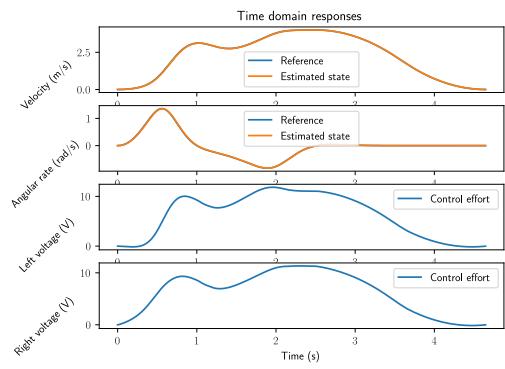


Figure 8.25: Ramsete controller's left/right velocity reference tracker response

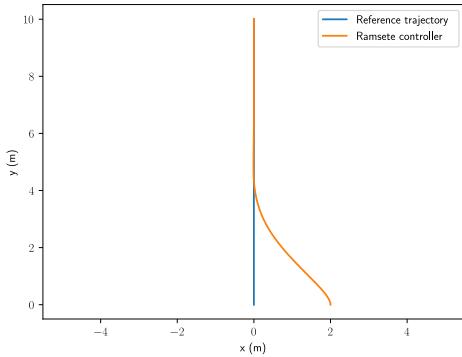


Figure 8.26: Ramsete controller response with velocity / angular velocity reference tracker ($b = 2$, $\zeta = 0.7$)

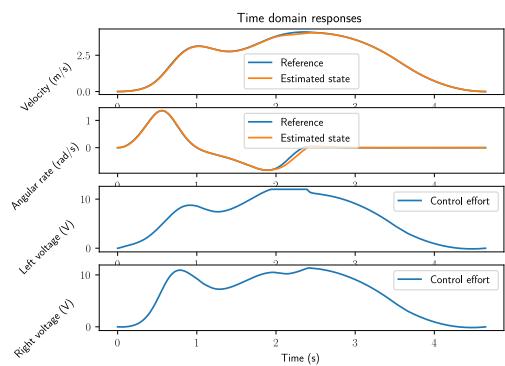


Figure 8.27: Ramsete controller's velocity / angular velocity reference tracker response

8.7 Linear time-varying unicycle controller (cascaded)

One can also create a linear time-varying controller with a cascaded control architecture like Ramsete. This section will derive a locally optimal replacement for Ramsete.

The change in global pose for a unicycle is defined by the following three equations.

$$\begin{aligned}\dot{x} &= v \cos \theta \\ \dot{y} &= v \sin \theta \\ \dot{\theta} &= \omega\end{aligned}$$

Here's the model as a vector function where $\mathbf{x} = [x \ y \ \theta]^T$ and $\mathbf{u} = [v \ \omega]^T$.

$$f(\mathbf{x}, \mathbf{u}) = \begin{bmatrix} v \cos \theta \\ v \sin \theta \\ \omega \end{bmatrix} \quad (8.43)$$

To create an LQR, we need to linearize this.

$$\frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} = \begin{bmatrix} 0 & 0 & -v \sin \theta \\ 0 & 0 & v \cos \theta \\ 0 & 0 & 0 \end{bmatrix}$$

$$\frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}} = \begin{bmatrix} \cos \theta & 0 \\ \sin \theta & 0 \\ 0 & 1 \end{bmatrix}$$

Therefore,

Theorem 8.7.1 — Linear time-varying unicycle state-space model.

$$\begin{aligned} \dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du} \\ \mathbf{x} &= \begin{bmatrix} x \\ y \\ \theta \end{bmatrix} \quad \mathbf{y} = [\theta] \quad \mathbf{u} = \begin{bmatrix} v \\ \omega \end{bmatrix} \\ \mathbf{A} &= \begin{bmatrix} 0 & 0 & -v \sin \theta \\ 0 & 0 & v \cos \theta \\ 0 & 0 & 0 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} \cos \theta & 0 \\ \sin \theta & 0 \\ 0 & 1 \end{bmatrix} \\ \mathbf{C} &= [0 \ 0 \ 1] \quad \mathbf{D} = \mathbf{0}_{1 \times 2} \end{aligned} \tag{8.44}$$

The LQR gain should be recomputed around the current operating point regularly due to the high nonlinearity of this system. A less computationally expensive controller will be developed next.

8.7.1 Explicit time-varying control law

As with the LTV differential drive, we will fit a vector function of the states to the LTV unicycle controller gains. Fortunately, the formulation of the controller dealing with cross-track error only requires linearization around different values of v rather than v and θ ; only a two-dimensional function is needed rather than a three-dimensional one. See figures 8.28 through 8.30 for plots of each controller gain for a range of velocities.

With the exception of the x gain plot, all functions are a variation of a square root. $v = 0$ and $v = 1$ were used to scale each function to produce a close approximation. The sign function¹⁰ is used for symmetry around the origin in the regression for y .

¹⁰The sign function is defined as follows:

$$\text{sgn}(x) = \begin{cases} -1 & x < 0 \\ 0 & x = 0 \\ 1 & x > 0 \end{cases}$$

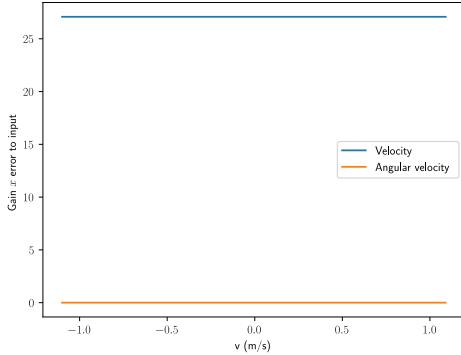


Figure 8.28: Linear time-varying unicycle controller cascaded LQR gain regression (x)

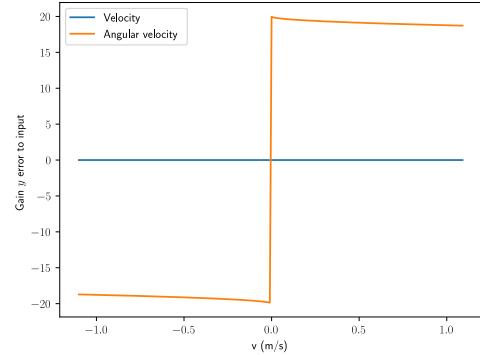


Figure 8.29: Linear time-varying unicycle controller cascaded LQR gain regression (y)

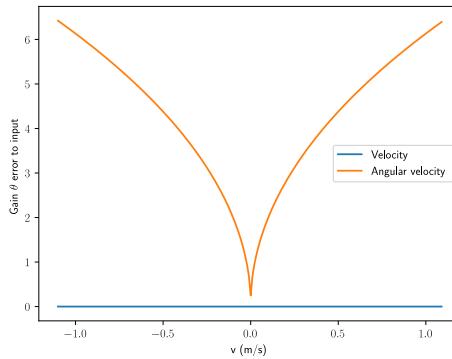


Figure 8.30: Linear time-varying unicycle controller cascaded LQR gain regression (θ)

Theorem 8.7.2 — Linear time-varying unicycle controller. The following \mathbf{A} and \mathbf{B} matrices of a continuous system are used to compute the LQR.

$$\mathbf{A} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & v \\ 0 & 0 & 0 \end{bmatrix} \quad \mathbf{B} = \begin{bmatrix} 1 & 0 \\ 0 & 0 \\ 0 & 1 \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} x \\ y \\ \theta \end{bmatrix} \quad \mathbf{u} = \begin{bmatrix} v \\ \omega \end{bmatrix} \quad (8.45)$$

The locally optimal controller for this model is $\mathbf{u} = \mathbf{K}(v)(\mathbf{r} - \mathbf{x})$ where

$$\mathbf{K}(v) = \begin{bmatrix} k_x & 0 & 0 \\ 0 & k_y(v) \operatorname{sgn}(v) & k_\theta \sqrt{|v|} \end{bmatrix} \quad (8.46)$$

$$k_y(v) = k_{y,0} + (k_{y,1} - k_{y,0}) \sqrt{|v|} \quad (8.47)$$

Using \mathbf{K} computed via LQR at $v = 0$

$$k_x = \mathbf{K}_{1,1} \quad k_{y,0} = \mathbf{K}_{2,2}$$

Using \mathbf{K} computed via LQR at $v = 1$

$$k_{y,1} = \mathbf{K}_{2,2} \quad k_\theta = \mathbf{K}_{2,3}$$

This page intentionally left blank

Estimation and localization

9	Stochastic control theory	93
9.1	Terminology	
9.2	State observers	
9.3	Introduction to probability	
9.4	Linear stochastic systems	
9.5	Two-sensor problem	
9.6	Kalman filter	
9.7	Kalman smoother	
9.8	Extended Kalman filter	
9.9	Unscented Kalman filter	
9.10	Multiple model adaptive estimation	
10	Pose estimation	123
10.1	Euler integration	
10.2	Pose exponential	
10.3	Pose correction	

This page intentionally left blank



9. Stochastic control theory

Stochastic control theory is a subfield of control theory that deals with the existence of uncertainty either in observations or in noise that drives the evolution of a [system](#). We assign probability distributions to this random noise and aim to achieve a desired control task despite the presence of this uncertainty.

Stochastic optimal control is concerned with doing this with minimum cost defined by some cost functional, like we did with LQR earlier. First, we'll cover the basics of probability and how we represent linear stochastic [systems](#) in state-space representation. Then, we'll derive an optimal estimator using this knowledge, the Kalman filter, and demonstrate creative applications of the Kalman filter theory.

9.1 Terminology

First, we should provide definitions for terms that have specific meanings in this field.

A causal system is one that uses only past information. A noncausal system also uses information from the future. A filter is a causal system that *filters* information through a probabilistic model to produce an estimate of a desired quantity that can't be measured directly. A smoother is a noncausal system, so it uses information from before and after the current state to produce a better estimate.

9.2 State observers

State [observers](#) are used to estimate [states](#) which cannot be measured directly. This can be due to noisy measurements or the [state](#) not being measurable (a hidden [state](#)). This information can be used for [localization](#), which is the process of using external measurements to determine an [agent](#)'s pose¹, or orientation in the world.

One type of [state](#) estimator is LQE. “LQE” stands for “Linear-Quadratic Estimator”. Similar to LQR, it places the estimator poles such that it minimizes the sum of squares of the estimation [error](#).

¹An agent is a system-agnostic term for independent controlled actors like robots or aircraft.

The Luenberger [observer](#) and Kalman filter are examples of these, where the latter chooses the pole locations optimally based on the [model](#) and measurement uncertainties.

Computer vision can also be used for [localization](#). By extracting features from an image taken by the [agent](#)'s camera, like a retroreflective target in FRC, and comparing them to known dimensions, one can determine where the [agent](#)'s camera would have to be to see that image. This can be used to correct our [state](#) estimate in the same way we do with an encoder or gyroscope.

9.2.1 Luenberger observer

We'll introduce the Luenberger observer first to demonstrate the general form of a state estimator and some of their properties.

Theorem 9.2.1 — Luenberger observer.

$$\dot{\hat{\mathbf{x}}} = \mathbf{A}\hat{\mathbf{x}} + \mathbf{B}\mathbf{u} + \mathbf{L}(\mathbf{y} - \hat{\mathbf{y}}) \quad (9.1)$$

$$\hat{\mathbf{y}} = \mathbf{C}\hat{\mathbf{x}} + \mathbf{D}\mathbf{u} \quad (9.2)$$

$$\hat{\mathbf{x}}_{k+1} = \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{y}_k - \hat{\mathbf{y}}_k) \quad (9.3)$$

$$\hat{\mathbf{y}}_k = \mathbf{C}\hat{\mathbf{x}}_k + \mathbf{D}\mathbf{u}_k \quad (9.4)$$

A	system matrix	$\hat{\mathbf{x}}$	state estimate vector
B	input matrix	\mathbf{u}	input vector
C	output matrix	\mathbf{y}	output vector
D	feedthrough matrix	$\hat{\mathbf{y}}$	output estimate vector
L	estimator gain matrix		

Matrix	Rows × Columns	Matrix	Rows × Columns
A	states × states	$\hat{\mathbf{x}}$	states × 1
B	states × inputs	\mathbf{u}	inputs × 1
C	outputs × states	\mathbf{y}	outputs × 1
D	outputs × inputs	$\hat{\mathbf{y}}$	outputs × 1
L	states × outputs		

Table 9.1: Luenberger observer matrix dimensions

Variables denoted with a hat are estimates of the corresponding variable. For example, $\hat{\mathbf{x}}$ is the estimate of the true [state](#) \mathbf{x} .

Notice that a Luenberger [observer](#) has an extra term in the [state](#) evolution equation. This term uses the difference between the estimated [outputs](#) and measured [outputs](#) to steer the estimated [state](#) toward the true [state](#). Large values of \mathbf{L} trust the measurements more while small values trust the [model](#) more.

- R** Using an estimator forfeits the performance guarantees from earlier, but the responses are still generally very good if the process and measurement noises are small enough. See John Doyle's paper *Guaranteed Margins for LQG Regulators* for a proof.

A Luenberger [observer](#) combines the prediction and update steps of an estimator. To run them separately, use the equations in theorem 9.2.2 instead.

Theorem 9.2.2 — Luenberger observer with separate predict/update.

Predict step

$$\hat{\mathbf{x}}_{k+1}^- = \mathbf{A}\hat{\mathbf{x}}_k^- + \mathbf{B}\mathbf{u}_k \quad (9.5)$$

Update step

$$\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{A}^{-1}\mathbf{L}(\mathbf{y}_k - \hat{\mathbf{y}}_k) \quad (9.6)$$

$$\hat{\mathbf{y}}_k = \mathbf{C}\hat{\mathbf{x}}_k^- \quad (9.7)$$

See appendix H.3.1 for a derivation.

Eigenvalues of closed-loop observer

The eigenvalues of the system matrix can be used to determine whether a [state observer](#)'s estimate will converge to the true [state](#).

Plugging equation (9.4) into equation (9.3) gives

$$\begin{aligned} \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{y}_k - \hat{\mathbf{y}}_k) \\ \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{y}_k - (\mathbf{C}\hat{\mathbf{x}}_k + \mathbf{D}\mathbf{u}_k)) \\ \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{y}_k - \mathbf{C}\hat{\mathbf{x}}_k - \mathbf{D}\mathbf{u}_k) \end{aligned}$$

Plugging in equation (5.4) gives

$$\begin{aligned} \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}((\mathbf{C}\mathbf{x}_k + \mathbf{D}\mathbf{u}_k) - \mathbf{C}\hat{\mathbf{x}}_k - \mathbf{D}\mathbf{u}_k) \\ \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{C}\mathbf{x}_k + \mathbf{D}\mathbf{u}_k - \mathbf{C}\hat{\mathbf{x}}_k - \mathbf{D}\mathbf{u}_k) \\ \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{C}\mathbf{x}_k - \mathbf{C}\hat{\mathbf{x}}_k) \\ \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}\mathbf{C}(\mathbf{x}_k - \hat{\mathbf{x}}_k) \end{aligned}$$

Let $E_k = \mathbf{x}_k - \hat{\mathbf{x}}_k$ be the [error](#) in the estimate $\hat{\mathbf{x}}_k$.

$$\hat{\mathbf{x}}_{k+1} = \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}\mathbf{C}\mathbf{E}_k$$

Subtracting this from equation (5.3) gives

$$\begin{aligned} \mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\mathbf{x}_k + \mathbf{B}\mathbf{u}_k - (\mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}\mathbf{C}\mathbf{E}_k) \\ \mathbf{E}_{k+1} &= \mathbf{A}\mathbf{x}_k + \mathbf{B}\mathbf{u}_k - (\mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}\mathbf{C}\mathbf{E}_k) \\ \mathbf{E}_{k+1} &= \mathbf{A}\mathbf{x}_k + \mathbf{B}\mathbf{u}_k - \mathbf{A}\hat{\mathbf{x}}_k - \mathbf{B}\mathbf{u}_k - \mathbf{L}\mathbf{C}\mathbf{E}_k \\ \mathbf{E}_{k+1} &= \mathbf{A}\mathbf{x}_k - \mathbf{A}\hat{\mathbf{x}}_k - \mathbf{L}\mathbf{C}\mathbf{E}_k \\ \mathbf{E}_{k+1} &= \mathbf{A}(\mathbf{x}_k - \hat{\mathbf{x}}_k) - \mathbf{L}\mathbf{C}\mathbf{E}_k \\ \mathbf{E}_{k+1} &= \mathbf{A}\mathbf{E}_k - \mathbf{L}\mathbf{C}\mathbf{E}_k \\ \mathbf{E}_{k+1} &= (\mathbf{A} - \mathbf{L}\mathbf{C})\mathbf{E}_k \end{aligned} \quad (9.8)$$

For equation (9.8) to have a bounded output, the eigenvalues of $\mathbf{A} - \mathbf{LC}$ must be within the unit circle. These eigenvalues represent how fast the estimator converges to the true **state** of the given **model**. A fast estimator converges quickly while a slow estimator avoids amplifying noise in the measurements used to produce a **state** estimate.

As stated before, the controller and estimator are dual problems. Controller gains can be found assuming perfect estimator (i.e., perfect knowledge of all **states**). Estimator gains can be found assuming an accurate **model** and a controller with perfect **tracking**.

The effect of noise can be seen if it is modeled **stochastically** as

$$\hat{\mathbf{x}}_{k+1} = \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}((\mathbf{y}_k + \nu_k) - \hat{\mathbf{y}}_k)$$

where ν_k is the measurement noise. Rearranging this equation yields

$$\begin{aligned}\hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{y}_k - \hat{\mathbf{y}}_k + \nu_k) \\ \hat{\mathbf{x}}_{k+1} &= \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k + \mathbf{L}(\mathbf{y}_k - \hat{\mathbf{y}}_k) + \mathbf{L}\nu_k\end{aligned}$$

As \mathbf{L} increases, the measurement noise is amplified.

9.3 Introduction to probability

Now we'll begin establishing probability concepts we need to describe and manipulate stochastic systems.

9.3.1 Random variables

A random variable is a variable whose values are the outcomes of a random phenomenon, like dice rolls or noisy process measurements. As such, a random variable is defined as a function that maps the outcomes of an unpredictable process to numerical quantities. A particular output of this function is called a sample. The sample space is the set of possible values taken by the random variable.

A probability density function (PDF) is a function of the random variable whose value at a given sample (measured value) in the sample space (the range of possible measured values) is the probability of that sample occurring. The area under the function over a range gives the probability that the sample falls within that range. Let x be a random variable, and let $p(x)$ denote the probability density function of x . The probability that the value of x will be in the interval $x \in [x_1, x_1 + dx]$ is $p(x_1) dx$. In other words, the probability is the area under the PDF within the region $[x_1, x_1 + dx]$ (see figure 9.1).

A probability of zero means that the sample will not occur and a probability of one means that the sample will always occur. Probability density functions require that no probabilities are negative and that the sum of all probabilities is 1. If the probabilities sum to 1, that means one of those outcomes *must* happen. In other words,

$$p(x) \geq 0, \int_{-\infty}^{\infty} p(x) dx = 1$$

or given that the probability of a given sample is greater than or equal to zero, the sum of probabilities for all possible input values is equal to one.

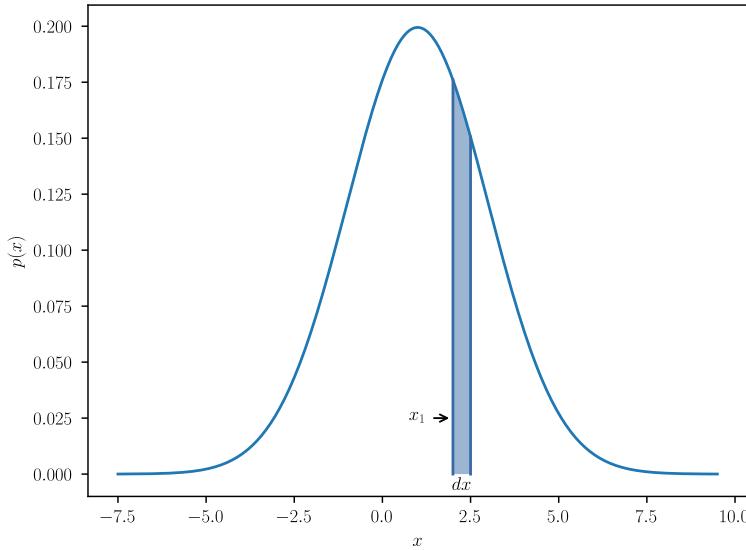


Figure 9.1: Probability density function

9.3.2 Expected value

Expected value or expectation is a weighted average of the values the PDF can produce where the weight for each value is the probability of that value occurring. This can be written mathematically as

$$\bar{x} = E[x] = \int_{-\infty}^{\infty} x p(x) dx$$

The expectation can be applied to random functions as well as random variables.

$$E[f(x)] = \int_{-\infty}^{\infty} f(x) p(x) dx$$

The mean of a random variable is denoted by an overbar (e.g., \bar{x}) pronounced x-bar. The expectation of the difference between a random variable and its mean $x - \bar{x}$ converges to zero. In other words, the expectation of a random variable is its mean. Here's a proof.

$$\begin{aligned} E[x - \bar{x}] &= \int_{-\infty}^{\infty} (x - \bar{x}) p(x) dx \\ E[x - \bar{x}] &= \int_{-\infty}^{\infty} x p(x) dx - \int_{-\infty}^{\infty} \bar{x} p(x) dx \\ E[x - \bar{x}] &= \int_{-\infty}^{\infty} x p(x) dx - \bar{x} \int_{-\infty}^{\infty} p(x) dx \\ E[x - \bar{x}] &= \bar{x} - \bar{x} \cdot 1 \\ E[x - \bar{x}] &= 0 \end{aligned}$$

9.3.3 Variance

Informally, variance is a measure of how far the outcome of a random variable deviates from its mean. Later, we will use variance to quantify how confident we are in the estimate of a random variable;

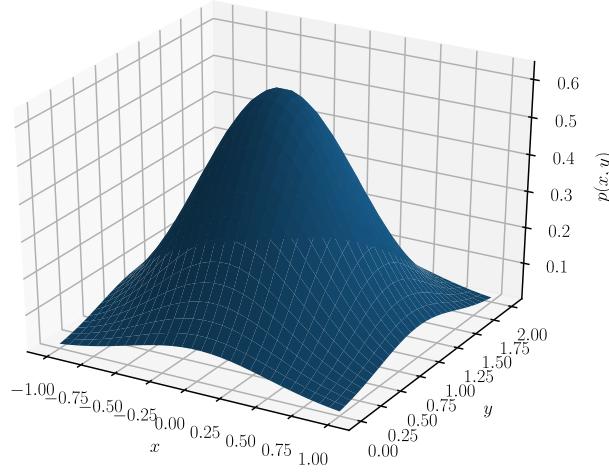


Figure 9.2: Joint probability density function

we can't know the true value of that variable without randomness, but we can give a bound on its randomness.

$$\text{var}(x) = \sigma^2 = E[(x - \bar{x})^2] = \int_{-\infty}^{\infty} (x - \bar{x})^2 p(x) dx$$

The standard deviation is the square root of the variance.

$$\text{std}[x] = \sigma = \sqrt{\text{var}(x)}$$

9.3.4 Joint probability density functions

Probability density functions can also include more than one variable. Let x and y are random variables. The joint probability density function $p(x, y)$ defines the probability $p(x, y) dx dy$, so that x and y are in the intervals $x \in [x, x + dx]$, $y \in [y, y + dy]$. In other words, the probability is the volume under a region of the PDF manifold (see figure 9.2 for an example of a joint PDF).

Joint probability density functions also require that no probabilities are negative and that the sum of all probabilities is 1.

$$p(x, y) \geq 0, \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} p(x, y) dx dy = 1$$

The expected values for joint PDFs are as follows.

$$\begin{aligned} E[x] &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} x dx dy \\ E[y] &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} y dx dy \\ E[f(x, y)] &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) dx dy \end{aligned}$$

The variance of a joint PDF measures how a variable correlates with itself (we'll cover variances with respect to other variables shortly).

$$\text{var}(x) = \Sigma_{xx} = E[(x - \bar{x})^2] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (x - \bar{x})^2 p(x, y) dx dy$$

$$\text{var}(y) = \Sigma_{yy} = E[(y - \bar{y})^2] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (y - \bar{y})^2 p(x, y) dx dy$$

9.3.5 Covariance

A covariance is a measurement of how a variable correlates with another. If they vary in the same direction, the covariance increases. If they vary in opposite directions, the covariance decreases.

$$\text{cov}(x, y) = \Sigma_{xy} = E[(x - \bar{x})(y - \bar{y})] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} (x - \bar{x})(y - \bar{y}) p(x, y) dx dy$$

9.3.6 Correlation

Two random variables are correlated if the result of one random variable affects the result of another. Correlation is defined as

$$\rho(x, y) = \frac{\Sigma_{xy}}{\sqrt{\Sigma_{xx}\Sigma_{yy}}}, |\rho(x, y)| \leq 1$$

So two variable's correlation is defined as their covariance over the geometric mean of their variances. Uncorrelated sources have a covariance of zero.

9.3.7 Independence

Two random variables are independent if the following relation is true.

$$p(x, y) = p(x) p(y)$$

This means that the values of x do not correlate with the values of y . That is, the outcome of one random variable does not affect another's outcome. If we assume independence,

$$E[xy] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} xy p(x, y) dx dy$$

$$E[xy] = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} xy p(x) p(y) dx dy$$

$$E[xy] = \int_{-\infty}^{\infty} x p(x) dx \int_{-\infty}^{\infty} y p(y) dy$$

$$E[xy] = E[x]E[y]$$

$$E[xy] = \bar{x}\bar{y}$$

$$\text{cov}(x, y) = E[(x - \bar{x})(y - \bar{y})]$$

$$\text{cov}(x, y) = E[(x - \bar{x})]E[(y - \bar{y})]$$

$$\text{cov}(x, y) = 0 \cdot 0$$

Therefore, the covariance Σ_{xy} is zero, as expected. Furthermore, $\rho(x, y) = 0$, which means they are uncorrelated.

9.3.8 Marginal probability density functions

Given two random variables x and y whose joint distribution is known, the marginal PDF $p(x)$ expresses the probability of x averaged over information about y . In other words, it's the PDF of x when y is unknown. This is calculated by integrating the joint PDF over y .

$$p(x) = \int_{-\infty}^{\infty} p(x, y) dy$$

9.3.9 Conditional probability density functions

Let us assume that we know the joint PDF $p(x, y)$ and the exact value for y . The conditional PDF gives the probability of x in the interval $[x, x + dx]$ for the given value y .

If $p(x, y)$ is known, then we also know $p(x, y = y^*)$. However, note that the latter is not the conditional density $p(x|y^*)$, instead

$$C(y^*) = \int_{-\infty}^{\infty} p(x, y = y^*) dx$$

$$p(x|y^*) = \frac{1}{C(y^*)} p(x, y = y^*)$$

The scale factor $\frac{1}{C(y^*)}$ is used to scale the area under the PDF to 1.

9.3.10 Bayes's rule

Bayes's rule is used to determine the probability of an event based on prior knowledge of conditions related to the event.

$$p(x, y) = p(x|y) p(y) = p(y|x) p(x)$$

If x and y are independent, then $p(x|y) = p(x)$, $p(y|x) = p(y)$, and $p(x, y) = p(x) p(y)$.

9.3.11 Conditional expectation

The concept of expectation can also be applied to conditional PDFs. This allows us to determine what the mean of a variable is given prior knowledge of other variables.

$$E[x|y] = \int_{-\infty}^{\infty} x p(x|y) dx = f(y), E[x|y] \neq E[x]$$

$$E[y|x] = \int_{-\infty}^{\infty} y p(y|x) dy = f(x), E[y|x] \neq E[y]$$

9.3.12 Conditional variances

$$\text{var}(x|y) = E[(x - E[x|y])^2|y]$$

$$\text{var}(x|y) = \int_{-\infty}^{\infty} (x - E[x|y])^2 p(x|y) dx$$

9.3.13 Random vectors

Now we will extend the probability concepts discussed so far to vectors where each element has a PDF.

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

The elements of \mathbf{x} are scalar variables jointly distributed with a joint density $p(x_1, x_2, \dots, x_n)$. The expectation is

$$\begin{aligned} E[\mathbf{x}] &= \bar{\mathbf{x}} = \int_{-\infty}^{\infty} \mathbf{x} p(\mathbf{x}) d\mathbf{x} \\ E[\mathbf{x}] &= \begin{bmatrix} E[x_1] \\ E[x_2] \\ \vdots \\ E[x_n] \end{bmatrix} \\ E[x_i] &= \int_{-\infty}^{\infty} \dots \int_{-\infty}^{\infty} x_i p(x_1, x_2, \dots, x_n) dx_1 \dots dx_n \\ E[f(\mathbf{x})] &= \int_{-\infty}^{\infty} f(\mathbf{x}) p(\mathbf{x}) d\mathbf{x} \end{aligned}$$

9.3.14 Covariance matrix

The covariance matrix for a random vector $\mathbf{x} \in \mathbb{R}^n$ is

$$\begin{aligned} \Sigma &= \text{cov}(\mathbf{x}, \mathbf{x}) = E[(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T] \\ \Sigma &= \begin{bmatrix} \text{cov}(x_1, x_1) & \text{cov}(x_1, x_2) & \dots & \text{cov}(x_1, x_n) \\ \text{cov}(x_2, x_1) & \text{cov}(x_2, x_2) & \dots & \text{cov}(x_1, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ \text{cov}(x_n, x_1) & \text{cov}(x_n, x_2) & \dots & \text{cov}(x_n, x_n) \end{bmatrix} \end{aligned}$$

This $n \times n$ matrix is symmetric and positive semidefinite. A positive semidefinite matrix satisfies the relation that for any $\mathbf{v} \in \mathbb{R}^n$ for which $\mathbf{v} \neq 0$, $\mathbf{v}^T \Sigma \mathbf{v} \geq 0$. In other words, the eigenvalues of Σ are all greater than or equal to zero.

9.3.15 Relations for independent random vectors

First, independent vectors imply linearity from $p(\mathbf{x}, \mathbf{y}) = p(\mathbf{x}) p(\mathbf{y})$.

$$\begin{aligned} E[\mathbf{Ax} + \mathbf{By}] &= \mathbf{A}E[\mathbf{x}] + \mathbf{B}E[\mathbf{y}] \\ E[\mathbf{Ax} + \mathbf{By}] &= \mathbf{A}\bar{\mathbf{x}} + \mathbf{B}\bar{\mathbf{y}} \end{aligned}$$

Second, independent vectors being uncorrelated means their covariance is zero.

$$\begin{aligned} \Sigma_{\mathbf{xy}} &= \text{cov}(\mathbf{x}, \mathbf{y}) \\ \Sigma_{\mathbf{xy}} &= E[(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{y} - \bar{\mathbf{y}})^T] \\ \Sigma_{\mathbf{xy}} &= E[\mathbf{xy}^T] - E[\mathbf{x}\bar{\mathbf{y}}^T] - E[\bar{\mathbf{x}}\mathbf{y}^T] + E[\bar{\mathbf{x}}\bar{\mathbf{y}}^T] \\ \Sigma_{\mathbf{xy}} &= E[\mathbf{xy}^T] - E[\mathbf{x}]\bar{\mathbf{y}}^T - \bar{\mathbf{x}}E[\mathbf{y}^T] + \bar{\mathbf{x}}\bar{\mathbf{y}}^T \\ \Sigma_{\mathbf{xy}} &= E[\mathbf{xy}^T] - \bar{\mathbf{x}}\bar{\mathbf{y}}^T - \bar{\mathbf{x}}\bar{\mathbf{y}}^T + \bar{\mathbf{x}}\bar{\mathbf{y}}^T \\ \Sigma_{\mathbf{xy}} &= E[\mathbf{xy}^T] - \bar{\mathbf{x}}\bar{\mathbf{y}}^T \end{aligned} \tag{9.9}$$

Now, compute $E[\mathbf{xy}^T]$.

$$E[\mathbf{xy}^T] = \int_X \int_Y \mathbf{xy}^T p(\mathbf{x}) p(\mathbf{y}) d\mathbf{x} d\mathbf{y}^T$$

Factor out constants from the inner integral. This includes variables which are held constant for each inner integral evaluation.

$$E[\mathbf{x}\mathbf{y}^T] = \int_X p(\mathbf{x}) \mathbf{x} d\mathbf{x} \int_Y p(\mathbf{y}) \mathbf{y}^T d\mathbf{y}^T$$

Each of these integrals is just the expected value of their respective integration variable.

$$E[\mathbf{x}\mathbf{y}^T] = \bar{\mathbf{x}}\bar{\mathbf{y}}^T \quad (9.10)$$

Substitute equation (9.10) into equation (9.9).

$$\begin{aligned}\Sigma_{\mathbf{x}\mathbf{y}} &= (\bar{\mathbf{x}}\bar{\mathbf{y}}^T) - \bar{\mathbf{x}}\bar{\mathbf{y}}^T \\ \Sigma_{\mathbf{x}\mathbf{y}} &= 0\end{aligned}$$

Using these results, we can compute the covariance of $\mathbf{z} = \mathbf{Ax} + \mathbf{By}$.

$$\begin{aligned}\Sigma_z &= \text{cov}(\mathbf{z}, \mathbf{z}) \\ \Sigma_z &= E[(\mathbf{z} - \bar{\mathbf{z}})(\mathbf{z} - \bar{\mathbf{z}})^T] \\ \Sigma_z &= E[(\mathbf{Ax} + \mathbf{By} - \mathbf{A}\bar{\mathbf{x}} - \mathbf{B}\bar{\mathbf{y}})(\mathbf{Ax} + \mathbf{By} - \mathbf{A}\bar{\mathbf{x}} - \mathbf{B}\bar{\mathbf{y}})^T] \\ \Sigma_z &= E[(\mathbf{A}(\mathbf{x} - \bar{\mathbf{x}}) + \mathbf{B}(\mathbf{y} - \bar{\mathbf{y}}))(\mathbf{A}(\mathbf{x} - \bar{\mathbf{x}}) + \mathbf{B}(\mathbf{y} - \bar{\mathbf{y}}))^T] \\ \Sigma_z &= E[(\mathbf{A}(\mathbf{x} - \bar{\mathbf{x}}) + \mathbf{B}(\mathbf{y} - \bar{\mathbf{y}}))((\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{A}^T + (\mathbf{y} - \bar{\mathbf{y}})^T \mathbf{B}^T)] \\ \Sigma_z &= E[\mathbf{A}(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{A}^T + \mathbf{A}(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{y} - \bar{\mathbf{y}})^T \mathbf{B}^T + \\ &\quad \mathbf{B}(\mathbf{y} - \bar{\mathbf{y}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{A}^T + \mathbf{B}(\mathbf{y} - \bar{\mathbf{y}})(\mathbf{y} - \bar{\mathbf{y}})^T \mathbf{B}^T]\end{aligned}$$

Since \mathbf{x} and \mathbf{y} are independent, $\Sigma_{xy} = 0$ and the cross terms cancel out.

$$\begin{aligned}\Sigma_z &= E[\mathbf{A}(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{A}^T + 0 + 0 + \mathbf{B}(\mathbf{y} - \bar{\mathbf{y}})(\mathbf{y} - \bar{\mathbf{y}})^T \mathbf{B}^T] \\ \Sigma_z &= E[\mathbf{A}(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{A}^T] + E[\mathbf{B}(\mathbf{y} - \bar{\mathbf{y}})(\mathbf{y} - \bar{\mathbf{y}})^T \mathbf{B}^T] \\ \Sigma_z &= \mathbf{A}E[(\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T] \mathbf{A}^T + \mathbf{B}E[(\mathbf{y} - \bar{\mathbf{y}})(\mathbf{y} - \bar{\mathbf{y}})^T] \mathbf{B}^T\end{aligned}$$

Recall that $\Sigma_x = \text{cov}(\mathbf{x}, \mathbf{x})$ and $\Sigma_y = \text{cov}(\mathbf{y}, \mathbf{y})$.

$$\Sigma_z = \mathbf{A}\Sigma_x\mathbf{A}^T + \mathbf{B}\Sigma_y\mathbf{B}^T$$

9.3.16 Gaussian random variables

A Gaussian random variable has the following properties:

$$\begin{aligned}E[x] &= \bar{x} \\ \text{var}(x) &= \sigma^2 \\ p(x) &= \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\bar{x})^2}{2\sigma^2}}\end{aligned}$$

While we could use any random variable to represent a random process, we use the Gaussian random variable a lot in probability theory due to the central limit theorem.

Definition 9.3.1 — Central limit theorem. When independent random variables are added, their properly normalized sum tends toward a normal distribution (a Gaussian distribution or “bell curve”).

This is the case even if the original variables themselves are not normally distributed. The theorem is a key concept in probability theory because it implies that probabilistic and statistical methods that work for normal distributions can be applicable to many problems involving other types of distributions.

For example, suppose that a sample is obtained containing a large number of independent observations, and that the arithmetic mean of the observed values is computed. The central limit theorem says that the computed values of the mean will tend toward being distributed according to a normal distribution.

9.4 Linear stochastic systems

Given the following stochastic system

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{Ax}_k + \mathbf{Bu}_k + \mathbf{\Gamma w}_k \\ \mathbf{y}_k &= \mathbf{Cx}_k + \mathbf{Du}_k + \mathbf{v}_k\end{aligned}$$

where \mathbf{w}_k is the process noise and \mathbf{v}_k is the measurement noise,

$$\begin{aligned}E[\mathbf{w}_k] &= 0 \\ E[\mathbf{w}_k \mathbf{w}_k^T] &= \mathbf{Q}_k \\ E[\mathbf{v}_k] &= 0 \\ E[\mathbf{v}_k \mathbf{v}_k^T] &= \mathbf{R}_k\end{aligned}$$

where \mathbf{Q}_k is the process noise covariance matrix and \mathbf{R}_k is the measurement noise covariance matrix. We assume the noise samples are independent, so $E[\mathbf{w}_k \mathbf{w}_j^T] = 0$ and $E[\mathbf{v}_k \mathbf{v}_j^T] = 0$ where $k \neq j$. Furthermore, process noise samples are independent from measurement noise samples.

We'll compute the expectation of these equations and their covariance matrices, which we'll use later for deriving the Kalman filter.

9.4.1 State vector expectation evolution

First, we will compute how the expectation of the system state evolves.

$$\begin{aligned}E[\mathbf{x}_{k+1}] &= E[\mathbf{Ax}_k + \mathbf{Bu}_k + \mathbf{\Gamma w}_k] \\ E[\mathbf{x}_{k+1}] &= E[\mathbf{Ax}_k] + E[\mathbf{Bu}_k] + E[\mathbf{\Gamma w}_k] \\ E[\mathbf{x}_{k+1}] &= \mathbf{A}E[\mathbf{x}_k] + \mathbf{B}E[\mathbf{u}_k] + \mathbf{\Gamma}E[\mathbf{w}_k] \\ E[\mathbf{x}_{k+1}] &= \mathbf{A}E[\mathbf{x}_k] + \mathbf{B}\mathbf{u}_k + 0 \\ \bar{\mathbf{x}}_{k+1} &= \mathbf{A}\bar{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k\end{aligned}$$

9.4.2 Error covariance matrix evolution

Now, we will use this to compute how the error covariance matrix \mathbf{P} evolves.

$$\mathbf{x}_{k+1} - \bar{\mathbf{x}}_{k+1} = \mathbf{Ax}_k + \mathbf{Bu}_k + \mathbf{\Gamma w}_k - (\mathbf{A}\bar{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k)$$

$$\mathbf{x}_{k+1} - \bar{\mathbf{x}}_{k+1} = \mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{\Gamma} \mathbf{w}_k$$

$$E[(\mathbf{x}_{k+1} - \bar{\mathbf{x}}_{k+1})(\mathbf{x}_{k+1} - \bar{\mathbf{x}}_{k+1})^T] = E[(\mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{\Gamma} \mathbf{w}_k)(\mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{\Gamma} \mathbf{w}_k)^T]$$

$$\mathbf{P}_{k+1} = E[(\mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{\Gamma} \mathbf{w}_k)(\mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{\Gamma} \mathbf{w}_k)^T]$$

$$\begin{aligned} \mathbf{P}_{k+1} &= E[(\mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k)(\mathbf{x}_k - \bar{\mathbf{x}}_k)^T \mathbf{A}^T] + E[\mathbf{A}(\mathbf{x}_k - \bar{\mathbf{x}}_k) \mathbf{w}_k^T \mathbf{\Gamma}^T] + \\ &\quad E[\mathbf{\Gamma} \mathbf{w}_k (\mathbf{x}_k - \bar{\mathbf{x}}_k)^T \mathbf{A}^T] + E[\mathbf{\Gamma} \mathbf{w}_k \mathbf{w}_k^T \mathbf{\Gamma}^T] \end{aligned}$$

$$\begin{aligned} \mathbf{P}_{k+1} &= \mathbf{A} E[(\mathbf{x}_k - \bar{\mathbf{x}}_k)(\mathbf{x}_k - \bar{\mathbf{x}}_k)^T] \mathbf{A}^T + \mathbf{A} E[(\mathbf{x}_k - \bar{\mathbf{x}}_k) \mathbf{w}_k^T] \mathbf{\Gamma}^T + \\ &\quad \mathbf{\Gamma} E[\mathbf{w}_k (\mathbf{x}_k - \bar{\mathbf{x}}_k)^T] \mathbf{A}^T + \mathbf{\Gamma} E[\mathbf{w}_k \mathbf{w}_k^T] \mathbf{\Gamma}^T \end{aligned}$$

$$\mathbf{P}_{k+1} = \mathbf{A} \mathbf{P}_k \mathbf{A}^T + \mathbf{A} E[(\mathbf{x}_k - \bar{\mathbf{x}}_k) \mathbf{w}_k^T] \mathbf{\Gamma}^T + \mathbf{\Gamma} E[\mathbf{w}_k (\mathbf{x}_k - \bar{\mathbf{x}}_k)^T] \mathbf{A}^T + \mathbf{\Gamma} \mathbf{Q}_k \mathbf{\Gamma}_k^T$$

Since the error and noise are independent, the cross terms are zero.

$$\mathbf{P}_{k+1} = \mathbf{A} \mathbf{P}_k \mathbf{A}^T + \mathbf{A} \underbrace{E[(\mathbf{x}_k - \bar{\mathbf{x}}_k) \mathbf{w}_k^T]}_0 \mathbf{\Gamma}^T + \mathbf{\Gamma} \underbrace{E[\mathbf{w}_k (\mathbf{x}_k - \bar{\mathbf{x}}_k)^T]}_0 \mathbf{A}^T + \mathbf{\Gamma} \mathbf{Q}_k \mathbf{\Gamma}_k^T$$

$$\mathbf{P}_{k+1} = \mathbf{A} \mathbf{P}_k \mathbf{A}^T + \mathbf{\Gamma} \mathbf{Q}_k \mathbf{\Gamma}^T$$

9.4.3 Measurement vector expectation

Next, we will compute the expectation of the output \mathbf{y} .

$$E[\mathbf{y}_k] = E[\mathbf{C} \mathbf{x}_k + \mathbf{D} \mathbf{u}_k + \mathbf{v}_k]$$

$$E[\mathbf{y}_k] = \mathbf{C} E[\mathbf{x}_k] + \mathbf{D} \mathbf{u}_k + 0$$

$$\bar{\mathbf{y}}_k = \mathbf{C} \bar{\mathbf{x}}_k + \mathbf{D} \mathbf{u}_k$$

9.4.4 Measurement covariance matrix

Now, we will use this to compute how the measurement covariance matrix \mathbf{S} evolves.

$$\mathbf{y}_k - \bar{\mathbf{y}}_k = \mathbf{C} \mathbf{x}_k + \mathbf{D} \mathbf{u}_k + \mathbf{v}_k - (\mathbf{C} \bar{\mathbf{x}}_k + \mathbf{D} \mathbf{u}_k)$$

$$\mathbf{y}_k - \bar{\mathbf{y}}_k = \mathbf{C}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{v}_k$$

$$E[(\mathbf{y}_k - \bar{\mathbf{y}}_k)(\mathbf{y}_k - \bar{\mathbf{y}}_k)^T] = E[(\mathbf{C}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{v}_k)(\mathbf{C}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{v}_k)^T]$$

$$\mathbf{S}_k = E[(\mathbf{C}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{v}_k)(\mathbf{C}(\mathbf{x}_k - \bar{\mathbf{x}}_k) + \mathbf{v}_k)^T]$$

$$\mathbf{S}_k = E[(\mathbf{C}(\mathbf{x}_k - \bar{\mathbf{x}}_k)(\mathbf{x}_k - \bar{\mathbf{x}}_k)^T \mathbf{C}^T] + E[\mathbf{v}_k \mathbf{v}_k^T]$$

$$\mathbf{S}_k = \mathbf{C} E[(\mathbf{x}_k - \bar{\mathbf{x}}_k)(\mathbf{x}_k - \bar{\mathbf{x}}_k)^T] \mathbf{C}^T + \mathbf{R}_k$$

$$\mathbf{S}_k = \mathbf{C} \mathbf{P}_k \mathbf{C}^T + \mathbf{R}_k$$

9.5 Two-sensor problem

We'll skip the probability derivations here, but given two data points with associated variances represented by Gaussian distributions, the information can be optimally combined into a third Gaussian distribution with its own mean value and variance. The expected value of x given a measurement z_1 is

$$E[x|z_1] = \mu = \frac{\sigma_0^2}{\sigma_0^2 + \sigma^2} z_1 + \frac{\sigma^2}{\sigma_0^2 + \sigma^2} x_0 \quad (9.11)$$

The variance of x given z_1 is

$$E[(x - \mu)^2 | z_1] = \frac{\sigma^2 \sigma_0^2}{\sigma_0^2 + \sigma^2} \quad (9.12)$$

The expected value, which is also the maximum likelihood value, is the linear combination of the prior expected (maximum likelihood) value and the measurement. The expected value is a reasonable estimator of x .

$$\begin{aligned} \hat{x} &= E[x | z_1] = \frac{\sigma_0^2}{\sigma_0^2 + \sigma^2} z_1 + \frac{\sigma^2}{\sigma_0^2 + \sigma^2} x_0 \\ \hat{x} &= w_1 z_1 + w_2 x_0 \end{aligned} \quad (9.13)$$

Note that the weights w_1 and w_2 sum to 1. When the prior (i.e., prior knowledge of [state](#)) is uninformative (a large variance)

$$w_1 = \lim_{\sigma_0^2 \rightarrow 0} \frac{\sigma_0^2}{\sigma_0^2 + \sigma^2} = 0 \quad (9.14)$$

$$w_2 = \lim_{\sigma_0^2 \rightarrow 0} \frac{\sigma^2}{\sigma_0^2 + \sigma^2} = 1 \quad (9.15)$$

and $\hat{x} = z_1$. That is, the weight is on the observations and the estimate is equal to the measurement.

Let us assume we have a [model](#) providing an almost exact prior for x . In that case, σ_0^2 approaches 0 and

$$w_1 = \lim_{\sigma_0^2 \rightarrow 0} \frac{\sigma_0^2}{\sigma_0^2 + \sigma^2} = 1 \quad (9.16)$$

$$w_2 = \lim_{\sigma_0^2 \rightarrow 0} \frac{\sigma^2}{\sigma_0^2 + \sigma^2} = 0 \quad (9.17)$$

The Kalman filter uses this optimal fusion as the basis for its operation.

9.6 Kalman filter

So far, we've derived equations for updating the expected value and state covariance without measurements and how to incorporate measurements into an initial [state](#) optimally. Now, we'll combine these concepts to produce an estimator which minimizes the error covariance for linear [systems](#).

9.6.1 Derivations

Given the *a posteriori* update equation $\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{K}_{k+1}(\mathbf{y}_{k+1} - \hat{\mathbf{y}}_{k+1})$, we want to find the value of \mathbf{K}_{k+1} that minimizes the *a posteriori* estimate covariance (the error covariance) because this minimizes the estimation error.

***a posteriori* estimate covariance update equation**

The following is the definition of the *a posteriori* estimate covariance matrix.

$$\mathbf{P}_{k+1}^+ = \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^+)$$

Substitute in the *a posteriori* update equation and expand the measurement equations.

$$\mathbf{P}_{k+1}^+ = \text{cov}(\mathbf{x}_{k+1} - (\hat{\mathbf{x}}_{k+1}^- + \mathbf{K}_{k+1}(\mathbf{y}_{k+1} - \hat{\mathbf{y}}_{k+1})))$$

$$\begin{aligned}
\mathbf{P}_{k+1}^+ &= \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \mathbf{K}_{k+1}(\mathbf{y}_{k+1} - \hat{\mathbf{y}}_{k+1})) \\
\mathbf{P}_{k+1}^+ &= \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \\
&\quad \mathbf{K}_{k+1}((\mathbf{C}_{k+1}\mathbf{x}_{k+1} + \mathbf{D}_{k+1}\mathbf{u}_{k+1} + \mathbf{v}_{k+1}) - (\mathbf{C}_{k+1}\hat{\mathbf{x}}_{k+1}^- + \mathbf{D}_{k+1}\mathbf{u}_{k+1}))) \\
\mathbf{P}_{k+1}^+ &= \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \\
&\quad \mathbf{K}_{k+1}(\mathbf{C}_{k+1}\mathbf{x}_{k+1} + \mathbf{D}_{k+1}\mathbf{u}_{k+1} + \mathbf{v}_{k+1} - \mathbf{C}_{k+1}\hat{\mathbf{x}}_{k+1}^- - \mathbf{D}_{k+1}\mathbf{u}_{k+1}))
\end{aligned}$$

Reorder terms.

$$\begin{aligned}
\mathbf{P}_{k+1}^+ &= \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \\
&\quad \mathbf{K}_{k+1}(\mathbf{C}_{k+1}\mathbf{x}_{k+1} - \mathbf{C}_{k+1}\hat{\mathbf{x}}_{k+1}^- + \mathbf{D}_{k+1}\mathbf{u}_{k+1} - \mathbf{D}_{k+1}\mathbf{u}_{k+1} + \mathbf{v}_{k+1}))
\end{aligned}$$

The $\mathbf{D}_{k+1}\mathbf{u}_{k+1}$ terms cancel.

$$\mathbf{P}_{k+1}^+ = \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \mathbf{K}_{k+1}(\mathbf{C}_{k+1}\mathbf{x}_{k+1} - \mathbf{C}_{k+1}\hat{\mathbf{x}}_{k+1}^- + \mathbf{v}_{k+1}))$$

Distribute \mathbf{K}_{k+1} to \mathbf{v}_{k+1} .

$$\mathbf{P}_{k+1}^+ = \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \mathbf{K}_{k+1}(\mathbf{C}_{k+1}\mathbf{x}_{k+1} - \mathbf{C}_{k+1}\hat{\mathbf{x}}_{k+1}^-) - \mathbf{K}_{k+1}\mathbf{v}_{k+1})$$

Factor out \mathbf{C}_{k+1} .

$$\mathbf{P}_{k+1}^+ = \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^- - \mathbf{K}_{k+1}\mathbf{C}_{k+1}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^-) - \mathbf{K}_{k+1}\mathbf{v}_{k+1})$$

Factor out $\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^-$ to the right.

$$\mathbf{P}_{k+1}^+ = \text{cov}((\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^-) - \mathbf{K}_{k+1}\mathbf{v}_{k+1})$$

Covariance is a linear operator, so it can be applied to each term separately. Covariance squares terms internally, so the negative sign on $\mathbf{K}_{k+1}\mathbf{v}_{k+1}$ is removed.

$$\mathbf{P}_{k+1}^+ = \text{cov}((\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^-)) + \text{cov}(\mathbf{K}_{k+1}\mathbf{v}_{k+1})$$

Now just evaluate the covariances.

$$\begin{aligned}
\mathbf{P}_{k+1}^+ &= (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1}) \text{cov}(\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^-)(\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})^T + \mathbf{K}_{k+1} \text{cov}(\mathbf{v}_{k+1}) \mathbf{K}_{k+1}^T \\
\mathbf{P}_{k+1}^+ &= (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1}) \mathbf{P}_{k+1}^- (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})^T + \mathbf{K}_{k+1} \mathbf{R}_{k+1} \mathbf{K}_{k+1}^T
\end{aligned}$$

Finding the optimal Kalman gain

The error in the *a posteriori state* estimation is $\mathbf{x}_{k+1} - \hat{\mathbf{x}}_{k+1}^-$. We want to minimize the expected value of the square of the magnitude of this vector. This is equivalent to minimizing the trace of the a posteriori estimate covariance matrix \mathbf{P}_{k+1}^+ .

We'll start with the equation for \mathbf{P}_{k+1}^+ .

$$\mathbf{P}_{k+1}^+ = (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1}) \mathbf{P}_{k+1}^- (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})^T + \mathbf{K}_{k+1} \mathbf{R}_{k+1} \mathbf{K}_{k+1}^T$$

We're going to expand the equation for \mathbf{P}_{k+1}^+ and collect terms. First, multiply in \mathbf{P}_{k+1}^- .

$$\mathbf{P}_{k+1}^+ = (\mathbf{P}_{k+1}^- - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-)(\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})^T + \mathbf{K}_{k+1} \mathbf{R}_{k+1} \mathbf{K}_{k+1}^T$$

Transpose each term in $\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1}$. \mathbf{I} is symmetric, so its transpose is dropped.

$$\mathbf{P}_{k+1}^+ = (\mathbf{P}_{k+1}^- - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-)(\mathbf{I} - \mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T) + \mathbf{K}_{k+1}\mathbf{R}_{k+1}\mathbf{K}_{k+1}^T$$

Multiply in $\mathbf{I} - \mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T$.

$$\mathbf{P}_{k+1}^+ = \mathbf{P}_{k+1}^-(\mathbf{I} - \mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T) - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-(\mathbf{I} - \mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T) + \mathbf{K}_{k+1}\mathbf{R}_{k+1}\mathbf{K}_{k+1}^T$$

Expand terms.

$$\begin{aligned} \mathbf{P}_{k+1}^+ &= \mathbf{P}_{k+1}^- - \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^- + \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T \\ &\quad \mathbf{K}_{k+1}\mathbf{R}_{k+1}\mathbf{K}_{k+1}^T \end{aligned}$$

Factor out \mathbf{K}_{k+1} and \mathbf{K}_{k+1}^T .

$$\begin{aligned} \mathbf{P}_{k+1}^+ &= \mathbf{P}_{k+1}^- - \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^- + \\ &\quad \mathbf{K}_{k+1}(\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T + \mathbf{R}_{k+1})\mathbf{K}_{k+1}^T \\ \mathbf{P}_{k+1}^+ &= \mathbf{P}_{k+1}^- - \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^- + \mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T \end{aligned} \quad (9.18)$$

Now take the trace.

$$\text{tr}(\mathbf{P}_{k+1}^+) = \text{tr}(\mathbf{P}_{k+1}^-) - \text{tr}(\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T) - \text{tr}(\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-) + \text{tr}(\mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T)$$

Transpose one of the terms twice.

$$\text{tr}(\mathbf{P}_{k+1}^+) = \text{tr}(\mathbf{P}_{k+1}^-) - \text{tr}((\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-)^T) - \text{tr}(\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-) + \text{tr}(\mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T)$$

\mathbf{P}_{k+1}^- is symmetric, so we can drop the transpose.

$$\text{tr}(\mathbf{P}_{k+1}^+) = \text{tr}(\mathbf{P}_{k+1}^-) - \text{tr}((\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-)^T) - \text{tr}(\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-) + \text{tr}(\mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T)$$

The trace of a matrix is equal to the trace of its transpose since the elements used in the trace are on the diagonal.

$$\begin{aligned} \text{tr}(\mathbf{P}_{k+1}^+) &= \text{tr}(\mathbf{P}_{k+1}^-) - \text{tr}(\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-) - \text{tr}(\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-) + \text{tr}(\mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T) \\ \text{tr}(\mathbf{P}_{k+1}^+) &= \text{tr}(\mathbf{P}_{k+1}^-) - 2\text{tr}(\mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-) + \text{tr}(\mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T) \end{aligned}$$

Given theorems 9.6.1 and 9.6.2

Theorem 9.6.1 $\frac{\partial}{\partial \mathbf{A}} \text{tr}(\mathbf{A}\mathbf{B}\mathbf{A}^T) = 2\mathbf{A}\mathbf{B}$ where \mathbf{B} is symmetric.

Theorem 9.6.2 $\frac{\partial}{\partial \mathbf{A}} \text{tr}(\mathbf{A}\mathbf{C}) = \mathbf{C}^T$

find the minimum of the trace of \mathbf{P}_{k+1}^+ by taking the partial derivative with respect to \mathbf{K} and setting the result to $\mathbf{0}$.

$$\begin{aligned}\frac{\partial \text{tr}(\mathbf{P}_{k+1}^+)}{\partial \mathbf{K}} &= \mathbf{0} - 2(\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-)^T + 2\mathbf{K}_{k+1}\mathbf{S}_{k+1} \\ \frac{\partial \text{tr}(\mathbf{P}_{k+1}^+)}{\partial \mathbf{K}} &= -2\mathbf{P}_{k+1}^{-T}\mathbf{C}_{k+1}^T + 2\mathbf{K}_{k+1}\mathbf{S}_{k+1} \\ \frac{\partial \text{tr}(\mathbf{P}_{k+1}^+)}{\partial \mathbf{K}} &= -2\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T + 2\mathbf{K}_{k+1}\mathbf{S}_{k+1} \\ \mathbf{0} &= -2\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T + 2\mathbf{K}_{k+1}\mathbf{S}_{k+1} \\ 2\mathbf{K}_{k+1}\mathbf{S}_{k+1} &= 2\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T \\ \mathbf{K}_{k+1}\mathbf{S}_{k+1} &= \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T \\ \mathbf{K}_{k+1} &= \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{S}_{k+1}^{-1}\end{aligned}$$

This is the optimal Kalman gain.

Simplified *a priori* estimate covariance update equation

If the optimal Kalman gain is used, the *a posteriori* estimate covariance matrix update equation can be simplified. First, we'll manipulate the equation for the optimal Kalman gain.

$$\begin{aligned}\mathbf{K}_{k+1} &= \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{S}_{k+1}^{-1} \\ \mathbf{K}_{k+1}\mathbf{S}_{k+1} &= \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T \\ \mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T &= \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T\end{aligned}$$

Now we'll substitute it into equation (9.18).

$$\begin{aligned}\mathbf{P}_{k+1}^+ &= \mathbf{P}_{k+1}^- - \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^- + \mathbf{K}_{k+1}\mathbf{S}_{k+1}\mathbf{K}_{k+1}^T \\ \mathbf{P}_{k+1}^+ &= \mathbf{P}_{k+1}^- - \mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^- + (\mathbf{P}_{k+1}^-\mathbf{C}_{k+1}^T\mathbf{K}_{k+1}^T) \\ \mathbf{P}_{k+1}^+ &= \mathbf{P}_{k+1}^- - \mathbf{K}_{k+1}\mathbf{C}_{k+1}\mathbf{P}_{k+1}^-\end{aligned}$$

Factor out \mathbf{P}_{k+1}^- to the right.

$$\mathbf{P}_{k+1}^+ = (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C}_{k+1})\mathbf{P}_{k+1}^-$$

9.6.2 Predict and update equations

Now that we've derived all the pieces we need, we can finally write all the equations for a Kalman filter. Theorem 9.6.3 shows the predict and update steps for a Kalman filter at the k^{th} timestep.

Theorem 9.6.3 — Kalman filter.

Predict step

$$\hat{\mathbf{x}}_{k+1}^- = \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k \quad (9.19)$$

$$\mathbf{P}_{k+1}^- = \mathbf{A}\mathbf{P}_k^-\mathbf{A}^T + \mathbf{Q}\mathbf{Q}^T \quad (9.20)$$

Update step

$$\mathbf{K}_{k+1} = \mathbf{P}_{k+1}^-\mathbf{C}^T(\mathbf{C}\mathbf{P}_{k+1}^-\mathbf{C}^T + \mathbf{R})^{-1} \quad (9.21)$$

$$\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{K}_{k+1}(\mathbf{y}_{k+1} - \mathbf{C}\hat{\mathbf{x}}_{k+1}^- - \mathbf{D}\mathbf{u}_{k+1}) \quad (9.22)$$

$$\mathbf{P}_{k+1}^+ = (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C})\mathbf{P}_{k+1}^- \quad (9.23)$$

A	system matrix	$\hat{\mathbf{x}}$	state estimate vector
B	input matrix	\mathbf{u}	input vector
C	output matrix	\mathbf{y}	output vector
D	feedthrough matrix	\mathbf{Q}	process noise intensity vector
P	error covariance matrix	\mathbf{Q}	process noise covariance matrix
K	Kalman gain matrix	\mathbf{R}	measurement noise covariance matrix

where a superscript of minus denotes *a priori* and plus denotes *a posteriori* estimate (before and after update respectively).

C, D, Q, and R from the equations derived earlier are made constants here.

R To implement a discrete time Kalman filter from a continuous model, the model and continuous time **Q** and **R** matrices can be [discretized](#) using theorem 6.7.1.

Matrix	Rows × Columns	Matrix	Rows × Columns
A	states × states	$\hat{\mathbf{x}}$	states × 1
B	states × inputs	\mathbf{u}	inputs × 1
C	outputs × states	\mathbf{y}	outputs × 1
D	outputs × inputs	\mathbf{Q}	states × 1
P	states × states	\mathbf{Q}	states × states
K	states × outputs	\mathbf{R}	outputs × outputs

Table 9.2: Kalman filter matrix dimensions

Unknown **states** in a Kalman filter are generally represented by a Wiener (pronounced VEE-ner) process². This process has the property that its variance increases linearly with time t .

²Explaining why we use the Wiener process would require going much more in depth into stochastic processes and Itô calculus, which is outside the scope of this book.

9.6.3 Setup

Equations to model

The following example [system](#) will be used to describe how to define and initialize the matrices for a Kalman filter.

A robot is between two parallel walls. It starts driving from one wall to the other at a velocity of $0.8\text{cm}/\text{s}$ and uses ultrasonic sensors to provide noisy measurements of the distances to the walls in front of and behind it. To estimate the distance between the walls, we will define three [states](#): robot position, robot velocity, and distance between the walls.

$$x_{k+1} = x_k + v_k \Delta T \quad (9.24)$$

$$v_{k+1} = v_k \quad (9.25)$$

$$x_{k+1}^w = x_k^w \quad (9.26)$$

This can be converted to the following state-space [model](#).

$$\mathbf{x}_k = \begin{bmatrix} x_k \\ v_k \\ x_k^w \end{bmatrix} \quad (9.27)$$

$$\mathbf{x}_{k+1} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \mathbf{x}_k + \begin{bmatrix} 0 \\ 0.8 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ 0.1 \\ 0 \end{bmatrix} w_k \quad (9.28)$$

where the Gaussian random variable w_k has a mean of 0 and a variance of 1. The observation [model](#) is

$$\mathbf{y}_k = \begin{bmatrix} 1 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix} \mathbf{x}_k + \theta_k \quad (9.29)$$

where the covariance matrix of Gaussian measurement noise θ is a 2×2 matrix with both diagonals 10cm^2 .

The [state](#) vector is usually initialized using the first measurement or two. The covariance matrix entries are assigned by calculating the covariance of the expressions used when assigning the state vector. Let $k = 2$.

$$\mathbf{Q} = [1] \quad (9.30)$$

$$\mathbf{R} = \begin{bmatrix} 10 & 0 \\ 0 & 10 \end{bmatrix} \quad (9.31)$$

$$\hat{\mathbf{x}} = \begin{bmatrix} \mathbf{y}_{k,1} \\ (\mathbf{y}_{k,1} - \mathbf{y}_{k-1,1})/dt \\ \mathbf{y}_{k,1} + \mathbf{y}_{k,2} \end{bmatrix} \quad (9.32)$$

$$\mathbf{P} = \begin{bmatrix} 10 & 10/dt & 10 \\ 10/dt & 20/dt^2 & 10/dt \\ 10 & 10/dt & 20 \end{bmatrix} \quad (9.33)$$

Initial conditions

To fill in the \mathbf{P} matrix, we calculate the covariance of each combination of [state](#) variables. The resulting value is a measure of how much those variables are correlated. Due to how the covariance calculation works out, the covariance between two variables is the sum of the variance of matching

terms which aren't constants multiplied by any constants the two have. If no terms match, the variables are uncorrelated and the covariance is zero.

In \mathbf{P}_{11} , the terms in \mathbf{x}_1 correlate with itself. Therefore, \mathbf{P}_{11} is \mathbf{x}_1 's variance, or $\mathbf{P}_{11} = 10$. For \mathbf{P}_{21} , One term correlates between \mathbf{x}_1 and \mathbf{x}_2 , so $\mathbf{P}_{21} = \frac{10}{dt}$. The constants from each are simply multiplied together. For \mathbf{P}_{22} , both measurements are correlated, so the variances add together. Therefore, $\mathbf{P}_{22} = \frac{20}{dt^2}$. It continues in this fashion until the matrix is filled up. Order doesn't matter for correlation, so the matrix is symmetric.

Selection of priors

Choosing good priors is important for a well performing filter, even if little information is known. This applies to both the measurement noise and the noise [model](#). The act of giving a [state](#) variable a large variance means you know something about the [system](#). Namely, you aren't sure whether your initial guess is close to the true [state](#). If you make a guess and specify a small variance, you are telling the filter that you are very confident in your guess. If that guess is incorrect, it will take the filter a long time to move away from your guess to the true value.

Covariance selection

While one could assume no correlation between the [state](#) variables and set the covariance matrix entries to zero, this may not reflect reality. The Kalman filter is still guaranteed to converge to the steady-state covariance after an infinite time, but it will take longer than otherwise.

Noise model selection

We typically use a Gaussian distribution for the noise [model](#) because the sum of many independent random variables produces a normal distribution by the central limit theorem. Kalman filters only require that the noise is zero-mean. If the true value has an equal probability of being anywhere within a certain range, use a uniform distribution instead. Each of these communicates information regarding what you know about a system in addition to what you do not.

Process noise and measurement noise covariance selection

Recall that the process noise covariance is \mathbf{Q} and the measurement noise covariance is \mathbf{R} . To tune the elements of these, it can be helpful to take a collection of measurements, then run the Kalman filter on them offline to evaluate its performance.

The diagonal elements of \mathbf{R} are the variances of each measurement, which can be easily determined from the offline measurements. The diagonal elements of \mathbf{Q} are the variances of each [state](#). They represent how much each [state](#) is expected to deviate from the [model](#).

Selecting \mathbf{Q} is more difficult. If the data is trusted too much over the model, one risks overfitting the data. One should balance estimating any hidden [states](#) sufficiently with actually filtering out the noise.

Modeling other noise colors

The Kalman filter assumes a [model](#) with zero-mean white noise. If the [model](#) is incomplete in some way, whether it's missing dynamics or assumes an incorrect noise [model](#), the residual $\tilde{\mathbf{y}} = \mathbf{y} - \mathbf{C}\hat{\mathbf{x}}$ over time will have probability characteristics not indicative of white noise (e.g., it isn't zero-mean).

To handle other colors of noise in a Kalman filter, define that color of noise in terms of white noise and augment the [model](#) with it.

9.6.4 Simulation

Figure 9.3 shows the [state](#) estimates and measurements of the Kalman filter over time. Figure 9.4 shows the position estimate and variance over time. Figure 9.5 shows the wall position estimate

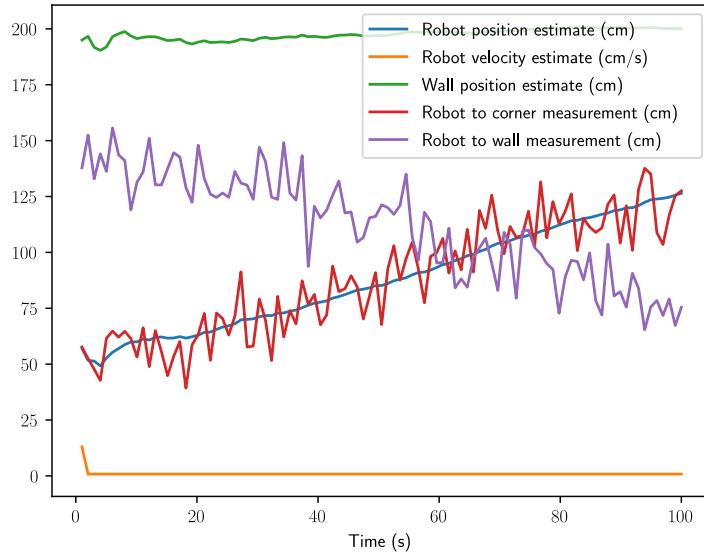


Figure 9.3: State estimates and measurements with Kalman filter

and variance over time. Notice how the variances decrease over time as the filter gathers more measurements. This means that the filter becomes more confident in its [state](#) estimates.

The final precisions in estimating the position of the robot and the wall are the square roots of the corresponding elements in the covariance matrix. That is, 0.5188 m and 0.4491 m respectively. They are smaller than the precision of the raw measurements, $\sqrt{10} = 3.1623 \text{ m}$. As expected, combining the information from several measurements produces a better estimate than any one measurement alone.

9.6.5 Kalman filter as Luenberger observer

A Kalman filter can be represented as a Luenberger [observer](#) by letting $\mathbf{L} = \mathbf{A}\mathbf{K}_k$ (see appendix H.3). The Luenberger observer has a constant observer gain matrix \mathbf{L} , so the steady-state Kalman gain is used to calculate it. We will demonstrate how to find this shortly.

Kalman filter theory provides a way to place the poles of the Luenberger observer optimally in the same way we placed the poles of the controller optimally with LQR. The eigenvalues of the Kalman filter are

$$\text{eig}(\mathbf{A}(\mathbf{I} - \mathbf{K}_k \mathbf{C})) \quad (9.34)$$

Steady-state Kalman gain

One may have noticed that the error covariance matrix can be updated independently of the rest of the [model](#). The error covariance matrix tends toward a steady-state value, and this matrix can be obtained via the discrete algebraic Riccati equation. This can then be used to compute a steady-state Kalman gain.

Snippet 9.1 computes the steady-state matrices for a Kalman filter.

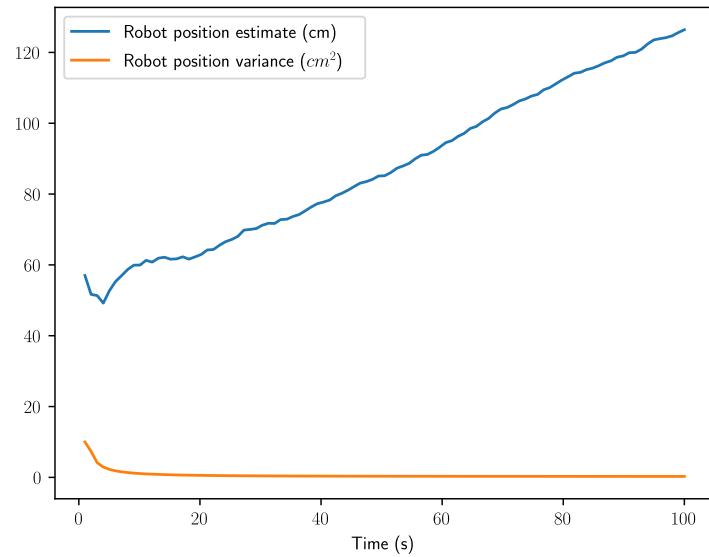


Figure 9.4: Robot position estimate and variance with Kalman filter

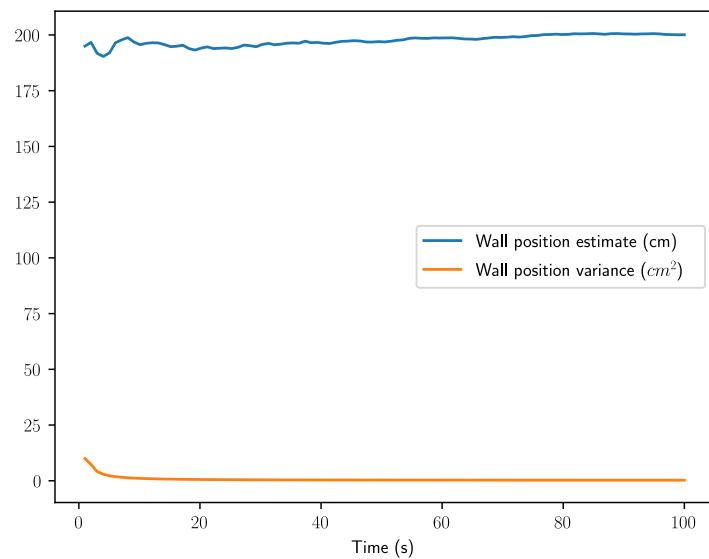


Figure 9.5: Wall position estimate and variance with Kalman filter

```

import control as ct
import numpy as np
import scipy as sp

def kalmd(sys, Q, R):
    """Solves for the steady state kalman gain and error covariance matrices.

    Keyword arguments:
    sys -- discrete state-space model
    Q -- process noise covariance matrix
    R -- measurement noise covariance matrix

    Returns:
    Kalman gain, error covariance matrix.
    """
    m = sys.A.shape[0]

    observability_rank = np.linalg.matrix_rank(ct.obsv(sys.A, sys.C))
    if observability_rank != m:
        print(
            "Warning: Observability of %d != %d, unobservable state"
            % (observability_rank, m)
        )

    # Compute the steady state covariance matrix
    P_prior = sp.linalg.solve_discrete_are(a=sys.A.T, b=sys.C.T, q=Q, r=R)
    S = sys.C * P_prior * sys.C.T + R
    K = P_prior * sys.C.T * np.linalg.inv(S)
    P = (np.eye(m) - K * sys.C) * P_prior

    return K, P

```

Snippet 9.1. Steady-state Kalman gain and error covariance matrices calculation in Python

9.7 Kalman smoother

The Kalman filter uses the data up to the current time to produce an optimal estimate of the system state. If data beyond the current time is available, it can be ran through a Kalman smoother to produce a better estimate. This is done by recording measurements, then applying the smoother to it offline.

The Kalman smoother does a forward pass on the available data, then a backward pass through the system dynamics so it takes into account the data before and after the current time. This produces state variances that are lower than that of a Kalman filter.

9.7.1 Derivations

9.7.2 State update equation

Let $\hat{\mathbf{x}}(t)$ be the state estimate from the forward pass based on samples 0 to t and $\hat{\mathbf{x}}_b(t)$ be the state estimate from the backward pass based on samples T to t .

$$\hat{\mathbf{x}}(t|T) = \mathbf{A}\hat{\mathbf{x}}(t) + \mathbf{A}'\hat{\mathbf{x}}_b(t) \quad (9.35)$$

where \mathbf{A} and \mathbf{A}' are weighting factors.

$$\hat{\mathbf{x}}(t|T) = \mathbf{A}\hat{\mathbf{x}}(t) + \mathbf{A}\tilde{\mathbf{x}}(t) + \mathbf{A}'\mathbf{x}(t) + \mathbf{A}'\tilde{\mathbf{x}}_b(t)$$

where $\tilde{\mathbf{x}}(t)$ represents the error in the forward state estimate and $\tilde{\mathbf{x}}_b(t)$ represents the error in the backward state estimate.

$$\begin{aligned}\mathbf{x}(t) + \tilde{\mathbf{x}}(t|T) &= \mathbf{A}\mathbf{x}(t) + \mathbf{A}\tilde{\mathbf{x}}(t) + \mathbf{A}'\mathbf{x}(t) + \mathbf{A}'\tilde{\mathbf{x}}_b(t) \\ \tilde{\mathbf{x}}(t|T) &= \mathbf{A}\mathbf{x}(t) - \mathbf{x}(t) + \mathbf{A}\tilde{\mathbf{x}}(t) + \mathbf{A}'\mathbf{x}(t) + \mathbf{A}'\tilde{\mathbf{x}}_b(t)\end{aligned}$$

Factor out $\mathbf{x}(t)$.

$$\tilde{\mathbf{x}}(t|T) = (\mathbf{A} + \mathbf{A}' - \mathbf{I})\mathbf{x}(t) + \mathbf{A}\tilde{\mathbf{x}}(t) + \mathbf{A}'\tilde{\mathbf{x}}_b(t)$$

For unbiased filtering errors such as $\tilde{\mathbf{x}}(t)$ and $\tilde{\mathbf{x}}_b(t)$, we want to have an unbiased smoothing error. Therefore, we set $\mathbf{A} + \mathbf{A}' - \mathbf{I}$ to zero. This yields $\mathbf{A}' = \mathbf{I} - \mathbf{A}$, so

$$\tilde{\mathbf{x}}(t|T) = \mathbf{A}\tilde{\mathbf{x}}(t) + (\mathbf{I} - \mathbf{A})\tilde{\mathbf{x}}_b(t)$$

9.7.3 Error covariance equation

Next, find the error covariance.

$$\mathbf{P}(t|T) = \mathbf{A}\mathbf{P}(t)\mathbf{A}^T + (\mathbf{I} - \mathbf{A})\mathbf{P}_b(t)(\mathbf{I} - \mathbf{A})^T \quad (9.36)$$

Find the minimum of the trace of $\mathbf{P}(t|T)$ by taking the partial derivative with respect to \mathbf{A} and setting the result to $\mathbf{0}$ ((t) has been dropped from covariance matrices for clarity).

$$\begin{aligned}\mathbf{0} &= 2\mathbf{A}\mathbf{P} + 2(\mathbf{I} - \mathbf{A})\mathbf{P}_b(-\mathbf{I}) \\ \mathbf{0} &= 2\mathbf{A}\mathbf{P} - 2(\mathbf{I} - \mathbf{A})\mathbf{P}_b \\ \mathbf{0} &= \mathbf{A}\mathbf{P} - (\mathbf{I} - \mathbf{A})\mathbf{P}_b \\ \mathbf{0} &= \mathbf{A}\mathbf{P} - (\mathbf{P}_b - \mathbf{A}\mathbf{P}_b) \\ \mathbf{0} &= \mathbf{A}\mathbf{P} - \mathbf{P}_b + \mathbf{A}\mathbf{P}_b \\ \mathbf{0} &= \mathbf{A}(\mathbf{P} + \mathbf{P}_b) - \mathbf{P}_b \\ \mathbf{A}(\mathbf{P} + \mathbf{P}_b) &= \mathbf{P}_b \\ \mathbf{A} &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1} \quad (9.37)\end{aligned}$$

$$\begin{aligned}\mathbf{I} - \mathbf{A} &= \mathbf{I} - \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1} \\ \mathbf{I} - \mathbf{A} &= (\mathbf{P} + \mathbf{P}_b)(\mathbf{P} + \mathbf{P}_b)^{-1} - \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1} \\ \mathbf{I} - \mathbf{A} &= (\mathbf{P} + \mathbf{P}_b - \mathbf{P}_b)(\mathbf{P} + \mathbf{P}_b)^{-1} \\ \mathbf{I} - \mathbf{A} &= \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1} \quad (9.38)\end{aligned}$$

$$\begin{aligned}\mathbf{A}^T &= (\mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1})^T \\ \mathbf{A}^T &= ((\mathbf{P} + \mathbf{P}_b)^{-1})^T \mathbf{P}_b^T \\ \mathbf{A}^T &= ((\mathbf{P} + \mathbf{P}_b)^T)^{-1} \mathbf{P}_b^T \\ \mathbf{A}^T &= (\mathbf{P}^T + \mathbf{P}_b^T)^{-1} \mathbf{P}_b^T\end{aligned}$$

Covariance matrices are symmetric, so

$$\mathbf{A}^T = (\mathbf{P} + \mathbf{P}_b)^{-1} \mathbf{P}_b \quad (9.39)$$

$$\begin{aligned}
(\mathbf{I} - \mathbf{A})^T &= (\mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1})^T \\
(\mathbf{I} - \mathbf{A})^T &= ((\mathbf{P} + \mathbf{P}_b)^{-1})^T \mathbf{P}^T \\
(\mathbf{I} - \mathbf{A})^T &= ((\mathbf{P} + \mathbf{P}_b)^T)^{-1} \mathbf{P}^T \\
(\mathbf{I} - \mathbf{A})^T &= (\mathbf{P}^T + \mathbf{P}_b^T)^{-1} \mathbf{P}^T
\end{aligned}$$

Covariance matrices are symmetric, so

$$(\mathbf{I} - \mathbf{A})^T = (\mathbf{P} + \mathbf{P}_b)^{-1} \mathbf{P} \quad (9.40)$$

Now starting from equation (9.36), substitute in equations (9.37) through (9.40).

$$\begin{aligned}
\mathbf{P}(t|T) &= \mathbf{A}\mathbf{P}\mathbf{A}^T + (\mathbf{I} - \mathbf{A})\mathbf{P}_b(\mathbf{I} - \mathbf{A})^T \\
\mathbf{P}(t|T) &= (\mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1})\mathbf{P}((\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}_b) + (\mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1})\mathbf{P}_b((\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}) \\
\mathbf{P}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}_b + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}
\end{aligned}$$

Apply theorem 9.7.1 to the right sides of each term to combine them.

Theorem 9.7.1 $\mathbf{AB} = (\mathbf{B}^{-1}\mathbf{A}^{-1})^{-1}$

$$\begin{aligned}
\mathbf{P}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}(\mathbf{P}_b^{-1}(\mathbf{P} + \mathbf{P}_b))^{-1} + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}_b(\mathbf{P}^{-1}(\mathbf{P} + \mathbf{P}_b))^{-1} \\
\mathbf{P}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}(\mathbf{P}_b^{-1}\mathbf{P} + \mathbf{I})^{-1} + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}_b(\mathbf{I} + \mathbf{P}^{-1}\mathbf{P}_b)^{-1}
\end{aligned}$$

Apply theorem 9.7.1 to the right sides of each term again.

$$\begin{aligned}
\mathbf{P}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}((\mathbf{P}_b^{-1}\mathbf{P} + \mathbf{I})\mathbf{P}^{-1})^{-1} + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}((\mathbf{I} + \mathbf{P}^{-1}\mathbf{P}_b)\mathbf{P}_b^{-1})^{-1} \\
\mathbf{P}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}(\mathbf{P}_b^{-1} + \mathbf{P}^{-1})^{-1} + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}(\mathbf{P}_b^{-1} + \mathbf{P}^{-1})^{-1} \\
\mathbf{P}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1} + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1}
\end{aligned}$$

Factor out $(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1}$ to the right.

$$\mathbf{P}(t|T) = (\mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1} + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1})(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1}$$

Factor out $(\mathbf{P} + \mathbf{P}_b)^{-1}$ to the right.

$$\begin{aligned}
\mathbf{P}(t|T) &= (\mathbf{P}_b + \mathbf{P})(\mathbf{P} + \mathbf{P}_b)^{-1}(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1} \\
\mathbf{P}(t|T) &= (\mathbf{P} + \mathbf{P}_b)(\mathbf{P} + \mathbf{P}_b)^{-1}(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1} \\
\mathbf{P}(t|T) &= I(\mathbf{P}^{-1} + \mathbf{P}_b^{-1})^{-1} \\
\mathbf{P}(t|T) &= (\mathbf{P}(t)^{-1} + \mathbf{P}_b(t)^{-1})^{-1}
\end{aligned} \quad (9.41)$$

9.7.4 Optimal estimate

Now find the optimal estimate $\hat{\mathbf{x}}(t|T)$ starting from equation (9.35).

$$\begin{aligned}
\hat{\mathbf{x}}(t|T) &= \mathbf{A}\hat{\mathbf{x}}(t) + \mathbf{A}'\hat{\mathbf{x}}_b(t) \\
\hat{\mathbf{x}}(t|T) &= \mathbf{A}\hat{\mathbf{x}}(t) + (\mathbf{I} - \mathbf{A})\hat{\mathbf{x}}_b(t)
\end{aligned}$$

$$\begin{aligned}\hat{\mathbf{x}}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}\hat{\mathbf{x}}(t) + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}\hat{\mathbf{x}}_b(t) \\ \hat{\mathbf{x}}(t|T) &= \mathbf{P}_b(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + \mathbf{P}(\mathbf{P} + \mathbf{P}_b)^{-1}\mathbf{P}_b\mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t)\end{aligned}$$

Apply theorem 9.7.1.

$$\begin{aligned}\hat{\mathbf{x}}(t|T) &= \mathbf{P}_b(\mathbf{P}^{-1}(\mathbf{P} + \mathbf{P}_b))^{-1}\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + \mathbf{P}(\mathbf{P}_b^{-1}(\mathbf{P} + \mathbf{P}_b))^{-1}\mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t) \\ \hat{\mathbf{x}}(t|T) &= \mathbf{P}_b(\mathbf{I} + \mathbf{P}^{-1}\mathbf{P}_b)^{-1}\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + \mathbf{P}(\mathbf{P}_b^{-1}\mathbf{P} + \mathbf{I})^{-1}\mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t)\end{aligned}$$

Apply theorem 9.7.1 again.

$$\begin{aligned}\hat{\mathbf{x}}(t|T) &= ((\mathbf{I} + \mathbf{P}^{-1}\mathbf{P}_b)\mathbf{P}_b^{-1})^{-1}\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + ((\mathbf{P}_b^{-1}\mathbf{P} + \mathbf{I})\mathbf{P}^{-1})^{-1}\mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t) \\ \hat{\mathbf{x}}(t|T) &= (\mathbf{P}_b^{-1} + \mathbf{P}^{-1})^{-1}\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + (\mathbf{P}_b^{-1} + \mathbf{P}^{-1})^{-1}\mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t)\end{aligned}$$

Factor out $(\mathbf{P}_b^{-1} + \mathbf{P}^{-1})^{-1}$ to the left.

$$\hat{\mathbf{x}}(t|T) = (\mathbf{P}_b^{-1} + \mathbf{P}^{-1})^{-1}(\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + \mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t))$$

Substitute in equation (9.41).

$$\begin{aligned}\hat{\mathbf{x}}(t|T) &= \mathbf{P}(t|T)(\mathbf{P}^{-1}\hat{\mathbf{x}}(t) + \mathbf{P}_b^{-1}\hat{\mathbf{x}}_b(t)) \\ \hat{\mathbf{x}}(t|T) &= \mathbf{P}(t|T)(\mathbf{P}^{-1}(t)\hat{\mathbf{x}}(t) + \mathbf{P}_b^{-1}(t)\hat{\mathbf{x}}_b(t))\end{aligned}\tag{9.42}$$

9.7.5 Predict and update equations

One first does a forward pass with the typical Kalman filter equations and stores the results. Then one can use the Rauch-Tung-Striebel (RTS) algorithm to do the backward pass (see theorem 9.7.2).

Theorem 9.7.2 shows the predict and update steps for the forward and backward passes for a Kalman smoother at the k^{th} timestep.

Theorem 9.7.2 — Kalman smoother.

Forward predict step

$$\hat{\mathbf{x}}_{k+1}^- = \mathbf{A}\hat{\mathbf{x}}_k + \mathbf{B}\mathbf{u}_k\tag{9.43}$$

$$\mathbf{P}_{k+1}^- = \mathbf{A}\mathbf{P}_k^-\mathbf{A}^T + \mathbf{\Gamma}\mathbf{Q}\mathbf{\Gamma}^T\tag{9.44}$$

Forward update step

$$\mathbf{K}_{k+1} = \mathbf{P}_{k+1}^-\mathbf{C}^T(\mathbf{C}\mathbf{P}_{k+1}^-\mathbf{C}^T + \mathbf{R})^{-1}\tag{9.45}$$

$$\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{K}_{k+1}(\mathbf{y}_{k+1} - \mathbf{C}\hat{\mathbf{x}}_{k+1}^- - \mathbf{D}\mathbf{u}_{k+1})\tag{9.46}$$

$$\mathbf{P}_{k+1}^+ = (\mathbf{I} - \mathbf{K}_{k+1}\mathbf{C})\mathbf{P}_{k+1}^-\tag{9.47}$$

Backward update step

$$\mathbf{K}_k = \mathbf{P}_k^+\mathbf{A}_k^T(\mathbf{P}_{k+1}^-)^{-1}\tag{9.48}$$

$$\hat{\mathbf{x}}_{k|N} = \hat{\mathbf{x}}_k^+ + \mathbf{K}_k(\hat{\mathbf{x}}_{k+1|N} - \hat{\mathbf{x}}_{k+1}^-)\tag{9.49}$$

$$\mathbf{P}_{k|N} = \mathbf{P}_k^+ + \mathbf{K}_k(\mathbf{P}_{k+1|N} - \mathbf{P}_{k+1}^-)\mathbf{K}_k^T\tag{9.50}$$

Backward initial conditions

$$\hat{\mathbf{x}}_{N|N} = \hat{\mathbf{x}}_N^+\tag{9.51}$$

$$\mathbf{P}_{N|N} = \mathbf{P}_N^+\tag{9.52}$$

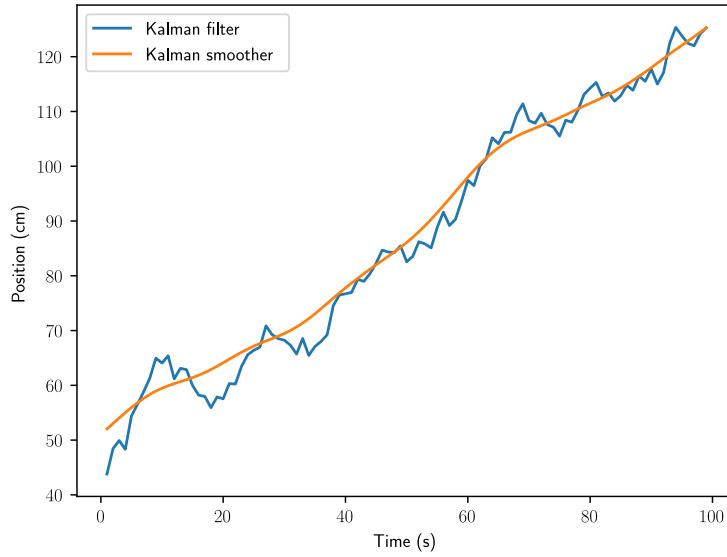


Figure 9.6: Robot position with Kalman smoother

9.7.6 Example

We will modify the robot model so that instead of a velocity of 0.8cm/s with random noise, the velocity is modeled as a random walk from the current velocity.

$$\mathbf{x}_k = \begin{bmatrix} x_k \\ v_k \\ x_k^w \end{bmatrix} \quad (9.53)$$

$$\mathbf{x}_{k+1} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \mathbf{x}_k + \begin{bmatrix} 0 \\ 0.1 \\ 0 \end{bmatrix} w_k \quad (9.54)$$

We will use the same observation model as before.

Using the same data from subsection 9.6.4, figures 9.6, 9.7, and 9.8 show the improved state estimates and figure 9.9 shows the improved robot position covariance with a Kalman smoother.

Notice how the wall position produced by the smoother is a constant. This is because that state has no dynamics, so the final estimate from the Kalman filter is already the best estimate.

See Roger Labbe's book *Kalman and Bayesian Filters in Python* for more on smoothing³.

³<https://github.com/rlabbe/Kalman-and-Bayesian-Filters-in-Python/blob/master/13-Smoothing.ipynb>

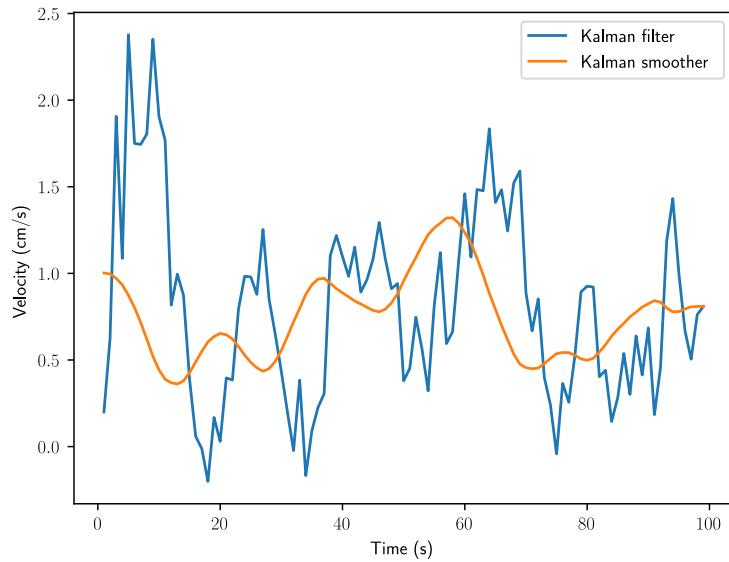


Figure 9.7: Robot velocity with Kalman smoother

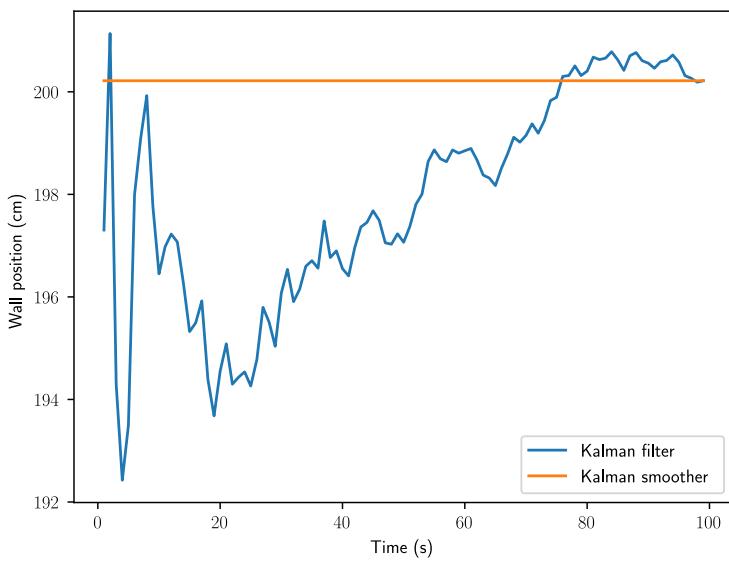


Figure 9.8: Wall position with Kalman smoother

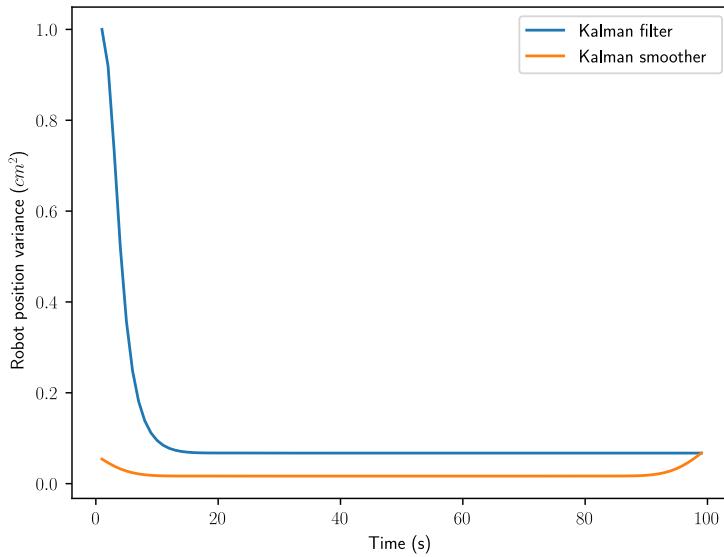


Figure 9.9: Robot position variance with Kalman smoother

9.8 Extended Kalman filter

In this book, we have covered the Kalman filter, which is the optimal unbiased estimator for linear [systems](#). It isn't optimal for nonlinear [systems](#), but several extensions to it have been developed to make it more accurate.

The extended Kalman filter [linearizes](#) the matrices used during the prediction step. **A**, **B**, **C**, and **D** are [linearized](#) as follows:

$$\mathbf{A} \approx \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} \quad \mathbf{B} \approx \frac{\partial f(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}} \quad \mathbf{C} \approx \frac{\partial h(\mathbf{x}, \mathbf{u})}{\partial \mathbf{x}} \quad \mathbf{D} \approx \frac{\partial h(\mathbf{x}, \mathbf{u})}{\partial \mathbf{u}}$$

From there, the continuous Kalman filter equations are used like normal to compute the error covariance matrix **P** and Kalman gain matrix. The [state](#) estimate update can still use the function $h(\mathbf{x})$ for accuracy.

$$\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{K}_{k+1}(\mathbf{y}_{k+1} - h(\hat{\mathbf{x}}_{k+1}^-))$$

9.9 Unscented Kalman filter

In this book, we have covered the Kalman filter, which is the optimal unbiased estimator for linear [systems](#). It isn't optimal for nonlinear [systems](#), but several extensions to it have been developed to make it more accurate.

The unscented Kalman filter propagates carefully chosen points called sigma points through the non-linear model to obtain an estimate of the true covariance (as opposed to a linearized version of it). We recommend reading Roger Labbe's book *Kalman and Bayesian Filters in Python* for more on unscented Kalman filters⁴.

⁴<https://github.com/rlabbe/Kalman-and-Bayesian-Filters-in-Python/blob/master/10-Unscented-Kalman-Filter.ipynb>

The original paper on the unscented Kalman filter is also an option [18]. The equations for van der Merwe's sigma point algorithm are in [31]. Here's a paper on a quaternion-based Unscented Kalman filter for orientation tracking [21].

9.10 Multiple model adaptive estimation

Multiple model adaptive estimation (MMAE) runs multiple Kalman filters with different [models](#) on the same data. The Kalman filter with the lowest residual has the highest likelihood of accurately reflecting reality. This can be used to detect certain [system states](#) like an aircraft engine failing without needing to invest in costly sensors to determine this directly.

For example, say you have three Kalman filters: one for turning left, one for turning right, and one for going straight. If the [control input](#) is attempting to fly the plane straight and the Kalman filter for going left has the lowest residual, the aircraft's left engine probably failed.

See Roger Labbe's book *Kalman and Bayesian Filters in Python* for more on MMAE⁵.

⁵MMAE section of <https://github.com/rlabbe/Kalman-and-Bayesian-Filters-in-Python/blob/master/14-Adaptive-Filtering.ipynb>

This page intentionally left blank

10. Pose estimation

Pose is a term for the orientation of an [agent](#) (a system with a controller). The plant usually includes the pose in its state vector. We'll cover several methods for estimating an agent's pose from local measurements such as encoders and gyroscope heading.

10.1 Euler integration

The simplest way to perform pose estimation via dead reckoning (that is, no direct measurements of the pose are used) is to integrate the velocity in each orthogonal direction over time. In two dimensions, one could use

$$\begin{aligned}x_{k+1} &= x_k + v_k \cos \theta_k T \\y_{k+1} &= y_k + v_k \sin \theta_k T \\\theta_{k+1} &= \theta_{gyro,k+1}\end{aligned}$$

where T is the sample period. This odometry approach assumes that the robot follows a straight path between samples (that is, $\omega = 0$ at all but the sample times).

10.2 Pose exponential

We can obtain a more accurate approximation of the pose by including first order dynamics for the heading θ . To provide a rationale for the math we're about to do, we need to cover some aspects of group theory.

10.2.1 What is a group?

In mathematics, a group is a set equipped with a binary operation (an operation with two arguments) that combines any two elements (of the set) to form a third element in such a way that four conditions called *group axioms* are satisfied: closure, associativity, identity, and invertibility.

Closure means that the result is in the same set as the arguments.

Associativity means that within an expression containing two or more occurrences in a row of the same associative operator, the order in which the operations are performed does not matter as long as the sequence of the operands is not changed. In other words, different groupings of the operations produces the same result.

Identity, or an identity element, is a special type of element of a set with respect to a binary operation on that set, which leaves any element of the set unchanged when combined with it. For example, the additive identity of the set of integers is zero, which means that any integer summed with zero produces that integer.

Invertibility means there is an element that can “undo” the effect of combination with another given element. For integers and the addition operator, the inverse element would be the negation.

10.2.2 What is a pose?

To develop what a pose is in group theory, we need to define a few key groups. $SO(2)$ is the special orthogonal group in dimension 2. They represent a 2D rotation.

$SE(2)$ is the special euclidean group in dimension 2. They represent a 2D rotation and a 2D translation, which we call a 2D pose. In other words, pose is an element of $SE(2)$.

10.2.3 What is a twist?

A 2D twist is an element of the tangent space of $SE(2)$ (like the tangential distance traveled by the robot along an arc in $SE(2)$). We use the “pose exponential” to map a twist (an element of the tangent space) to an element of $SE(2)$. In other words, we map a twist to a pose.

We call it a pose exponential because it’s an exponential map onto a pose. The term exponential is used because the solution for integrating a change in something (like a slope or tangent) over time is usually an exponential. For example, $\frac{dx}{dt} = ax$ has the solution $x = x_0 e^{at}$.

We use the pose exponential to take encoder measurement deltas and gyro angle deltas (which are in the tangent space and are thus a twist) and turn them into a change in pose. This gets added to the pose from the last update.

10.2.4 Derivation

We can obtain a more accurate approximation of the pose than Euler integration by including first order dynamics for the heading θ .

$$\mathbf{x} = \begin{bmatrix} x \\ y \\ \theta \end{bmatrix}$$

v_x , v_y , and ω are the x and y velocities of the robot within its local coordinate frame, which will be treated as constants.



There are two coordinate frames used here: robot and global. A superscript on the left side of a matrix denotes the coordinate frame in which that matrix is represented. The robot’s coordinate frame is denoted by R and the global coordinate frame is denoted by G .

In the robot frame (the tangent space)

$$\begin{aligned} {}^R dx &= {}^R v_x dt \\ {}^R dy &= {}^R v_y dt \end{aligned}$$

$${}^R d\theta = {}^R \omega dt$$

To transform this into the global frame SE(2), we apply a counterclockwise rotation matrix where θ changes over time.

$$\begin{aligned} {}^G \begin{bmatrix} dx \\ dy \\ d\theta \end{bmatrix} &= \begin{bmatrix} \cos \theta(t) & -\sin \theta(t) & 0 \\ \sin \theta(t) & \cos \theta(t) & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} dt \\ {}^G \begin{bmatrix} dx \\ dy \\ d\theta \end{bmatrix} &= \begin{bmatrix} \cos \omega t & -\sin \omega t & 0 \\ \sin \omega t & \cos \omega t & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} dt \end{aligned}$$

Now, integrate the matrix equation (matrices are integrated element-wise). This derivation heavily utilizes the integration method described in section 11.2.1.

$$\begin{aligned} {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \left[\begin{bmatrix} \frac{\sin \omega t}{\omega} & \frac{\cos \omega t}{\omega} & 0 \\ -\frac{\cos \omega t}{\omega} & \frac{\sin \omega t}{\omega} & 0 \\ 0 & 0 & t \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \right]_0^t \\ {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \left[\begin{bmatrix} \frac{\sin \omega t}{\omega} & \frac{\cos \omega t - 1}{\omega} & 0 \\ \frac{1 - \cos \omega t}{\omega} & \frac{\sin \omega t}{\omega} & 0 \\ 0 & 0 & t \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \right] \end{aligned}$$

This equation assumes a starting orientation of $\theta = 0$. For nonzero starting orientations, we can apply a counterclockwise rotation by θ .

$${}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{bmatrix} \frac{\sin \omega t}{\omega} & \frac{\cos \omega t - 1}{\omega} & 0 \\ \frac{1 - \cos \omega t}{\omega} & \frac{\sin \omega t}{\omega} & 0 \\ 0 & 0 & t \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \right] \quad (10.1)$$

If we factor out a t , we can use change in pose between updates instead of velocities.

$$\begin{aligned} {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{bmatrix} \frac{\sin \omega t}{\omega} & \frac{\cos \omega t - 1}{\omega} & 0 \\ \frac{1 - \cos \omega t}{\omega} & \frac{\sin \omega t}{\omega} & 0 \\ 0 & 0 & t \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \right] \\ {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{bmatrix} \frac{\sin \omega t}{\omega t} & \frac{\cos \omega t - 1}{\omega t} & 0 \\ \frac{1 - \cos \omega t}{\omega t} & \frac{\sin \omega t}{\omega t} & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \right] t \\ {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \left[\begin{bmatrix} \frac{\sin \omega t}{\omega t} & \frac{\cos \omega t - 1}{\omega t} & 0 \\ \frac{1 - \cos \omega t}{\omega t} & \frac{\sin \omega t}{\omega t} & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} v_x t \\ v_y t \\ \omega t \end{bmatrix} \right] \\ {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} \frac{\sin \Delta \theta}{\Delta \theta} & \frac{\cos \Delta \theta - 1}{\Delta \theta} & 0 \\ \frac{1 - \cos \Delta \theta}{\Delta \theta} & \frac{\sin \Delta \theta}{\Delta \theta} & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} \quad (10.2) \end{aligned}$$

The vector ${}^R [\Delta x \ \Delta y \ \Delta \theta]^T$ is a twist because it's an element of the tangent space (the robot's local coordinate frame).

Equation (10.1) used the current velocity and projected the model forward to the next timestep (into the future). As such, the prediction must be done after any controller calculations are performed. With equation (10.2), the locally measured pose deltas can only be measured using past samples, so the model update must be performed before any controller calculations to advance the model to the current timestep.

When the robot is traveling on a straight trajectory ($\omega = 0$), some expressions within the equation above are indeterminate. We can approximate these with Taylor series expansions.

$$\begin{aligned}\frac{\sin \omega t}{\omega} &= t - \frac{t^3 \omega^2}{6} + \dots \\ \frac{\sin \omega t}{\omega} &\sim t - \frac{t^3 \omega^2}{6} \\ \frac{\cos \omega t - 1}{\omega} &= -\frac{t^2 \omega}{2} + \frac{t^4 \omega^3}{4} - \dots \\ \frac{\cos \omega t - 1}{\omega} &\sim -\frac{t^2 \omega}{2} \\ \frac{1 - \cos \omega t}{\omega} &= \frac{t^2 \omega}{2} - \frac{t^4 \omega^3}{4} + \dots \\ \frac{1 - \cos \omega t}{\omega} &\sim \frac{t^2 \omega}{2}\end{aligned}$$

If we let $\omega = 0$, we should get the standard kinematic equations like $x = vt$ with a rotation applied to them.

$$\begin{aligned}\frac{\sin \omega t}{\omega} &\sim t - \frac{t^3 \cdot 0^2}{6} = t \\ \frac{\cos \omega t - 1}{\omega} &\sim -\frac{t^2 \cdot 0}{2} = 0 \\ \frac{1 - \cos \omega t}{\omega} &\sim \frac{t^2 \cdot 0}{2} = 0\end{aligned}$$

Now substitute these into equation (10.3).

$$\begin{aligned}{}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} t & 0 & 0 \\ 0 & t & 0 \\ 0 & 0 & t \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \\ {}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} &= \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} v_x t \\ v_y t \\ \omega t \end{bmatrix}\end{aligned}$$

As expected, the equations simplify to the first order case with a rotation matrix applied to the velocities in the robot's local coordinate frame.

Differential drive robots have $v_y = 0$ since they only move in the direction of the current heading, which is along the x-axis. Therefore,

$${}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} v_x t \\ 0 \\ \omega t \end{bmatrix}$$

$${}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} = \begin{bmatrix} v_x t \cos \theta \\ v_y t \sin \theta \\ \omega t \end{bmatrix}$$

Theorem 10.2.1 — Pose exponential.

$${}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \frac{\sin \omega t}{\omega} & \frac{\cos \omega t - 1}{\omega} & 0 \\ \frac{1 - \cos \omega t}{\omega} & \frac{\sin \omega t}{\omega} & 0 \\ 0 & 0 & t \end{bmatrix} {}^R \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \quad (10.3)$$

or

$${}^G \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} \frac{\sin \Delta \theta}{\Delta \theta} & \frac{\cos \Delta \theta - 1}{\Delta \theta} & 0 \\ \frac{1 - \cos \Delta \theta}{\Delta \theta} & \frac{\sin \Delta \theta}{\Delta \theta} & 0 \\ 0 & 0 & 1 \end{bmatrix} {}^R \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta \theta \end{bmatrix} \quad (10.4)$$

where G denotes global coordinate frame and R denotes robot's coordinate frame.

For sufficiently small ω :

$$\frac{\sin \omega t}{\omega} = t - \frac{t^3 \omega^2}{6} \quad \frac{\cos \omega t - 1}{\omega} = -\frac{t^2 \omega}{2} \quad \frac{1 - \cos \omega t}{\omega} = \frac{t^2 \omega}{2} \quad (10.5)$$

$$\frac{\sin \omega t}{\omega t} = 1 - \frac{t^2 \omega^2}{6} \quad \frac{\cos \omega t - 1}{\omega t} = -\frac{t \omega}{2} \quad \frac{1 - \cos \omega t}{\omega t} = \frac{t \omega}{2} \quad (10.6)$$

Δx	change in pose's x	v_x	velocity along x -axis
Δy	change in pose's y	v_y	velocity along y -axis
$\Delta \theta$	change in pose's θ	ω	angular velocity
t	Time since last pose update	θ	starting angle in global coordinate frame

This change in pose can be added directly to the previous pose estimate to update it.

Figure 10.1 shows the error in the global pose coordinates over time for the simpler odometry method compared to the method using twists (uses the Ramsete controller from subsection 8.6).

The highest error is 4cm in x over a 10m path starting with a 2m displacement. The difference would be even more noticeable for paths with higher curvatures and longer durations. One mitigation is using a smaller update period. However, there are bigger sources of error like turning scrub on skid steer robots that should be dealt with before odometry numerical accuracy.

10.2.5 Lie groups

While we avoided the topic in our explanation, pose is what's known as a Lie group (a group that is also a differentiable manifold). There's a lot of mathematical and controls results developed around Lie groups, so we're mentioning the connection here in case you want to search the Internet for more information.

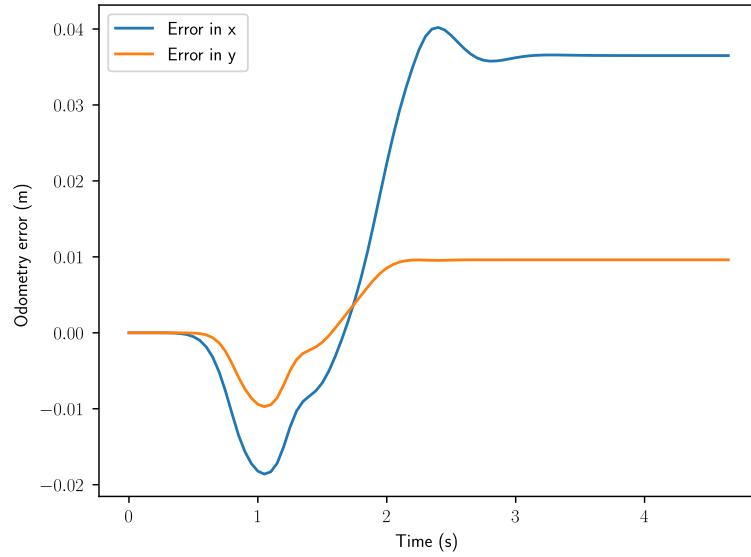


Figure 10.1: Odometry error compared to method using twists ($dt = 0.05s$)

10.3 Pose correction

The previous methods for pose estimation have assumed no direct pose measurements are available. At best, only heading is available. To augment any of them with corrections from full pose measurements, an Extended Kalman filter (section 9.8) or Unscented Kalman filter (section 9.9) can be used.

IV

System modeling

11	Calculus	131
11.1	Derivatives	
11.2	Integrals	
11.3	Tables	
12	Dynamics	135
12.1	Linear motion	
12.2	Angular motion	
12.3	Vectors	
12.4	Curvilinear motion	
13	Model examples	145
13.1	DC brushed motor	
13.2	Elevator	
13.3	Flywheel	
13.4	Single-jointed arm	
13.5	Pendulum	
13.6	Differential drive	
14	System identification	161
14.1	1 DOF mechanism model	
14.2	Drivetrain velocity model	

This page intentionally left blank

11. Calculus

This book uses derivatives and integrals occasionally to represent small changes in values over small changes in time and the infinitesimal sum of values over time respectively. The formulas and tables presented here are all you'll need to carry through with the math in later chapters.

If you are interested in more information after reading this chapter, 3Blue1Brown does a fantastic job of introducing them in his *Essence of calculus* video series [12]. We recommend watching videos 1 through 3 and 7 through 11 from that playlist for a solid foundation. The Taylor series (presented in video 11) will be used in chapter 6.

11.1 Derivatives

Derivatives are expressions for the slope of a curve at arbitrary points. Common notations for this operation on a function like $f(x)$ include

Leibniz notation	Lagrange notation
$\frac{d}{dx} f(x)$	$f'(x)$
$\frac{d^2}{dx^2} f(x)$	$f''(x)$
$\frac{d^3}{dx^3} f(x)$	$f'''(x)$
$\frac{d^4}{dx^4} f(x)$	$f^{(4)}(x)$
$\frac{d^n}{dx^n} f(x)$	$f^{(n)}(x)$

Table 11.1: Notation for derivatives of $f(x)$

Lagrange notation is usually voiced as “ f prime of x ”, “ f double-prime of x ”, etc.

11.1.1 Power rule

$$f(x) = x^n$$

$$f'(x) = nx^{n-1}$$

11.1.2 Product rule

This is for taking the derivative of the product of two expressions.

$$h(x) = f(x)g(x)$$

$$h'(x) = f'(x)g(x) + f(x)g'(x)$$

11.1.3 Chain rule

This is for taking the derivative of nested expressions.

$$h(x) = f(g(x))$$

$$h'(x) = f'(g(x)) \cdot g'(x)$$

For example

$$h(x) = (3x + 2)^5$$

$$h'(x) = 5(3x + 2)^4 \cdot (3)$$

$$h'(x) = 15(3x + 2)^4$$

11.2 Integrals

The integral is the inverse operation of the derivative and calculates the area under a curve. Here is an example of one based on table 11.2.

$$\int e^{at} dt$$

$$\frac{1}{a}e^{at} + C$$

The arbitrary constant C is needed because when you take a derivative, constants are discarded because vertical offsets don't affect the slope. When performing the inverse operation, we don't have enough information to determine the constant.

However, we can provide bounds for the integration.

$$\int_0^t e^{at} dt$$

$$\left(\frac{1}{a}e^{at} + C \right) \Big|_0^t$$

$$\left(\frac{1}{a}e^{at} + C \right) - \left(\frac{1}{a}e^{a \cdot 0} + C \right)$$

$$\left(\frac{1}{a}e^{at} + C \right) - \left(\frac{1}{a} + C \right)$$

$$\frac{1}{a}e^{at} + C - \frac{1}{a} - C$$

$$\frac{1}{a}e^{at} - \frac{1}{a}$$

When we do this, the constant cancels out.

11.2.1 Change of variables

Change of variables substitutes an expression with a single variable to make the calculation more straightforward. Here's an example of integration which utilizes it.

$$\int \cos \omega t \, dt$$

Let $u = \omega t$.

$$\begin{aligned} du &= \omega \, dt \\ dt &= \frac{1}{\omega} \, du \end{aligned}$$

Now substitute the expressions for u and dt in.

$$\begin{aligned} &\int \cos u \frac{1}{\omega} \, du \\ &\frac{1}{\omega} \int \cos u \, du \\ &\frac{1}{\omega} \sin u + C \\ &\frac{1}{\omega} \sin \omega t + C \end{aligned}$$

11.3 Tables

11.3.1 Common derivatives and integrals

$\int f(x) \, dx$	$f(x)$	$f'(x)$
ax	a	0
$\frac{1}{2}ax^2$	ax	a
$\frac{1}{a+1}x^{a+1}$	x^a	ax^{a-1}
$\frac{1}{a}e^{ax}$	e^{ax}	ae^{ax}
$-\cos(x)$	$\sin(x)$	$\cos(x)$
$\sin(x)$	$\cos(x)$	$-\sin(x)$
$\cos(x)$	$-\sin(x)$	$-\cos(x)$
$-\sin(x)$	$-\cos(x)$	$\sin(x)$

Table 11.2: Common derivatives and integrals

This page intentionally left blank



12. Dynamics

12.1 Linear motion

$$\sum F = ma$$

where $\sum F$ is the sum of all forces applied to an object in Newtons, m is the mass of the object in kg , and a is the net acceleration of the object in $\frac{m}{s^2}$.

$$x(t) = x_0 + v_0 t + \frac{1}{2} a t^2$$

where $x(t)$ is an object's position at time t , x_0 is the initial position, v_0 is the initial velocity, and a is the acceleration.

12.2 Angular motion

$$\sum \tau = I\alpha$$

where $\sum \tau$ is the sum of all torques applied to an object in Newton-meters, I is the moment of inertia of the object in $kg \cdot m^2$ (also called the rotational mass), and α is the net angular acceleration of the object in $\frac{rad}{s^2}$.

$$\theta(t) = \theta_0 + \omega_0 t + \frac{1}{2} \alpha t^2$$

where $\theta(t)$ is an object's angle at time t , θ_0 is the initial angle, ω_0 is the initial angular velocity, and α is the angular acceleration.

12.3 Vectors

Vectors are quantities with a magnitude and a direction. Vectors in three-dimensional space have a coordinate for each spatial direction x , y , and z . Let's take the vector $\vec{a} = \langle 1, 2, 3 \rangle$. \vec{a} is a three-dimensional vector that describes a movement 1 unit in the x direction, 2 units in the y direction, and 3 units in the z direction.

We define \hat{i} , \hat{j} , and \hat{k} as vectors that represent the fundamental movements one can make in the three-dimensional space: 1 unit of movement in the x direction, 1 unit of movement y direction, and 1 unit of movement in the z direction respectively. These three vectors form a *basis* of three-dimensional space because copies of them can be added together to reach any point in three-dimensional space.

$$\hat{i} = \langle 1, 0, 0 \rangle$$

$$\hat{j} = \langle 0, 1, 0 \rangle$$

$$\hat{k} = \langle 0, 0, 1 \rangle$$

We can also write the vector \vec{a} in terms of these basis vectors.

$$\vec{a} = 1\hat{i} + 2\hat{j} + 3\hat{k}$$

12.3.1 Basic vector operations

We will now show this is equivalent to the original notation through some vector mathematics. First, we'll substitute in the values for the basis vectors.

$$\vec{a} = 1\langle 1, 0, 0 \rangle + 2\langle 0, 1, 0 \rangle + 3\langle 0, 0, 1 \rangle$$

Scalars are multiplied component-wise with vectors.

$$\vec{a} = \langle 1, 0, 0 \rangle + \langle 0, 2, 0 \rangle + \langle 0, 0, 3 \rangle$$

Vectors are added by summing each of their components.

$$\vec{a} = \langle 1, 2, 3 \rangle$$

12.3.2 Cross product

The cross product is denoted by \times . The cross product of the basis vectors are computed as follows.

$$\begin{aligned}\hat{i} \times \hat{j} &= \hat{k} \\ \hat{j} \times \hat{k} &= \hat{i} \\ \hat{k} \times \hat{i} &= \hat{j}\end{aligned}$$

They proceed in a cyclic fashion through i , j , and k . If a vector is crossed with itself, it produces the zero vector (a scalar zero for each coordinate). The cross product of the basis vectors in the opposite order progress backwards and include a negative sign.

$$\begin{aligned}\hat{i} \times \hat{k} &= -\hat{j} \\ \hat{k} \times \hat{j} &= -\hat{i} \\ \hat{j} \times \hat{i} &= -\hat{k}\end{aligned}$$

Given vectors $\vec{u} = a\hat{i} + b\hat{j} + c\hat{k}$ and $\vec{v} = d\hat{i} + e\hat{j} + f\hat{k}$, $\vec{u} \times \vec{v}$ is computed using the distributive property.

$$\begin{aligned}
 \vec{u} \times \vec{v} &= (a\hat{i} + b\hat{j} + c\hat{k}) \times (d\hat{i} + e\hat{j} + f\hat{k}) \\
 \vec{u} \times \vec{v} &= ad(\hat{i} \times \hat{i}) + ae(\hat{i} \times \hat{j}) + af(\hat{i} \times \hat{k}) + \\
 &\quad bd(\hat{j} \times \hat{i}) + be(\hat{j} \times \hat{j}) + bf(\hat{j} \times \hat{k}) + \\
 &\quad cd(\hat{k} \times \hat{i}) + ce(\hat{k} \times \hat{j}) + cf(\hat{k} \times \hat{k}) \\
 \vec{u} \times \vec{v} &= ae\hat{k} + af(-\hat{j}) + bd(-\hat{k}) + b\hat{f}\hat{i} + cd\hat{j} + ce(-\hat{i}) \\
 \vec{u} \times \vec{v} &= ae\hat{k} - af\hat{j} - bd\hat{k} + b\hat{f}\hat{i} + cd\hat{j} - ce\hat{i} \\
 \vec{u} \times \vec{v} &= (bf - ce)\hat{i} + (cd - af)\hat{j} + (ae - bd)\hat{k}
 \end{aligned}$$

12.4 Curvilinear motion

For these derivations, we'll assume positive x (\hat{i}) is forward, positive y (\hat{j}) is to the left, positive z (\hat{k}) is up, and the robot is facing in the x direction. This axes convention is known as North-West-Up (NWU), and is shown in figure 12.1.

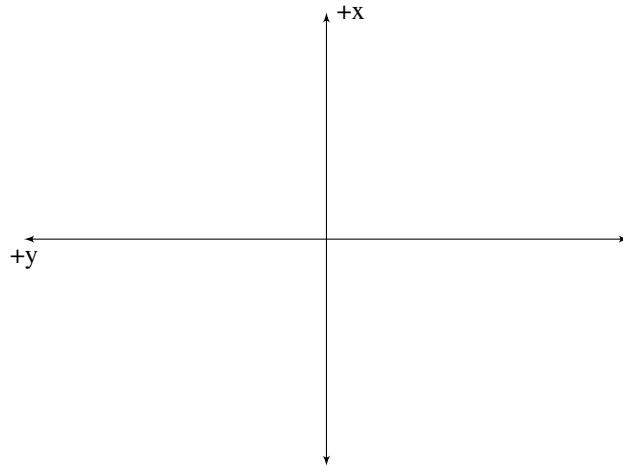


Figure 12.1: 2D projection of North-West-Up (NWU) axes convention. The positive z axis is pointed out of the page toward the reader.

The main equation we'll need is the following.

$$\vec{v}_B = \vec{v}_A + \omega_A \times \vec{r}_{B|A}$$

where \vec{v}_B is the velocity vector at point B, \vec{v}_A is the velocity vector at point A, ω_A is the angular velocity vector at point A, and $\vec{r}_{B|A}$ is the distance vector from point A to point B (also described as the “distance to B relative to A”).

12.4.1 Differential drive

A differential drive has two wheels, one on each side, separated by some distance $2r_b$. The forces they generate when moving forward are shown in figure 12.2.

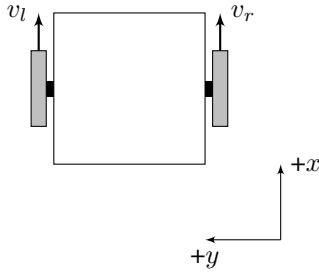


Figure 12.2: Differential drive free body diagram

Inverse kinematics

The mapping from v and ω to the left and right wheel velocities v_l and v_r is derived as follows. Let \vec{v}_c be the velocity vector of the center of rotation, \vec{v}_l be the velocity vector of the left wheel, \vec{v}_r be the velocity vector of the right wheel, r_b is the distance from the center of rotation to each wheel, and ω is the counterclockwise turning rate around the center of rotation.

Once we have the vector equation representing the wheel's velocity, we'll project it onto the wheel direction vector using the dot product.

First, we'll derive v_l .

$$\begin{aligned}\vec{v}_l &= v_c \hat{i} + \omega \hat{k} \times r_b \hat{j} \\ \vec{v}_l &= v_c \hat{i} - \omega r_b \hat{i} \\ \vec{v}_l &= (v_c - \omega r_b) \hat{i}\end{aligned}$$

Now, project this vector onto the left wheel, which is pointed in the \hat{i} direction.

$$v_l = (v_c - \omega r_b) \hat{i} \cdot \frac{\hat{i}}{\|\hat{i}\|}$$

The magnitude of \hat{i} is 1, so the denominator cancels.

$$\begin{aligned}v_l &= (v_c - \omega r_b) \hat{i} \cdot \hat{i} \\ v_l &= v_c - \omega r_b\end{aligned} \tag{12.1}$$

Next, we'll derive v_r .

$$\begin{aligned}\vec{v}_r &= v_c \hat{i} + \omega \hat{k} \times r_b \hat{j} \\ \vec{v}_r &= v_c \hat{i} + \omega r_b \hat{i} \\ \vec{v}_r &= (v_c + \omega r_b) \hat{i}\end{aligned}$$

Now, project this vector onto the right wheel, which is pointed in the \hat{i} direction.

$$v_r = (v_c + \omega r_b) \hat{i} \cdot \frac{\hat{i}}{\|\hat{i}\|}$$

The magnitude of \hat{i} is 1, so the denominator cancels.

$$v_r = (v_c + \omega r_b) \hat{i} \cdot \hat{i}$$

$$v_r = v_c + \omega r_b \quad (12.2)$$

So the two inverse kinematic equations are as follows.

$$v_l = v_c - \omega r_b \quad (12.3)$$

$$v_r = v_c + \omega r_b \quad (12.4)$$

Forward kinematics

The mapping from the left and right wheel velocities v_l and v_r to v and ω is derived as follows.

$$\begin{aligned} v_r &= v_c + \omega r_b \\ v_c &= v_r - \omega r_b \end{aligned} \quad (12.5)$$

Substitute equation (12.5) equation for v_l .

$$\begin{aligned} v_l &= v_c - \omega r_b \\ v_l &= (v_r - \omega r_b) - \omega r_b \\ v_l &= v_r - 2\omega r_b \\ 2\omega r_b &= v_r - v_l \\ \omega &= \frac{v_r - v_l}{2r_b} \end{aligned}$$

Substitute this back into equation (12.5).

$$\begin{aligned} v_c &= v_r - \omega r_b \\ v_c &= v_r - \left(\frac{v_r - v_l}{2r_b} \right) r_b \\ v_c &= v_r - \frac{v_r - v_l}{2} \\ v_c &= v_r - \frac{v_r}{2} + \frac{v_l}{2} \\ v_c &= \frac{v_r + v_l}{2} \end{aligned}$$

So the two forward kinematic equations are as follows.

$$v_c = \frac{v_r + v_l}{2} \quad (12.6)$$

$$\omega = \frac{v_r - v_l}{2r_b} \quad (12.7)$$

12.4.2 Mecanum drive

A mecanum drive has four wheels, one on each corner of a rectangular chassis. The wheels have rollers offset at 45 degrees (whether it's clockwise or not varies per wheel). The forces they generate when moving forward are shown in figure 12.3.

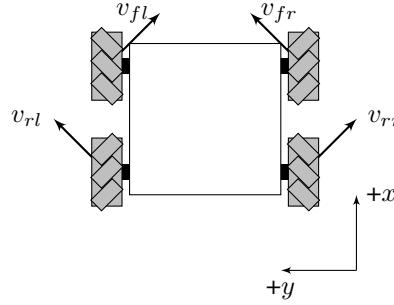


Figure 12.3: Mecanum drive free body diagram

Inverse kinematics

First, we'll derive the front-left wheel kinematics.

$$\begin{aligned}\vec{v}_{fl} &= v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{fl_x} \hat{i} + r_{fl_y} \hat{j}) \\ \vec{v}_{fl} &= v_x \hat{i} + v_y \hat{j} + \omega r_{fl_x} \hat{j} - \omega r_{fl_y} \hat{i} \\ \vec{v}_{fl} &= (v_x - \omega r_{fl_y}) \hat{i} + (v_y + \omega r_{fl_x}) \hat{j}\end{aligned}$$

Project the front-left wheel onto its wheel vector.

$$\begin{aligned}v_{fl} &= ((v_x - \omega r_{fl_y}) \hat{i} + (v_y + \omega r_{fl_x}) \hat{j}) \cdot \frac{\hat{i} - \hat{j}}{\sqrt{2}} \\ v_{fl} &= ((v_x - \omega r_{fl_y}) - (v_y + \omega r_{fl_x})) \frac{1}{\sqrt{2}} \\ v_{fl} &= (v_x - \omega r_{fl_y} - v_y - \omega r_{fl_x}) \frac{1}{\sqrt{2}} \\ v_{fl} &= (v_x - v_y - \omega r_{fl_y} - \omega r_{fl_x}) \frac{1}{\sqrt{2}} \\ v_{fl} &= (v_x - v_y - \omega(r_{fl_x} + r_{fl_y})) \frac{1}{\sqrt{2}} \\ v_{fl} &= \frac{1}{\sqrt{2}} v_x - \frac{1}{\sqrt{2}} v_y - \frac{1}{\sqrt{2}} \omega(r_{fl_x} + r_{fl_y})\end{aligned}\tag{12.8}$$

Next, we'll derive the front-right wheel kinematics.

$$\begin{aligned}\vec{v}_{fr} &= v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{fr_x} \hat{i} + r_{fr_y} \hat{j}) \\ \vec{v}_{fr} &= v_x \hat{i} + v_y \hat{j} + \omega r_{fr_x} \hat{i} - \omega r_{fr_y} \hat{j} \\ \vec{v}_{fr} &= (v_x - \omega r_{fr_y}) \hat{i} + (v_y + \omega r_{fr_x}) \hat{j}\end{aligned}$$

Project the front-right wheel onto its wheel vector.

$$\begin{aligned}v_{fr} &= ((v_x - \omega r_{fr_y}) \hat{i} + (v_y + \omega r_{fr_x}) \hat{j}) \cdot (\hat{i} + \hat{j}) \frac{1}{\sqrt{2}} \\ v_{fr} &= ((v_x - \omega r_{fr_y}) + (v_y + \omega r_{fr_x})) \frac{1}{\sqrt{2}} \\ v_{fr} &= (v_x - \omega r_{fr_y} + v_y + \omega r_{fr_x}) \frac{1}{\sqrt{2}}\end{aligned}$$

$$\begin{aligned}
v_{fr} &= (v_x + v_y - \omega r_{fr_y} + \omega r_{fr_x}) \frac{1}{\sqrt{2}} \\
v_{fr} &= (v_x + v_y + \omega(r_{fr_x} - r_{fr_y})) \frac{1}{\sqrt{2}} \\
v_{fr} &= \frac{1}{\sqrt{2}}v_x + \frac{1}{\sqrt{2}}v_y + \frac{1}{\sqrt{2}}\omega(r_{fr_x} - r_{fr_y})
\end{aligned} \tag{12.9}$$

Next, we'll derive the rear-left wheel kinematics.

$$\begin{aligned}
\vec{v}_{rl} &= v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{rl_x} \hat{i} + r_{rl_y} \hat{j}) \\
\vec{v}_{rl} &= v_x \hat{i} + v_y \hat{j} + \omega r_{rl_x} \hat{j} - \omega r_{rl_y} \hat{i} \\
\vec{v}_{rl} &= (v_x - \omega r_{rl_y}) \hat{i} + (v_y + \omega r_{rl_x}) \hat{j}
\end{aligned}$$

Project the rear-left wheel onto its wheel vector.

$$\begin{aligned}
v_{rl} &= ((v_x - \omega r_{rl_y}) \hat{i} + (v_y + \omega r_{rl_x}) \hat{j}) \cdot (\hat{i} + \hat{j}) \frac{1}{\sqrt{2}} \\
v_{rl} &= ((v_x - \omega r_{rl_y}) + (v_y + \omega r_{rl_x})) \frac{1}{\sqrt{2}} \\
v_{rl} &= (v_x - \omega r_{rl_y} + v_y + \omega r_{rl_x}) \frac{1}{\sqrt{2}} \\
v_{rl} &= (v_x + v_y - \omega r_{rl_y} + \omega r_{rl_x}) \frac{1}{\sqrt{2}} \\
v_{rl} &= (v_x + v_y + \omega(r_{rl_x} - r_{rl_y})) \frac{1}{\sqrt{2}} \\
v_{rl} &= \frac{1}{\sqrt{2}}v_x + \frac{1}{\sqrt{2}}v_y + \frac{1}{\sqrt{2}}\omega(r_{rl_x} - r_{rl_y})
\end{aligned} \tag{12.10}$$

Next, we'll derive the rear-right wheel kinematics.

$$\begin{aligned}
\vec{v}_{rr} &= v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{rr_x} \hat{i} + r_{rr_y} \hat{j}) \\
\vec{v}_{rr} &= v_x \hat{i} + v_y \hat{j} + \omega r_{rr_x} \hat{j} - \omega r_{rr_y} \hat{i} \\
\vec{v}_{rr} &= (v_x - \omega r_{rr_y}) \hat{i} + (v_y + \omega r_{rr_x}) \hat{j}
\end{aligned}$$

Project the rear-right wheel onto its wheel vector.

$$\begin{aligned}
v_{rr} &= ((v_x - \omega r_{rr_y}) \hat{i} + (v_y + \omega r_{rr_x}) \hat{j}) \cdot \frac{\hat{i} - \hat{j}}{\sqrt{2}} \\
v_{rr} &= ((v_x - \omega r_{rr_y}) - (v_y + \omega r_{rr_x})) \frac{1}{\sqrt{2}} \\
v_{rr} &= (v_x - \omega r_{rr_y} - v_y - \omega r_{rr_x}) \frac{1}{\sqrt{2}} \\
v_{rr} &= (v_x - v_y - \omega r_{rr_y} - \omega r_{rr_x}) \frac{1}{\sqrt{2}} \\
v_{rr} &= (v_x - v_y - \omega(r_{rr_x} + r_{rr_y})) \frac{1}{\sqrt{2}} \\
v_{rr} &= \frac{1}{\sqrt{2}}v_x - \frac{1}{\sqrt{2}}v_y - \frac{1}{\sqrt{2}}\omega(r_{rr_x} + r_{rr_y})
\end{aligned} \tag{12.11}$$

This gives the following inverse kinematic equations.

$$\begin{aligned} v_{fl} &= \frac{1}{\sqrt{2}}v_x - \frac{1}{\sqrt{2}}v_y - \frac{1}{\sqrt{2}}\omega(r_{fl_x} + r_{fl_y}) \\ v_{fr} &= \frac{1}{\sqrt{2}}v_x + \frac{1}{\sqrt{2}}v_y + \frac{1}{\sqrt{2}}\omega(r_{fr_x} - r_{fr_y}) \\ v_{rl} &= \frac{1}{\sqrt{2}}v_x + \frac{1}{\sqrt{2}}v_y + \frac{1}{\sqrt{2}}\omega(r_{rl_x} - r_{rl_y}) \\ v_{rr} &= \frac{1}{\sqrt{2}}v_x - \frac{1}{\sqrt{2}}v_y - \frac{1}{\sqrt{2}}\omega(r_{rr_x} + r_{rr_y}) \end{aligned}$$

Now we'll factor them out into matrices.

$$\begin{aligned} \begin{bmatrix} v_{fl} \\ v_{fr} \\ v_{rl} \\ v_{rr} \end{bmatrix} &= \begin{bmatrix} \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}}(r_{fl_x} + r_{fl_y}) \\ \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}}(r_{fr_x} - r_{fr_y}) \\ \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}}(r_{rl_x} - r_{rl_y}) \\ \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}}(r_{rr_x} + r_{rr_y}) \end{bmatrix} \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \\ \begin{bmatrix} v_{fl} \\ v_{fr} \\ v_{rl} \\ v_{rr} \end{bmatrix} &= \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & -1 & -(r_{fl_x} + r_{fl_y}) \\ 1 & 1 & (r_{fr_x} - r_{fr_y}) \\ 1 & 1 & (r_{rl_x} - r_{rl_y}) \\ 1 & -1 & -(r_{rr_x} + r_{rr_y}) \end{bmatrix} \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \end{aligned} \quad (12.12)$$

Forward kinematics

Let \mathbf{M} be the 4×3 inverse kinematics matrix above including the $\frac{1}{\sqrt{2}}$ factor. The forward kinematics are

$$\begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} = \mathbf{M}^\dagger \begin{bmatrix} v_{fl} \\ v_{fr} \\ v_{rl} \\ v_{rr} \end{bmatrix} \quad (12.13)$$

where \mathbf{M}^\dagger is the pseudoinverse of \mathbf{M} .

12.4.3 Swerve drive

A swerve drive has an arbitrary number of wheels which can rotate in place independent of the chassis. The forces they generate are shown in figure 12.4.

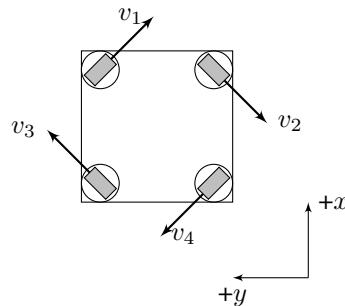


Figure 12.4: Swerve drive free body diagram

Inverse kinematics

$$\vec{v}_{wheel1} = \vec{v}_{robot} + \vec{\omega}_{robot} \times \vec{r}_{robot2wheel1}$$

$$\vec{v}_{wheel2} = \vec{v}_{robot} + \vec{\omega}_{robot} \times \vec{r}_{robot2wheel2}$$

$$\vec{v}_{wheel3} = \vec{v}_{robot} + \vec{\omega}_{robot} \times \vec{r}_{robot2wheel3}$$

$$\vec{v}_{wheel4} = \vec{v}_{robot} + \vec{\omega}_{robot} \times \vec{r}_{robot2wheel4}$$

where \vec{v} is linear velocity vector, $\vec{\omega}$ is angular velocity vector, \vec{r} is distance vector from the center of rotation to the wheel, $\vec{v}_{robot} = v_x \hat{i} + v_y \hat{j}$, and $\vec{r}_{robot2wheel} = r_x \hat{i} + r_y \hat{j}$.

$$\vec{v}_1 = v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{1x} \hat{i} + r_{1y} \hat{j})$$

$$\vec{v}_2 = v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{2x} \hat{i} + r_{2y} \hat{j})$$

$$\vec{v}_3 = v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{3x} \hat{i} + r_{3y} \hat{j})$$

$$\vec{v}_4 = v_x \hat{i} + v_y \hat{j} + \omega \hat{k} \times (r_{4x} \hat{i} + r_{4y} \hat{j})$$

$$\vec{v}_1 = v_x \hat{i} + v_y \hat{j} + (\omega r_{1x} \hat{j} - \omega r_{1y} \hat{i})$$

$$\vec{v}_2 = v_x \hat{i} + v_y \hat{j} + (\omega r_{2x} \hat{j} - \omega r_{2y} \hat{i})$$

$$\vec{v}_3 = v_x \hat{i} + v_y \hat{j} + (\omega r_{3x} \hat{j} - \omega r_{3y} \hat{i})$$

$$\vec{v}_4 = v_x \hat{i} + v_y \hat{j} + (\omega r_{4x} \hat{j} - \omega r_{4y} \hat{i})$$

$$\vec{v}_1 = v_x \hat{i} + v_y \hat{j} + \omega r_{1x} \hat{j} - \omega r_{1y} \hat{i}$$

$$\vec{v}_2 = v_x \hat{i} + v_y \hat{j} + \omega r_{2x} \hat{j} - \omega r_{2y} \hat{i}$$

$$\vec{v}_3 = v_x \hat{i} + v_y \hat{j} + \omega r_{3x} \hat{j} - \omega r_{3y} \hat{i}$$

$$\vec{v}_4 = v_x \hat{i} + v_y \hat{j} + \omega r_{4x} \hat{j} - \omega r_{4y} \hat{i}$$

$$\vec{v}_1 = v_x \hat{i} - \omega r_{1y} \hat{i} + v_y \hat{j} + \omega r_{1x} \hat{j}$$

$$\vec{v}_2 = v_x \hat{i} - \omega r_{2y} \hat{i} + v_y \hat{j} + \omega r_{2x} \hat{j}$$

$$\vec{v}_3 = v_x \hat{i} - \omega r_{3y} \hat{i} + v_y \hat{j} + \omega r_{3x} \hat{j}$$

$$\vec{v}_4 = v_x \hat{i} - \omega r_{4y} \hat{i} + v_y \hat{j} + \omega r_{4x} \hat{j}$$

$$\vec{v}_1 = (v_x - \omega r_{1y}) \hat{i} + (v_y + \omega r_{1x}) \hat{j}$$

$$\vec{v}_2 = (v_x - \omega r_{2y}) \hat{i} + (v_y + \omega r_{2x}) \hat{j}$$

$$\vec{v}_3 = (v_x - \omega r_{3y}) \hat{i} + (v_y + \omega r_{3x}) \hat{j}$$

$$\vec{v}_4 = (v_x - \omega r_{4y}) \hat{i} + (v_y + \omega r_{4x}) \hat{j}$$

Separate the \hat{i} components into independent equations.

$$v_{1x} = v_x - \omega r_{1y}$$

$$v_{2x} = v_x - \omega r_{2y}$$

$$v_{3x} = v_x - \omega r_{3y}$$

$$v_{4x} = v_x - \omega r_{4y}$$

Separate the \mathbf{j} -hat components into independent equations.

$$\begin{aligned} v_{1y} &= v_y + \omega r_{1x} \\ v_{2y} &= v_y + \omega r_{2x} \\ v_{3y} &= v_y + \omega r_{3x} \\ v_{4y} &= v_y + \omega r_{4x} \end{aligned}$$

Now we'll factor them out into matrices.

$$\begin{bmatrix} v_{1x} \\ v_{2x} \\ v_{3x} \\ v_{4x} \\ v_{1y} \\ v_{2y} \\ v_{3y} \\ v_{4y} \end{bmatrix} = \begin{bmatrix} 1 & 0 & -r_{1y} \\ 1 & 0 & -r_{2y} \\ 1 & 0 & -r_{3y} \\ 1 & 0 & -r_{4y} \\ 0 & 1 & r_{1x} \\ 0 & 1 & r_{2x} \\ 0 & 1 & r_{3x} \\ 0 & 1 & r_{4x} \end{bmatrix} \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix}$$

Rearrange the rows so the x and y components are in adjacent rows.

$$\begin{bmatrix} v_{1x} \\ v_{1y} \\ v_{2x} \\ v_{2y} \\ v_{3x} \\ v_{3y} \\ v_{4x} \\ v_{4y} \end{bmatrix} = \begin{bmatrix} 1 & 0 & -r_{1y} \\ 0 & 1 & r_{1x} \\ 1 & 0 & -r_{2y} \\ 0 & 1 & r_{2x} \\ 1 & 0 & -r_{3y} \\ 0 & 1 & r_{3x} \\ 1 & 0 & -r_{4y} \\ 0 & 1 & r_{4x} \end{bmatrix} \begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} \quad (12.14)$$

To convert from swerve module x and y velocity components to a velocity and heading, use the Pythagorean theorem and arctangent respectively, as shown below.

$$v = \sqrt{v_x^2 + v_y^2} \quad (12.15)$$

$$\theta = \tan^{-1} \left(\frac{v_y}{v_x} \right) \quad (12.16)$$

Forward kinematics

Let \mathbf{M} be the 8×3 inverse kinematics matrix above. The forward kinematics are

$$\begin{bmatrix} v_x \\ v_y \\ \omega \end{bmatrix} = \mathbf{M}^\dagger \begin{bmatrix} v_{1x} \\ v_{1y} \\ v_{2x} \\ v_{2y} \\ v_{3x} \\ v_{3y} \\ v_{4x} \\ v_{4y} \end{bmatrix} \quad (12.17)$$

where \mathbf{M}^\dagger is the pseudoinverse of \mathbf{M} .

13. Model examples

A **model** is a set of differential equations describing how the **system** behaves over time. There are two common approaches for developing them.

1. Collecting data on the physical system's behavior and performing **system** identification with it.
2. Using physics to derive the **system**'s model from first principles.

We'll use the second approach in this book.

The **models** derived here should cover most types of motion seen on an FRC robot. Furthermore, they can be easily tweaked to describe many types of mechanisms just by pattern-matching. There's only so many ways to hook up a mass to a motor in FRC. The flywheel **model** can be used for spinning mechanisms, the elevator **model** can be used for spinning mechanisms transformed to linear motion, and the single-jointed arm **model** can be used for rotating servo mechanisms (it's just the flywheel **model** augmented with a position **state**).

These **models** assume all motor controllers driving DC brushed motors are set to brake mode instead of coast mode. Brake mode behaves the same as coast mode except where the applied voltage is zero. In brake mode, the motor leads are shorted together to prevent movement. In coast mode, the motor leads are an open circuit.

13.1 DC brushed motor

We will be deriving a first-order **model** for a DC brushed motor. A second-order **model** would include the inductance of the motor windings as well, but we're assuming the time constant of the inductor is small enough that its affect on the **model** behavior is negligible for FRC use cases (see subsection 3.3.1 for a demonstration of this for a real DC brushed motor).

The first-order **model** will only require numbers from the motor's datasheet. The second-order **model** would require measuring the motor inductance as well, which generally isn't in the datasheet. It can be difficult to measure accurately without the right equipment.

13.1.1 Equations of motion

The circuit for a DC brushed motor is shown in figure 13.1.

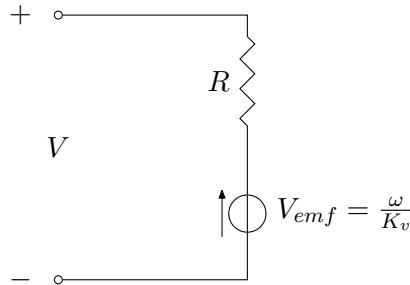


Figure 13.1: DC brushed motor circuit

V is the voltage applied to the motor, I is the current through the motor in Amps, R is the resistance across the motor in Ohms, ω is the angular velocity of the motor in radians per second, and K_v is the angular velocity constant in radians per second per Volt. This circuit reflects the following relation.

$$V = IR + \frac{\omega}{K_v} \quad (13.1)$$

The mechanical relation for a DC brushed motor is

$$\tau = K_t I \quad (13.2)$$

where τ is the torque produced by the motor in Newton-meters and K_t is the torque constant in Newton-meters per Amp. Therefore

$$I = \frac{\tau}{K_t}$$

Substitute this into equation (13.1).

$$V = \frac{\tau}{K_t} R + \frac{\omega}{K_v} \quad (13.3)$$

13.1.2 Calculating constants

A typical motor's datasheet should include graphs of the motor's measured torque and current for different angular velocities for a given voltage applied to the motor. Figure 13.2 is an example. Data for the most common motors in FRC can be found at <https://motors.vex.com>.

Torque constant K_t

$$\begin{aligned} \tau &= K_t I \\ K_t &= \frac{\tau}{I} \\ K_t &= \frac{\tau_{stall}}{I_{stall}} \end{aligned} \quad (13.4)$$

where τ_{stall} is the stall torque and I_{stall} is the stall current of the motor from its datasheet.

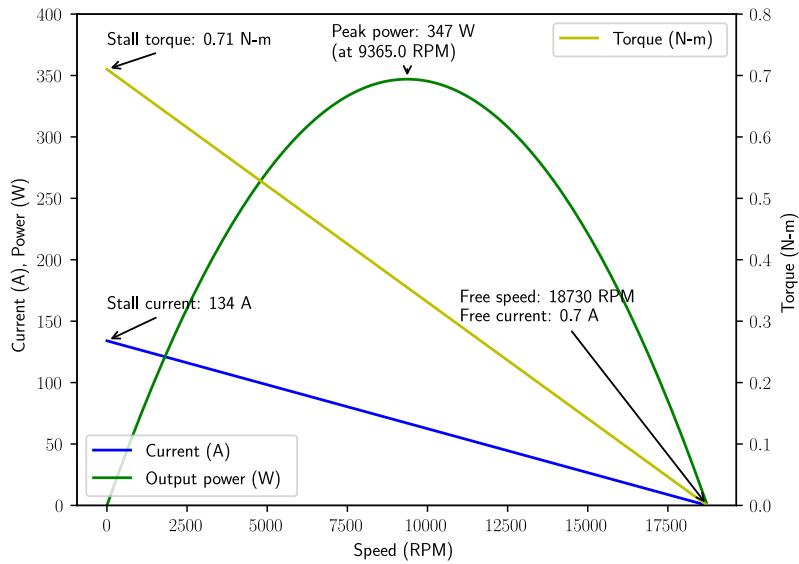


Figure 13.2: Example motor datasheet for 775pro

Resistance R

Recall equation (13.1).

$$V = IR + \frac{\omega}{K_v}$$

When the motor is stalled, $\omega = 0$.

$$\begin{aligned} V &= I_{stall}R \\ R &= \frac{V}{I_{stall}} \end{aligned} \tag{13.5}$$

where I_{stall} is the stall current of the motor and V is the voltage applied to the motor at stall.

Angular velocity constant K_v

Recall equation (13.1).

$$\begin{aligned} V &= IR + \frac{\omega}{K_v} \\ V - IR &= \frac{\omega}{K_v} \\ K_v &= \frac{\omega}{V - IR} \end{aligned}$$

When the motor is spinning under no load

$$K_v = \frac{\omega_{free}}{V - I_{free}R} \tag{13.6}$$

where ω_{free} is the angular velocity of the motor under no load (also known as the free speed), and V is the voltage applied to the motor when it's spinning at ω_{free} , and I_{free} is the current drawn by the motor under no load.

If several identical motors are being used in one gearbox for a mechanism, multiply the stall torque by the number of motors.

13.1.3 Current limiting

Current limiting of a DC brushed motor reduces the maximum input voltage to avoid exceeding a current threshold. Predictive current limiting uses a projected estimate of the current, so the voltage is reduced before the current threshold is exceeded. Reactive current limiting uses an actual current measurement, so the voltage is reduced after the current threshold is exceeded.

The following pseudocode demonstrates each type of current limiting.

```

# Normal feedback control
V = K @ (r - x)

# Calculations for predictive current limiting
omega = angular_velocity_measurement
I = V / R - omega / (Kv * R)

# Calculations for reactive current limiting
I = current_measurement
omega = Kv * V - I * R * Kv # or can be angular velocity measurement

# If predicted/actual current above max, limit current by reducing voltage
if I > I_max:
    V = I_max * R + omega / Kv

```

Snippet 13.1. Limits current of DC motor to I_{max}

13.2 Elevator

13.2.1 Equations of motion

This elevator consists of a DC brushed motor attached to a pulley that drives a mass up or down.

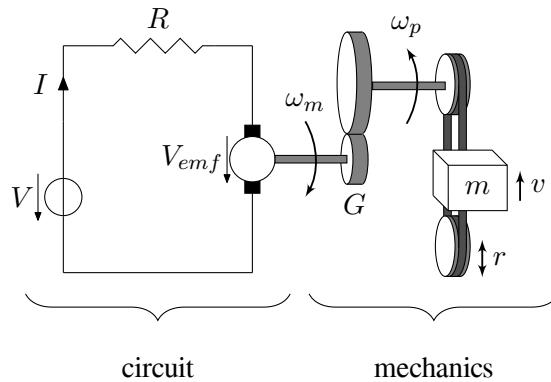


Figure 13.3: Elevator system diagram

Gear ratios are written as output over input, so G is greater than one in figure 13.3.

We want to derive an equation for the carriage acceleration a (derivative of v) given an input voltage V , which we can integrate to get carriage velocity and position.

First, we'll find a torque to substitute into the equation for a DC brushed motor. Based on figure 13.3

$$\tau_m G = \tau_p \quad (13.7)$$

where G is the gear ratio between the motor and the pulley and τ_p is the torque produced by the pulley.

$$rF_m = \tau_p \quad (13.8)$$

where r is the radius of the pulley. Substitute equation (13.7) into τ_m in the DC brushed motor equation (13.3).

$$\begin{aligned} V &= \frac{\tau_p}{K_t} R + \frac{\omega_m}{K_v} \\ V &= \frac{\tau_p}{GK_t} R + \frac{\omega_m}{K_v} \end{aligned}$$

Substitute in equation (13.8) for τ_p .

$$V = \frac{rF_m}{GK_t} R + \frac{\omega_m}{K_v} \quad (13.9)$$

The angular velocity of the motor armature ω_m is

$$\omega_m = G\omega_p \quad (13.10)$$

where ω_p is the angular velocity of the pulley. The velocity of the mass (the elevator carriage) is

$$\begin{aligned} v &= r\omega_p \\ \omega_p &= \frac{v}{r} \end{aligned} \quad (13.11)$$

Substitute equation (13.11) into equation (13.10).

$$\omega_m = G \frac{v}{r} \quad (13.12)$$

Substitute equation (13.12) into equation (13.9) for ω_m .

$$\begin{aligned} V &= \frac{rF_m}{GK_t} R + \frac{G \frac{v}{r}}{K_v} \\ V &= \frac{RrF_m}{GK_t} + \frac{G}{rK_v} v \end{aligned}$$

Solve for F_m .

$$\begin{aligned} \frac{RrF_m}{GK_t} &= V - \frac{G}{rK_v} v \\ F_m &= \left(V - \frac{G}{rK_v} v \right) \frac{GK_t}{Rr} \\ F_m &= \frac{GK_t}{Rr} V - \frac{G^2 K_t}{Rr^2 K_v} v \end{aligned} \quad (13.13)$$

$$\sum F = ma \quad (13.14)$$

where $\sum F$ is the sum of forces applied to the elevator carriage, m is the mass of the elevator carriage in kilograms, and a is the acceleration of the elevator carriage.

$$ma = F_m$$

- R** Gravity is not part of the modeled dynamics because it complicates the state-space model and the controller will behave well enough without it.

$$\begin{aligned} ma &= \left(\frac{GK_t}{Rr} V - \frac{G^2 K_t}{Rr^2 K_v} v \right) \\ a &= \frac{GK_t}{Rrm} V - \frac{G^2 K_t}{Rr^2 m K_v} v \end{aligned} \quad (13.15)$$

This model will be converted to state-space notation in section 8.1.

13.3 Flywheel

13.3.1 Equations of motion

This flywheel consists of a DC brushed motor attached to a spinning mass of non-negligible moment of inertia.

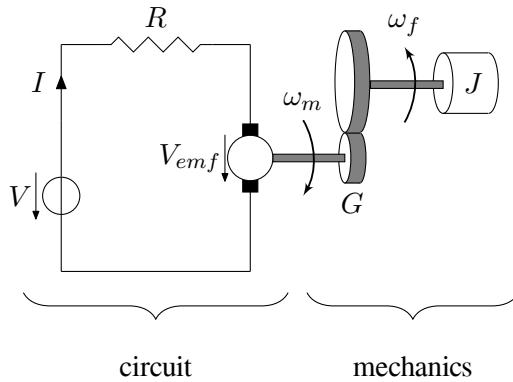


Figure 13.4: Flywheel system diagram

Gear ratios are written as output over input, so G is greater than one in figure 13.4.

We want to derive an equation for the flywheel angular acceleration $\dot{\omega}_f$ given an input voltage V , which we can integrate to get flywheel angular velocity.

We will start with the equation derived earlier for a DC brushed motor, equation (13.3).

$$V = \frac{\tau_m}{K_t} R + \frac{\omega_m}{K_v}$$

Solve for the angular acceleration. First, we'll rearrange the terms because from inspection, V is the model input, ω_m is the state, and τ_m contains the angular acceleration.

$$V = \frac{R}{K_t} \tau_m + \frac{1}{K_v} \omega_m$$

Solve for τ_m .

$$\begin{aligned} V &= \frac{R}{K_t} \tau_m + \frac{1}{K_v} \omega_m \\ \frac{R}{K_t} \tau_m &= V - \frac{1}{K_v} \omega_m \end{aligned}$$

$$\tau_m = \frac{K_t}{R}V - \frac{K_t}{K_v R} \omega_m$$

Since $\tau_m G = \tau_f$ and $\omega_m = G\omega_f$

$$\begin{aligned} \left(\frac{\tau_f}{G}\right) &= \frac{K_t}{R}V - \frac{K_t}{K_v R}(G\omega_f) \\ \frac{\tau_f}{G} &= \frac{K_t}{R}V - \frac{GK_t}{K_v R}\omega_f \\ \tau_f &= \frac{GK_t}{R}V - \frac{G^2K_t}{K_v R}\omega_f \end{aligned} \quad (13.16)$$

The torque applied to the flywheel is defined as

$$\tau_f = J\dot{\omega}_f \quad (13.17)$$

where J is the moment of inertia of the flywheel and $\dot{\omega}_f$ is the angular acceleration. Substitute equation (13.17) into equation (13.16).

$$\begin{aligned} (J\dot{\omega}_f) &= \frac{GK_t}{R}V - \frac{G^2K_t}{K_v R}\omega_f \\ \dot{\omega}_f &= \frac{GK_t}{RJ}V - \frac{G^2K_t}{K_v R J}\omega_f \\ \dot{\omega}_f &= -\frac{G^2K_t}{K_v R J}\omega_f + \frac{GK_t}{RJ}V \end{aligned}$$

We'll relabel ω_f as ω for convenience.

$$\dot{\omega} = -\frac{G^2K_t}{K_v R J}\omega + \frac{GK_t}{RJ}V \quad (13.18)$$

This model will be converted to state-space notation in section 8.2.

13.3.2 Calculating constants

Moment of inertia J

Given the simplicity of this mechanism, it may be easier to compute this value theoretically using material properties in CAD. A procedure for measuring it experimentally is presented below.

First, rearrange equation (13.18) into the form $y = mx + b$ such that J is in the numerator of m .

$$\begin{aligned} \dot{\omega} &= -\frac{G^2K_t}{K_v R J}\omega + \frac{GK_t}{RJ}V \\ J\dot{\omega} &= -\frac{G^2K_t}{K_v R}\omega + \frac{GK_t}{R}V \end{aligned}$$

Multiply by $\frac{K_v R}{G^2 K_t}$ on both sides.

$$\begin{aligned} \frac{JK_v R}{G^2 K_t} \dot{\omega} &= -\omega + \frac{GK_t}{R} \cdot \frac{K_v R}{G^2 K_t} V \\ \frac{JK_v R}{G^2 K_t} \dot{\omega} &= -\omega + \frac{K_v}{G} V \end{aligned}$$

$$\omega = -\frac{JK_vR}{G^2K_t}\dot{\omega} + \frac{K_v}{G}V \quad (13.19)$$

The test procedure is as follows.

1. Run the flywheel forward at a constant voltage. Record the angular velocity over time.
2. Compute the angular acceleration from the angular velocity data as the difference between each sample divided by the time between them.
3. Perform a linear regression of angular velocity versus angular acceleration. The slope of this line has the form $-\frac{JK_vR}{G^2K_t}$ as per equation (13.19).
4. Multiply the slope by $-\frac{G^2K_t}{K_vR}$ to obtain a least squares estimate of J .

13.4 Single-jointed arm

13.4.1 Equations of motion

This single-jointed arm consists of a DC brushed motor attached to a pulley that spins a straight bar in pitch.

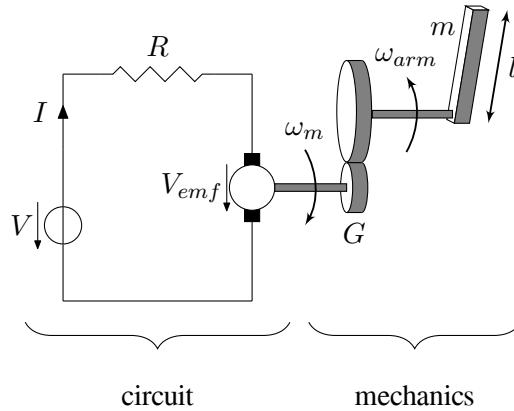


Figure 13.5: Single-jointed arm system diagram

Gear ratios are written as output over input, so G is greater than one in figure 13.5.

We want to derive an equation for the arm angular acceleration ω_{arm} given an input voltage V , which we can integrate to get arm angular velocity and angle.

We will start with the equation derived earlier for a DC brushed motor, equation (13.3).

$$V = \frac{\tau_m}{K_t}R + \frac{\omega_m}{K_v}$$

Solve for the angular acceleration. First, we'll rearrange the terms because from inspection, V is the model input, ω_m is the state, and τ_m contains the angular acceleration.

$$V = \frac{R}{K_t}\tau_m + \frac{1}{K_v}\omega_m$$

Solve for τ_m .

$$V = \frac{R}{K_t}\tau_m + \frac{1}{K_v}\omega_m$$

$$\begin{aligned}\frac{R}{K_t}\tau_m &= V - \frac{1}{K_v}\omega_m \\ \tau_m &= \frac{K_t}{R}V - \frac{K_t}{K_vR}\omega_m\end{aligned}$$

Since $\tau_m G = \tau_{arm}$ and $\omega_m = G\omega_{arm}$

$$\begin{aligned}\left(\frac{\tau_{arm}}{G}\right) &= \frac{K_t}{R}V - \frac{K_t}{K_vR}(G\omega_{arm}) \\ \frac{\tau_{arm}}{G} &= \frac{K_t}{R}V - \frac{GK_t}{K_vR}\omega_{arm} \\ \tau_{arm} &= \frac{GK_t}{R}V - \frac{G^2K_t}{K_vR}\omega_{arm}\end{aligned}\tag{13.20}$$

The torque applied to the arm is defined as

$$\tau_{arm} = J\dot{\omega}_{arm}\tag{13.21}$$

where J is the moment of inertia of the arm and $\dot{\omega}_{arm}$ is the angular acceleration. Substitute equation (13.21) into equation (13.20).

$$\begin{aligned}(J\dot{\omega}_{arm}) &= \frac{GK_t}{R}V - \frac{G^2K_t}{K_vR}\omega_{arm} \\ \dot{\omega}_{arm} &= -\frac{G^2K_t}{K_vRJ}\omega_{arm} + \frac{GK_t}{RJ}V\end{aligned}$$

We'll relabel ω_{arm} as ω for convenience.

$$\dot{\omega} = -\frac{G^2K_t}{K_vRJ}\omega + \frac{GK_t}{RJ}V\tag{13.22}$$

This model will be converted to state-space notation in section 8.3.

13.4.2 Calculating constants

Moment of inertia J

Given the simplicity of this mechanism, it may be easier to compute this value theoretically using material properties in CAD. J can also be approximated as the moment of inertia of a thin rod rotating around one end. Therefore

$$J = \frac{1}{3}ml^2\tag{13.23}$$

where m is the mass of the arm and l is the length of the arm. Otherwise, a procedure for measuring it experimentally is presented below.

First, rearrange equation (13.22) into the form $y = mx + b$ such that J is in the numerator of m .

$$\begin{aligned}\dot{\omega} &= -\frac{G^2K_t}{K_vRJ}\omega + \frac{GK_t}{RJ}V \\ J\dot{\omega} &= -\frac{G^2K_t}{K_vR}\omega + \frac{GK_t}{R}V\end{aligned}$$

Multiply by $\frac{K_v R}{G^2 K_t}$ on both sides.

$$\begin{aligned}\frac{JK_v R}{G^2 K_t} \dot{\omega} &= -\omega + \frac{GK_t}{R} \cdot \frac{K_v R}{G^2 K_t} V \\ \frac{JK_v R}{G^2 K_t} \dot{\omega} &= -\omega + \frac{K_v}{G} V \\ \omega &= -\frac{JK_v R}{G^2 K_t} \dot{\omega} + \frac{K_v}{G} V\end{aligned}\quad (13.24)$$

The test procedure is as follows.

1. Orient the arm such that its axis of rotation is aligned with gravity (i.e., the arm is on its side). This avoids gravity affecting the measurements.
2. Run the arm forward at a constant voltage. Record the angular velocity over time.
3. Compute the angular acceleration from the angular velocity data as the difference between each sample divided by the time between them.
4. Perform a linear regression of angular velocity versus angular acceleration. The slope of this line has the form $-\frac{JK_v R}{G^2 K_t}$ as per equation (13.24).
5. Multiply the slope by $-\frac{G^2 K_t}{K_v R}$ to obtain a least squares estimate of J .

13.5 Pendulum

Kinematics and dynamics are a rather large topics, so for now, we'll just focus on the basics required for working with the [models](#) in this book. We'll derive the same [model](#), a pendulum, using three approaches: sum of forces, sum of torques, and conservation of energy.

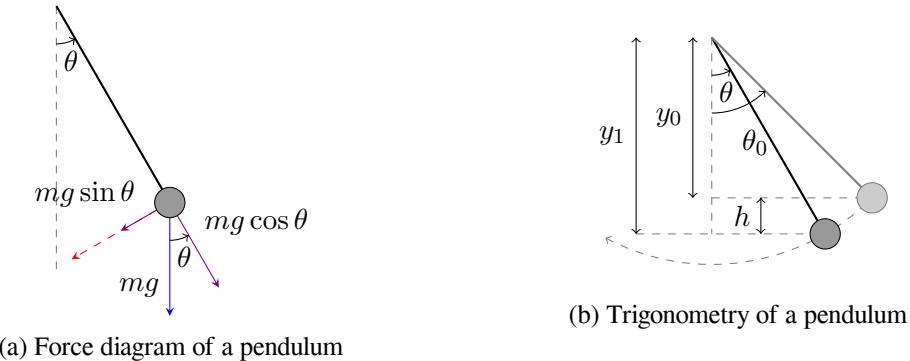


Figure 13.6: Pendulum force diagrams

13.5.1 Force derivation

Consider figure 13.6a, which shows the forces acting on a pendulum.

Note that the path of the pendulum sweeps out an arc of a circle. The angle θ is measured in radians. The blue arrow is the gravitational force acting on the bob, and the violet arrows are that same force resolved into components parallel and perpendicular to the bob's instantaneous motion. The direction of the bob's instantaneous velocity always points along the red axis, which is considered the tangential axis because its direction is always tangent to the circle. Consider Newton's second law

$$F = ma$$

where F is the sum of forces on the object, m is mass, and a is the acceleration. Because we are only concerned with changes in speed, and because the bob is forced to stay in a circular path, we apply Newton's equation to the tangential axis only. The short violet arrow represents the component of the gravitational force in the tangential axis, and trigonometry can be used to determine its magnitude. Therefore

$$\begin{aligned}-mg \sin \theta &= ma \\ a &= -g \sin \theta\end{aligned}$$

where g is the acceleration due to gravity near the surface of the earth. The negative sign on the right hand side implies that θ and a always point in opposite directions. This makes sense because when a pendulum swings further to the left, we would expect it to accelerate back toward the right.

This linear acceleration a along the red axis can be related to the change in angle θ by the arc length formulas; s is arc length and l is the length of the pendulum.

$$\begin{aligned}s &= l\theta \\ v &= \frac{ds}{dt} = l \frac{d\theta}{dt} \\ a &= \frac{d^2s}{dt^2} = l \frac{d^2\theta}{dt^2}\end{aligned}\tag{13.25}$$

Therefore

$$\begin{aligned}l \frac{d^2\theta}{dt^2} &= -g \sin \theta \\ \frac{d^2\theta}{dt^2} &= -\frac{g}{l} \sin \theta \\ \ddot{\theta} &= -\frac{g}{l} \sin \theta\end{aligned}$$

13.5.2 Torque derivation

The equation can be obtained using two definitions for torque.

$$\tau = \mathbf{r} \times \mathbf{F}$$

First start by defining the torque on the pendulum bob using the force due to gravity.

$$\tau = \mathbf{l} \times \mathbf{F}_g$$

where \mathbf{l} is the length vector of the pendulum and \mathbf{F}_g is the force due to gravity.

For now just consider the magnitude of the torque on the pendulum.

$$|\tau| = -mgl \sin \theta$$

where m is the mass of the pendulum, g is the acceleration due to gravity, l is the length of the pendulum and θ is the angle between the length vector and the force due to gravity.

Next rewrite the angular momentum.

$$\mathbf{L} = \mathbf{r} \times \mathbf{p} = m\mathbf{r} \times (\omega \times \mathbf{r})$$

Again just consider the magnitude of the angular momentum.

$$\begin{aligned} |\mathbf{L}| &= mr^2\omega \\ |\mathbf{L}| &= ml^2 \frac{d\theta}{dt} \\ \frac{d}{dt} |\mathbf{L}| &= ml^2 \frac{d^2\theta}{dt^2} \end{aligned}$$

According to $\tau = \frac{d\mathbf{L}}{dt}$, we can just compare the magnitudes.

$$\begin{aligned} -mgl \sin \theta &= ml^2 \frac{d^2\theta}{dt^2} \\ -\frac{g}{l} \sin \theta &= \frac{d^2\theta}{dt^2} \\ \ddot{\theta} &= -\frac{g}{l} \sin \theta \end{aligned}$$

which is the same result from force analysis.

13.5.3 Energy derivation

The equation can also be obtained via the conservation of mechanical energy principle: any object falling a vertical distance h would acquire kinetic energy equal to that which it lost to the fall. In other words, gravitational potential energy is converted into kinetic energy. Change in potential energy is given by

$$\Delta U = mgh$$

The change in kinetic energy (body started from rest) is given by

$$\Delta K = \frac{1}{2}mv^2$$

Since no energy is lost, the gain in one must be equal to the loss in the other

$$\frac{1}{2}mv^2 = mgh$$

The change in velocity for a given change in height can be expressed as

$$v = \sqrt{2gh}$$

Using equation (13.25), this equation can be rewritten in terms of $\frac{d\theta}{dt}$.

$$\begin{aligned} v &= l \frac{d\theta}{dt} = \sqrt{2gh} \\ \frac{d\theta}{dt} &= \frac{2gh}{l} \end{aligned} \tag{13.26}$$

where h is the vertical distance the pendulum fell. Look at figure 13.6b, which presents the trigonometry of a pendulum. If the pendulum starts its swing from some initial angle θ_0 , then y_0 , the vertical distance from the pivot point, is given by

$$y_0 = l \cos \theta_0$$

Similarly for y_1 , we have

$$y_1 = l \cos \theta$$

Then h is the difference of the two

$$h = l(\cos \theta - \cos \theta_0)$$

Substituting this into equation (13.26) gives

$$\frac{d\theta}{dt} = \sqrt{\frac{2g}{l}(\cos \theta - \cos \theta_0)}$$

This equation is known as the first integral of motion. It gives the velocity in terms of the location and includes an integration constant related to the initial displacement (θ_0). We can differentiate by applying the chain rule with respect to time. Doing so gives the acceleration.

$$\begin{aligned} \frac{d}{dt} \frac{d\theta}{dt} &= \frac{d}{dt} \sqrt{\frac{2g}{l}(\cos \theta - \cos \theta_0)} \\ \frac{d^2\theta}{dt^2} &= \frac{1}{2} \frac{-\frac{2g}{l} \sin \theta}{\sqrt{\frac{2g}{l}(\cos \theta - \cos \theta_0)}} \frac{d\theta}{dt} \\ \frac{d^2\theta}{dt^2} &= \frac{1}{2} \frac{-\frac{2g}{l} \sin \theta}{\sqrt{\frac{2g}{l}(\cos \theta - \cos \theta_0)}} \sqrt{\frac{2g}{l}(\cos \theta - \cos \theta_0)} \\ \frac{d^2\theta}{dt^2} &= -\frac{g}{l} \sin \theta \\ \ddot{\theta} &= -\frac{g}{l} \sin \theta \end{aligned}$$

which is the same result from force analysis.

13.6 Differential drive

13.6.1 Equations of motion

This drivetrain consists of two DC brushed motors per side which are chained together on their respective sides and drive wheels which are assumed to be massless.

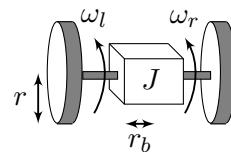


Figure 13.7: Differential drive system diagram

We want to derive equations for the accelerations of the left and right sides of the robot \dot{v}_l and \dot{v}_r given left and right input voltages V_l and V_r .

From equation (13.16) of the flywheel model derivations

$$\tau = \frac{GK_t}{R}V - \frac{G^2K_t}{K_vR}\omega \quad (13.27)$$

where τ is the torque applied by one wheel of the differential drive, G is the gear ratio of the differential drive, K_t is the torque constant of the motor, R is the resistance of the motor, and K_v is the angular velocity constant. Since $\tau = rF$ and $\omega = \frac{v}{r}$ where v is the velocity of a given drive side along the ground and r is the drive wheel radius

$$\begin{aligned}(rF) &= \frac{GK_t}{R}V - \frac{G^2K_t}{K_vR} \left(\frac{v}{r}\right) \\ rF &= \frac{GK_t}{R}V - \frac{G^2K_t}{K_vRr}v \\ F &= \frac{GK_t}{Rr}V - \frac{G^2K_t}{K_vRr^2}v \\ F &= -\frac{G^2K_t}{K_vRr^2}v + \frac{GK_t}{Rr}V\end{aligned}$$

Therefore, for each side of the robot

$$\begin{aligned}F_l &= -\frac{G_l^2K_t}{K_vRr^2}v_l + \frac{G_lK_t}{Rr}V_l \\ F_r &= -\frac{G_r^2K_t}{K_vRr^2}v_r + \frac{G_rK_t}{Rr}V_r\end{aligned}$$

where the l and r subscripts denote the side of the robot to which each variable corresponds.

Let $C_1 = -\frac{G_l^2K_t}{K_vRr^2}$, $C_2 = \frac{G_lK_t}{Rr}$, $C_3 = -\frac{G_r^2K_t}{K_vRr^2}$, and $C_4 = \frac{G_rK_t}{Rr}$.

$$F_l = C_1v_l + C_2V_l \quad (13.28)$$

$$F_r = C_3v_r + C_4V_r \quad (13.29)$$

First, find the sum of forces.

$$\begin{aligned}\sum F &= ma \\ F_l + F_r &= m\dot{v} \\ F_l + F_r &= m\frac{\dot{v}_l + \dot{v}_r}{2} \\ \frac{2}{m}(F_l + F_r) &= \dot{v}_l + \dot{v}_r \\ \dot{v}_l &= \frac{2}{m}(F_l + F_r) - \dot{v}_r\end{aligned} \quad (13.30)$$

Next, find the sum of torques.

$$\begin{aligned}\sum \tau &= J\dot{\omega} \\ \tau_l + \tau_r &= J \left(\frac{\dot{v}_r - \dot{v}_l}{2r_b} \right)\end{aligned}$$

where r_b is the radius of the differential drive.

$$(-r_bF_l) + (r_bF_r) = J \frac{\dot{v}_r - \dot{v}_l}{2r_b}$$

$$\begin{aligned}
-r_b F_l + r_b F_r &= \frac{J}{2r_b}(\dot{v}_r - \dot{v}_l) \\
-F_l + F_r &= \frac{J}{2r_b^2}(\dot{v}_r - \dot{v}_l) \\
\frac{2r_b^2}{J}(-F_l + F_r) &= \dot{v}_r - \dot{v}_l \\
\dot{v}_r &= \dot{v}_l + \frac{2r_b^2}{J}(-F_l + F_r)
\end{aligned}$$

Substitute in equation (13.30) for \dot{v}_l to obtain an expression for \dot{v}_r .

$$\begin{aligned}
\dot{v}_r &= \left(\frac{2}{m}(F_l + F_r) - \dot{v}_r \right) + \frac{2r_b^2}{J}(-F_l + F_r) \\
2\dot{v}_r &= \frac{2}{m}(F_l + F_r) + \frac{2r_b^2}{J}(-F_l + F_r) \\
\dot{v}_r &= \frac{1}{m}(F_l + F_r) + \frac{r_b^2}{J}(-F_l + F_r) \tag{13.31}
\end{aligned}$$

$$\begin{aligned}
\dot{v}_r &= \frac{1}{m}F_l + \frac{1}{m}F_r - \frac{r_b^2}{J}F_l + \frac{r_b^2}{J}F_r \\
\dot{v}_r &= \left(\frac{1}{m} - \frac{r_b^2}{J} \right)F_l + \left(\frac{1}{m} + \frac{r_b^2}{J} \right)F_r \tag{13.32}
\end{aligned}$$

Substitute equation (13.31) back into equation (13.30) to obtain an expression for \dot{v}_l .

$$\begin{aligned}
\dot{v}_l &= \frac{2}{m}(F_l + F_r) - \left(\frac{1}{m}(F_l + F_r) + \frac{r_b^2}{J}(-F_l + F_r) \right) \\
\dot{v}_l &= \frac{1}{m}(F_l + F_r) - \frac{r_b^2}{J}(-F_l + F_r) \\
\dot{v}_l &= \frac{1}{m}(F_l + F_r) + \frac{r_b^2}{J}(F_l - F_r) \\
\dot{v}_l &= \frac{1}{m}F_l + \frac{1}{m}F_r + \frac{r_b^2}{J}F_l - \frac{r_b^2}{J}F_r \\
\dot{v}_l &= \left(\frac{1}{m} + \frac{r_b^2}{J} \right)F_l + \left(\frac{1}{m} - \frac{r_b^2}{J} \right)F_r \tag{13.33}
\end{aligned}$$

Now, plug the expressions for F_l and F_r into equation (13.32).

$$\begin{aligned}
\dot{v}_r &= \left(\frac{1}{m} - \frac{r_b^2}{J} \right)F_l + \left(\frac{1}{m} + \frac{r_b^2}{J} \right)F_r \\
\dot{v}_r &= \left(\frac{1}{m} - \frac{r_b^2}{J} \right)(C_1v_l + C_2V_l) + \left(\frac{1}{m} + \frac{r_b^2}{J} \right)(C_3v_r + C_4V_r) \tag{13.34}
\end{aligned}$$

Now, plug the expressions for F_l and F_r into equation (13.33).

$$\begin{aligned}
\dot{v}_l &= \left(\frac{1}{m} + \frac{r_b^2}{J} \right)F_l + \left(\frac{1}{m} - \frac{r_b^2}{J} \right)F_r \\
\dot{v}_l &= \left(\frac{1}{m} + \frac{r_b^2}{J} \right)(C_1v_l + C_2V_l) + \left(\frac{1}{m} - \frac{r_b^2}{J} \right)(C_3v_r + C_4V_r) \tag{13.35}
\end{aligned}$$

This model will be converted to state-space notation in section 8.5.

13.6.2 Calculating constants

Moment of inertia J

We'll use empirical measurements of linear and angular velocity to determine J . First, we'll derive the equation required to perform a linear regression using velocity test data.

$$\begin{aligned}\tau_1 &= \mathbf{r} \times \mathbf{F} \\ \tau_1 &= rma\end{aligned}$$

where τ_1 is the torque applied by a drive motor during only linear acceleration, r is the wheel radius, m is the robot mass, and a is the linear acceleration.

$$\tau_2 = I\alpha$$

where τ_2 is the torque applied by a drive motor during only angular acceleration, I is the moment of inertia (same as J), and α is the angular acceleration. If a constant voltage is applied during both the linear and angular acceleration tests, $\tau_1 = \tau_2$. Therefore,

$$rma = I\alpha$$

Integrate with respect to time.

$$\begin{aligned}rmv + C_1 &= I\omega + C_2 \\ rmv &= I\omega + C_3 \\ v &= \frac{I}{rm}\omega + C_3\end{aligned}\tag{13.36}$$

where v is linear velocity and ω is angular velocity. C_1 , C_2 , and C_3 are arbitrary constants of integration that won't be needed. The test procedure is as follows.

1. Run the drivetrain forward at a constant voltage. Record the linear velocity over time using encoders.
2. Rotate the drivetrain around its center by applying the same voltage as the linear acceleration test with the motors driving in opposite directions. Record the angular velocity over time using a gyroscope.
3. Perform a linear regression of linear velocity versus angular velocity. The slope of this line has the form $\frac{I}{rm}$ as per equation (13.36).
4. Multiply the slope by rm to obtain a least squares estimate of I .

14. System identification

For now, this chapter just covers how to obtain state-space models for common mechanisms from a feedforward model derived from performance data.

14.1 1 DOF mechanism model

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$$

\mathbf{x} = velocity $\dot{\mathbf{x}}$ = acceleration \mathbf{u} = voltage

If $\mathbf{u} = K_v v$, then $\mathbf{x} = v$ and $\dot{\mathbf{x}} = 0$. Therefore

$$0 = \mathbf{Av} + \mathbf{B}(K_v v) \quad (14.1)$$

If $\mathbf{u} = K_v v + K_a a$, then $\mathbf{x} = v$ and $\dot{\mathbf{x}} = a$. Therefore

$$a = \mathbf{Av} + \mathbf{B}(K_v v + K_a a) \quad (14.2)$$

Subtract equation (14.1) from equation (14.2).

$$\begin{aligned} a &= \mathbf{B}(K_a a) \\ 1 &= \mathbf{B}K_a \\ \mathbf{B} &= \frac{1}{K_a} \end{aligned}$$

Substitute \mathbf{B} back into (14.1) to obtain \mathbf{A} .

$$\begin{aligned} \mathbf{0} &= \mathbf{Av} + \left(\frac{1}{K_a} \right) (K_v v) \\ \mathbf{0} &= \mathbf{Av} + \frac{K_v}{K_a} v \\ -\frac{K_v}{K_a} v &= \mathbf{Av} \\ \mathbf{A} &= -\frac{K_v}{K_a} \end{aligned}$$

A model with position and velocity states would be

Theorem 14.1.1 — 1 DOF mechanism position model.

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$$

$$\mathbf{x} = \begin{bmatrix} \text{position} \\ \text{velocity} \end{bmatrix} \quad \mathbf{u} = [\text{voltage}]$$

$$\dot{\mathbf{x}} = \begin{bmatrix} 0 & 1 \\ 0 & -\frac{K_v}{K_a} \end{bmatrix} \mathbf{x} + \begin{bmatrix} 0 \\ \frac{1}{K_a} \end{bmatrix} \mathbf{u} \quad (14.3)$$

14.2 Drivetrain velocity model

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$$

If $\mathbf{u} = \begin{bmatrix} K_{v1}v \\ K_{v1}v \end{bmatrix}$, then $\mathbf{x} = \begin{bmatrix} v \\ v \end{bmatrix}$ and $\dot{\mathbf{x}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$.

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} A_1 & A_2 \\ A_2 & A_1 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} B_1 & B_2 \\ B_2 & B_1 \end{bmatrix} \begin{bmatrix} K_{v1}v \\ K_{v1}v \end{bmatrix}$$

Since the column vectors contain the same element, the elements in the second row can be rearranged.

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} A_1 & A_2 \\ A_1 & A_2 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} B_1 & B_2 \\ B_1 & B_2 \end{bmatrix} \begin{bmatrix} K_{v1}v \\ K_{v1}v \end{bmatrix}$$

Since the rows are linearly dependent, we can use just one of them.

$$\mathbf{0} = [A_1 \ A_2] v + [B_1 \ B_2] K_{v1}v$$

$$\mathbf{0} = [v \ v \ K_{v1}v \ K_{v1}v] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

$$\mathbf{0} = [1 \ 1 \ K_{v1} \ K_{v1}] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

If $\mathbf{u} = [K_{v1}v + K_{a1}a]$, then $\mathbf{x} = \begin{bmatrix} v \\ v \end{bmatrix}$ and $\dot{\mathbf{x}} = \begin{bmatrix} a \\ a \end{bmatrix}$.

$$\begin{bmatrix} a \\ a \end{bmatrix} = \begin{bmatrix} A_1 & A_2 \\ A_2 & A_1 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} B_1 & B_2 \\ B_2 & B_1 \end{bmatrix} \begin{bmatrix} K_{v1}v + K_{a1}a \\ K_{v1}v + K_{a1}a \end{bmatrix}$$

Since the column vectors contain the same element, the elements in the second row can be rearranged.

$$\begin{bmatrix} a \\ a \end{bmatrix} = \begin{bmatrix} A_1 & A_2 \\ A_1 & A_2 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} B_1 & B_2 \\ B_1 & B_2 \end{bmatrix} \begin{bmatrix} K_{v1}v + K_{a1}a \\ K_{v1}v + K_{a1}a \end{bmatrix}$$

Since the rows are linearly dependent, we can use just one of them.

$$a = [A_1 \ A_2] v + [B_1 \ B_2] [K_{v1}v + K_{a1}a]$$

$$a = [v \ v \ K_{v1}v + K_{a1}a \ K_{v1} + K_{a1}a] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

If $\mathbf{u} = \begin{bmatrix} -K_{v2}v \\ K_{v2}v \end{bmatrix}$, then $\mathbf{x} = \begin{bmatrix} -v \\ v \end{bmatrix}$ and $\dot{\mathbf{x}} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$.

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} A_1 & A_2 \\ A_2 & A_1 \end{bmatrix} \begin{bmatrix} -v \\ v \end{bmatrix} + \begin{bmatrix} B_1 & B_2 \\ B_2 & B_1 \end{bmatrix} \begin{bmatrix} -K_{v2}v \\ K_{v2}v \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} -A_1 & A_2 \\ -A_2 & A_1 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} -B_1 & B_2 \\ -B_2 & B_1 \end{bmatrix} \begin{bmatrix} K_{v2}v \\ K_{v2}v \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} -A_1 & A_2 \\ A_1 & -A_2 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} -B_1 & B_2 \\ B_1 & -B_2 \end{bmatrix} \begin{bmatrix} K_{v2}v \\ K_{v2}v \end{bmatrix}$$

Since the column vectors contain the same element, the elements in the second row can be rearranged.

$$\begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} -A_1 & A_2 \\ -A_1 & A_2 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} -B_1 & B_2 \\ -B_1 & B_2 \end{bmatrix} \begin{bmatrix} K_{v2}v \\ K_{v2}v \end{bmatrix}$$

Since the rows are linearly dependent, we can use just one of them.

$$0 = [-A_1 \ A_2] v + [-B_1 \ B_2] K_{v2}v$$

$$0 = -vA_1 + vA_2 - K_{v2}vB_1 + K_{v2}vB_2$$

$$0 = [-v \ v \ -K_{v2}v \ K_{v2}v] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

$$0 = [-1 \ 1 \ -K_{v2} \ K_{v2}] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

If $\mathbf{u} = \begin{bmatrix} -K_{v2}v - K_{a2}a \\ K_{v2}v + K_{a2}a \end{bmatrix}$, then $\mathbf{x} = \begin{bmatrix} -v \\ v \end{bmatrix}$ and $\dot{\mathbf{x}} = \begin{bmatrix} -a \\ a \end{bmatrix}$.

$$\begin{bmatrix} -a \\ a \end{bmatrix} = \begin{bmatrix} A_1 & A_2 \\ A_2 & A_1 \end{bmatrix} \begin{bmatrix} -v \\ v \end{bmatrix} + \begin{bmatrix} B_1 & B_2 \\ B_2 & B_1 \end{bmatrix} \begin{bmatrix} -K_{v2}v - K_{a2}a \\ K_{v2}v + K_{a2}a \end{bmatrix}$$

$$\begin{bmatrix} -a \\ a \end{bmatrix} = \begin{bmatrix} -A_1 & A_2 \\ -A_2 & A_1 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} -B_1 & B_2 \\ -B_2 & B_1 \end{bmatrix} \begin{bmatrix} K_{v2}v + K_{a2}a \\ K_{v2}v + K_{a2}a \end{bmatrix}$$

$$\begin{bmatrix} -a \\ a \end{bmatrix} = \begin{bmatrix} -A_1 & A_2 \\ A_1 & -A_2 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} -B_1 & B_2 \\ B_1 & -B_2 \end{bmatrix} \begin{bmatrix} K_{v2}v + K_{a2}a \\ K_{v2}v + K_{a2}a \end{bmatrix}$$

Since the column vectors contain the same element, the elements in the second row can be rearranged.

$$\begin{bmatrix} -a \\ -a \end{bmatrix} = \begin{bmatrix} -A_1 & A_2 \\ -A_1 & A_2 \end{bmatrix} \begin{bmatrix} v \\ v \end{bmatrix} + \begin{bmatrix} -B_1 & B_2 \\ -B_1 & B_2 \end{bmatrix} \begin{bmatrix} K_{v2}v + K_{a2}a \\ K_{v2}v + K_{a2}a \end{bmatrix}$$

Since the rows are linearly dependent, we can use just one of them.

$$-a = [-A_1 \ A_2] v + [-B_1 \ B_2] [K_{v2}v + K_{a2}a]$$

$$-a = -vA_1 + vA_2 - (K_{v2}v + K_{a2}a)B_1 + K_{v2}v + K_{a2}a)B_2$$

$$-a = [-v \ v \ -(K_{v2}v + K_{a2}a) \ K_{v2}v + K_{a2}a] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

$$a = [v \ -v \ K_{v2}v + K_{a2}a \ -(K_{v2}v + K_{a2}a)] \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

Now stack the rows.

$$\begin{bmatrix} 0 \\ a \\ 0 \\ a \end{bmatrix} = \begin{bmatrix} 1 & 1 & K_{v1} & K_{v1} \\ v & v & K_{v1}v + K_{a1}a & K_{v1}v + K_{a1}a \\ -1 & 1 & -K_{v2} & K_{v2} \\ v & -v & K_{v2}v + K_{a2}a & -(K_{v2}v + K_{a2}a) \end{bmatrix} \begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix}$$

Solve for matrix elements with Wolfram Alpha. Let $b = K_{v1}$, $c = K_{a1}$, $d = K_{v2}$, and $f = K_{a2}$.

inverse of $\{\{1, 1, b, b\}, \{v, v, b v + c a, b v + c a\}, \{-1, 1, -d, d\}, \{v, -v, d v + f a, -(d v + f a)\}\} * \{\{0\}, \{a\}, \{0\}, \{a\}\}$

$$\begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix} = \frac{1}{2cf} \begin{bmatrix} -cd - bf \\ cd - bf \\ c + f \\ f - c \end{bmatrix}$$

$$\begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix} = \frac{1}{2K_{a1}K_{a2}} \begin{bmatrix} -K_{a1}K_{v2} - K_{v1}K_{a2} \\ K_{a1}K_{v2} - K_{v1}K_{a2} \\ K_{a1} + K_{a2} \\ K_{a2} - K_{a1} \end{bmatrix}$$

To summarize,

Theorem 14.2.1 — Drivetrain velocity model.

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$$

$$\mathbf{x} = \begin{bmatrix} \text{left velocity} \\ \text{right velocity} \end{bmatrix} \quad \mathbf{u} = \begin{bmatrix} \text{left voltage} \\ \text{right voltage} \end{bmatrix}$$

$$\dot{\mathbf{x}} = \begin{bmatrix} A_1 & A_2 \\ A_2 & A_1 \end{bmatrix} \mathbf{x} + \begin{bmatrix} B_1 & B_2 \\ B_2 & B_1 \end{bmatrix} \mathbf{u}$$

$$\begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix} = \frac{1}{2K_{a,linear}K_{a,angular}} \begin{bmatrix} -K_{a,linear}K_{v,angular} - K_{v,linear}K_{a,angular} \\ K_{a,linear}K_{v,angular} - K_{v,linear}K_{a,angular} \\ K_{a,linear} + K_{a,angular} \\ K_{a,angular} - K_{a,linear} \end{bmatrix} \quad (14.4)$$

If K_v and K_a are the same for both the linear and angular cases, it devolves to the one-dimensional case. This means the left and right sides are decoupled.

$$\begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix} = \frac{1}{2K_aK_a} \begin{bmatrix} -K_aK_v - K_vK_a \\ K_aK_v - K_vK_a \\ K_a + K_a \\ K_a - K_a \end{bmatrix}$$

$$\begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix} = \frac{1}{2K_a K_a} \begin{bmatrix} -2K_v K_a \\ 0 \\ 2K_a \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} A_1 \\ A_2 \\ B_1 \\ B_2 \end{bmatrix} = \begin{bmatrix} -\frac{K_v}{K_a} \\ 0 \\ \frac{1}{K_a} \\ 0 \end{bmatrix}$$

V

Motion planning

15	Motion profiles	169
15.1	1 DOF motion profiles	
15.2	2 DOF motion profiles	
16	Trajectory optimization	173

This page intentionally left blank



15. Motion profiles

If smooth, predictable motion of a [system](#) over time is desired, it's best to only change a [system's](#) reference as fast as the [system](#) is able to physically move. Motion profiles, also known as trajectories, are used for this purpose. For multi-state [systems](#), each [state](#) is given its own trajectory. Since these [states](#) are usually position and velocity, they share different derivatives of the same profile.

15.1 1 DOF motion profiles

For one degree of freedom (1 DOF) point-to-point movements in FRC, the most commonly used profiles are the trapezoidal profile (figure 15.1) and the S-curve profile (figure 15.2). These profiles accelerate the [system](#) to a maximum velocity from rest, then decelerate it later such that the final acceleration velocity, are zero at the moment the [system](#) arrives at the desired location.

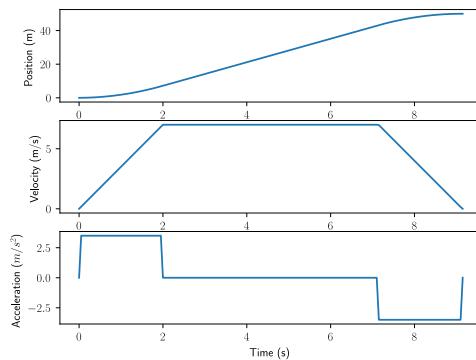


Figure 15.1: Trapezoidal profile

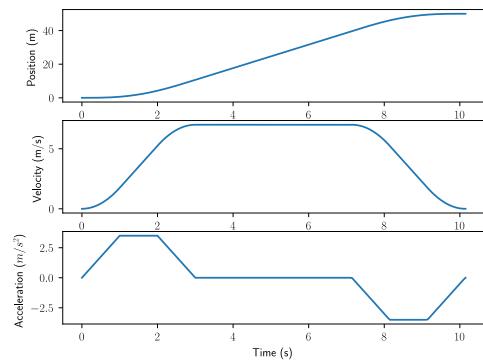


Figure 15.2: S-curve profile

These profiles are given their names based on the shape of their velocity trajectory. The trapezoidal profile has a velocity trajectory shaped like a trapezoid and the S-curve profile has a velocity trajectory

shaped like an S-curve.

In the context of a point-to-point move, a full S-curve consists of seven distinct phases of motion. Phase I starts moving the [system](#) from rest at a linearly increasing acceleration until it reaches the maximum acceleration. In phase II, the profile accelerates at this maximum acceleration rate until it must start decreasing as it approaches the maximum velocity. This occurs in phase III when the acceleration linearly decreases until it reaches zero. In phase IV, the velocity is constant until deceleration begins, at which point the profile decelerates in a manner symmetric to phases I, II and III.

A trapezoidal profile, on the other hand, has three phases. It is a subset of an S-curve profile, having only the phases corresponding to phase II of the S-curve profile (constant acceleration), phase IV (constant velocity), and phase VI (constant deceleration). This reduced number of phases underscores the difference between these two profiles: the S-curve profile has extra motion phases which transition between periods of acceleration and periods of nonacceleration; the trapezoidal profile has instantaneous transitions between these phases. This can be seen in the acceleration graphs of the corresponding velocity profiles for these two profile types.

15.1.1 Jerk

The motion characteristic that defines the change in acceleration, or transitional period, is known as "jerk". Jerk is defined as the rate of change of acceleration with time. In a trapezoidal profile, the jerk (change in acceleration) is infinite at the phase transitions, while in the S-curve profile the jerk is a constant value, spreading the change in acceleration over a period of time.

From figures 15.1 and 15.2, we can see S-curve profiles are smoother than trapezoidal profiles. Why, however, do the S-curve profile result in less load oscillation? For a given load, the higher the jerk, the greater the amount of unwanted vibration energy will be generated, and the broader the frequency spectrum of the vibration's energy will be.

This means that more rapid changes in acceleration induce more powerful vibrations, and more vibrational modes will be excited. Because vibrational energy is absorbed in the system mechanics, it may cause an increase in [settling time](#) or reduced accuracy if the vibration frequency matches resonances in the mechanical system.

15.1.2 Profile selection

Since trapezoidal profiles spend their time at full acceleration or full deceleration, they are, from the standpoint of profile execution, faster than S-curve profiles. However, if this "all on"/"all off" approach causes an increase in settling time, the advantage is lost. Often, only a small amount of "S" (transition between acceleration and no acceleration) can substantially reduce induced vibration. Therefore to optimize throughput, the S-curve profile must be tuned for each a given load and given desired transfer speed.

What S-curve form is right for a given [system](#)? On an application by application basis, the specific choice of the form of the S-curve will depend on the mechanical nature of the [system](#) and the desired performance specifications. For example, in medical applications which involve liquid transfers that should not be jostled, it would be appropriate to choose a profile with no phase II and VI segment at all. Instead the acceleration transitions would be spread out as far as possible, thereby maximizing smoothness.

In other applications involving high speed pick and place, overall transfer speed is most important, so a good choice might be an S-curve with transition phases (phases I, III, V, and VII) that are five to fifteen percent of phase II and VI. In this case, the S-curve profile will add a small amount of time to

the overall transfer time. However, the reduced load oscillation at the end of the move considerably decreases the total effective transfer time. Trial and error using a motion measurement system is generally the best way to determine the right amount of “S” because modeling high frequency dynamics is difficult to do accurately.

Another consideration is whether that “S” segment will actually lead to smoother control of the [system](#). If the high frequency dynamics at play are negligible, one can use the simpler trapezoidal profile.

15.1.3 Profile equations

The trapezoidal profile uses the following equations.

$$x(t) = x_0 + v_0 t + \frac{1}{2} a t^2$$

$$v(t) = v_0 + a t$$

where $x(t)$ is the position at time t , x_0 is the initial position, v_0 is the initial velocity, and a is the acceleration at time t . The S-curve profile equations also include jerk.

$$x(t) = x_0 + v_0 t + \frac{1}{2} a t^2 + \frac{1}{6} j t^3$$

$$v(t) = v_0 + a t + \frac{1}{2} j t^2$$

$$a(t) = a_0 + j t$$

where j is the jerk at time t , $a(t)$ is the acceleration at time t , and a_0 is the initial acceleration.

More derivations are required to determine when to start and stop the different profile phases. The derivations for a trapezoid profile are in appendix H.4. We’re omitting an S-curve profile derivation because the continuous acceleration over time that S-curves provide isn’t required for FRC mechanisms.

15.1.4 Other profile types

The ultimate goal of any profile is to match the profile’s motion characteristics to the desired application. Trapezoidal and S-curve profiles work well when the [system](#)’s torque response curve is fairly flat. In other words, when the output torque does not vary that much over the range of velocities the [system](#) will be experiencing. This is true for most servo motor systems, whether DC brushed or DC brushless.

Step motors, however, do not have flat torque/speed curves. Torque output is nonlinear, sometimes has a large drop at a location called the “mid-range instability”, and generally drops off at higher velocities.

Mid-range instability occurs at the step frequency when the motor’s natural resonance frequency matches the current step rate. To address mid-range instability, the most common technique is to use a nonzero starting velocity. This means that the profile instantly “jumps” to a programmed velocity upon initial acceleration, and while decelerating. While crude, this technique sometimes provides better results than a smooth ramp for zero, particularly for [systems](#) that do not use a microstepping drive technique.

To address torque drop-off at higher velocities, a parabolic profile can be used. The corresponding acceleration curve has the smallest acceleration when the velocity is highest. This is a good match for stepper motors because there is less torque available at higher speeds. However, notice that starting

and ending accelerations are very high, and there is no “S” phase where the acceleration smoothly transitions to zero. If load oscillation is a problem, parabolic profiles may not work as well as an S-curve despite the fact that a standard S-curve profile is not optimized for a stepper motor from the standpoint of the torque/speed curve.

15.1.5 Further reading

FRC teams 254 and 971 gave a talk at FIRST World Championships in 2015 about motion profiles [1].

15.2 2 DOF motion profiles

In FRC, two degree of freedom (2 DOF) point-to-point movements are almost always within the context of drivetrains where the two degrees of freedom are the x and y axes.

A *path* is a set of (x, y) points for the drivetrain to follow. A *trajectory* is a path that includes both the states (e.g., position and velocity) and control inputs (e.g., voltage) of the drivetrain as functions of time.

Currently, the most common form of multidimensional trajectory planning in FRC is based on polynomial splines. The splines determine the path of the trajectory, and the path is parameterized by a trapezoidal motion profile to create a trajectory. *A Dive into WPILib Trajectories* by Declan Freeman-Gleason introduces 2 DOF trajectory planning in FRC using polynomial splines [16].

Planning Motion Trajectories for Mobile Robots Using Splines by Christoph Sprunk provides a more general treatment of spline-based trajectory generation [26].



16. Trajectory optimization

To contextualize some definitions related to trajectory optimization, let's use the example of a drivetrain in FRC. A *path* is a set of (x, y) points for the *system* to follow. A *trajectory* is a path that includes both the states (e.g., position and velocity) and control inputs (e.g., voltage) of the *system* as functions of time. *Trajectory optimization* refers to a set of methods for finding the best choice of trajectory (for some mathematical definition of “best”) by selecting states and inputs of the *system* as functions of time.

Matthew Kelly has an approachable introduction to trajectory optimization on his website [20].

This page intentionally left blank

Appendices

A	Simplifying block diagrams	177
A.1	Cascaded blocks	
A.2	Combining blocks in parallel	
A.3	Removing a block from a feedforward loop	
A.4	Eliminating a feedback loop	
A.5	Removing a block from a feedback loop	
B	Laplace domain analysis	181
B.1	Projections	
B.2	Fourier transform	
B.3	Laplace transform	
B.4	Laplace transform definition	
B.5	Case study: steady-state error	
B.6	Case study: flywheel PID control	
C	Robust control	195
C.1	Gain margin and phase margin	
D	Installing Python packages	197
D.1	Windows instructions	
D.2	Linux instructions	
E	Linear-quadratic regulator	199
E.1	Derivation	
E.2	Output feedback	
E.3	Implicit model following	
F	QR-weighted linear plant inversion	207
F.1	Necessary theorems	
F.2	Setup	
F.3	Minimization	
G	Steady-state feedforward	211
G.1	Continuous case	
G.2	Discrete case	
G.3	Deriving steady-state input	
H	Derivations	215
H.1	Transfer function in feedback	
H.2	Zero-order hold for state-space	
H.3	Kalman filter as Luenberger observer	
H.4	Trapezoidal motion profile	
	Glossary	220
	Bibliography	221
	Online	
	Miscellaneous	
	Index	223

This page intentionally left blank

A. Simplifying block diagrams

A.1 Cascaded blocks

$$Y = (P_1 P_2)X \quad (\text{A.1})$$

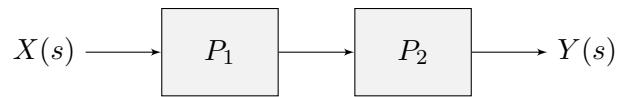


Figure A.1: Cascaded blocks

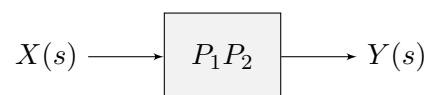


Figure A.2: Simplified cascaded blocks

A.2 Combining blocks in parallel

$$Y = P_1 X \pm P_2 X \quad (\text{A.2})$$

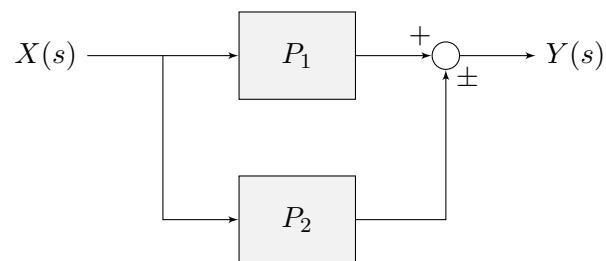


Figure A.3: Parallel blocks

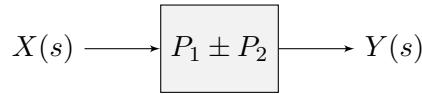


Figure A.4: Simplified parallel blocks

A.3 Removing a block from a feedforward loop

$$Y = P_1 X \pm P_2 X \quad (\text{A.3})$$

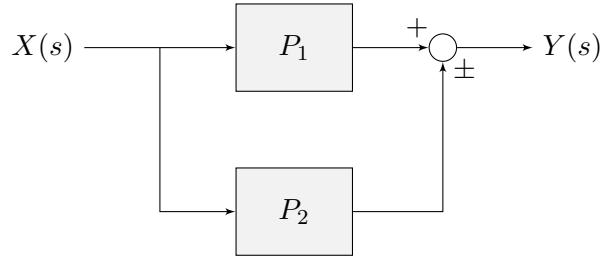


Figure A.5: Feedforward loop

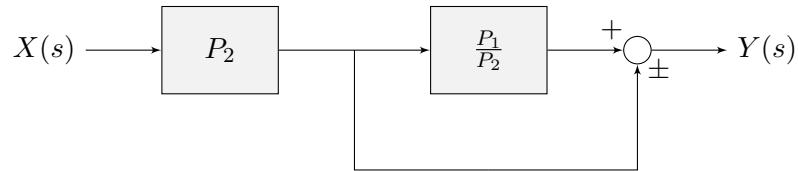


Figure A.6: Transformed feedforward loop

A.4 Eliminating a feedback loop

$$Y = P_1(X \mp P_2 Y) \quad (\text{A.4})$$

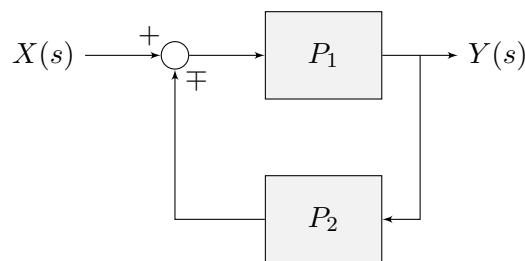


Figure A.7: Feedback loop

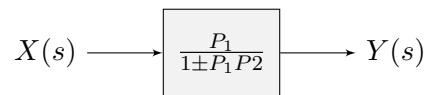


Figure A.8: Simplified feedback loop

A.5 Removing a block from a feedback loop

$$Y = P_1(X \mp P_2 Y) \quad (\text{A.5})$$

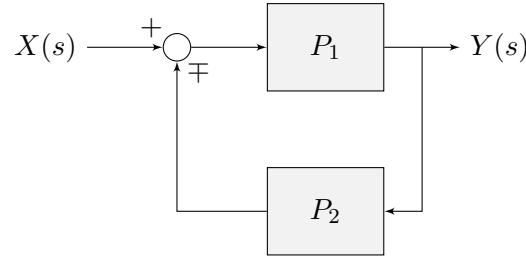


Figure A.9: Feedback loop

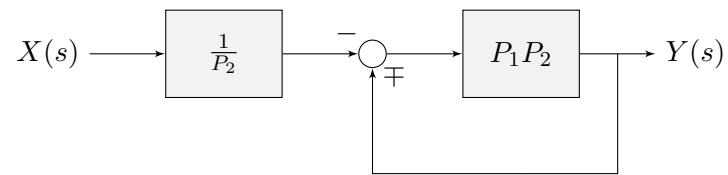


Figure A.10: Transformed feedback loop

This page intentionally left blank

B. Laplace domain analysis

This appendix uses Laplace transforms and transfer functions to analyze properties of control systems like [steady-state error](#).

These case studies cover various aspects of PID control using the algebraic approach of transfer functions. For this, we'll be using equation (B.1), the transfer function for a PID controller.

$$K(s) = K_p + \frac{K_i}{s} + K_d s \quad (\text{B.1})$$

First, we need to define what Laplace transforms and transfer functions are, which is rooted in the concept of orthogonal projections.

B.1 Projections

Consider a two-dimensional Euclidean space \mathbb{R}^2 shown in figure B.1 (each \mathbb{R} is a dimension whose domain is the set of real numbers, so \mathbb{R}^2 is the standard x-y plane).

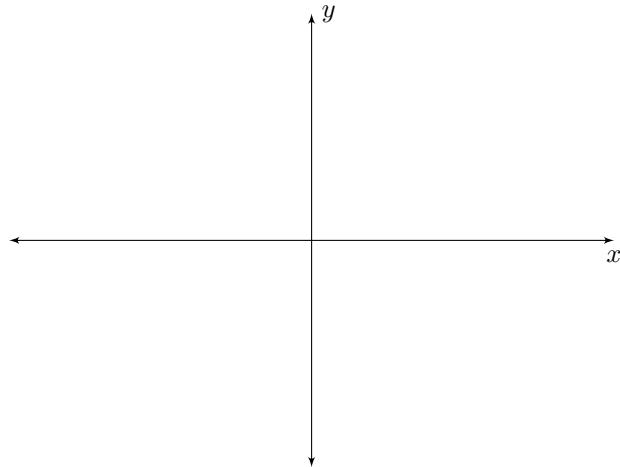
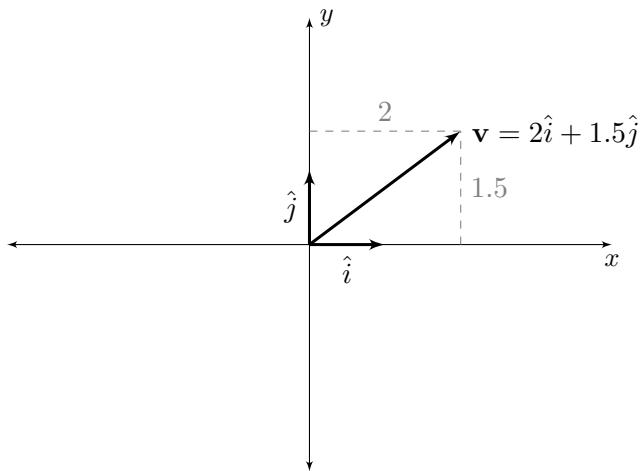
Ordinarily, we notate points in this plane by their components in the set of basis vectors $\{\hat{i}, \hat{j}\}$, where \hat{i} (pronounced i-hat) is the unit vector in the positive x direction and \hat{j} is the unit vector in the positive y direction. Figure B.2 shows an example vector \mathbf{v} in this basis.

How do we find the coordinates of \mathbf{v} in this basis mathematically? As long as the basis is *orthogonal* (i.e., the basis vectors are at right angles to each other), we simply take the *orthogonal projection* of \mathbf{v} onto \hat{i} and \hat{j} . Intuitively, this means finding “the amount of \mathbf{v} that points in the direction of \hat{i} or \hat{j} ”. Note that a set of orthogonal vectors have a dot product of zero with respect to each other.

More formally, we can calculate projections with the dot product - the projection of \mathbf{v} onto any other vector \mathbf{w} is as follows.

$$\text{proj}_{\mathbf{w}} \mathbf{v} = \frac{\mathbf{v} \cdot \mathbf{w}}{|\mathbf{w}|}$$

Since \hat{i} and \hat{j} are *unit vectors*, their magnitudes are 1 so the coordinates of \mathbf{v} are $\mathbf{v} \cdot \hat{i}$ and $\mathbf{v} \cdot \hat{j}$.

Figure B.1: Euclidean space \mathbb{R}^2 Figure B.2: \mathbf{v} with basis set $\{\hat{i}, \hat{j}\}$

We can use this same process to find the coordinates of \mathbf{v} in *any* orthogonal basis. For example, imagine the basis $\{\hat{i} + \hat{j}, \hat{i} - \hat{j}\}$ - the coordinates in this basis are given by $\frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}}$ and $\frac{\mathbf{v} \cdot (\hat{i} - \hat{j})}{\sqrt{2}}$. Let's "unwrap" the formula for dot product and look a bit more closely.

$$\frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}} = \frac{1}{\sqrt{2}} \sum_{i=0}^n \mathbf{v}_i (\hat{i} + \hat{j})_i$$

where the subscript i denotes which component of each vector and n is the total number of components. To change coordinates, we expanded both \mathbf{v} and $\hat{i} + \hat{j}$ in a basis, multiplied their components, and added them up. Here's a more concrete example. Let $\mathbf{v} = 2\hat{i} + 1.5\hat{j}$ from figure B.2. First, we'll project \mathbf{v} onto the $\hat{i} + \hat{j}$ basis vector.

$$\begin{aligned} \frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}} &= \frac{1}{\sqrt{2}} (2\hat{i} \cdot \hat{i} + 1.5\hat{j} \cdot \hat{j}) \\ \frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}} &= \frac{1}{\sqrt{2}} (2 + 1.5) \\ \frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}} &= \frac{3.5}{\sqrt{2}} \end{aligned}$$

$$\frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}} = \frac{3.5\sqrt{2}}{2}$$

$$\frac{\mathbf{v} \cdot (\hat{i} + \hat{j})}{\sqrt{2}} = 1.75\sqrt{2}$$

Next, we'll project \mathbf{v} onto the $\hat{i} - \hat{j}$ basis vector.

$$\frac{\mathbf{v} \cdot (\hat{i} - \hat{j})}{\sqrt{2}} = \frac{1}{\sqrt{2}}(2\hat{i} \cdot \hat{i} - 1.5\hat{j} \cdot \hat{j})$$

$$\frac{\mathbf{v} \cdot (\hat{i} - \hat{j})}{\sqrt{2}} = \frac{1}{\sqrt{2}}(2 - 1.5)$$

$$\frac{\mathbf{v} \cdot (\hat{i} - \hat{j})}{\sqrt{2}} = \frac{0.5}{\sqrt{2}}$$

$$\frac{\mathbf{v} \cdot (\hat{i} - \hat{j})}{\sqrt{2}} = \frac{0.5\sqrt{2}}{2}$$

$$\frac{\mathbf{v} \cdot (\hat{i} - \hat{j})}{\sqrt{2}} = 0.25\sqrt{2}$$

Figure B.3 shows this result geometrically with respect to the basis $\{\hat{i} + \hat{j}, \hat{i} - \hat{j}\}$.

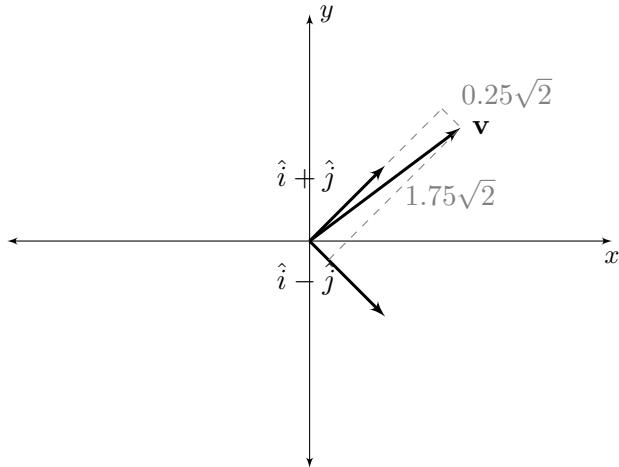


Figure B.3: \mathbf{v} with basis $\{\hat{i} + \hat{j}, \hat{i} - \hat{j}\}$

The previous example was only a change of coordinates in a finite-dimensional vector space. However, as we will see, the core idea does not change much when we move to more complicated structures. Observe the formula for the Fourier transform.

$$\hat{f}(\xi) = \int_{-\infty}^{\infty} f(x) e^{-2\pi i x \xi} dx \text{ where } \xi \in \mathbb{R}$$

This is fundamentally the same formula we had before. $f(x)$ has taken the place of v_n , $e^{-2\pi i x \xi}$ has taken the place of $(\hat{i} + \hat{j})_i$, and the sum over i has turned into an integral over dx , but the underlying concept is the same. To change coordinates in a *function space*, we simply take the orthogonal projection onto our new basis *functions*. In the case of the Fourier transform, the function basis is the

family of functions of the form $f(x) = e^{-2\pi ix\xi}$ for $\xi \in \mathbb{R}$. Since these functions are oscillatory at a frequency determined by ξ , we can think of this as a “frequency basis”.



Watch the “Abstract vector spaces” video from 3Blue1Brown’s *Essence of linear algebra* series (17 minutes) [2] for a more geometric introduction to using functions as a basis.

Now, the Laplace transform is somewhat more complicated - as it turns out, the Fourier basis is orthogonal, so the analogy to the simpler vector space holds almost-precisely. The Laplace transform is *not* orthogonal, so we can’t interpret it *strictly* as a change of coordinates in the traditional sense. However, the intuition is the same: we are taking the orthogonal projection of our original function onto the functions of our new basis set.

$$F(s) = \int_0^\infty f(t)e^{-st} dt, \text{ where } s \in \mathbb{C}$$

Here, it becomes obvious that the Laplace transform is a *generalization* of the Fourier transform in that the basis family is strictly larger (we have allowed the “frequency” parameter to take *complex* values, as opposed to merely *real* values). As a result, the Laplace basis contains functions that grow and decay, while the Fourier basis does not.

B.2 Fourier transform

The Fourier transform decomposes a function of time into its component frequencies. Each of these frequencies is part of what’s called a *basis*. These basis waveforms can be multiplied by their respective contribution amount and summed to produce the original signal (this weighted sum is called a linear combination). In other words, the Fourier transform provides a way for us to determine, given some signal, what frequencies can we add together and in what amounts to produce the original signal.

Think of an Fmajor4 chord which has the notes F_4 (349.23 Hz), A_4 (440 Hz), and C_4 (261.63 Hz). The waveform over time looks like figure B.4.

Notice how this complex waveform can be represented just by three frequencies. They show up as Dirac delta functions¹ in the frequency domain with the area underneath them equal to their contribution (see figure B.5).

¹The Dirac delta function is zero everywhere except at the origin. The nonzero region has an infinitesimal width and has a height such that the area within that region is 1.

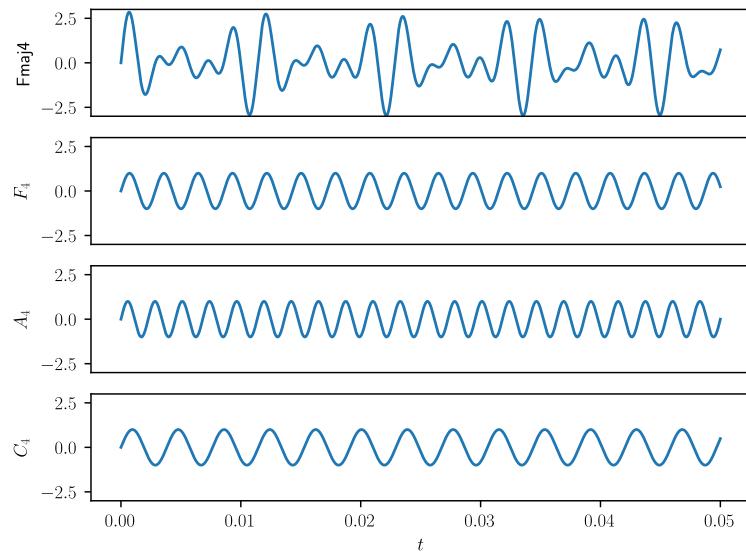


Figure B.4: Frequency decomposition of Fmajor4 chord

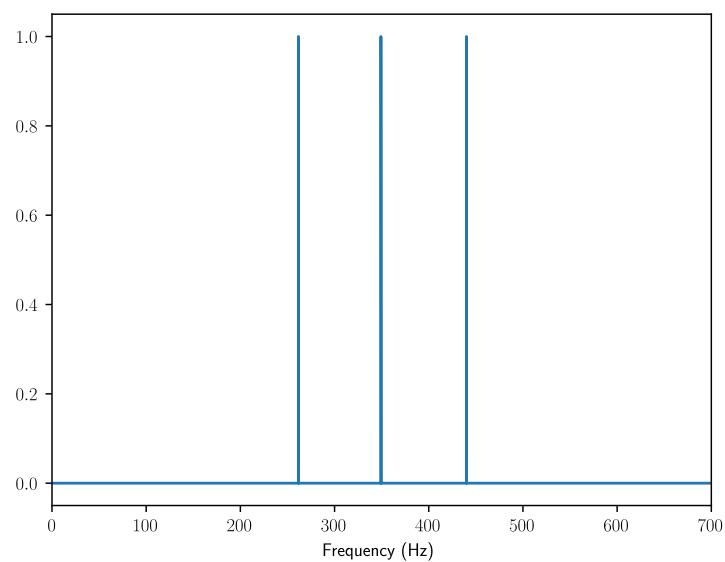


Figure B.5: Fourier transform of Fmajor4 chord

B.3 Laplace transform

The Laplace domain is a generalization of the frequency domain that has the frequency ($j\omega$) on the imaginary y-axis and a real number on the x-axis, yielding a two-dimensional coordinate system. We represent coordinates in this space as a complex number $s = \sigma + j\omega$. The real part σ corresponds to the x-axis and the imaginary part $j\omega$ corresponds to the y-axis (see figure B.6).

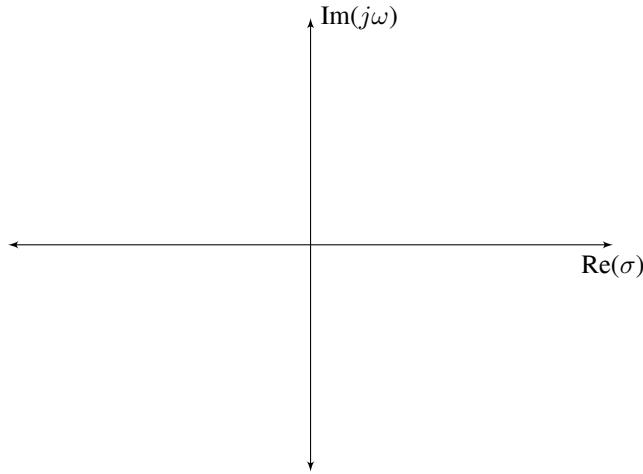


Figure B.6: Laplace domain

To extend our analogy of each coordinate being represented by some basis, we now have the y coordinate representing the oscillation frequency of the [system response](#) (the frequency domain) and also the x coordinate representing the speed at which that oscillation decays and the [system](#) converges to zero (i.e., a decaying exponential). Figure 3.2 shows this for various points.

If we move the component frequencies in the Fmajor4 chord example parallel to the real axis to $\sigma = -25$, the resulting time domain response attenuates according to the decaying exponential e^{-25t} (see figure B.7).

Note that this explanation as a basis isn't exact because the Laplace basis isn't orthogonal (that is, the x and y coordinates affect each other and have cross-talk). In the frequency domain, we had a basis of sine waves that we represented as delta functions in the frequency domain. Each frequency contribution was independent of the others. In the Laplace domain, this is not the case; a pure exponential is $\frac{1}{s-a}$ (a rational function where a is a real number) instead of a delta function. This function is nonzero at points that aren't actually frequencies present in the time domain. Figure B.8 demonstrates this, which shows the Laplace transform of the Fmajor4 chord plotted in 3D.

Notice how the values of the function around each component frequency decrease according to $\frac{1}{\sqrt{x^2+y^2}}$ in the x and y directions (in just the x direction, it would be $\frac{1}{x}$).

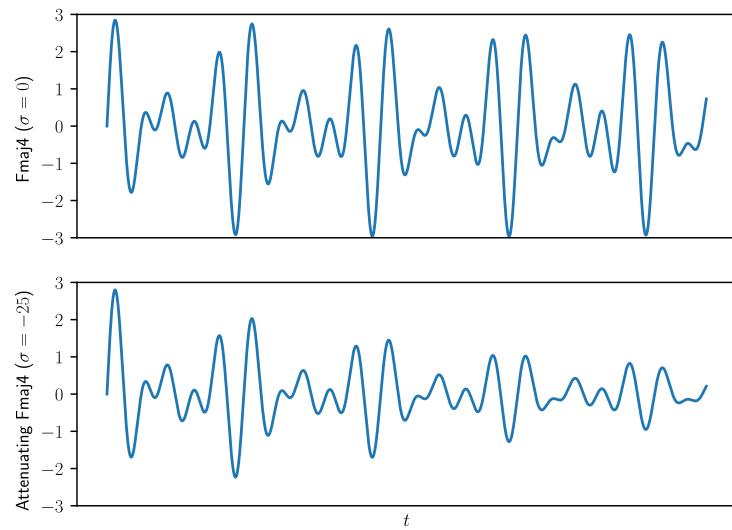
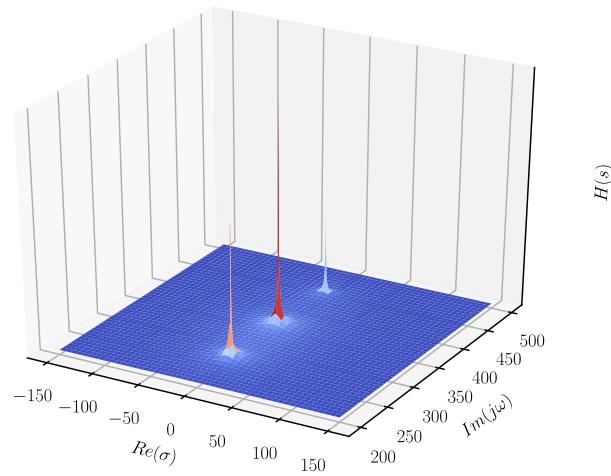
Figure B.7: Fmajor4 chord at $\sigma = 0$ and $\sigma = -25$ 

Figure B.8: Laplace transform of Fmajor4 chord plotted in 3D

B.4 Laplace transform definition

The Laplace transform of a function $f(t)$ is defined as

$$\mathcal{L}\{f(t)\} = F(s) = \int_0^\infty f(t)e^{-st} dt$$

We won't be computing any Laplace transforms by hand using this formula (everyone in the real world looks these up in a table anyway). Common Laplace transforms (assuming zero initial conditions) are shown in table B.1. Of particular note are the Laplace transforms for the derivative, unit step², and exponential decay. We can see that a derivative is equivalent to multiplying by s , and an integral is equivalent to multiplying by $\frac{1}{s}$.

	Time domain	Laplace domain
Linearity	$a f(t) + b g(t)$	$a F(s) + b G(s)$
Convolution	$(f * g)(t)$	$F(s) G(s)$
Derivative	$f'(t)$	$s F(s)$
n^{th} derivative	$f^{(n)}(t)$	$s^n F(s)$
Unit step	$u(t)$	$\frac{1}{s}$
Ramp	$t u(t)$	$\frac{1}{s^2}$
Exponential decay	$e^{-\alpha t} u(t)$	$\frac{1}{s+\alpha}$

Table B.1: Common Laplace transforms and Laplace transform properties with zero initial conditions

B.5 Case study: steady-state error

To demonstrate the problem of [steady-state error](#), we will use a DC brushed motor controlled by a velocity PID controller. A DC brushed motor has a transfer function from voltage (V) to angular velocity ($\dot{\theta}$) of

$$G(s) = \frac{\dot{\Theta}(s)}{V(s)} = \frac{K}{(Js+b)(Ls+R) + K^2} \quad (\text{B.2})$$

First, we'll try controlling it with a P controller defined as

$$K(s) = K_p$$

When these are in unity feedback, the transfer function from the input voltage to the error is

$$\begin{aligned} \frac{E(s)}{V(s)} &= \frac{1}{1 + K(s)G(s)} \\ E(s) &= \frac{1}{1 + K(s)G(s)} V(s) \\ E(s) &= \frac{1}{1 + (K_p) \left(\frac{K}{(Js+b)(Ls+R) + K^2} \right)} V(s) \end{aligned}$$

²The unit step $u(t)$ is defined as 0 for $t < 0$ and 1 for $t \geq 0$.

$$E(s) = \frac{1}{1 + \frac{K_p K}{(Js+b)(Ls+R)+K^2}} V(s)$$

The steady-state of a transfer function can be found via

$$\lim_{s \rightarrow 0} sH(s) \quad (\text{B.3})$$

since steady-state has an input frequency of zero.

$$\begin{aligned} e_{ss} &= \lim_{s \rightarrow 0} sE(s) \\ e_{ss} &= \lim_{s \rightarrow 0} s \frac{1}{1 + \frac{K_p K}{(Js+b)(Ls+R)+K^2}} V(s) \\ e_{ss} &= \lim_{s \rightarrow 0} s \frac{1}{1 + \frac{K_p K}{(Js+b)(Ls+R)+K^2}} \frac{1}{s} \\ e_{ss} &= \lim_{s \rightarrow 0} \frac{1}{1 + \frac{K_p K}{(Js+b)(Ls+R)+K^2}} \\ e_{ss} &= \frac{1}{1 + \frac{K_p K}{(J(0)+b)(L(0)+R)+K^2}} \\ e_{ss} &= \frac{1}{1 + \frac{K_p K}{bR+K^2}} \end{aligned} \quad (\text{B.4})$$

Notice that the steady-state error is nonzero. To fix this, an integrator must be included in the controller.

$$K(s) = K_p + \frac{K_i}{s}$$

The same steady-state calculations are performed as before with the new controller.

$$\begin{aligned} \frac{E(s)}{V(s)} &= \frac{1}{1 + K(s)G(s)} \\ E(s) &= \frac{1}{1 + K(s)G(s)} V(s) \\ E(s) &= \frac{1}{1 + \left(K_p + \frac{K_i}{s}\right) \left(\frac{K}{(Js+b)(Ls+R)+K^2}\right)} \left(\frac{1}{s}\right) \\ e_{ss} &= \lim_{s \rightarrow 0} s \frac{1}{1 + \left(K_p + \frac{K_i}{s}\right) \left(\frac{K}{(Js+b)(Ls+R)+K^2}\right)} \left(\frac{1}{s}\right) \\ e_{ss} &= \lim_{s \rightarrow 0} \frac{1}{1 + \left(K_p + \frac{K_i}{s}\right) \left(\frac{K}{(Js+b)(Ls+R)+K^2}\right)} \\ e_{ss} &= \lim_{s \rightarrow 0} \frac{1}{1 + \left(K_p + \frac{K_i}{s}\right) \left(\frac{K}{(Js+b)(Ls+R)+K^2}\right)} \frac{s}{s} \\ e_{ss} &= \lim_{s \rightarrow 0} \frac{s}{s + (K_p s + K_i) \left(\frac{K}{(Js+b)(Ls+R)+K^2}\right)} \\ e_{ss} &= \frac{0}{0 + (K_p(0) + K_i) \left(\frac{K}{(J(0)+b)(L(0)+R)+K^2}\right)} \end{aligned}$$

$$e_{ss} = \frac{0}{K_i \frac{K}{bR+K^2}}$$

The denominator is nonzero, so $e_{ss} = 0$. Therefore, an integrator is required to eliminate steady-state error in all cases for this model.

It should be noted that e_{ss} in equation (B.4) approaches zero for $K_p = \infty$. This is known as a bang-bang controller. In practice, an infinite switching frequency cannot be achieved, but it may be close enough for some performance specifications.

B.6 Case study: flywheel PID control

PID controllers typically control voltage to a motor in FRC independent of the equations of motion of that motor. For position PID control, large values of K_p can lead to overshoot and K_d is commonly used to reduce overshoots. Let's consider a flywheel controlled with a standard PID controller. Why wouldn't K_d provide damping for velocity overshoots in this case?

PID control is designed to control second-order and first-order systems well. It can be used to control a lot of things, but struggles when given higher order systems. It has three degrees of freedom. Two are used to place the two poles of the system, and the third is used to remove steady-state error. With higher order systems like a one input, seven state system, there aren't enough degrees of freedom to place the system's poles in desired locations. This will result in poor control.

The math for PID doesn't assume voltage, a motor, etc. It defines an output based on derivatives and integrals of its input. We happen to use it for motors because it actually works pretty well for it because motors are second-order systems.

The following math will be in continuous time, but the same ideas apply to discrete time. This is all assuming a velocity controller.

Our simple motor model hooked up to a mass is

$$V = IR + \frac{\omega}{K_v} \quad (B.5)$$

$$\tau = IK_t \quad (B.6)$$

$$\tau = J \frac{d\omega}{dt} \quad (B.7)$$

For an explanation of where these equations come from, read section 13.1.

First, we'll solve for $\frac{d\omega}{dt}$ in terms of V .

Substitute equation (B.6) into equation (B.5).

$$\begin{aligned} V &= IR + \frac{\omega}{K_v} \\ V &= \left(\frac{\tau}{K_t} \right) R + \frac{\omega}{K_v} \end{aligned}$$

Substitute in equation (B.7).

$$V = \frac{\left(J \frac{d\omega}{dt} \right)}{K_t} R + \frac{\omega}{K_v}$$

Solve for $\frac{d\omega}{dt}$.

$$\begin{aligned} V &= \frac{J \frac{d\omega}{dt}}{K_t} R + \frac{\omega}{K_v} \\ V - \frac{\omega}{K_v} &= \frac{J \frac{d\omega}{dt}}{K_t} R \\ \frac{d\omega}{dt} &= \frac{K_t}{JR} \left(V - \frac{\omega}{K_v} \right) \\ \frac{d\omega}{dt} &= -\frac{K_t}{JR K_v} \omega + \frac{K_t}{JR} V \end{aligned}$$

Now take the Laplace transform. Because the Laplace transform is a linear operator, we can take the Laplace transform of each term individually. Based on table B.1, $\frac{d\omega}{dt}$ becomes $s\omega$ and $\omega(t)$ and $V(t)$ become $\omega(s)$ and $V(s)$ respectively (the parenthetical notation has been dropped for clarity).

$$s\omega = -\frac{K_t}{JR K_v} \omega + \frac{K_t}{JR} V \quad (\text{B.8})$$

Solve for the transfer function $H(s) = \frac{\omega}{V}$.

$$\begin{aligned} s\omega &= -\frac{K_t}{JR K_v} \omega + \frac{K_t}{JR} V \\ \left(s + \frac{K_t}{JR K_v} \right) \omega &= \frac{K_t}{JR} V \\ \frac{\omega}{V} &= \frac{\frac{K_t}{JR}}{s + \frac{K_t}{JR K_v}} \end{aligned}$$

That gives us a pole at $-\frac{K_t}{JR K_v}$, which is actually stable. Notice that there is only one pole.

First, we'll use a simple P loop.

$$V = K_p(\omega_{goal} - \omega)$$

Substitute this controller into equation (B.8).

$$s\omega = -\frac{K_t}{JR K_v} \omega + \frac{K_t}{JR} K_p (\omega_{goal} - \omega)$$

Solve for the transfer function $H(s) = \frac{\omega}{\omega_{goal}}$.

$$\begin{aligned} s\omega &= -\frac{K_t}{JR K_v} \omega + \frac{K_t K_p}{JR} (\omega_{goal} - \omega) \\ s\omega &= -\frac{K_t}{JR K_v} \omega + \frac{K_t K_p}{JR} \omega_{goal} - \frac{K_t K_p}{JR} \omega \\ \left(s + \frac{K_t}{JR K_v} + \frac{K_t K_p}{JR} \right) \omega &= \frac{K_t K_p}{JR} \omega_{goal} \\ \frac{\omega}{\omega_{goal}} &= \frac{\frac{K_t K_p}{JR}}{\left(s + \frac{K_t}{JR K_v} + \frac{K_t K_p}{JR} \right)} \end{aligned}$$

This has a pole at $-\left(\frac{K_t}{JRK_v} + \frac{K_tK_p}{JR}\right)$. Assuming that that quantity is negative (i.e., we are stable), that pole corresponds to a time constant of $\frac{1}{\frac{K_t}{JRK_v} + \frac{K_tK_p}{JR}}$.

As can be seen above, a flywheel has a single pole. It therefore only needs a single pole controller to place all of its poles anywhere.



This analysis assumes that the motor is well coupled to the mass and that the time constant of the inductor is small enough that it doesn't factor into the motor equations. In Austin Schuh's experience with 971's robots, these are pretty good assumptions.

Next, we'll try a PD loop. (This will use a perfect derivative, but anyone following along closely already knows that we can't really take a derivative here, so the math will need to be updated at some point. We could switch to discrete time and pick a differentiation method, or pick some other way of modeling the derivative.)

$$V = K_p(\omega_{goal} - \omega) + K_d s(\omega_{goal} - \omega)$$

Substitute this controller into equation (B.8).

$$\begin{aligned} s\omega &= -\frac{K_t}{JRK_v}\omega + \frac{K_t}{JR}(K_p(\omega_{goal} - \omega) + K_d s(\omega_{goal} - \omega)) \\ s\omega &= -\frac{K_t}{JRK_v}\omega + \frac{K_tK_p}{JR}(\omega_{goal} - \omega) + \frac{K_tK_d s}{JR}(\omega_{goal} - \omega) \\ s\omega &= -\frac{K_t}{JRK_v}\omega + \frac{K_tK_p}{JR}\omega_{goal} - \frac{K_tK_p}{JR}\omega + \frac{K_tK_d s}{JR}\omega_{goal} - \frac{K_tK_d s}{JR}\omega \end{aligned}$$

Collect the common terms on separate sides and refactor.

$$\begin{aligned} s\omega + \frac{K_tK_d s}{JR}\omega + \frac{K_t}{JRK_v}\omega + \frac{K_tK_p}{JR}\omega &= \frac{K_tK_p}{JR}\omega_{goal} + \frac{K_tK_d s}{JR}\omega_{goal} \\ \left(s\left(1 + \frac{K_tK_d}{JR}\right) + \frac{K_t}{JRK_v} + \frac{K_tK_p}{JR}\right)\omega &= \frac{K_t}{JR}(K_p + K_d s)\omega_{goal} \\ \frac{\omega}{\omega_{goal}} &= \frac{\frac{K_t}{JR}(K_p + K_d s)}{\left(s\left(1 + \frac{K_tK_d}{JR}\right) + \frac{K_t}{JRK_v} + \frac{K_tK_p}{JR}\right)} \end{aligned}$$

So, we added a zero at $-\frac{K_p}{K_d}$ and moved our pole to $-\frac{\frac{K_t}{JRK_v} + \frac{K_tK_p}{JR}}{1 + \frac{K_tK_d}{JR}}$. This isn't progress. We've added more complexity to our [system](#) and, practically speaking, gotten nothing good in return. Zeroes should be avoided if at all possible because they amplify unwanted high frequency modes of the [system](#) and are noisier the faster the [system](#) is sampled. At least this is a stable zero, but it's still undesirable.

In summary, derivative doesn't help on an ideal flywheel. K_d may compensate for unmodeled dynamics such as accelerating projectiles slowing the flywheel down, but that effect may also increase

recovery time; K_d drives the acceleration to zero in the undesired case of negative acceleration as well as well as the actually desired case of positive acceleration.

We'll cover a superior compensation method much later in subsection 5.11.2 that avoids zeroes in the controller, doesn't act against the desired control action, and facilitates better [tracking](#).

This page intentionally left blank

C. Robust control

C.1 Gain margin and phase margin

One generally needs to learn about Bode plots and Nyquist plots to truly understand gain and phase margin and their origins, but those plots are large topics unto themselves. Since we won't be using either of them for controller design, we'll just cover what gain and phase margin are in a general sense and how they are used.

Gain margin and phase margin are two metrics for measuring a [system](#)'s relative stability. Gain and phase margin are the amounts by which the closed-loop gain and phase can be varied respectively before the [system](#) becomes unstable. In a sense, they are safety margins for when unmodeled dynamics affect the [system response](#).

For a more thorough explanation of gain and phase margin, watch Brian Douglas's video on them [15]. He has other videos too on classical control methods like Bode and Nyquist plots that we recommend.

This page intentionally left blank

D. Installing Python packages

D.1 Windows instructions

To install Python, download the installer for Python 3.5 or higher from <https://www.python.org/downloads/> and run it.

To install Python packages, run `py -3 -m pip install pkg` via cmd.exe or Powershell where `pkg` should be the name of the package. Packages can be upgraded with `py -3 -m pip install --user --upgrade pkg`.

D.2 Linux instructions

To install Python, install the appropriate packages from table D.1 using your system's package manager.

Debian/Ubuntu	Arch Linux
<code>python3</code>	<code>python</code>
<code>python3-pip</code>	<code>python-pip</code>

Table D.1: Required system packages

To install Python packages, run `pip3 install --user pkg` where `pkg` should be the name of the package. Using `--user` makes installation not require root privileges. Packages can be upgraded with `pip3 install --user --upgrade pkg`.

This page intentionally left blank

E. Linear-quadratic regulator

This appendix will go into more detail on the linear-quadratic regulator's derivation and interesting applications.

E.1 Derivation

Let a continuous time linear system be defined as

$$\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu} \quad (\text{E.1})$$

with the cost function

$$J = \int_0^{\infty} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \mathbf{Q} & \mathbf{N} \\ \mathbf{N}^T & \mathbf{R} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt$$

where J represents a trade-off between state excursion and control effort with the weighting factors \mathbf{Q} , \mathbf{R} , and \mathbf{N} . \mathbf{Q} is the weight matrix for error, \mathbf{R} is the weight matrix for control effort, and \mathbf{N} is a cross weight matrix between error and control effort. \mathbf{N} is commonly utilized when penalizing the output in addition to the state and input.

$$\begin{aligned} J &= \int_0^{\infty} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \mathbf{Qx} + \mathbf{Nu} \\ \mathbf{N}^T \mathbf{x} + \mathbf{Ru} \end{bmatrix} dt \\ J &= \int_0^{\infty} [\mathbf{x}^T \quad \mathbf{u}^T] \begin{bmatrix} \mathbf{Qx} + \mathbf{Nu} \\ \mathbf{N}^T \mathbf{x} + \mathbf{Ru} \end{bmatrix} dt \\ J &= \int_0^{\infty} (\mathbf{x}^T (\mathbf{Qx} + \mathbf{Nu}) + \mathbf{u}^T (\mathbf{N}^T \mathbf{x} + \mathbf{Ru})) dt \end{aligned}$$

$$\begin{aligned}
J &= \int_0^\infty (\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{x}^T \mathbf{N} \mathbf{u} + \mathbf{u}^T \mathbf{N}^T \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt \\
J &= \int_0^\infty \left(\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{x}^T \mathbf{N} \mathbf{u} + (\mathbf{x}^T \mathbf{N} \mathbf{u})^T + \mathbf{u}^T \mathbf{R} \mathbf{u} \right) dt \\
J &= \int_0^\infty (\mathbf{x}^T \mathbf{Q} \mathbf{x} + 2\mathbf{x}^T \mathbf{N} \mathbf{u} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt \\
J &= \int_0^\infty (\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u} + 2\mathbf{x}^T \mathbf{N} \mathbf{u}) dt
\end{aligned}$$

The feedback [control law](#) which minimizes J subject to the constraint $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ is

$$\mathbf{u} = -\mathbf{K}\mathbf{x}$$

where \mathbf{K} is given by

$$\mathbf{K} = \mathbf{R}^{-1} (\mathbf{B}^T \mathbf{S} + \mathbf{N}^T)$$

and \mathbf{S} is found by solving the continuous time algebraic Riccati equation defined as

$$\mathbf{A}^T \mathbf{S} + \mathbf{S} \mathbf{A} - (\mathbf{S} \mathbf{B} + \mathbf{N}) \mathbf{R}^{-1} (\mathbf{B}^T \mathbf{S} + \mathbf{N}^T) + \mathbf{Q} = 0$$

or alternatively

$$\mathcal{A}^T \mathbf{S} + \mathbf{S} \mathcal{A} - \mathbf{S} \mathbf{B} \mathbf{R}^{-1} \mathbf{B}^T \mathbf{S} + \mathbf{Q} = 0$$

with

$$\begin{aligned}
\mathcal{A} &= \mathbf{A} - \mathbf{B} \mathbf{R}^{-1} \mathbf{N}^T \\
\mathcal{Q} &= \mathbf{Q} - \mathbf{N} \mathbf{R}^{-1} \mathbf{N}^T
\end{aligned}$$

If there is no cross-correlation between [error](#) and [control effort](#), \mathbf{N} is a zero matrix and the cost function simplifies to

$$J = \int_0^\infty (\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

The feedback [control law](#) which minimizes this J subject to $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ is

$$\mathbf{u} = -\mathbf{K}\mathbf{x}$$

where \mathbf{K} is given by

$$\mathbf{K} = \mathbf{R}^{-1} \mathbf{B}^T \mathbf{S}$$

and \mathbf{S} is found by solving the continuous time algebraic Riccati equation defined as

$$\mathbf{A}^T \mathbf{S} + \mathbf{S} \mathbf{A} - \mathbf{S} \mathbf{B} \mathbf{R}^{-1} \mathbf{B}^T \mathbf{S} + \mathbf{Q} = 0$$

The discrete time LQR [controller](#) is computed via a slightly different cost function, constraint, and resulting algebraic Riccati equation. Snippet E.1 computes the optimal infinite horizon, discrete time LQR [controller](#).

```

import control as ct
import numpy as np
import scipy as sp

def lqr(*args, **kwargs):
    """Solves for the optimal linear-quadratic regulator (LQR).

    For a continuous system:

    .. math:: \dot{x} = A * x + B * u
    .. math:: J = \int_0^\infty (x^T Q x + u^T R u + 2 x^T N u) dt

    For a discrete system:

    .. math:: \dot{x}(n+1) = A x(n) + B u(n)
    .. math:: J = \sum_0^\infty (x^T Q x + u^T R u + 2 x^T N u) \Delta T

    Keyword arguments:
    sys -- StateSpace object representing a linear system.
    Q -- numpy.array(states x states), state cost matrix.
    R -- numpy.array(inputs x inputs), control effort cost matrix.
    N -- numpy.array(states x inputs), cross weight matrix.

    Returns:
    K -- numpy.array(states x inputs), controller gain matrix.
    """
    sys = args[0]
    Q = args[1]
    R = args[2]
    if len(args) == 4:
        N = args[3]
    else:
        N = np.zeros((sys.A.shape[0], sys.B.shape[1]))

    m = sys.A.shape[0]

    controllability_rank = np.linalg.matrix_rank(ct.ctrb(sys.A, sys.B))
    if controllability_rank != m:
        print(
            "Warning: Controllability of %d != %d, uncontrollable state"
            % (controllability_rank, m)
        )

    if sys.dt == None:
        P = sp.linalg.solve_continuous_are(a=sys.A, b=sys.B, q=Q, r=R, s=N)
        return np.linalg.solve(R, sys.B.T @ P + N.T)
    else:
        P = sp.linalg.solve_discrete_are(a=sys.A, b=sys.B, q=Q, r=R, s=N)
        return np.linalg.solve(R + sys.B.T @ P @ sys.B, sys.B.T @ P @ sys.A +
                               N.T)

```

Snippet E.1. Infinite horizon, discrete time LQR computation in Python

Other formulations of LQR for finite horizon and discrete time can be seen on Wikipedia [14].

MIT OpenCourseWare has a rigorous proof of the results shown above [30].

E.2 Output feedback

LQR is normally used for state feedback on

$$\begin{aligned}\dot{\mathbf{x}} &= \mathbf{Ax} + \mathbf{Bu} \\ \mathbf{y} &= \mathbf{Cx} + \mathbf{Du}\end{aligned}$$

with the cost functional

$$J = \int_0^\infty (\mathbf{x}^T \mathbf{Q} \mathbf{x} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

However, we may not know how to select costs for some of the states, or we don't care what certain internal states are doing. Output feedback can accommodate this. Not only can we make our output contain a subset of states, but we can use any other cost metric we can think of as long as it's representable as a linear combination of the states and inputs¹.

For output feedback, we want to minimize the following cost functional.

$$J = \int_0^\infty (\mathbf{y}^T \mathbf{Q} \mathbf{y} + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

Substitute in the expression for \mathbf{y} .

$$J = \int_0^\infty ((\mathbf{Cx} + \mathbf{Du})^T \mathbf{Q} (\mathbf{Cx} + \mathbf{Du}) + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

Apply the transpose to the left-hand side of the \mathbf{Q} term.

$$J = \int_0^\infty ((\mathbf{x}^T \mathbf{C}^T + \mathbf{u}^T \mathbf{D}^T) \mathbf{Q} (\mathbf{Cx} + \mathbf{Du}) + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

Factor out $\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T$ from the left side and $\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}$ from the right side of each term.

$$\begin{aligned}J &= \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \mathbf{C}^T \\ \mathbf{D}^T \end{bmatrix} \mathbf{Q} [\mathbf{C} \quad \mathbf{D}] \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} + \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt \\ J &= \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \left(\begin{bmatrix} \mathbf{C}^T \\ \mathbf{D}^T \end{bmatrix} \mathbf{Q} [\mathbf{C} \quad \mathbf{D}] + \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \right) \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt\end{aligned}$$

Multiply in \mathbf{Q} .

$$J = \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \left(\begin{bmatrix} \mathbf{C}^T \mathbf{Q} \\ \mathbf{D}^T \mathbf{Q} \end{bmatrix} [\mathbf{C} \quad \mathbf{D}] + \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \right) \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt$$

Multiply matrices in the left term together.

$$J = \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \left(\begin{bmatrix} \mathbf{C}^T \mathbf{Q} \mathbf{C} & \mathbf{C}^T \mathbf{Q} \mathbf{D} \\ \mathbf{D}^T \mathbf{Q} \mathbf{C} & \mathbf{D}^T \mathbf{Q} \mathbf{D} \end{bmatrix} + \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \right) \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt$$

¹We'll see this later on in section E.3 when we define the cost metric as deviation from the behavior of another model.

Add the terms together.

$$J = \int_0^\infty \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \underbrace{\mathbf{C}^T \mathbf{Q} \mathbf{C}}_{\mathbf{Q}} & \underbrace{\mathbf{C}^T \mathbf{Q} \mathbf{D}}_{\mathbf{N}} \\ \underbrace{\mathbf{D}^T \mathbf{Q} \mathbf{C}}_{\mathbf{N}^T} & \underbrace{\mathbf{D}^T \mathbf{Q} \mathbf{D} + \mathbf{R}}_{\mathbf{R}} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt \quad (\text{E.2})$$

Thus, output feedback can be defined as the following optimization problem.

Theorem E.2.1 — Linear-quadratic regulator with output feedback.

$$\min_{\mathbf{u}} \int_0^\infty \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \underbrace{\mathbf{C}^T \mathbf{Q} \mathbf{C}}_{\mathbf{Q}} & \underbrace{\mathbf{C}^T \mathbf{Q} \mathbf{D}}_{\mathbf{N}} \\ \underbrace{\mathbf{D}^T \mathbf{Q} \mathbf{C}}_{\mathbf{N}^T} & \underbrace{\mathbf{D}^T \mathbf{Q} \mathbf{D} + \mathbf{R}}_{\mathbf{R}} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt$$

subject to $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$ (E.3)

The optimal control policy \mathbf{u}^* is $\mathbf{K}(\mathbf{r} - \mathbf{y})$ where \mathbf{r} is the desired output and \mathbf{y} is the measured output defined as $\mathbf{y} = \mathbf{C}\mathbf{x} + \mathbf{D}\mathbf{u}$. \mathbf{K} can be computed via the typical LQR equations based on the algebraic Riccati equation.

If the output is only the whole state vector, then $\mathbf{C} = \mathbf{I}$, $\mathbf{D} = \mathbf{0}$, and the cost functional simplifies to that of state feedback LQR.

$$J = \int_0^\infty \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \mathbf{Q} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt$$

E.3 Implicit model following

If we want to design a feedback controller that erases the dynamics of our system and makes it behave like some other system, we can use *implicit model following*. This is used on the Blackhawk helicopter at NASA Ames research center when they want to make it fly like experimental aircraft (within the limits of the helicopter's actuators, of course).

Let the original system dynamics be

$$\begin{aligned} \dot{\mathbf{x}} &= \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u} \\ \mathbf{y} &= \mathbf{C}\mathbf{x} \end{aligned}$$

and the desired system dynamics be

$$\dot{\mathbf{z}} = \mathbf{A}_{ref}\mathbf{z}$$

$$\begin{aligned} \dot{\mathbf{y}} &= \mathbf{C}\dot{\mathbf{x}} \\ \dot{\mathbf{y}} &= \mathbf{C}(\mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}) \\ \dot{\mathbf{y}} &= \mathbf{C}\mathbf{A}\mathbf{x} + \mathbf{C}\mathbf{B}\mathbf{u} \end{aligned}$$

We want to minimize the following cost functional.

$$J = \int_0^\infty ((\dot{\mathbf{y}} - \dot{\mathbf{z}})^T \mathbf{Q} (\dot{\mathbf{y}} - \dot{\mathbf{z}}) + \mathbf{u}^T \mathbf{R} \mathbf{u}) dt$$

We'll be measuring the desired system's state, so let $\mathbf{y} = \mathbf{z}$.

$$\begin{aligned}\dot{\mathbf{z}} &= \mathbf{A}_{ref}\mathbf{y} \\ \dot{\mathbf{z}} &= \mathbf{A}_{ref}\mathbf{Cx}\end{aligned}$$

Therefore,

$$\begin{aligned}\dot{\mathbf{y}} - \dot{\mathbf{z}} &= \mathbf{CAx} + \mathbf{CBu} - (\mathbf{A}_{ref}\mathbf{Cx}) \\ \dot{\mathbf{y}} - \dot{\mathbf{z}} &= (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})\mathbf{x} + \mathbf{CBu}\end{aligned}$$

Substitute this into the cost functional.

$$\begin{aligned}J &= \int_0^\infty ((\dot{\mathbf{y}} - \dot{\mathbf{z}})^T \mathbf{Q}(\dot{\mathbf{y}} - \dot{\mathbf{z}}) + \mathbf{u}^T \mathbf{Ru}) dt \\ J &= \int_0^\infty (((\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})\mathbf{x} + \mathbf{CBu})^T \mathbf{Q}((\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})\mathbf{x} + \mathbf{CBu}) + \mathbf{u}^T \mathbf{Ru}) dt\end{aligned}$$

Apply the transpose to the left-hand side of the \mathbf{Q} term.

$$J = \int_0^\infty ((\mathbf{x}^T (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T + \mathbf{u}^T (\mathbf{CB})^T) \mathbf{Q}((\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})\mathbf{x} + \mathbf{CBu}) + \mathbf{u}^T \mathbf{Ru}) dt$$

Factor out $\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T$ from the left side and $\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}$ from the right side of each term.

$$\begin{aligned}J &= \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \\ (\mathbf{CB})^T \end{bmatrix} \mathbf{Q} [\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C} \quad \mathbf{CB}] \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} + \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt \\ J &= \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \left(\begin{bmatrix} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \\ (\mathbf{CB})^T \end{bmatrix} \mathbf{Q} [\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C} \quad \mathbf{CB}] + \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \right) \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt\end{aligned}$$

Multiply in \mathbf{Q} .

$$J = \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \left(\begin{bmatrix} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q} \\ (\mathbf{CB})^T \mathbf{Q} \end{bmatrix} [\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C} \quad \mathbf{CB}] + \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \right) \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt$$

Multiply matrices in the left term together.

$$J = \int_0^\infty \left(\begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \left(\begin{bmatrix} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C}) & (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q} (\mathbf{CB}) \\ (\mathbf{CB})^T \mathbf{Q} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C}) & (\mathbf{CB})^T \mathbf{Q} (\mathbf{CB}) \end{bmatrix} + \begin{bmatrix} \mathbf{0} & \mathbf{0} \\ \mathbf{0} & \mathbf{R} \end{bmatrix} \right) \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} \right) dt$$

Add the terms together.

$$J = \int_0^\infty \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \underbrace{(\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})}_{\mathbf{Q}} & \underbrace{(\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q} (\mathbf{CB})}_{\mathbf{N}} \\ \underbrace{(\mathbf{CB})^T \mathbf{Q} (\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})}_{\mathbf{N}^T} & \underbrace{(\mathbf{CB})^T \mathbf{Q} (\mathbf{CB}) + \mathbf{R}}_{\mathbf{R}} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt \quad (\text{E.4})$$

Thus, implicit model following can be defined as the following optimization problem.

Theorem E.3.1 — Implicit model following.

$$\min_{\mathbf{u}} \int_0^\infty \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \underbrace{(\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q}(\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})}_{\mathbf{Q}} & \underbrace{(\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})^T \mathbf{Q}(\mathbf{CB})}_{\mathbf{N}} \\ \underbrace{(\mathbf{CB})^T \mathbf{Q}(\mathbf{CA} - \mathbf{A}_{ref}\mathbf{C})}_{\mathbf{N}^T} & \underbrace{(\mathbf{CB})^T \mathbf{Q}(\mathbf{CB}) + \mathbf{R}}_{\mathbf{R}} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt$$

subject to $\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}$ (E.5)

The optimal control policy \mathbf{u}^* is $-\mathbf{Kx}$. \mathbf{K} can be computed via the typical LQR equations based on the algebraic Riccati equation.

The control law $\mathbf{u}_{imf} = -\mathbf{Kx}$ makes $\dot{\mathbf{x}} = \mathbf{Ax} + \mathbf{Bu}_{imf}$ match the open-loop response of $\dot{\mathbf{z}} = \mathbf{A}_{ref}\mathbf{z}$.

If the original and desired system have the same states, then $\mathbf{C} = \mathbf{I}$ and the cost functional simplifies to

$$J = \int_0^\infty \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix}^T \begin{bmatrix} \underbrace{(\mathbf{A} - \mathbf{A}_{ref})^T \mathbf{Q}(\mathbf{A} - \mathbf{A}_{ref})}_{\mathbf{Q}} & \underbrace{(\mathbf{A} - \mathbf{A}_{ref})^T \mathbf{QB}}_{\mathbf{N}} \\ \underbrace{\mathbf{B}^T \mathbf{Q}(\mathbf{A} - \mathbf{A}_{ref})}_{\mathbf{N}^T} & \underbrace{\mathbf{B}^T \mathbf{QB} + \mathbf{R}}_{\mathbf{R}} \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ \mathbf{u} \end{bmatrix} dt (E.6)$$

This page intentionally left blank

F. QR-weighted linear plant inversion

F.1 Necessary theorems

The following theorem and corollary will be needed to derive the QR-weighted linear plant inversion equation.

Theorem F.1.1 $\frac{\partial(\mathbf{Ax}+\mathbf{b})^T \mathbf{C}(\mathbf{Dx}+\mathbf{e})}{\partial \mathbf{x}} = \mathbf{A}^T \mathbf{C}(\mathbf{Dx} + \mathbf{e}) + \mathbf{D}^T \mathbf{C}^T(\mathbf{Ax} + \mathbf{b})$

Corollary F.1.2 $\frac{\partial(\mathbf{Ax}+\mathbf{b})^T \mathbf{C}(\mathbf{Ax}+\mathbf{b})}{\partial \mathbf{x}} = 2\mathbf{A}^T \mathbf{C}(\mathbf{Ax} + \mathbf{b})$ where \mathbf{C} is symmetric.

Proof:

$$\begin{aligned} \frac{\partial(\mathbf{Ax}+\mathbf{b})^T \mathbf{C}(\mathbf{Ax}+\mathbf{b})}{\partial \mathbf{x}} &= \mathbf{A}^T \mathbf{C}(\mathbf{Ax} + \mathbf{b}) + \mathbf{A}^T \mathbf{C}^T(\mathbf{Ax} + \mathbf{b}) \\ \frac{\partial(\mathbf{Ax}+\mathbf{b})^T \mathbf{C}(\mathbf{Ax}+\mathbf{b})}{\partial \mathbf{x}} &= (\mathbf{A}^T \mathbf{C} + \mathbf{A}^T \mathbf{C}^T)(\mathbf{Ax} + \mathbf{b}) \end{aligned}$$

\mathbf{C} is symmetric, so

$$\begin{aligned} \frac{\partial(\mathbf{Ax}+\mathbf{b})^T \mathbf{C}(\mathbf{Ax}+\mathbf{b})}{\partial \mathbf{x}} &= (\mathbf{A}^T \mathbf{C} + \mathbf{A}^T \mathbf{C})(\mathbf{Ax} + \mathbf{b}) \\ \frac{\partial(\mathbf{Ax}+\mathbf{b})^T \mathbf{C}(\mathbf{Ax}+\mathbf{b})}{\partial \mathbf{x}} &= 2\mathbf{A}^T \mathbf{C}(\mathbf{Ax} + \mathbf{b}) \end{aligned}$$

F.2 Setup

Let's start with the equation for the [reference](#) dynamics

$$\mathbf{r}_{k+1} = \mathbf{A}\mathbf{r}_k + \mathbf{B}\mathbf{u}_k$$

where \mathbf{u}_k is the feedforward input. Note that this feedforward equation does not and should not take into account any feedback terms. We want to find the optimal \mathbf{u}_k such that we minimize the [tracking](#) error between \mathbf{r}_{k+1} and \mathbf{r}_k .

$$\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k = \mathbf{B}\mathbf{u}_k$$

To solve for \mathbf{u}_k , we need to take the inverse of the nonsquare matrix \mathbf{B} . This isn't possible, but we can find the pseudoinverse given some constraints on the [state tracking](#) error and [control effort](#). To find the optimal solution for these sorts of trade-offs, one can define a cost function and attempt to minimize it. To do this, we'll first solve the expression for 0.

$$\mathbf{0} = \mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)$$

This expression will be the [state tracking](#) cost we use in our cost function.

Our cost function will use an H_2 norm with \mathbf{Q} as the [state](#) cost matrix with dimensionality $states \times states$ and \mathbf{R} as the [control input](#) cost matrix with dimensionality $inputs \times inputs$.

$$\mathbf{J} = (\mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k))^T \mathbf{Q} (\mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)) + \mathbf{u}_k^T \mathbf{R} \mathbf{u}_k$$

F.3 Minimization

Given theorem 5.10.1 and corollary F.1.2, find the minimum of \mathbf{J} by taking the partial derivative with respect to \mathbf{u}_k and setting the result to 0.

$$\begin{aligned} \frac{\partial \mathbf{J}}{\partial \mathbf{u}_k} &= 2\mathbf{B}^T \mathbf{Q} (\mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)) + 2\mathbf{R}\mathbf{u}_k \\ \mathbf{0} &= 2\mathbf{B}^T \mathbf{Q} (\mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)) + 2\mathbf{R}\mathbf{u}_k \\ \mathbf{0} &= \mathbf{B}^T \mathbf{Q} (\mathbf{B}\mathbf{u}_k - (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)) + \mathbf{R}\mathbf{u}_k \\ \mathbf{0} &= \mathbf{B}^T \mathbf{Q} \mathbf{B}\mathbf{u}_k - \mathbf{B}^T \mathbf{Q} (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) + \mathbf{R}\mathbf{u}_k \\ \mathbf{B}^T \mathbf{Q} \mathbf{B}\mathbf{u}_k + \mathbf{R}\mathbf{u}_k &= \mathbf{B}^T \mathbf{Q} (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \\ (\mathbf{B}^T \mathbf{Q} \mathbf{B} + \mathbf{R})\mathbf{u}_k &= \mathbf{B}^T \mathbf{Q} (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \\ \mathbf{u}_k &= (\mathbf{B}^T \mathbf{Q} \mathbf{B} + \mathbf{R})^{-1} \mathbf{B}^T \mathbf{Q} (\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k) \end{aligned}$$

Theorem F.3.1 — QR-weighted linear plant inversion. Given the discrete model $\mathbf{x}_{k+1} = \mathbf{A}\mathbf{x}_k + \mathbf{B}\mathbf{u}_k$, the plant inversion feedforward is

$$\mathbf{u}_k = \mathbf{K}_{ff}(\mathbf{r}_{k+1} - \mathbf{A}\mathbf{r}_k)$$

where $\mathbf{K}_{ff} = (\mathbf{B}^T \mathbf{Q} \mathbf{B} + \mathbf{R})^{-1} \mathbf{B}^T \mathbf{Q}$, \mathbf{r}_{k+1} is the reference at the next timestep, and \mathbf{r}_k is the reference at the current timestep.

Figure F.1 shows [plant](#) inversion applied to a second-order CIM motor model.

[Plant](#) inversion isn't as effective with both \mathbf{Q} and \mathbf{R} cost because the \mathbf{R} matrix penalized [control effort](#). The [reference tracking](#) with no cost matrices is much better.

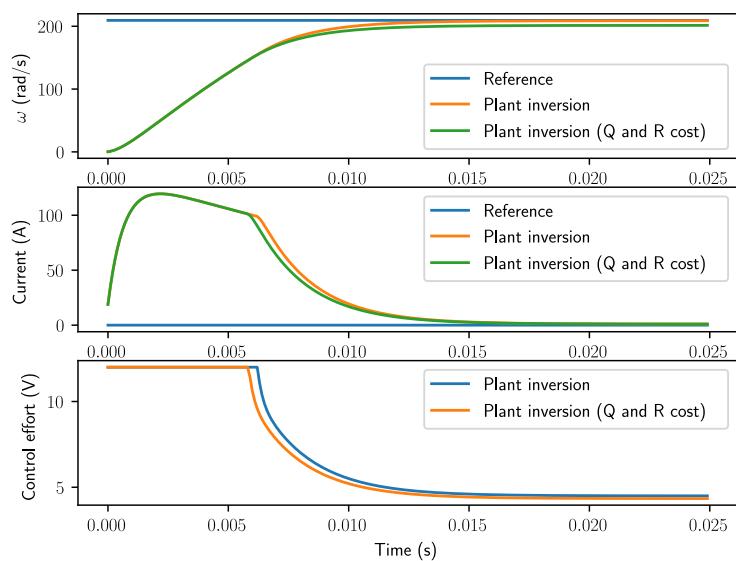


Figure F.1: Second-order CIM motor response with plant inversion

This page intentionally left blank

G. Steady-state feedforward

Steady-state feedforwards apply the [control effort](#) required to keep a [system](#) at the [reference](#) if it is no longer moving (i.e., the [system](#) is at steady-state). The first steady-state feedforward converts desired [outputs](#) to desired [states](#).

$$\mathbf{x}_c = \mathbf{N}_x \mathbf{y}_c$$

\mathbf{N}_x converts desired [outputs](#) \mathbf{y}_c to desired [states](#) \mathbf{x}_c (also known as \mathbf{r}). For steady-state, that is

$$\mathbf{x}_{ss} = \mathbf{N}_x \mathbf{y}_{ss} \quad (G.1)$$

The second steady-state feedforward converts the desired [outputs](#) \mathbf{y} to the [control input](#) required at steady-state.

$$\mathbf{u}_c = \mathbf{N}_u \mathbf{y}_c$$

\mathbf{N}_u converts the desired [outputs](#) \mathbf{y} to the [control input](#) \mathbf{u} required at steady-state. For steady-state, that is

$$\mathbf{u}_{ss} = \mathbf{N}_u \mathbf{y}_{ss} \quad (G.2)$$

G.1 Continuous case

To find the [control input](#) required at steady-state, set equation (5.1) to zero.

$$\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}\mathbf{u}$$

$$\mathbf{y} = \mathbf{C}\mathbf{x} + \mathbf{D}\mathbf{u}$$

$$\mathbf{0} = \mathbf{A}\mathbf{x}_{ss} + \mathbf{B}\mathbf{u}_{ss}$$

$$\mathbf{y}_{ss} = \mathbf{C}\mathbf{x}_{ss} + \mathbf{D}\mathbf{u}_{ss}$$

$$\mathbf{0} = \mathbf{A}\mathbf{N}_x \mathbf{y}_{ss} + \mathbf{B}\mathbf{N}_u \mathbf{y}_{ss}$$

$$\mathbf{y}_{ss} = \mathbf{C}\mathbf{N}_x\mathbf{y}_{ss} + \mathbf{D}\mathbf{N}_u\mathbf{y}_{ss}$$

$$\begin{aligned}\begin{bmatrix} \mathbf{0} \\ \mathbf{y}_{ss} \end{bmatrix} &= \begin{bmatrix} \mathbf{A}\mathbf{N}_x + \mathbf{B}\mathbf{N}_u \\ \mathbf{C}\mathbf{N}_x + \mathbf{D}\mathbf{N}_u \end{bmatrix} \mathbf{y}_{ss} \\ \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} &= \begin{bmatrix} \mathbf{A}\mathbf{N}_x + \mathbf{B}\mathbf{N}_u \\ \mathbf{C}\mathbf{N}_x + \mathbf{D}\mathbf{N}_u \end{bmatrix} \\ \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} &= \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix} \begin{bmatrix} \mathbf{N}_x \\ \mathbf{N}_u \end{bmatrix} \\ \begin{bmatrix} \mathbf{N}_x \\ \mathbf{N}_u \end{bmatrix} &= \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}^\dagger \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix}\end{aligned}$$

where † is the Moore-Penrose pseudoinverse.

G.2 Discrete case

Now, we'll do the same thing for the discrete system. To find the control input required at steady-state, set equation (5.3) to zero.

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{A}\mathbf{x}_k + \mathbf{B}\mathbf{u}_k \\ \mathbf{y}_k &= \mathbf{C}\mathbf{x}_k + \mathbf{D}\mathbf{u}_k\end{aligned}$$

$$\begin{aligned}\mathbf{x}_{ss} &= \mathbf{A}\mathbf{x}_{ss} + \mathbf{B}\mathbf{u}_{ss} \\ \mathbf{y}_{ss} &= \mathbf{C}\mathbf{x}_{ss} + \mathbf{D}\mathbf{u}_{ss}\end{aligned}$$

$$\begin{aligned}\mathbf{0} &= (\mathbf{A} - \mathbf{I})\mathbf{x}_{ss} + \mathbf{B}\mathbf{u}_{ss} \\ \mathbf{y}_{ss} &= \mathbf{C}\mathbf{x}_{ss} + \mathbf{D}\mathbf{u}_{ss}\end{aligned}$$

$$\begin{aligned}\mathbf{0} &= (\mathbf{A} - \mathbf{I})\mathbf{N}_x\mathbf{y}_{ss} + \mathbf{B}\mathbf{N}_u\mathbf{y}_{ss} \\ \mathbf{y}_{ss} &= \mathbf{C}\mathbf{N}_x\mathbf{y}_{ss} + \mathbf{D}\mathbf{N}_u\mathbf{y}_{ss}\end{aligned}$$

$$\begin{aligned}\begin{bmatrix} \mathbf{0} \\ \mathbf{y}_{ss} \end{bmatrix} &= \begin{bmatrix} (\mathbf{A} - \mathbf{I})\mathbf{N}_x + \mathbf{B}\mathbf{N}_u \\ \mathbf{C}\mathbf{N}_x + \mathbf{D}\mathbf{N}_u \end{bmatrix} \mathbf{y}_{ss} \\ \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} &= \begin{bmatrix} (\mathbf{A} - \mathbf{I})\mathbf{N}_x + \mathbf{B}\mathbf{N}_u \\ \mathbf{C}\mathbf{N}_x + \mathbf{D}\mathbf{N}_u \end{bmatrix} \\ \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} &= \begin{bmatrix} \mathbf{A} - \mathbf{I} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix} \begin{bmatrix} \mathbf{N}_x \\ \mathbf{N}_u \end{bmatrix} \\ \begin{bmatrix} \mathbf{N}_x \\ \mathbf{N}_u \end{bmatrix} &= \begin{bmatrix} \mathbf{A} - \mathbf{I} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}^\dagger \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix}\end{aligned}$$

where † is the Moore-Penrose pseudoinverse.

G.3 Deriving steady-state input

Now, we'll find an expression that uses \mathbf{N}_x and \mathbf{N}_u to convert the reference \mathbf{r} to a control input feedforward \mathbf{u}_{ff} . Let's start with equation (G.1).

$$\begin{aligned}\mathbf{x}_{ss} &= \mathbf{N}_x \mathbf{y}_{ss} \\ \mathbf{N}_x^\dagger \mathbf{x}_{ss} &= \mathbf{y}_{ss}\end{aligned}$$

Now substitute this into equation (G.2).

$$\begin{aligned}\mathbf{u}_{ss} &= \mathbf{N}_u \mathbf{y}_{ss} \\ \mathbf{u}_{ss} &= \mathbf{N}_u (\mathbf{N}_x^\dagger \mathbf{x}_{ss}) \\ \mathbf{u}_{ss} &= \mathbf{N}_u \mathbf{N}_x^\dagger \mathbf{x}_{ss}\end{aligned}$$

\mathbf{u}_{ss} and \mathbf{x}_{ss} are also known as \mathbf{u}_{ff} and \mathbf{r} respectively.

$$\mathbf{u}_{ff} = \mathbf{N}_u \mathbf{N}_x^\dagger \mathbf{r}$$

So all together, we get theorem G.3.1.

Theorem G.3.1 — Steady-state feedforward.

Continuous:

$$\begin{bmatrix} \mathbf{N}_x \\ \mathbf{N}_u \end{bmatrix} = \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}^\dagger \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} \quad (\text{G.3})$$

Discrete:

$$\begin{bmatrix} \mathbf{N}_x \\ \mathbf{N}_u \end{bmatrix} = \begin{bmatrix} \mathbf{A} - \mathbf{I} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}^\dagger \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} \quad (\text{G.4})$$

$$\mathbf{u}_{ff} = \mathbf{N}_u \mathbf{N}_x^\dagger \mathbf{r} \quad (\text{G.5})$$

In the augmented matrix, \mathbf{B} should contain one column corresponding to an actuator and \mathbf{C} should contain one row whose [output](#) will be driven by that actuator. More than one actuator or output can be included in the computation at once, but the result won't be the same as if they were computed independently and summed afterward.

After computing the feedforward for each actuator-output pair, the respective collections of \mathbf{N}_x and \mathbf{N}_u matrices can be summed to produce the combined feedforward.

If the augmented matrix in theorem G.3.1 is square (number of [inputs](#) = number of [outputs](#)), the normal matrix inverse can be used instead.

G.3.1 Case study: second-order CIM motor model

Each feedforward implementation has advantages. The steady-state feedforward allows using specific actuators to maintain the [reference](#) at steady-state. [Plant](#) inversion doesn't support this, but can be used for [reference tracking](#) as well with the same tuning parameters as LQR design. Figure G.1 shows both types of feedforwards applied to a second-order CIM motor model.

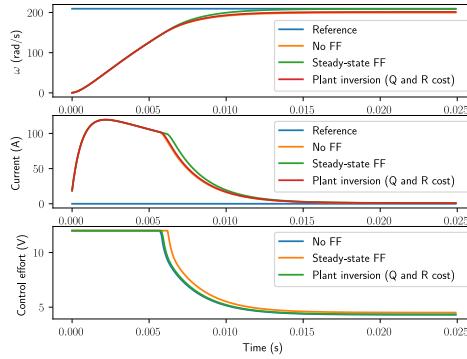


Figure G.1: Second-order CIM motor response with various feedforwards

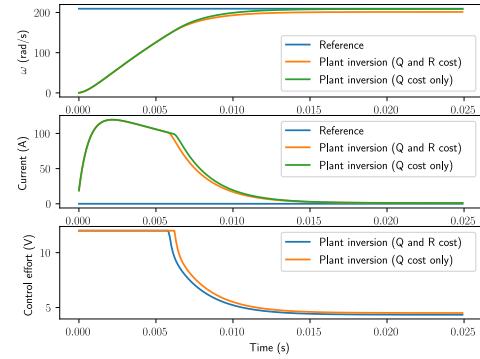


Figure G.2: Second-order CIM motor response with plant inversions

Plant inversion isn't as effective in figure G.1 because the **R** matrix penalized control effort. If the **R** cost matrix is removed from the plant inversion calculation, the reference tracking is much better (see figure G.2).

H. Derivations

H.1 Transfer function in feedback

Given the feedback network in figure H.1, find an expression for $Y(s)$.

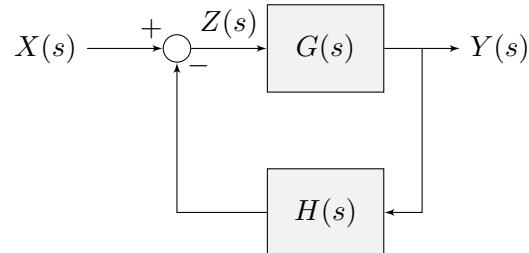


Figure H.1: Closed-loop block diagram

$$\begin{aligned}
 Y(s) &= Z(s)G(s) \\
 Z(s) &= X(s) - Y(s)H(s) \\
 X(s) &= Z(s) + Y(s)H(s) \\
 X(s) &= Z(s) + Z(s)G(s)H(s) \\
 \frac{Y(s)}{X(s)} &= \frac{Z(s)G(s)}{Z(s) + Z(s)G(s)H(s)} \\
 \frac{Y(s)}{X(s)} &= \frac{G(s)}{1 + G(s)H(s)}
 \end{aligned} \tag{H.1}$$

A more general form is

$$\frac{Y(s)}{X(s)} = \frac{G(s)}{1 \mp G(s)H(s)} \tag{H.2}$$

where positive feedback uses the top sign and negative feedback uses the bottom sign.

H.2 Zero-order hold for state-space

Starting with the continuous model

$$\dot{\mathbf{x}}(t) = \mathbf{A}\mathbf{x}(t) + \mathbf{B}\mathbf{u}(t)$$

by premultiplying the model by $e^{-\mathbf{A}t}$, we get

$$\begin{aligned} e^{-\mathbf{A}t}\dot{\mathbf{x}}(t) &= e^{-\mathbf{A}t}\mathbf{A}\mathbf{x}(t) + e^{-\mathbf{A}t}\mathbf{B}\mathbf{u}(t) \\ e^{-\mathbf{A}t}\dot{\mathbf{x}}(t) - e^{-\mathbf{A}t}\mathbf{A}\mathbf{x}(t) &= e^{-\mathbf{A}t}\mathbf{B}\mathbf{u}(t) \end{aligned}$$

The derivative of the matrix exponential is

$$\frac{d}{dt}e^{\mathbf{A}t} = \mathbf{A}e^{\mathbf{A}t} = e^{\mathbf{A}t}\mathbf{A}$$

so we recognize the previous equation as

$$\frac{d}{dt}(e^{-\mathbf{A}t}\mathbf{x}(t)) = e^{-\mathbf{A}t}\mathbf{B}\mathbf{u}(t)$$

By integrating this equation, we get

$$\begin{aligned} e^{-\mathbf{A}t}\mathbf{x}(t) - e^0\mathbf{x}(0) &= \int_0^t e^{-\mathbf{A}\tau}\mathbf{B}\mathbf{u}(\tau) d\tau \\ \mathbf{x}(t) &= e^{\mathbf{A}t}\mathbf{x}(0) + \int_0^t e^{\mathbf{A}(t-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \end{aligned}$$

which is an analytical solution to the continuous model. Now we want to discretize it.

$$\begin{aligned} \mathbf{x}_k &\stackrel{def}{=} \mathbf{x}(kT) \\ \mathbf{x}_k &= e^{\mathbf{A}kT}\mathbf{x}(0) + \int_0^{kT} e^{\mathbf{A}(kT-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}(k+1)T}\mathbf{x}(0) + \int_0^{(k+1)T} e^{\mathbf{A}((k+1)T-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}(k+1)T}\mathbf{x}(0) + \int_0^{kT} e^{\mathbf{A}((k+1)T-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau + \int_{kT}^{(k+1)T} e^{\mathbf{A}((k+1)T-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}(k+1)T}\mathbf{x}(0) + \int_0^{kT} e^{\mathbf{A}((k+1)T-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau + \int_{kT}^{(k+1)T} e^{\mathbf{A}(kT+T-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}T} \underbrace{\left(e^{\mathbf{A}kT}\mathbf{x}(0) + \int_0^{kT} e^{\mathbf{A}(kT-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \right)}_{\mathbf{x}_k} + \int_{kT}^{(k+1)T} e^{\mathbf{A}(kT+T-\tau)}\mathbf{B}\mathbf{u}(\tau) d\tau \end{aligned}$$

We assume that \mathbf{u} is constant during each timestep, so it can be pulled out of the integral.

$$\mathbf{x}_{k+1} = e^{\mathbf{A}T}\mathbf{x}_k + \left(\int_{kT}^{(k+1)T} e^{\mathbf{A}(kT+T-\tau)} d\tau \right) \mathbf{B}\mathbf{u}_k$$

The second term can be simplified by substituting it with the function $v(\tau) = kT + T - \tau$. Note that $d\tau = -dv$.

$$\begin{aligned}\mathbf{x}_{k+1} &= e^{\mathbf{A}T} \mathbf{x}_k - \left(\int_{v(kT)}^{v((k+1)T)} e^{\mathbf{A}v} dv \right) \mathbf{B} \mathbf{u}_k \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}T} \mathbf{x}_k - \left(\int_T^0 e^{\mathbf{A}v} dv \right) \mathbf{B} \mathbf{u}_k \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}T} \mathbf{x}_k + \left(\int_0^T e^{\mathbf{A}v} dv \right) \mathbf{B} \mathbf{u}_k \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}T} \mathbf{x}_k + \mathbf{A}^{-1} e^{\mathbf{A}v} \Big|_0^T \mathbf{B} \mathbf{u}_k \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}T} \mathbf{x}_k + \mathbf{A}^{-1} (e^{\mathbf{A}T} - e^{\mathbf{A}0}) \mathbf{B} \mathbf{u}_k \\ \mathbf{x}_{k+1} &= e^{\mathbf{A}T} \mathbf{x}_k + \mathbf{A}^{-1} (e^{\mathbf{A}T} - \mathbf{I}) \mathbf{B} \mathbf{u}_k\end{aligned}$$

which is an exact solution to the [discretization](#) problem.

H.3 Kalman filter as Luenberger observer

A Luenberger [observer](#) is defined as

$$\hat{\mathbf{x}}_{k+1}^+ = \mathbf{A} \hat{\mathbf{x}}_k^- + \mathbf{B} \mathbf{u}_k + \mathbf{L} (\mathbf{y}_k - \hat{\mathbf{y}}_k) \quad (\text{H.3})$$

$$\hat{\mathbf{y}}_k = \mathbf{C} \hat{\mathbf{x}}_k^- \quad (\text{H.4})$$

where a superscript of minus denotes *a priori* and plus denotes *a posteriori* estimate. Combining equation (H.3) and equation (H.4) gives

$$\hat{\mathbf{x}}_{k+1}^+ = \mathbf{A} \hat{\mathbf{x}}_k^- + \mathbf{B} \mathbf{u}_k + \mathbf{L} (\mathbf{y}_k - \mathbf{C} \hat{\mathbf{x}}_k^-) \quad (\text{H.5})$$

The following is a Kalman filter that considers the current update step and the next predict step together rather than the current predict step and current update step.

Update step

$$\mathbf{K}_k = \mathbf{P}_k^- \mathbf{C}^T (\mathbf{C} \mathbf{P}_k^- \mathbf{C}^T + \mathbf{R})^{-1} \quad (\text{H.6})$$

$$\hat{\mathbf{x}}_k^+ = \hat{\mathbf{x}}_k^- + \mathbf{K}_k (\mathbf{y}_k - \mathbf{C} \hat{\mathbf{x}}_k^-) \quad (\text{H.7})$$

$$\mathbf{P}_k^+ = (\mathbf{I} - \mathbf{K}_k \mathbf{C}) \mathbf{P}_k^- \quad (\text{H.8})$$

Predict step

$$\hat{\mathbf{x}}_{k+1}^+ = \mathbf{A} \hat{\mathbf{x}}_k^+ + \mathbf{B} \mathbf{u}_k \quad (\text{H.9})$$

$$\mathbf{P}_{k+1}^- = \mathbf{A} \mathbf{P}_k^+ \mathbf{A}^T + \mathbf{\Gamma} \mathbf{Q} \mathbf{\Gamma}^T \quad (\text{H.10})$$

Substitute equation (H.7) into equation (H.9).

$$\begin{aligned}\hat{\mathbf{x}}_{k+1}^+ &= \mathbf{A} (\hat{\mathbf{x}}_k^- + \mathbf{K}_k (\mathbf{y}_k - \mathbf{C} \hat{\mathbf{x}}_k^-)) + \mathbf{B} \mathbf{u}_k \\ \hat{\mathbf{x}}_{k+1}^+ &= \mathbf{A} \hat{\mathbf{x}}_k^- + \mathbf{A} \mathbf{K}_k (\mathbf{y}_k - \mathbf{C} \hat{\mathbf{x}}_k^-) + \mathbf{B} \mathbf{u}_k \\ \hat{\mathbf{x}}_{k+1}^+ &= \mathbf{A} \hat{\mathbf{x}}_k^- + \mathbf{B} \mathbf{u}_k + \mathbf{A} \mathbf{K}_k (\mathbf{y}_k - \mathbf{C} \hat{\mathbf{x}}_k^-)\end{aligned}$$

Let $\mathbf{L} = \mathbf{A} \mathbf{K}_k$.

$$\hat{\mathbf{x}}_{k+1}^+ = \mathbf{A} \hat{\mathbf{x}}_k^- + \mathbf{B} \mathbf{u}_k + \mathbf{L} (\mathbf{y}_k - \mathbf{C} \hat{\mathbf{x}}_k^-) \quad (\text{H.11})$$

which matches equation (H.5). Therefore, the eigenvalues of the Kalman filter [observer](#) can be obtained by

$$\begin{aligned} & \text{eig}(\mathbf{A} - \mathbf{LC}) \\ & \text{eig}(\mathbf{A} - (\mathbf{AK}_k)(\mathbf{C})) \\ & \text{eig}(\mathbf{A}(\mathbf{I} - \mathbf{K}_k\mathbf{C})) \end{aligned} \quad (\text{H.12})$$

H.3.1 Luenberger observer with separate prediction and update

To run a Luenberger [observer](#) with separate prediction and update steps, substitute the relationship between the Luenberger [observer](#) and Kalman filter matrices derived above into the Kalman filter equations.

Appendix H.3 shows that $\mathbf{L} = \mathbf{AK}_k$. Since \mathbf{L} and \mathbf{A} are constant, one must assume \mathbf{K}_k has reached steady-state. Then, $\mathbf{K} = \mathbf{A}^{-1}\mathbf{L}$. Substitute this into the Kalman filter update equation.

$$\begin{aligned} \hat{\mathbf{x}}_{k+1}^+ &= \hat{\mathbf{x}}_{k+1}^- + \mathbf{K}(\mathbf{y}_{k+1} - \mathbf{C}\hat{\mathbf{x}}_{k+1}^-) \\ \hat{\mathbf{x}}_{k+1}^+ &= \hat{\mathbf{x}}_{k+1}^- + \mathbf{A}^{-1}\mathbf{L}(\mathbf{y}_{k+1} - \mathbf{C}\hat{\mathbf{x}}_{k+1}^-) \end{aligned}$$

Substitute in equation (9.4).

$$\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{A}^{-1}\mathbf{L}(\mathbf{y}_{k+1} - \hat{\mathbf{y}}_{k+1})$$

The predict step is the same as the Kalman filter's. Therefore, a Luenberger [observer](#) run with prediction and update steps is written as follows.

Predict step

$$\hat{\mathbf{x}}_{k+1}^- = \mathbf{A}\hat{\mathbf{x}}_k^- + \mathbf{B}\mathbf{u}_k \quad (\text{H.13})$$

Update step

$$\hat{\mathbf{x}}_{k+1}^+ = \hat{\mathbf{x}}_{k+1}^- + \mathbf{A}^{-1}\mathbf{L}(\mathbf{y}_{k+1} - \hat{\mathbf{y}}_{k+1}) \quad (\text{H.14})$$

$$\hat{\mathbf{y}}_{k+1} = \mathbf{C}\hat{\mathbf{x}}_{k+1}^- \quad (\text{H.15})$$

H.4 Trapezoidal motion profile

Glossary

agent An independent actor being controlled through autonomy or human-in-the-loop (e.g., a robot, aircraft, etc.).

control effort A term describing how much force, pressure, etc. an actuator is exerting.

control input The input of a [plant](#) used for the purpose of controlling it.

control law Also known as control policy, is a mathematical formula used by the [controller](#) to determine the [input](#) u that is sent to the [plant](#). This control law is designed to drive the [system](#) from its current [state](#) to some other desired [state](#).

control system Monitors and controls the behavior of a [system](#).

controller Applies an [input](#) to a [plant](#) to bring about a desired [system state](#) by driving the difference between a [reference signal](#) and the [output](#) to zero.

discretization The process by which a continuous (e.g., analog) [system](#) or [controller](#) design is converted to discrete (e.g., digital).

disturbance An external force acting on a [system](#) that isn't included in the [system's model](#).

disturbance rejection The quality of a feedback control [system](#) to compensate for external forces to reach a desired [reference](#).

error [Reference](#) minus an [output](#) or [state](#).

feedback controller Used in positive or negative feedback with a [plant](#) to bring about a desired [system state](#) by driving the difference between a [reference signal](#) and the [output](#) to zero.

feedback gain The gain from the [output](#) to an earlier point in a [control system](#) diagram.

feedforward controller A [controller](#) that injects information about the [system's dynamics](#) (like a [model](#) does) or the desired movement. The feedforward handles parts of the control actions we already know must be applied to make a [system](#) track a [reference](#), then the feedback controller compensates for what we do not or cannot know about the [system's behavior](#) at runtime.

gain A proportional value that shows the relationship between the magnitude of an input signal to the magnitude of an output signal at steady-state.

gain margin See section C.1 on gain and phase margin.

impulse response The response of a [system](#) to the Dirac delta function.

input An input to the [plant](#) (hence the name) that can be used to change the [plant's state](#).

linearization A method by which a nonlinear [system](#)'s dynamics are approximated by a linear [system](#).

localization The process of using measurements of the environment to determine an [agent's pose](#).

model A set of mathematical equations that reflects some aspect of a physical [system](#)'s behavior.

noise immunity The quality of a [system](#) to have its performance or stability unaffected by noise in the [outputs](#) (see also: [robustness](#)).

observer In control theory, a [system](#) that estimates the internal [state](#) of a given real [system](#) from measurements of the [input](#) and [output](#) of the real system.

open-loop gain The gain directly from the [input](#) to the [output](#), ignoring loops.

output Measurements from sensors.

output-based control Controls the [system's state](#) via the [outputs](#).

overshoot The amount by which a [system's state](#) surpasses the [reference](#) after rising toward it.

phase margin See section C.1 on gain and phase margin.

plant The [system](#) or collection of actuators being controlled.

pose The orientation of an [agent](#) in the world, which is represented by all or part of the [agent's state](#).

process variable The term used to describe the [output](#) of a [plant](#) in the context of PID control.

realization In control theory, this is an implementation of a given input-output behavior as a state-space [model](#).

reference The desired state. This value is used as the reference point for a controller's error calculation.

regulator A [controller](#) that attempts to minimize the [error](#) from a constant [reference](#) in the presence of disturbances.

rise time The time a [system](#) takes to initially reach the [reference](#) after applying a [step input](#).

robustness The quality of a feedback [control system](#) to remain stable in response to disturbances and uncertainty.

setpoint The term used to describe the [reference](#) of a PID controller.

settling time The time a [system](#) takes to settle at the [reference](#) after a [step input](#) is applied.

state A characteristic of a [system](#) (e.g., velocity) that can be used to determine the [system's future behavior](#).

state feedback Uses [state](#) instead of [output](#) in feedback.

steady-state error Error after [system](#) reaches equilibrium.

step input A [system input](#) that is 0 for $t < 0$ and a constant greater than 0 for $t \geq 0$. A step input that is 1 for $t \geq 0$ is called a unit step input.

step response The response of a [system](#) to a [step input](#).

stochastic process A process whose [model](#) is partially or completely defined by random variables.

system A term encompassing a [plant](#) and its interaction with a [controller](#) and [observer](#), which is treated as a single entity. Mathematically speaking, a [system](#) maps [inputs](#) to [outputs](#) through a linear combination of [states](#).

system response The behavior of a [system](#) over time for a given [input](#).

time-invariant The [system's](#) fundamental response does not change over time.

tracking In control theory, the process of making the output of a [control system](#) follow the [reference](#).

unity feedback A feedback network in a [control system](#) diagram with a feedback gain of 1.

Bibliography

Online

- [14] Wikipedia Commons. *Linear-quadratic regulator*. URL: https://en.wikipedia.org/wiki/Linear-quadratic_regulator (visited on 03/24/2018) (cited on page 201).
- [16] Declan Freeman-Gleason. *A Dive into WPILib Trajectories*. 2020. URL: <https://pietroglyph.github.io/trajectory-presentation/> (cited on page 172).
- [17] Sean Humbert. *Why do we have to linearize around an equilibrium point?* URL: https://www.cds.caltech.edu/%7Emurray/courses/cds101/fa02/faq/02-10-09_linearization.html (visited on 07/12/2018) (cited on page 58).
- [18] Simon J. Julier and Jeffrey K. Uhlmann. *A New Extension of the Kalman Filter to Nonlinear Systems*. URL: https://www.cs.unc.edu/~welch/kalman/media/pdf/Julier1997_SPIE_KF.pdf (visited on 09/29/2019) (cited on page 121).
- [19] Kapitanyuk, Proskurnikov, and Cao. *A guiding vector field algorithm for path following control of nonholonomic mobile robots*. URL: <https://arxiv.org/pdf/1610.04391.pdf> (visited on 08/09/2018) (cited on page 60).
- [20] Matthew Kelly. *An Introduction to Trajectory Optimization: How to Do Your Own Direct Collocation*. URL: https://www.matthewpeterkelly.com/research/MatthewKelly_IntroTrajectoryOptimization_SIAM_Review_2017.pdf (visited on 08/27/2019) (cited on page 173).
- [21] Edgar Kraft. *A Quaternion-based Unscented Kalman Filter for Orientation Tracking*. URL: <https://kodlab.seas.upenn.edu/uploads/Arun/UKFpaper.pdf> (visited on 07/11/2018) (cited on page 121).
- [22] Charles F. Van Loan. *Computing Integrals Involving the Matrix Exponential*. URL: <https://www.cs.cornell.edu/cv/ResearchPDF/computing.integrals.involving.Matrix.Exp.pdf> (visited on 06/21/2018) (cited on page 55).
- [23] Andrew W. Long et al. *The Banana Distribution is Gaussian: A Localization Study with Exponential Coordinates*. 2008. URL: https://rpk.lcsr.jhu.edu/wp-content/uploads/2014/08/p34_Long12_The-Banana-Distribution.pdf (cited on page 93).

- [24] Luca, Oriolo, and Vendittelli. *Control of Wheeled Mobile Robots: An Experimental Overview*. URL: <https://www.dis.uniroma1.it/~labrob/pub/papers/Ramsete01.pdf> (visited on 08/09/2018) (cited on pages 82, 83).
- [25] MIT OpenCourseWare. *Linear Quadratic Regulator*. URL: <https://ocw.mit.edu/courses/mechanical-engineering/2-154-maneuvering-and-control-of-surface-and-underwater-vehicles-13-49-fall-2004/lecture-notes/lec19.pdf> (visited on 06/26/2018) (cited on page 35).
- [26] Christoph Sprunk. *Planning Motion Trajectories for Mobile Robots Using Splines*. 2008. URL: <http://www2.informatik.uni-freiburg.de/~lau/students/Sprunk2008.pdf> (cited on page 172).
- [29] Russ Tedrake. *Underactuated Robotics*. 2019. URL: <http://underactuated.mit.edu/underactuated.html> (cited on page 60).
- [30] Russ Tedrake. *Chapter 9. Linear Quadratic Regulators*. URL: <http://underactuated.mit.edu/lqr.html> (visited on 03/22/2020) (cited on page 201).
- [31] Eric A. Wan and Rudolph van der Merwe. *The Unscented Kalman Filter for Nonlinear Estimation*. URL: <https://www.seas.harvard.edu/courses/cs281/papers/unscented.pdf> (visited on 06/26/2018) (cited on page 121).

Misc

- [1] FRC team 254. *Motion Planning and Control in FRC*. 2015. URL: <https://www.youtube.com/watch?v=8319J1BEHwM> (cited on page 172).
- [2] 3Blue1Brown. *Abstract vector spaces*. 2016. URL: <https://www.youtube.com/watch?v=TgKwz5Ikpc8> (cited on page 184).
- [3] 3Blue1Brown. *Eigenvectors and eigenvalues*. 2016. URL: <https://www.youtube.com/watch?v=PFDu9oVAE-g> (cited on page 28).
- [4] 3Blue1Brown. *Essence of linear algebra preview*. 2016. URL: <https://www.youtube.com/watch?v=kjBOesZCoqc> (cited on page 27).
- [5] 3Blue1Brown. *Inverse matrices, column space, and null space*. 2016. URL: <https://www.youtube.com/watch?v=uQhTuRIWMxw> (cited on page 28).
- [6] 3Blue1Brown. *Linear combinations, span, and basis vectors*. 2016. URL: <https://www.youtube.com/watch?v=k7RM-ot2NWY> (cited on page 27).
- [7] 3Blue1Brown. *Linear transformations and matrices*. 2016. URL: <https://www.youtube.com/watch?v=kYB8IZa5AuE> (cited on page 27).
- [8] 3Blue1Brown. *Matrix multiplication as composition*. 2016. URL: <https://www.youtube.com/watch?v=XkY2DOUCWMU> (cited on page 27).
- [9] 3Blue1Brown. *Nonsquare matrices as transformations between dimensions*. 2016. URL: https://www.youtube.com/watch?v=v8VSDg_WQIA (cited on page 28).
- [10] 3Blue1Brown. *The determinant*. 2016. URL: <https://www.youtube.com/watch?v=Ip3X9LOh2dk> (cited on page 28).
- [11] 3Blue1Brown. *Vectors, what even are they?* 2016. URL: https://www.youtube.com/watch?v=fNk_zzaMoSs (cited on page 27).
- [12] 3Blue1Brown. *Essence of calculus*. 2017. URL: <https://www.youtube.com/playlist?list=PLZHQB0bOWTQDMsr9K-rj53DwVRMYO3t5Yr> (cited on page 131).
- [13] 3Blue1Brown. *Taylor series*. 2017. URL: <https://www.youtube.com/watch?v=3d6DsjIBzJ4> (cited on page 53).
- [15] Brian Douglas. *Gain and Phase Margins Explained!* 2015. URL: <https://www.youtube.com/watch?v=ThoA4amCAX4> (cited on page 195).
- [27] Zach Star. *What does the Laplace Transform really tell us? A visual explanation (plus applications)*. 2019. URL: <https://www.youtube.com/watch?v=n2y7n6jw5d0> (cited on page 19).
- [28] Russ Tedrake. *6.832 Underactuated Robotics*. 2019. URL: https://www.youtube.com/channel/UChfUOAhz7ynELF-s_1LPpWg/videos (cited on page 60).

Index

block diagrams, 8
 simplification, 177

controller design
 actuator saturation, 17
 controllability, 32
 linear time-varying control, 72
 linear-quadratic regulator, 34
 Bryson's rule, 35
 definition, 35
 implicit model following, 203
 output feedback, 202

observability, 32
 pole placement, 34

digital signal processing
 aliasing, 47
 Nyquist frequency, 47

discretization, 48
 backward Euler method, 48
 bilinear transform, 48
 forward Euler method, 48
 matrix exponential, 52
 Taylor series, 53
 zero-order hold, 54

feedforward
 linear plant inversion, 41
 steady-state feedforward, 213

FRC models
 DC brushed motor equations, 146
 differential drive equations, 70
 elevator equations, 61
 flywheel equations, 63
 single-jointed arm equations, 66

gain, 8

integral control
 input error estimation, 43
 plant augmentation, 43

linear algebra
 basis vectors, 27
 linear combination, 27
 vectors, 27

matrices
 determinant, 28
 eigenvalues, 28
 inverse, 28
 linear systems, 28
 linear transformation, 27
 multiplication, 27
 pseudoinverse, 28
 rank, 28
 transpose, 28

model augmentation

- of controller, 37
- of observer, 37
- of output, 38
- of plant, 37
- motion profiles
 - 1 DOF, 169
 - 2 DOF, 172
 - S-curve, 170
 - trajectory optimization, 173
 - trapezoidal, 170
- nonlinear control
 - extended Kalman filter, 120
 - linear time-varying control, 72
 - linearization, 57, 69
 - Lyapunov stability, 58
 - Unscented Kalman filter, 120
- optimal control
 - linear plant inversion, 41
 - linear time-varying control, 72
 - linear-quadratic regulator, 34
 - Bryson's rule, 35
 - definition, 35
- physics
 - conservation of energy, 156
 - sum of forces, 154
 - sum of torques, 155
- PID control, 11, 29, 190
- probability, 96
 - Bayes's rule, 100
 - central limit theorem, 103
 - conditional expectation, 100
 - conditional probability density functions, 100
 - conditional variances, 100
 - covariance, 99
- covariance matrix, 101
- expected value, 97
- marginal probability density functions, 100
- probability density function, 96
- probability density functions, 98
- random variables, 96
- variance, 97
- stability
 - eigenvalues, 34, 95
 - gain margin, 195
 - phase margin, 195
 - poles and zeroes, 20
 - root locus, 22
- state-space controllers
 - closed-loop, 33
 - open-loop, 31
- state-space observers
 - Extended Kalman filter, 120
 - Kalman filter
 - as Luenberger observer, 112
 - derivations, 103, 105
 - equations, 108
 - setup, 110
 - Kalman smoother, 114
 - derivations, 114
 - equations, 117
 - Luenberger observer, 94
 - multiple model adaptive estimation, 121
 - Unscented Kalman filter, 120
- steady-state error, 188
- stochastic
 - linear systems, 103
 - measurement noise, 103
 - process noise, 103
 - two-sensor problem, 104

